

Amazon.AWS-Solutions-Architect-Professional.v2021-01-26.q127

Exam Code:	AWS-Solutions-Architect-Professional
Exam Name:	AWS Certified Solutions Architect - Professional
Certification Provider:	Amazon
Free Question Number:	127
Version:	v2021-01-26
# of views:	1584
# of Questions views:	46197
https://www.freecram.com/torrent/Amazon.AWS-Solutions-Architect-Professional.v2021-01-26.q127.html	

NEW QUESTION: 1

A solutions architect is designing a publicly accessible web application that is on an Amazon CloudFront distribution with an Amazon S3 website endpoint as the origin. When the solution is deployed, the website returns an Error 403: Access Denied message. Which steps should the solutions architect take to correct the issue? (Select TWO.)

- A. Remove the origin access identity (OAI) from the CloudFront distribution
- B. Change the storage class from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA)
- C. Remove the S3 block public access option from the S3 bucket
- D. Disable S3 object versioning
- E. Remove the requester pays option from the S3 bucket

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 2

A financial services company receives a regular data feed from its credit card servicing partner. Approximately

5,000 records are sent every 15 minutes in plaintext delivered over HTTPS directly into an Amazon S3 bucket with server-side encryption. This feed contains sensitive credit card primary account number (PAN) data. The company needs to automatically mask the PAN before sending the data to another S3 bucket for additional internal processing. The company also needs to remove and merge specific fields and then transform the record into JSON format. Additionally, extra feeds are likely to be added in the future, so any design needs to be easily expandable. Which solutions will meet these requirements?

- A. Trigger an AWS Lambda function on file delivery that extracts each record and writes it to an Amazon SQS queue. Configure an AWS Fargate container application to automatically scale to a

single instance when the SQS queue contains messages Have the application process each record and transform the record into JSON format When the queue is empty send the results to another S3 bucket for internal processing and scale down the AWS Fargate instance

B. Trigger an AWS Lambda function on file delivery that extracts each record and writes it to an Amazon SQS queue Trigger another Lambda function when new messages arrive in the SQS queue to process the records, writing the results to a temporary location in Amazon S3 Trigger a final Lambda function once the SQS queue is empty to transform the records into JSON format and send the results to another S3 bucket for internal processing n

C. Create an AWS Glue crawler and custom classifier based on the data feed formats and build a table definition to match Trigger an AWS Lambda function on file delivery to start an AWS Glue ETL job to transform the entire record according to the processing and transformation requirements Define the output format as JSON Once complete have the ETL job send the results to another S3 bucket for internal processing

D. Create an AWS Glue crawler and custom classifier based upon the data feed formats and build a table definition to match Perform an Amazon Athena query on file delivery to start an Amazon EMR ETL job to transform the entire record according to the processing and transformation requirements Define the output format as JSON Once complete send the results to another S3 bucket for internal processing and scale down the EMR cluster

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 3

An AWS customer has a web application that runs on premises. The web application (etches data from a third party API that is behind a firewall. The third party accepts only one public CIDR block in each client's allow list The customer wants to migrate their web application to the AWS Cloud The application will be hosted on a set of Amazon EC2 instances behind an Application Load Balancer (ALB) in a VPC. The ALB is located in public subnets The EC2 instances are located in private subnets. NAT gateways provide internet access to the private subnets.

How should a solutions architect ensure that the web application can continue to call the third-party API after the migration?

A. Associate a block of customer owned public IP addresses to the VPC Enable public IP addressing for public subnets in the VPC

B. Create Elastic IP addresses from the block of customer owned IP addresses Assign the static Elastic IP addresses to the ALB

C. Register a block of customer-owned public IP addresses in the AWS account Create Elastic IP addresses from the address block and assign them to the NAT gateways in the VPC

D. Register a block of customer-owned public IP addresses in the AWS account Set up AWS Global Accelerator to use Elastic IP addresses from the address block Set the ALB as the accelerator endpoint

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 4

A company runs a Windows Server host in a public subnet that is configured to allow a team of administrators to connect over RDP to troubleshoot issues with hosts in a private subnet. The host must be available at all times outside of a scheduled maintenance window, and needs to receive the latest operating system updates within 3 days of release.

What should be done to manage the host with the LEAST amount of administrative effort?

A. Run the host in AWS OpsWorks Stacks. Use a Chef recipe to harden the AMI during instance launch.

Use an AWS Lambda scheduled event to run the Upgrade Operating System stack command to apply system updates.

B. Run the host in a single-instance AWS Elastic Beanstalk environment. Configure the environment with a custom AMI to use a hardened machine image from AWS Marketplace. Apply system updates with AWS Systems Manager Patch Manager.

C. Run the host on AWS WorkSpaces. Use Amazon WorkSpaces Application Manager (WAM) to harden the host. Configure Windows automatic updates to occur every 3 days.

D. Run the host in an Auto Scaling group with a minimum and maximum instance count of 1. Use a hardened machine image from AWS Marketplace. Apply system updates with AWS Systems Manager Patch Manager.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 5

A company is using AWS Organizations to manage multiple AWS accounts. For security purposes, the company requires the creation of an Amazon Simple Notification Service (Amazon SNS) topic that enables integration with a third-party alerting system in all the Organizations member accounts.

A solutions architect used an AWS CloudFormation template to create the SNS topic and stack sets to automate the deployment of CloudFormation stacks. Trusted access has been enabled in Organizations.

What should the solutions architect do to deploy the CloudFormation StackSets in all AWS accounts?

A. Create stacks in the Organizations master account. Use service-managed permissions. Set deployment options to deploy to the organization. Enable CloudFormation StackSets automatic deployment.

B. Create stacks in the Organizations member accounts. Use self-service permissions. Set deployment options to deploy to an organization. Enable the CloudFormation StackSets automatic deployment.

C. Create a stack set in the Organizations member accounts. Use service-managed permissions. Set deployment options to deploy to an organization. Use CloudFormation StackSets drift detection.

D. Create stacks in the Organization master account. Use service-managed permissions. Set deployment options to deploy to the organization. Enable CloudFormation StackSets drift detection.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 6

A company has several teams, and each team has their own Amazon RDS database that totals 100 TB. The company is building a data query platform for Business Intelligence Analysts to generate a weekly business report. The new system must run ad-hoc SQL queries. What is the MOST cost-effective solution?

- A.** Create an Amazon EMR cluster with enough core nodes. Run an Apache Spark job to copy data from the RDS databases to an Hadoop Distributed File System (HDFS). Use a local Apache Hive metastore to maintain the table definition. Use Spark SQL to run the query.
- B.** Use an AWS Glue crawler to crawl all the databases and create tables in the AWS Glue Data Catalog. Use an AWS Glue ETL Job to load data from the RDS databases to Amazon S3, and use Amazon Athena to run the queries.
- C.** Use an AWS Glue ETL job to copy all the RDS databases to a single Amazon Aurora PostgreSQL database. Run SQL queries on the Aurora PostgreSQL database.
- D.** Create a new Amazon Redshift cluster. Create an AWS Glue ETL job to copy data from the RDS databases to the Amazon Redshift cluster. Use Amazon Redshift to run the query.

Answer: **B** ([LEAVE A REPLY](#))

NEW QUESTION: 7

A retail company is running an application that stores invoice files in an Amazon S3 bucket and metadata about the files in an Amazon DynamoDB table. The S3 bucket and DynamoDB table are in us-east-1. The company wants to protect itself from data corruption and loss of connectivity to either Region. Which option meets these requirements?

- A.** Create a DynamoDB global table to replicate data between us-east-1 and eu-west-1. Enable continuous backup on the DynamoDB table in us-east-1. Enable versioning on the S3 bucket.
- B.** Create an AWS Lambda function triggered by Amazon CloudWatch Events to make regular backups of the DynamoDB table. Set up S3 cross-region replication from us-east-1 to eu-west-1. Set up MFA delete on the S3 bucket in us-east-1.
- C.** Create a DynamoDB global table to replicate data between us-east-1 and eu-west-1. Enable versioning on the S3 bucket. Implement strict ACLs on the S3 bucket.
- D.** Create a DynamoDB global table to replicate data between us-east-1 and eu-west-1. Enable continuous backup on the DynamoDB table in us-east-1. Set up S3 cross-region replication from us-east-1 to eu-west-1.

Answer: **D** ([LEAVE A REPLY](#))

Explanation

<https://aws.amazon.com/blogs/aws/new-cross-region-replication-for-amazon-s3/>

NEW QUESTION: 8

A company is migrating to the cloud. It wants to evaluate the configurations of virtual machines in its existing data center environment to ensure that it can size new Amazon EC2 instances

accurately. The company wants to collect metrics, such as CPU, memory, and disk utilization, and it needs an inventory of what processes are running on each instance. The company would also like to monitor network connections to map communications between servers.

Which would enable the collection of this data MOST cost effectively?

- A.** Enable AWS Application Discovery Service in the AWS Management Console and configure the corporate firewall to allow scans over a VPN.
- B.** Use AWS Application Discovery Service and deploy the data collection agent to each virtual machine in the data center.
- C.** Configure the Amazon CloudWatch agent on all servers within the local environment and publish metrics to Amazon CloudWatch Logs.
- D.** Use AWS Application Discovery Service and enable agentless discovery in the existing virtualization environment.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 9

A company is using an existing orchestration tool to manage thousands of Amazon EC2 instances. A recent penetration test found a vulnerability in the company's software stack. This vulnerability has prompted the company to perform a full evaluation of its current production environment. The analysis determined that the following vulnerabilities exist within the environment

- * Operating systems with outdated libraries and known vulnerabilities are being used in production
- * Relational databases hosted and managed by the company are running unsupported versions with known vulnerabilities
- * Data stored in databases is not encrypted

The solutions architect intends to use AWS Config to continuously audit and assess the compliance of the company's AWS resource configurations with the company's policies and guidelines. What additional steps will enable the company to secure its environments and track resources while adhering to best practices?

- A.** Install the AWS Systems Manager Agent on all existing instances using the company's current orchestration tool. Use the Systems Manager Run Command to execute a list of commands to upgrade software on each instance using operating system-specific tools. Enable AWS KMS encryption on all Amazon EBS volumes.
- B.** Create an AWS CloudFormation template for the EC2 instances. Use EC2 user data in the CloudFormation template to install the AWS Systems Manager Agent, and enable AWS KMS encryption on all Amazon EBS volumes. Have CloudFormation replace all running instances. Use Systems Manager Patch Manager to establish a patch baseline and deploy a Systems Manager Maintenance Windows task to execute AWS-RunPatchBaseline using the patch baseline.
- C.** Use AWS Application Discovery Service to evaluate all running EC2 instances. Use the AWS CLI to modify each instance, and use EC2 user data to install the AWS Systems Manager Agent.

during boot Schedule patching to run as a Systems Manager Maintenance Windows task Migrate all relational databases to Amazon RDS and enable AWS KMS encryption

D. Install the AWS Systems Manager Agent on all existing instances using the company's current orchestration tool Migrate all relational databases to Amazon RDS and enable AWS KMS encryption Use Systems Manager Patch Manager to establish a patch baseline and deploy a Systems Manager Maintenance Windows task to execute AWS-RunPatchBaseline using the patch baseline.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 10

A company is building an AWS landing zone and has asked a Solutions Architect to design a multi-account access strategy that will allow hundreds of users to use corporate credentials to access the AWS Console. The company is running a Microsoft Active Directory and users will use an AWS Direct Connect connection to connect to AWS. The company also wants to be able to federate to third-party services and providers, including custom applications.

Which solution meets the requirements by using the LEAST amount of management overhead?

- A.** Connect the Active Directory to AWS by using single sign-on and an Active Directory Federation Services (AD FS) with SAML 2.0, and then configure the identity Provider (IdP) system to use form-based authentication. Build the AD FS portal page with corporate branding, and integrate third-party applications that support SAML 2.0 as required.
- B.** Create a two-way Forest trust relationship between the on-premises Active Directory and the AWS Directory Service. Set up AWS Single Sign-On with AWS Organizations. Use single sign-on integrations for connections with third-party applications.
- C.** Configure single sign-on by connecting the on-premises Active Directory using the AWS Directory Service AD Connector. Enable federation to the AWS services and accounts by using the IAM applications and services linking function. Leverage third-party single sign-on as needed.
- D.** Connect the company's Active Directory to AWS by using AD FS and SAML 2.0. Configure the AD FS claim rule to leverage Regex and a common Active Directory naming convention for the security group to allow federation of all AWS accounts. Leverage third-party single sign-on as needed, and add it to the AD FS server.

Answer: D ([LEAVE A REPLY](#))

Explanation

<https://aws.amazon.com/blogs/security/aws-federated-authentication-with-active-directory-federation-services-ad>

NEW QUESTION: 11

A company plans to migrate to AWS. A solutions architect uses AWS Application Discovery Service over the fleet and discovers that there is an Oracle data warehouse and several PostgreSQL databases.

Which combination of migration patterns will reduce licensing costs and operational overhead? (Select TWO.)

- A. Lift and shift the PostgreSQL databases to Amazon EC2 using AWS DMS.
- B. Lift and shift the Oracle data warehouse to Amazon EC2 using AWS DMS.
- C. Migrate the Oracle data warehouse to Amazon Redshift using AWS SCT and AWS DMS
- D. Migrate the PostgreSQL databases to Amazon RDS for PostgreSQL using AWS DMS.
- E. Migrate the Oracle data warehouse to an Amazon EMR managed cluster using AWS DMS.

Answer: C,D ([LEAVE A REPLY](#))

NEW QUESTION: 12

A company wants to migrate its on-premises data center to the AWS Cloud. This includes thousands of virtualized Linux and Microsoft Windows servers, SAN storage, Java and PHP applications with MYSQL, and Oracle databases. There are many dependent services hosted either in the same data center or externally. The technical documentation is incomplete and outdated. A solution architect needs to understand the current environment and estimate the cloud resource costs after the migration.

Which tools or services should solution architect use to plan the cloud migration (Select THREE.)

- A. AWS x-Ray
- B. AWS SMS
- C. AWS Cloud Adoption Readiness Tool (CART)
- D. AWS Application Discovery Service
- E. Amazon Inspector
- F. AWS Migration Hub

Answer: C,D,F ([LEAVE A REPLY](#))

NEW QUESTION: 13

A solutions architect is designing a web application on AWS that requires 99.99% availability. The application will consist of a three-tier architecture that supports 300.000 web requests each minute when experiencing peak traffic. The application will use Amazon Route 53 for DNS resolution. Amazon CloudFront as the content delivery network (CDN), an Elastic Load Balancer for load balancing. Amazon EC2 Auto Scaling groups to scale the application tier, and Amazon Aurora MySQL as the backend database. The backend database load will average 90% reads and 10% writes. The company wants to build a cost-effective solution, but reliability is critical.

Which set of strategies should the solutions architect use?

- A. Build the application in a single AWS Region. Deploy the EC2 application layer to three Availability Zones using an Auto Scaling group with dynamic scaling based on request metrics. Use a Multi-AZ Amazon Aurora MySQL DB cluster with two Aurora Replicas. Each Aurora Replica must have enough capacity to support 50% of the peak read queries.
- B. Build the application in a single AWS Region. Deploy the EC2 application layer to two Availability Zones using an Auto Scaling group with a minimum desired capacity sufficient to process 300.000 requests each minute. Use a Multi-AZ Amazon Aurora MySQL DB cluster with one Aurora Replica.

The Aurora Replica must have enough capacity to support 50% of the peak read and write queries.

C. Build the application in two AWS Regions Deploy the EC2 application layer to two Availability Zones using an Auto Scaling group with dynamic scaling based on the request metrics in each Region. In the second Region, deploy an Amazon Aurora MySQL cross-Region replica. Use Amazon Route 53 to distribute traffic between Regions and configure failover if a Region becomes unavailable.

D. Build the application in a single AWS Region. Deploy the EC2 application layer to three Availability Zones using an Auto Scaling group with a minimum desired capacity sufficient to process 450,000 requests each minute. Use a Multi-AZ Amazon Aurora MySQL DB duster with two Aurora Replicas.

Each Aurora Replica must have enough capacity to support 100% of the peak read queries.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 14

A company has an application that uses Amazon EC2 instances in an Auto Scaling group. The Quality Assurance (QA) department needs to launch a large number of short-lived environments to test the application.

The application environments are currently launched by the Manager of the department using an AWS CloudFormation template. To launch the stack, the Manager uses a role with permission to use CloudFormation, EC2 and Auto Scaling APIs. The Manager wants to allow testers to launch their own environments, but does not want to grant broad permission to each user. Which set up would achieve these goals?

A. Upload the AWS CloudFormation template to Amazon S3. Give users in the QA department permission to assume the Manager's role and add a policy that restricts the permissions to the template and the resources it creates. Train users to launch the template from the CloudFormation console.

B. Create an AWS Service Catalog product from the environment template. Add a launch constraint to the product with the existing role. Give users in the QA department permission to use AWS Service Catalog APIs only. Train users to launch the templates from the AWS Service Catalog console.

C. Upload the AWS CloudFormation template to Amazon S3. Give users in the QA department permission to use CloudFormation and S3 APIs, with conditions that restrict the permission to the template and the resources it creates. Train users to launch the template from the CloudFormation console.

D. Create an AWS Elastic Beanstalk application from the environment template. Give users in the QA department permission to use Elastic Beanstalk permissions only. Train users to launch Elastic beanstalk environments with the Elastic Beanstalk CLI, passing the existing role to the environment as a service role.

Answer: B ([LEAVE A REPLY](#))

Explanation

<https://aws.amazon.com/blogs/mt/how-to-launch-secure-and-governed-aws-resources-with-aws-cloudformation->

NEW QUESTION: 15

A company is adding a new approved external vendor that only supports IPv6 connectivity. The company's backend systems sit in the private subnet of an Amazon VPC. The company uses a NAT gateway to allow these systems to communicate with external vendors over IPv4. Company policy requires systems that communicate with external vendors use a security group that limits access to only approved external vendors.

The virtual private cloud (VPC) uses the default network ACL.

The Systems Operator successfully assigns IPv6 addresses to each of the backend systems. The Systems Operator also updates the outbound security group to include the IPv6 CIDR of the external vendor (destination). The systems within the VPC are able to ping one another successfully over IPv6. However, these systems are unable to communicate with the external vendor.

What changes are required to enable communication with the external vendor?

- A.** Create an IPv6 NAT instance. Add a route for destination 0.0.0.0/0 pointing to the NAT instance.
- B.** Enable IPv6 on the NAT gateway. Add a route for destination ::/0 pointing to the NAT gateway.
- C.** Enable IPv6 on the internet gateway. Add a route for destination 0.0.0.0/0 pointing to the IGW.
- D.** Create an egress-only internet gateway. Add a route for destination ::/0 pointing to the gateway.

Answer: ([SHOW ANSWER](#))

Explanation

<https://docs.aws.amazon.com/vpc/latest/userguide/egress-only-internet-gateway.html>

NEW QUESTION: 16

A company runs a three-tier application in AWS. Users report that the application performance can vary greatly depending on the time of day and functionality being accessed.

The application includes the following components:

- * Eight t2.large front-end web servers that serve static content and proxy dynamic content from the application tier.
- * Four t2.large application servers.
- * One db.m4.large Amazon RDS MySQL Multi-AZ DB instance.

Operations has determined that the web and application tiers are network constrained.

Which of the following should cost effectively improve application performance? (Choose two.)

- A.** Replace web and app tiers with t2.xlarge instances
- B.** Use AWS Auto Scaling and m4.large instances for the web and application tiers
- C.** Convert the MySQL RDS instance to a self-managed MySQL cluster on Amazon EC2
- D.** Create an Amazon CloudFront distribution to cache content
- E.** Increase the size of the Amazon RDS instance to db.m4.xlarge

Answer: B,D ([LEAVE A REPLY](#))

Explanation

<https://aws.amazon.com/ec2/instance-types/>

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (216 Q&As Dumps, **30%OFF** Special Discount: **freecram**)

NEW QUESTION: 17

A company is creating a centralized logging service running on Amazon EC2 that will receive and analyze logs from hundreds of AWS accounts. AWS PrivateLink is being used to provide connectivity between the client services and the logging service.

In each AWS account with a client an interface endpoint has been created for the logging service and is available. The logging service running on EC2 instances with a Network Load Balancer (NLB) are deployed in different subnets. The clients are unable to submit logs using the VPC endpoint.

Which combination of steps should a solutions architect take to resolve this issue? (Select TWO)

- A.** Check the security group for the logging service running on the EC2 instances to ensure it allows ingress from the NLB subnets.
- B.** Check the security group for the NLB to ensure it allows ingress from the interface endpoint subnets
- C.** Check that the NACL is attached to the logging service subnet to allow communications to and from the NLB subnets Check that the NACL is attached to the NLB subnet to allow communications to and from the logging service subnets running on EC2 instances.
- D.** Check the security group for the logging service running on EC2 instances to ensure it allows ingress from the clients
- E.** Check that the NACL is attached to the logging service subnets to allow communications to and from the interface endpoint subnets Check that the NACL is attached to the interface endpoint subnet to allow communications to and from the logging service subnets running on EC2 instances.

Answer: A,E ([LEAVE A REPLY](#))

NEW QUESTION: 18

An online magazine will launch its latest edition this month. This edition will be the first to be distributed globally. The magazine's dynamic website currently uses an Application Load Balance

in front of the web tier, a fleet of Amazon EC2 instances for web and application servers, and Amazon Aurora MySQL. Portions of the website include static content and almost all traffic is read-only.

The magazine is exporting a significant spike in internet traffic when the new edition is launched. Optimal performance is a top priority for the week following the launch.

Which combination of steps should a solutions architect take to reduce system response times for a global audience? (Select Two.)

- A.** Migrate the database from Amazon Aurora to Amazon RDS for MySQL. Ensure all three of the application tiers—web, application, and database—are in private subnets.
- B.** Introduce Amazon Route 53 with latency-based routing and Amazon CloudFront distributions. Ensure the web and application tiers are each in Auto Scaling groups.
- C.** Use an Aurora global database for physical cross-Region replication. Use Amazon S3 with cross-Region replication for static content and resources. Deploy the web and application tiers in Regions across the world.
- D.** Ensure the web and application tiers are each in Auto Scaling groups. Introduce an AWS Direct Connect connection. Deploy the web and application tiers in Regions across the world.
- E.** Use logical cross-Region replication to replicate the Aurora MySQL database to a secondary Region. Replace the web servers with Amazon S3. Deploy S3 buckets in cross-Region replication mode.

Answer: B,C ([LEAVE A REPLY](#))

NEW QUESTION: 19

A company has released a new version of a website to target an audience in Asia and South America. The website's media assets are hosted on Amazon S3 and have an Amazon CloudFront distribution to improve end-user performance. However, users are having a poor login experience; the authentication service is only available in the us-east-1 AWS Region. How can the Solutions Architect improve the login experience and maintain high security and performance with minimal management overhead?

- A.** Replicate the setup in each new geography and use Amazon Route 53 geo-based routing to route traffic to the AWS Region closest to the users.
- B.** Use an Amazon Route 53 weighted routing policy to route traffic to the CloudFront distribution. Use CloudFront cached HTTP methods to improve the user login experience.
- C.** Use Amazon Lambda@Edge attached to the CloudFront viewer request trigger to authenticate and authorize users by maintaining a secure cookie token with a session expiry to improve the user experience in multiple geographies.
- D.** Replicate the setup in each geography and use Network Load Balancers to route traffic to the authentication service running in the closest region to users.

Answer: ([SHOW ANSWER](#)**)**

Explanation

There are several benefits to using Lambda@Edge for authorization operations. First, performance is improved by running the authorization function using Lambda@Edge closest to

the viewer, reducing latency and response time to the viewer request. The load on your origin servers is also reduced by offloading CPU-intensive operations such as verification of JSON Web Token (JWT) signatures. Finally, there are security benefits such as filtering out unauthorized requests before they reach your origin infrastructure.

<https://aws.amazon.com/blogs/networking-and-content-delivery/authorization-edge-how-to-use-lambda-edge-and->

NEW QUESTION: 20

A company hosts a game player-matching service on a public facing, physical, on-premises instance that all users are able to access over the internet. All traffic to the instance uses UDP. The company wants to migrate the service to AWS and provide a high level of security. A solutions architect needs to design a solution for the player-matching service using AWS. Which combination of steps should the solutions architect take to meet these requirements? (Select THREE)

- A. Enable AWS Shield Advanced on all public-facing resources.
- B. Configure a network ACL rule to block all non-UDP traffic. Associate the network ACL with the subnets that hold the load balancer instances.
- C. Use an Application Load Balancer (ALB) in front of the player-matching instance. Use a friendly DNS entry in Amazon Route 53 pointing to the ALB's internet-facing fully qualified domain name (FQDN).
- D. Use a Network Load Balancer (NLB) in front of the player-matching instance. Use a friendly DNS entry in Amazon Route 53 pointing to the NLB's Elastic IP address
- E. Use Amazon CloudFront with an Elastic Load Balancer as an origin.
- F. Define an AWS WAF rule to explicitly drop non-UDP traffic, and associate the rule with the load balancer. .

Answer: A,B,D ([LEAVE A REPLY](#))

NEW QUESTION: 21

A company runs an ordering system on AWS using Amazon SQS and AWS Lambda, with each order received as a JSON message. recently the company had a marketing event that led to a tenfold increase in orders. With this increase, the following undesired behaviors started in the ordering system:

- * Lambda failures while processing orders lead to queue backlogs.
- * The same orders have been processed multiple times.

A solutions Architect has been asked to solve the existing issues with the ordering system and add the following resiliency features:

- * Retain problematic orders for analysis.
- * Send notification if errors go beyond a threshold value.

How should the Solutions Architect meet these requirements?

- A. Receive multiple messages with each Lambda invocation, use long polling when receiving the messages, log the errors from the message processing code using Amazon CloudWatch Logs,

create a dead letter queue with AWS Lambda to capture failed invocations, create CloudWatch events on Lambda errors for notification.

B. Receive single messages with each Lambda invocation, put additional Lambda workers to poll the queue, delete messages after processing, increase the message timer for the messages, use Amazon CloudWatch Logs for messages that could not be processed, create a CloudWatch alarm on Lambda errors for notification.

C. Receive multiple messages with each Lambda invocation, add error handling to message processing code and delete messages after processing, increase the visibility timeout for the messages, create a delay queue for messages that could not be processed, create an Amazon CloudWatch metric on Lambda errors for notification.

D. Receive multiple messages with each Lambda invocation, add error handling to message processing code and delete messages after processing, increase the visibility timeout for the messages, create a dead letter queue for messages that could not be processed, create an Amazon CloudWatch alarm on Lambda errors for notification.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 22

A large company is migrating its entire IT portfolio to AWS. Each business unit in the company has a standalone AWS account that supports both development and test environments. New accounts to support production workloads will be needed soon.

The Finance department requires a centralized method for payment but must maintain visibility into each group's spending to allocate costs.

The Security team requires a centralized mechanism to control IAM usage in all the company's accounts.

What combination of the following options meet the company's needs with LEAST effort? (Choose two.)

A. Use a collection of parameterized AWS CloudFormation templates defining common IAM permissions that are launched into each account. Require all new and existing accounts to launch the appropriate stacks to enforce the least privilege model.

B. Use AWS Organizations to create a new organization from a chosen payer account and define an organizational unit hierarchy. Invite the existing accounts to join the organization and create new accounts using Organizations.

C. Require each business unit to use its own AWS accounts. Tag each AWS account appropriately and enable Cost Explorer to administer chargebacks.

D. Enable all features of AWS Organizations and establish appropriate service control policies that filter IAM permissions for sub-accounts.

E. Consolidate all of the company's AWS accounts into a single AWS account. Use tags for billing purposes and IAM's Access Advice feature to enforce the least privilege model.

Answer: B,D ([LEAVE A REPLY](#))

Explanation

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ce-what-is.html>

NEW QUESTION: 23

A company is currently in the design phase of an application that will need an RPO of less than 5 minutes and an RTO of less than 10 minutes. The solutions architecture team is forecasting that the database will store approximately 10 TB of data. As part of the design, they are looking for a database solution that will provide the company with the ability to fail over to a secondary Region. Which solution will meet these business requirements at the LOWEST cost?

- A.** Deploy an Amazon RDS instance with a cross-Region read replica in a secondary Region. In the event of a failure, promote the read replica to become the primary.
- B.** Deploy an Amazon Aurora DB cluster in the primary Region and another in a secondary Region. Use AWS DMS to keep the secondary Region in sync.
- C.** Deploy an Amazon Aurora DB cluster and take snapshots of the cluster every 5 minutes. Once a snapshot is complete, copy the snapshot to a secondary Region to serve as a backup in the event of a failure.
- D.** Deploy an Amazon RDS instance with a read replica in the same Region. In the event of a failure, promote the read replica to become the primary.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 24

A development team has created a new flight tracker application that provides near-real-time data to users. The application has a front end that consists of an Application Load Balancer (ALB) in front of two large Amazon EC2 instances in a single Availability Zone. Data is stored in a single Amazon RDS MySQL DB instance. An Amazon Route 53 DNS record points to the ALB.

Management wants the development team to improve the solution to achieve maximum reliability with the least amount of operational overhead.

Which set of actions should the team take?

- A.** Replace the DB instance with Amazon Aurora with Aurora Replicas. Deploy the application to multiple smaller EC2 instances across multiple Availability Zones in an Auto Scaling group behind an ALB.
- B.** Create RDS MySQL read replicas. Deploy the application to multiple AWS Regions. Use a Route 53 latency-based routing to route to the application.
- C.** Configure the DB instance as Multi-AZ. Deploy the application to two additional EC2 instances in different Availability Zones behind an ALB.
- D.** Replace the DB instance with Amazon DynamoDB global tables. Deploy the application in multiple AWS Regions. Use a Route 53 latency-based routing policy to route to the application.

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 25

A company is migrating its marketing website and content management system from an on-premises data center to AWS. The company wants the AWS application to be developed in a

VPC with Amazon EC2 instances used for the web servers and an Amazon RDS instance for the database.

The company has a runbook document that describes the installation process of the on-premises system. The company would like to base the AWS system on the processes referenced in the runbook document. The runbook document describes the installation and configuration of the operating systems, network settings, the website, and content management system software on the servers. After the migration is complete, the company wants to be able to make changes quickly to take advantage of other AWS features.

How can the application and environment be deployed and automated in AWS, while allowing for future changes?

- A.** Update the runbook to describe how to create the VPC, the EC2 instances, and the RDS instance for the application by using the AWS Console. Make sure that the rest of the steps in the runbook are updated to reflect any changes that may come from the AWS migration.
- B.** Write an AWS CloudFormation template that creates the VPC, the EC2 instances, and the RDS instance for the application. Ensure that the rest of the steps in the runbook are updated to reflect any changes that may come from the AWS migration.
- C.** Write a Python script that uses the AWS API to create the VPC, the EC2 instances, and the RDS instance for the application. Write shell scripts that implement the rest of the steps in the runbook. Have the Python script copy and run the shell scripts on the newly created instances to complete the installation.
- D.** Write an AWS CloudFormation template that creates the VPC, the EC2 instances, and the RDS instance for the application. Include EC2 user data in the AWS CloudFormation template to install and configure the software.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 26

A Company has a security event whereby an Amazon S3 bucket with sensitive information was made public.

Company policy is to never have public S3 objects, and the Compliance team must be informed immediately when any public objects are identified.

How can the presence of a public S3 object be detected, set to trigger alarm notifications, and automatically remediated in the future? (Choose two.)

- A.** Turn on object-level logging for Amazon S3. Turn on Amazon S3 event notifications to notify by using an Amazon SNS topic when a PutObject API call is made with a public-read permission.
- B.** Configure an Amazon CloudWatch Events rule that invokes an AWS Lambda function to secure the S3 bucket.
- C.** Use the S3 bucket permissions for AWS Trusted Advisor and configure a CloudWatch event to notify by using Amazon SNS.
- D.** Turn on object-level logging for Amazon S3. Configure a CloudWatch event to notify by using an SNS topic when a PutObject API call with public-read permission is detected in the AWS CloudTrail logs.

E. Schedule a recursive Lambda function to regularly change all object permissions inside the S3 bucket.

Answer: ([SHOW ANSWER](#))

Explanation

<https://aws.amazon.com/blogs/security/how-to-detect-and-automatically-remediate-unintended-permissions-in-am>

NEW QUESTION: 27

An enterprise company's data science team wants to provide a safe, cost-effective way to provide easy access to Amazon SageMaker. The data scientists have limited AWS knowledge and need to be able to launch a Jupyter notebook instance. The notebook instance needs to have a preconfigured AWS KMS key to encrypt data at rest on the machine learning storage volume without exposing the complex setup requirements.

Which approach will allow the company to set up a self-service mechanism for the data scientists to launch Jupyter notebooks in its AWS accounts with the LEAST amount of operational overhead?

A. Create an AWS CloudFormation template to launch a Jupyter notebook instance using the `AWS::SageMaker::NotebookInstance` resource type with a preconfigured KMS key. Add a user-friendly name to the CloudFormation template. Display the URL to the notebook using the Outputs section.

Distribute the CloudFormation template to the data scientists using a shared Amazon S3 bucket.

B. Create an AWS CloudFormation template to launch a Jupyter notebook instance using the `AWS::SageMaker::NotebookInstance` resource type with a preconfigured KMS key. Simplify the parameter names, such as the instance size, by mapping them to Small, Large, and X-Large using the Mappings section in CloudFormation. Display the URL to the notebook using the Outputs section, then upload the template into an AWS Service Catalog product in the data scientist's portfolio, and share it with the data scientist's IAM role.

C. Create an AWS CLI script that the data scientists can run locally. Provide step-by-step instructions about the parameters to be provided while executing the AWS CLI script to launch a Jupyter notebook with a preconfigured KMS key. Distribute the CLI script to the data scientists using a shared Amazon S3 bucket.

D. Create a serverless front end using a static Amazon S3 website to allow the data scientists to request a Jupyter notebook instance by filling out a form. Use Amazon API Gateway to receive requests from the S3 website and trigger a central AWS Lambda function to make an API call to Amazon SageMaker that launch a notebook instance with a preconfigured KMS key for the data scientists. Then call back to the front-end website to display the URL to the notebook instance.

Answer: A ([LEAVE A REPLY](#))

NEW QUESTION: 28

A Solutions Architect is designing a highly available and reliable solution for a cluster of Amazon EC2 instances.

The Solutions Architect must ensure that any EC2 instance within the cluster recovers automatically after a system failure. The solution must ensure that the recovered instance maintains the same IP address.

How can these requirements be met?

- A.** Create an AWS Lambda script to restart any EC2 instances that shut down unexpectedly.
- B.** Create an Auto Scaling group for each EC2 instance that has a minimum and maximum size of 1.
- C.** Create a new t2.micro instance to monitor the cluster instances. Configure the t2.micro instance to issue an `aws ec2 reboot-instances` command upon failure.
- D.** Create an Amazon CloudWatch alarm for the `StatusCheckFailed_System` metric, and then configure an EC2 action to recover the instance.

Answer: A ([LEAVE A REPLY](#))

Explanation

References: <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-recover.html>

NEW QUESTION: 29

A company must deploy multiple independent instances of an application. The front-end application is internet accessible. However, corporate policy stipulates that the backends are to be isolated from each other and the internet, yet accessible from a centralized administration server. The application setup should be automated to minimize the opportunity for mistakes as new instances are deployed.

Which option meets the requirements and MINIMIZES costs?

- A.** Use an AWS CloudFormation template to create identical IAM roles for each region. Use AWS CloudFormation StackSets to deploy each application instance by using parameters to customize for each instance, and use security groups to isolate each instance while permitting access to the central server.
- B.** Create each instance of the application IAM roles and resources in separate accounts by using AWS CloudFormation StackSets. Include a VPN connection to the VPN gateway of the central administration server.
- C.** Use the parameters of the AWS CloudFormation template to customize the deployment into separate accounts. Include a NAT gateway to allow communication back to the central administration server.
- D.** Duplicate the application IAM roles and resources in separate accounts by using a single CloudFormation template. Include VPC peering to connect the VPC of each application instance to a central VPC.

Answer: A ([LEAVE A REPLY](#))

NEW QUESTION: 30

An enterprise company is building an infrastructure services platform for its users. The company has the following requirements:

- * Provide least privilege access to users when launching AWS infrastructure so users cannot provision unapproved services
 - * Use a central account to manage the creation of infrastructure services
 - * Provide the ability to distribute infrastructure services to multiple accounts in AWS Organizations
 - * Provide the ability to enforce tags on any infrastructure that is started by users Which combination of actions using AWS services will meet these requirements? (Select THREE.)
- A.** Allow user IAM roles to have ServiceCatalogEndUserAccess permissions only Use an automation script to import the central portfolios to local AWS accounts, copy the TagOption assign users access and apply launch constraints
- B.** Develop infrastructure services using AWS Cloud Formation templates Add the templates to a central Amazon S3 bucket and add the-IAM roles or users that require access to the S3 bucket policy
- C.** Use the AWS Service Catalog TagOption Library to maintain a list of tags required by the company Apply the TagOption to AWS Service Catalog products or portfolios
- D.** Allow user IAM roles to have AWSCloudFormationFullAccess and AmazonS3ReadOnlyAccess permissions Add an Organizations SCP at the AWS account root user level to deny all services except AWS CloudFormation and Amazon S3.
- E.** Develop infrastructure services using AWS Cloud For matron templates Upload each template as an AWS Service Catalog product to portfolios created in a central AWS account Share these portfolios with the Organizations structure created for the company
- F.** Use the AWS CloudFormation Resource Tags property to enforce the application of tags to any CloudFormation templates that will be created for users

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 31

A North American company with headquarters on the East Coast is deploying a new web application running on Amazon EC2 in the us-east-1 Region. The application should dynamically scale to meet user demand and maintain resiliency. Additionally, the application must have disaster recover capabilities in an active-passive configuration with the us-west-1 Region.

Which steps should a solutions architect take after creating a VPC in the us-east-1 Region?

- A.** Create a VPC in the us-west-1 Region Use inter-Region VPC peering to connect both VPCs. Deploy an Application Load Balancer (ALB) that spans both VPCs. Deploy EC2 instances across multiple Availability Zones as part of an Auto Scaling group in each VPC served by the ALB. Create an Amazon Route 53 record that points to the ALB.
- B.** Create a VPC in the us-west-1 Region Use inter-Region VPC peering to connect both VPCs. Deploy an Application Load Balancer (ALB) spanning multiple Availability Zones (AZs) to the VPC in the us-east-1 Region Deploy EC2 instances across multiple AZs in each Region as part of an Auto Scaling group spanning both VPCs and served by the ALB.
- C.** Deploy an Application Load Balancer (ALB) spanning multiple Availability Zones (AZs) to the VPC in the us-east-1 Region. Deploy EC2 instances across multiple AZs as part of an Auto Scaling group served by the ALB Deploy the same solution to the us-west-1 Region Create

separate Amazon Route 53 records in each Region that point to the ALB in the Region. Use Route 53 health checks to provide high availability across both Regions.

D. Deploy an Application Load Balancer (ALB) spanning multiple Availability Zones (AZs) to the VPC in the us-east-1 Region. Deploy EC2 instances across multiple AZs as part of an Auto Scaling group served by the ALB. Deploy the same solution to the us-west-1 Region. Create an Amazon Route 53 record set with a failover routing policy and health checks enabled to provide high availability across both Regions.

Answer: ([SHOW ANSWER](#))

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (**216 Q&As Dumps**, **30%OFF Special Discount: freecram**)

NEW QUESTION: 32

A company has a large on-premises Apache Hadoop cluster with a 20 PB HDFS database. The cluster is growing every quarter by roughly 200 instances and 1 PB. The company's goals are to enable resiliency for its Hadoop data, limit the impact of losing cluster nodes, and significantly reduce costs. The current cluster runs

24/7 and supports a variety of analysis workloads, including interactive queries and batch processing.

Which solution would meet these requirements with the LEAST expense and down time?

A. Use AWS Snowmobile to migrate the existing cluster data to Amazon S3. Create a persistent Amazon EMR cluster initially sized to handle the interactive workload based on historical data from the on-premises cluster. Store the data on EMRFS. Minimize costs using Reserved Instances for master and core nodes and Spot Instances for task nodes, and auto scale task nodes based on Amazon CloudWatch metrics. Create job-specific, optimized clusters for batch workloads that are similarly optimized.

B. Use AWS Snowmobile to migrate the existing cluster data to Amazon S3. Create a persistent Amazon EMR cluster of similar size and configuration to the current cluster. Store the data on EMRFS.

Minimize costs by using Reserved Instances. As the workload grows each quarter, purchase additional Reserved Instances and add to the cluster.

C. Use AWS Snowball to migrate the existing cluster data to Amazon S3. Create a persistent Amazon EMR cluster initially sized to handle the interactive workloads based on historical data from the on-premises cluster. Store the on EMRFS. Minimize costs using Reserved Instances for

master and core nodes and Spot Instances for task nodes, and auto scale task nodes based on Amazon CloudWatch metrics. Create job-specific, optimized clusters for batch workloads that are similarly optimized.

D. Use AWS Direct Connect to migrate the existing cluster data to Amazon S3. Create a persistent Amazon EMR cluster initially sized to handle the interactive workload based on historical data from the on-premises cluster. Store the data on EMRFS. Minimize costs using Reserved Instances for master and core nodes and Spot Instances for task nodes, and auto scale task nodes based on Amazon CloudWatch metrics. Create job-specific, optimized clusters for batch workloads that are similarly optimized.

Answer: A (LEAVE A REPLY)

Explanation

Q: How should I choose between Snowmobile and Snowball?

To migrate large datasets of 10PB or more in a single location, you should use Snowmobile. For datasets less than 10PB or distributed in multiple locations, you should use Snowball. In addition, you should evaluate the amount of available bandwidth in your network backbone. If you have a high speed backbone with hundreds of Gb/s of spare throughput, then you can use Snowmobile to migrate the large datasets all at once. If you have limited bandwidth on your backbone, you should consider using multiple Snowballs to migrate the data incrementally.

NEW QUESTION: 33

A fitness tracking company serves users around the world, with its primary markets in North America and Asia. The company needs to design an infrastructure for its read heavy user authorization application with the following requirements:

- * Be resilient to problem with the application in any region.
- * Write to a database In a single Region.
- * Read from multiple regions.
- * Support resiliency across application tiers in each Region.
- * Support the relational database semantics reflected in the application.

Which combination of steps should a solution architect take? (Select TWO.)

A. Use an Amazon Route 53 geolocation routing combined with a failover routing policy.

B. Use an Amazon Route 53 geoproximity routing policy combined with a multivalue answer routing policy.

C. Deploy web, application, and MySQL database servers to Amazon EC2 instance in each Region. Set up the application so that reads and writes are local to the Region. Create snapshots of the web, application, and database servers and store the snapshots in an Amazon S3 bucket in both Regions. Set up cross- Region replication for database layer.

D. Set up web, application, and Amazon RDS for MySQL instances in each Region. Set up the application so that reads are local and writes are partitioned based on the user. Set up a Multi-AZ failover for the web, application, and database servers, Set up cross-Region replication for the database layer.

E. Set up active-active web and application servers in each Region. Deploy an Amazon Aurora global database with clusters in each Region. Set up the application to use the in-Region Aurora database endpoints. Create snapshots of the web application servers and store them in an Amazon S3 bucket in both Region.

Answer: A,D (LEAVE A REPLY)

NEW QUESTION: 34

A solution architect is migrating an existing workload to AWS Fargate. The task can only run in a private subnet within the VPC where there is no direct connectivity from outside the system to the application. When the Fargate task is launched, the task fails with the following error:

```
CannotPullContainerError: API error (500): Get https://111122223333.dkr.ecr.us-east-1.amazonaws.com/v2/: net/http: request canceled while waiting for connection
```

How should the solution architect correct this error?

A. Ensure the task is set to DISABLED for the auto-assign public IP setting when launching the task.

Configure a NAT gateway in the public subnet in the VPC to route requests to the internet

B. Ensure the task is set to ENABLED for the auto-assign public IP setting when launching the task.

C. Ensure the network mode is set to bridge in the Fargate task definition.

D. Ensure the task is set to DISABLED for the auto-assign public IP setting when launching the task.

Configure a NAT gateway in the private subnet in the VPC to route requests to the internet

Answer: A (LEAVE A REPLY)

NEW QUESTION: 35

A company is running multiple applications on Amazon EC2. Each application is deployed and managed by multiple business units. All applications are deployed on a single AWS account but on different virtual private clouds (VPCs). The company uses a separate VPC in the same account for test and development purposes.

Production applications suffered multiple outages when users accidentally terminated and modified resources that belonged to another business unit. A Solutions Architect has been asked to improve the availability of the company applications while allowing the Developers access to the resources they need.

Which option meets the requirements with the LEAST disruption?

A. Create an AWS account for each business unit. Move each business unit's instances to its own account and set up a federation to allow users to access their business unit's account.

B. Set up a federation to allow users to use their corporate credentials, and lock the users down to their own VPC. Use a network ACL to block each VPC from accessing other VPCs.

C. Implement a tagging policy based on business units. Create an IAM policy so that each user can terminate instances belonging to their own business units only.

D. Set up role-based access for each user and provide limited permissions based on individual roles and the services for which each user is responsible.

Answer: C ([LEAVE A REPLY](#))

Explanation

<https://aws.amazon.com/blogs/security/resource-level-permissions-for-ec2-controlling-management-access-on-sp>

NEW QUESTION: 36

A company is refactoring an existing web service that provides read and write access to structured data. The service must respond to short but significant spikes in the system load. The service must be fault tolerant across multiple AWS Regions.

Which actions should be taken to meet these requirements?

A. Store the data in Amazon Aurora global databases. Add Auto Scaling replicas to both Regions. Run the web service on Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer in each Region. Configure the instances to download the web service code in the user data. In Amazon Route 53, configure an alias record for the company's domain and a multi-value routing policy.

B. Store the data in Amazon DocumentDB. Create a single global Amazon CloudFront distribution with a custom origin built on edge-optimized Amazon API Gateway and AWS Lambda. Assign the company's domain as an alternate domain for the distribution. and configure Amazon Route 53 with an alias to the CloudFront distribution.

C. Store the data in an Amazon DynamoDB global table in two Regions using on-demand capacity mode. In both Regions, run the web service as Amazon ECS Fargate tasks in an Auto Scaling ECS service behind an Application Load Balancer (ALB). In Amazon Route 53, configure an alias record in the company's domain and a Route 53 latency-based routing policy with health checks to distribute traffic between the two ALBs.

D. Store the data in replicated Amazon S3 buckets in two Regions. Create an Amazon CloudFront distribution in each Region, with custom origins built on Amazon API Gateway and AWS Lambda launched in each Region. Assign the company's domain as an alternate domain for both distributions and configure Amazon Route 53 with a failover routing policy between them.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 37

A company has several Amazon EC2 instances in both public and private subnets within a VPC that is not connected to the corporate network. A security group associated with the EC2 instances allows the company to use the Windows remote desktop protocol (RDP) over the internet to access the instances. The security team has noticed connection attempts from unknown sources. The company wants to implement a more secure solution to access the EC2 instances.

Which strategy should a solutions architect implement?

- A.** Deploy a Linux bastion host with an Elastic IP address in the public subnet. Allow access to the bastion host from 0.0.0.0/0.
- B.** Deploy a Linux bastion host on the corporate network that has access to all instances in the VPC.
- C.** Deploy AWS Systems Manager Agent on the EC2 instances Access the EC2 instances using Session Manager, restricting access to users with permission.
- D.** Establish a Site-to-Site VPN connecting the corporate network to the VPC. Update the security groups to allow access from the corporate network only.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 38

A company's lease of a colocated storage facility will expire in 90 days. The company wants to move to AWS to avoid signing a contract extension. The company environment consists of 200 virtual machines and a NAS with 40 TB of data Most of the data is archival, yet instant access is required when data is requested Leadership wants to ensure minimal downtime during the migration Each virtual machine has a number of customized configurations. The company's existing 1Gbps network connection is mostly idle especially after business hours Which combination of steps should the company take to migrate to AWS while minimizing downtime and operational impact? (Select TWO)

- A.** Use AWS Storage Gateway to migrate the data to cloud-native storage
- B.** Use AWS Snowball to migrate the data
- C.** Use AWS SMS to migrate the virtual machines
- D.** Use new Amazon EC2 instances and reinstall all application code.
- E.** Use AWS SMS to copy the infrequently accessed data from the NAS

Answer: A,C ([LEAVE A REPLY](#))

NEW QUESTION: 39

A company is migrating its on-premises systems to AWS. The user environment consists of the following systems:

- * Windows and Linux virtual machines running on VMware.
- * Physical servers running Red Hat Enterprise Linux.

The company wants to be able to perform the following steps before migrating to AWS:

- * Identify dependencies between on-premises systems.
- * Group systems together into applications to build migration plans.
- * Review performance data using Amazon Athena to ensure that Amazon EC2 instances are right-sized.

How can these requirements be met?

- A.** Populate the AWS Application Discovery Service import template with information from an on-premises configuration management database (CMDB). Upload the completed import template to Amazon S3, then import the data into Application Discovery Service.

B. Install the AWS Application Discovery Service Discovery Agent on the physical on-pre-map servers.

Install the AWS Application Discovery Service Discovery Connector in VMware vCenter. Allow the Discovery Agent to collect data for a period of time.

C. Install the AWS Application Discovery Service Discovery Agent on each of the on-premises systems.

Allow the Discovery Agent to collect data for a period of time.

D. Install the AWS Application Discovery Service Discovery Connector on each of the on-premises systems and in VMware vCenter. Allow the Discovery Connector to collect data for one week.

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 40

A large multinational company runs a timesheet application on AWS that is used by staff across the world.

The application runs on Amazon EC2 instances in an Auto Scaling group behind an Elastic Load Balancing (ELB) load balancer, and stores in an Amazon RDS MySQL Multi-AZ database instance.

The CFO is concerned about the impact on the business if the application is not available. The application must not be down for more than two hours, but the solution must be as cost-effective as possible.

How should the Solutions Architect meet the CFO's requirements while minimizing data loss?

A. Configure a read replica in another region. Create an AWS CloudFormation template of the application infrastructure. When an issue occurs, promote the read replica and configure as an Amazon RDS Multi-AZ database instance and use the AWS CloudFormation template to create the environment in another region using the promoted Amazon RDS instance. Update the DNS record to point to the other region's ELB.

B. In another region, configure a read replica and create a copy of the infrastructure. When an issue occurs, promote the read replica and configure as an Amazon RDS Multi-AZ database instance. Update the DNS to point to the other region's ELB.

C. Configure a 1-day window of 60-minute snapshots of the Amazon RDS Multi-AZ database instance which is copied to another region. Create an AWS CloudFormation template of the application infrastructure that uses the latest copied snapshot. When an issue occurs, use the AWS CloudFormation template to create the environment in another region. Update the DNS record to point to the other region's ELB.

D. Configure a 1-day window of 60-minute snapshots of the Amazon RDS Multi-AZ database instance.

Create an AWS CloudFormation template of the application infrastructure that uses the latest snapshot.

When an issue occurs, use the AWS CloudFormation template to create the environment in another region. Update the DNS record to point to the other region's ELB.

Answer: A ([LEAVE A REPLY](#))

NEW QUESTION: 41

A company that is developing a mobile game is making game assets available in two AWS Regions. Game assets are served from a set of Amazon EC2 instances behind an Application Load Balancer (ALB) in each Region. The company requires game assets to be fetched from the closest Region. If game assets become unavailable in the closest Region, they should be fetched from the other Region.

What should a solutions architect do to meet these requirements?

- A.** Create an Amazon Route 53 health check for each ALB. Create a Route 53 failover routing record pointing to the two ALBs. Set the Evaluate Target Health value to Yes.
- B.** Create two Amazon CloudFront distributions, each with one ALB as the origin. Create an Amazon Route 53 failover routing record pointing to the two CloudFront distributions. Set the Evaluate Target Health value to Yes.
- C.** Create an Amazon Route 53 health check for each ALB. Create a Route 53 latency alias record pointing to the two ALBs. Set the Evaluate Target Health value to Yes.
- D.** Create an Amazon CloudFront distribution. Create an origin group with one origin for each ALB. Set one of the origins as primary.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 42

A government agency is building a forms submission portal using AWS to allow citizen to submit and retrieve sensitive documents. The solution was built using serverless architecture, with the front-end code developed using HTML and JavaScript and the backend architecture using Amazon API Gateway and Amazon S3.

The portal must meet the following security requirements:

- * Requests to the backend infrastructure should be allowed only if they originate from a specific country.
- * Requests to the backend infrastructure should prevent brute attacks from individual IP addresses by not allowing more than 3000 requests per minutes for 10 requests per seconds for each IP address.
- * All access attempts to the backend infrastructure must be logged.

Which steps should a solution architect take to meet these requirements? (Select Two)

- A.** Configure the AWS WAF web ACL to an Amazon CloudWatch Logs group. Configure API Gateway to log to an Amazon Cloudwatch Logs group
- B.** Create an AWS WAP web ACL with a custom condition that allows access attempts from the authorized country only, and a rate-based rule with a rate-based rule with rate limit 3000 requests per 5 minutes. Then associate the web ACL with the API Gateway API
- C.** Configure the AWS WAF web ACL to log to an Amazon Kinesis Data Firehose delivery with Amazon Elasticsearch Service (Amazon ES) as the destination. Configure API Gateway to log to an Amazon CloudWatch Logs group.

D. Configure the API Gateway API with a custom rule condition that allow APIs to be called from the authorized country only. Then enable default method throttling, setting the rate limit in 10 requests per seconds.

E. Configure Amazon Cloud with a geographical restriction that allows access attempts from the authorized country only, and a rate-based rule with a rate limit of 3000 requests per 5 minutes. Then Add the API Gateway API as a custom origin.

Answer: A,B (LEAVE A REPLY)

NEW QUESTION: 43

A company manages more than 200 separate internet-facing web applications. All of the applications are deployed to AWS in a single AWS Region The fully qualified domain names (FQDNs) of all of the applications are made available through HTTPS using Application Load Balancers (ALBs). The ALBs are configured to use public SSL/TLS certificates.

A Solutions Architect needs to migrate the web applications to a multi-region architecture. All HTTPS services should continue to work without interruption.

Which approach meets these requirements?

A. Request a certificate for each FQDN using AWS KMS. Associate the certificates with the ALBs in the primary AWS Region. Enable cross-region availability in AWS KMS for the certificates and associate the certificates with the ALBs in the secondary AWS Region.

B. Generate the key pairs and certificate requests for each FQDN using AWS KMS. Associate the certificates with the ALBs in both the primary and secondary AWS Regions.

C. Request a certificate for each FQDN using AWS Certificate Manager. Associate the certificates with the ALBs in both the primary and secondary AWS Regions.

D. Request certificates for each FQDN in both the primary and secondary AWS Regions using AWS Certificate Manager. Associate the certificates with the corresponding ALBs in each AWS Region.

Answer: D (LEAVE A REPLY)

Explanation

<https://docs.aws.amazon.com/acm/latest/userguide/acm-regions.html>

Certificates in ACM are regional resources. To use a certificate with Elastic Load Balancing for the same fully qualified domain name (FQDN) or set of FQDNs in more than one AWS region, you must request or import a certificate for each region. For certificates provided by ACM, this means you must revalidate each domain name in the certificate for each region. You cannot copy a certificate between regions.

NEW QUESTION: 44

What combination of steps could a Solutions Architect take to protect a web workload running on Amazon EC2 from DDoS and application layer attacks? (Select two.)

A. Put the EC2 instances behind a Network Load Balancer and configure AWS WAF on it.

B. Migrate the DNS to Amazon Route 53 and use AWS Shield

C. Put the EC2 instances in an Auto Scaling group and configure AWS WAF on it.

- D. Create and use an Amazon CloudFront distribution and configure AWS WAF on it.
- E. Create and use an internet gateway in the VPC and use AWS Shield.

Answer: B,D (LEAVE A REPLY)

Explanation

References: <https://aws.amazon.com/answers/networking/aws-ddos-attack-mitigation/>

NEW QUESTION: 45

A large financial company is deploying applications that consist of Amazon EC2 and Amazon RDS instances to the AWS Cloud using AWS Cloud Formation.

The CloudFormation stack has the following stack policy:

```
{
  "Statement" : [
    {
      "Effect" : "Allow",
      "Action" : ["Update:*"],
      "Principal" : "*",
      "Resource" : "*"
    }
  ]
}
```

The company wants to ensure that developers do not lose data by accidentally removing or replacing RDS instances when updating the Cloud Formation stack. Developers also still need to be able to modify or remove EC2 instances as needed. How should the company change the stack policy to meet these requirements?

- A. Modify the statement to specify "Effect" "Deny" "Action" ["Update Delete"] for all logical RDS resources
- B. Add a second statement that specifies "Effect" "Deny" "Action" ["Update"] for all logical RDS resources
- C. Modify the statement to specify "Effect" "Deny" "Action" ["Update *"] for all logical RDS resources
- D. Add a second statement that specifies "Effect" "Deny" "Action" ["Update Delete" "Update Replace"] for all logical RDS resources

Answer: B (LEAVE A REPLY)

NEW QUESTION: 46

A company with several AWS accounts is using AWS Organizations and service control policies (SCPs). An Administrator created the following SCP and has attached it to an organizational unit (OU) that contains AWS account 1111-1111-1111:

```

{
  "Version": "2012-10-27",
  "Statement": [
    {
      "Sid": "AllowsAllActions",
      "Effect": "Allow",
      "Action": "*",
      "Resource": "*"
    },
    {
      "Sid": "DenyCloudTrail",
      "Effect": "Deny",
      "Action": "cloudtrail:*",
      "Resource": "*"
    }
  ]
}

```

Developers working in account 1111-1111-1111 complain that they cannot create Amazon S3 buckets. How should the Administrator address this problem?

- A. Remove the SCP from account 1111-1111-1111.
- B. Instruct the Developers to add Amazon S3 permissions to their IAM entities.
- C. Remove the account from the OU, and attach the SCP directly to account 1111-1111-1111.
- D. Add s3:CreateBucket with "Allow" effect to the SCP.

Answer: ([SHOW ANSWER](#))

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (216 Q&As Dumps, **30%OFF** Special Discount: **freecram**)

NEW QUESTION: 47

A company has a web-based application deployed in the ap-south-east-2 Region behind an Application Load Balancer ALB). AWS Certificate Manager (ACM) has issued a TLS certificate for example.com. This certificate is deployed to the ALB. There is a record set in Amazon Route 53 for example.com associated to the ALB.

Due to increased load on the application, the company wants to use Amazon CloudFront. This transition cannot cause application downtime.

Which combination of actions can achieve this? (Choose Three.)

- A.** Update the ALB security group to allow access from the CloudFront Edge locations only.
 - B.** Create a new ACM certificate in the us-east-1 Region for example.com. Create a CloudFront distribution and use the ACM certificate in the us-east-1 Region. Set origin example.com as the custom origin.
 - C.** Create a CloudFront distribution and use the existing certificate associated with the ALB in the ap-southeast-2 Region. Set origin example.com as the custom origin.
 - D.** Create a new ACM certificate in the ap-southeast-2 Region for origin-example.com and example.com.
- Associate this certificate to the existing ALB. Add a DNS entry in Route 53 for origin.example.com associated with the existing ALB.
- E.** Create a new ACM certificate in the us-east-1 Region for example.com. Create a new ALB in the us-east-1 Region as the origin of the CloudFront distribution. Attach the security group associated with the ALB to the CloudFront distribution.
 - F.** Update Route 53 for example.com to the alias record of the CloudFront distribution.

Answer: A,B,D ([LEAVE A REPLY](#))

NEW QUESTION: 48

A company's service for video game recommendations has just gone viral. The company has new users from all over the world. The website for the service is hosted on a set of Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB). The website consists of static content with different resources being loaded depending on the device type. Users recently reported that the load time for the website has increased. Administrators are reporting high loads on the EC2 instances that host the service.

Which set of actions should a solutions architect take to improve response times?

- A.** Move content to Amazon S3. Create an Amazon CloudFront distribution to serve content out of the S3 bucket. Use Lambda@Edge to load different resources based on the User-Agent HTTP header.
- B.** Move content to Amazon S3. Create an Amazon CloudFront distribution to serve content out of the S3 bucket. Use the User-Agent HTTP header to load different content.
- C.** Create a separate ALB for each device type. Create one Auto Scaling group behind each ALB. Use Amazon Route 53 to route to different ALBs depending on the User-Agent HTTP header.
- D.** Create separate Auto Scaling groups based on device types. Switch to a Network Load Balancer (NLB). Use the User-Agent HTTP header in the NLB to route to a different set of EC2 instances.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 49

A company is deploying a public-facing global application on AWS using Amazon CloudFront. The application communicates with an external system. A solutions architect needs to ensure the data is secured during end-to-end transit and at rest. Which combination of steps will satisfy these requirements? (Select THREE)

- A.** Create a public certificate for the required domain in AWS Certificate Manager and deploy it to CloudFront an Application Load Balancer and Amazon EC2 instances
- B.** Communicate with the external system using plaintext and use the VPN to encrypt the data in transit
- C.** Use SSL or encrypt data while communicating with the external system using a VPN
- D.** Provision Amazon EBS encrypted volumes using AWS KMS
- E.** Acquire a public certificate from a third-party vendor and deploy it to CloudFront an Application Load Balancer and Amazon EC2 instances
- F.** Provision Amazon EBS encrypted volumes using AWS KMS and ensure explicit encryption of data when writing to Amazon EBS

Answer: A,C,F (LEAVE A REPLY)

NEW QUESTION: 50

An enterprise runs 103 line-of-business applications on virtual machines in an on-premises data center. Many of the applications are simple PHP, Java, or Ruby web applications, are no longer actively developed, and serve little traffic.

Which approach should be used to migrate these applications to AWS with the LOWEST infrastructure costs ?

- A.** Deploy the applications to single-instance AWS Elastic Beanstalk environments without a load balancer.
- B.** Use AWS SMS to create AMIs for each virtual machine and run them in Amazon EC2.
- C.** Convert each application to a Docker image and deploy to a small Amazon ECS cluster behind an Application Load Balancer.
- D.** Use VM Import/Export to create AMIs for each virtual machine and run them in single-instance AWS Elastic Beanstalk environments by configuring a custom image.

Answer: C (LEAVE A REPLY)

Explanation

<https://docs.aws.amazon.com/elasticbeanstalk/latest/dg/using-features-managing-env-types.html>

NEW QUESTION: 51

A company has a single AWS master billing account, which is the root of the AWS Organizations hierarchy.

The company has multiple AWS accounts within this hierarchy, all organized into organization units (OUs).

More OUS and AWS accounts will continue to be created as other parts of the business migrate applications to AWS. These business units may need to use different AWS services. The Security team is implementing the following requirements for all current and future AWS accounts.

- * Control policies must be applied across all accounts to prohibit AWS servers.
- * Exceptions to the control policies are allowed based on valid use cases.

Which solution will meet these requirements with minimal optional overhead?

A. Use an SCP in Organization to implement a deny list of AWS service. Apply this SCP at each OU level

. Leave the default AWS managed SCP at the root level For any specific executions for an OU, create a new SCP for that OU.

B. Use an SCP In organizations to implement a deny list of AWS service. Apply this SCP at the root level and each OU. Remove the default AWS managed SCP from the root level and all OU levels. For any specific exceptions, modify the SCP attached to that OU, and add the required AWS required services to the allow list.

C. Use an SCP in Organizations to implement a deny list of AWS servers. Apply this SCP at the level. For any specific exceptions for an OU, create a new SCP for that OU and add the required AWS services the allow list.

D. Use an SCP in Organizations to implement an allow list of AWS services. Apply this SCP at the root level. Remove the default AWS managed SCP from the root level and all OU levels. For any specific exceptions for an OU, modify the SCP attached to that OU, and add the required AWS services to the allow list.

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 52

A company has created an account for individual Development teams, resulting in a total of 200 accounts. All accounts have a single virtual private cloud (VPC) in a single region with multiple microservices running in Docker containers that need to communicate with microservices in other accounts. The Security team requirements state that these microservices must not traverse the public internet, and only certain internal services should be allowed to call other individual services. If there is any denied network traffic for a service, the Security team must be notified of any denied requests, including the source IP.

How can connectivity be established between services while meeting the security requirements?

A. Create a VPC peering connection between the VPCs. Use security groups on the instances to allow traffic from the security group IDs that are permitted to call the microservice. Apply network ACLs to and allow traffic from the local VPC and peered VPCs only. Within the task definition in Amazon ECS for each of the microservices, specify a log configuration by using the awslogs driver. Within Amazon CloudWatch Logs, create a metric filter and alarm off of the number of HTTP 403 responses. Create an alarm when the number of messages exceeds a threshold set by the Security team.

B. Ensure that no CIDR ranges are overlapping, and attach a virtual private gateway (VGW) to each VPC.

Provision an IPsec tunnel between each VGW and enable route propagation on the route table. Configure security groups on each service to allow the CIDR ranges of the VPCs on the other accounts.

Enable VPC Flow Logs, and use an Amazon CloudWatch Logs subscription filter for rejected traffic.

Create an IAM role and allow the Security team to call the AssumeRole action for each account.

C. Deploy a transit VPC by using third-party marketplace VPN appliances running on Amazon EC2, dynamically routed VPN connections between the VPN appliance, and the virtual private gateways (VGWs) attached to each VPC within the region. Adjust network ACLs to allow traffic from the local VPC only. Apply security groups to the microservices to allow traffic from the VPN appliances only.

Install the awslogs agent on each VPN appliance, and configure logs to forward to Amazon CloudWatch Logs in the security account for the Security team to access.

D. Create a Network Load Balancer (NLB) for each microservice. Attach the NLB to a PrivateLink endpoint service and whitelist the accounts that will be consuming this service. Create an interface endpoint in the consumer VPC and associate a security group that allows only the security group IDs of the services authorized to call the producer service. On the producer services, create security groups for each microservice and allow only the CIDR range the allowed services. Create VPC Flow Logs on each VPC to capture rejected traffic that will be delivered to an Amazon CloudWatch Logs group. Create a CloudWatch Logs subscription that streams the log data to a security account.

Answer: D ([LEAVE A REPLY](#))

Explanation

AWS PrivateLink provides private connectivity between VPCs, AWS services, and on-premises applications, securely on the Amazon network. AWS PrivateLink makes it easy to connect services across different accounts and VPCs to significantly simplify the network architecture. It seems like the next VPC peering.

<https://aws.amazon.com/privatelink/>

NEW QUESTION: 53

An enterprise company is using a multi-account AWS strategy. There are separate accounts for development, staging, and production workloads. To control costs and improve governance, the following requirements have been defined:

- * The company must be able to calculate the AWS costs for each project
- * The company must be able to calculate the AWS costs for each environment: development, staging, and production
- * Commonly deployed IT services must be centrally managed
- * Business units can deploy pre-approved IT services only
- * Usage of AWS resources in the development account must be limited

Which combination of actions should be taken to meet these requirements? (Select THREE.)

- A.** Configure SCPs in AWS Organizations to allow services available using AWS
- B.** Configure a billing alarm in Amazon CloudWatch.
- C.** Configure custom budgets and define thresholds using Cost Explorer
- D.** Configure AWS Trusted Advisor to obtain weekly emails with cost-saving estimates
- E.** Apply environment, cost center, and application name tags to all taggable resources
- F.** Create a portfolio for each business unit and add products to the portfolios using AWS CloudFormation in AWS Service Catalog

Answer: A,B,D ([LEAVE A REPLY](#))

NEW QUESTION: 54

A company has a mobile app with users in Europe. When the app is used, it downloads a configuration file that is device and app version-specific. The company has the following architecture:

- *Configuration files are stored in Amazon S3 in the eu-west-1 Region and served to the users using Amazon CloudFront.

- *Lambda@Edge is used to extract the device and version information from the app requests. It then updates the requests to load the correct configuration.

The company uses the configuration file load time as a key performance metric, and targets a response time of

100 ms or less. The app recently launched in the ap-southeast-2 Region, and the latency for requests from users in Australia is significantly above the 100 ms target. A solutions architect needs to recommend a solution.

Which solution will reduce latency for users in Australia?

- A.** Create an S3 bucket in the ap-southeast-2 Region. Use cross-Region replication to synchronize from the bucket in the eu-west-1 Region. Modify Lambda@Edge to access Amazon S3 in the Region that is closest to the user.
- B.** Create an S3 bucket in the ap-southeast-2 Region. Use cross-Region replication to synchronize from the bucket in the eu-west-1 Region. Create an Amazon Route 53 hosted zone with latency-based routing configured for both buckets. Modify Lambda@Edge to update the origin of the request to be the Route 53 hosted zone that is closest to the user.
- C.** Configure S3 Transfer Acceleration on the bucket. Add the Transfer Acceleration Edge endpoints for Australia and Europe as CloudFront origins. Modify Lambda@Edge to update the origin of the request to be the Transfer Acceleration endpoint in the Region that is closest to the user.
- D.** Configure S3 Transfer Acceleration on the bucket. Modify Lambda@Edge to access Amazon S3 using the Transfer Acceleration endpoint in the Region that is closest to the user.

Answer: A ([LEAVE A REPLY](#))

NEW QUESTION: 55

A public retail web application uses an Application Load Balancer (ALB) in front of Amazon EC2 instances running across multiple Availability Zones (AZs) in a Region backed by an Amazon RDS MySQL Multi-AZ deployment. Target group health checks are configured to use HTTP and pointed at the product catalog page.

Auto Scaling is configured to maintain the web fleet size based on the ALB health check.

Recently, the application experienced an outage. Auto Scaling continuously replaced the instances during the outage. A subsequent investigation determined that the web server metrics were within the normal range, but the database tier was experiencing high load, resulting in severely elevated query response times.

Which of the following changes together would remediate these issues while improving monitoring capabilities for the availability and functionality of the entire application stack for future growth? (Select TWO.)

A. Configure the target group health check to point at a simple HTML page instead of a product catalog page and the Amazon Route 53 health check against the product page to evaluate full application functionality. Configure Amazon CloudWatch alarms to notify administrators when the site fails.

B. Configure the target group health check to use a TCP check of the Amazon EC2 web server and the Amazon Route 53 health check against the product page to evaluate full application functionality.

Configure Amazon CloudWatch alarms to notify administrators when the site fails.

C. Configure an Amazon ElastiCache cluster and place it between the web application and RDS MySQL instances to reduce the load on the backend database tier.

D. Configure read replicas for Amazon RDS MySQL and use the single reader endpoint in the web application to reduce the load on the backend database tier.

E. Configure an Amazon CloudWatch alarm for Amazon RDS with an action to recover a high-load, impaired RDS instance in the database tier.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 56

A company that is new to AWS reports it has exhausted its service limits across several accounts that are on the Basic Support plan. The company would like to prevent this from happening in the future.

What is the MOST efficient way of monitoring and managing all service limits in the company's accounts?

A. Use Amazon CloudWatch and AWS Lambda to periodically calculate the limits across all linked accounts using AWS Trusted Advisor, provide notifications using Amazon SNS if the limits are close to exceeding the threshold.

B. Reach out to AWS Support to proactively increase the limits across all accounts. That way, the customer avoids creating and managing infrastructure just to raise the service limits.

C. Use Amazon CloudWatch and AWS Lambda to periodically calculate the limits across all linked accounts using AWS Trusted Advisor, programmatically increase the limits that are close to exceeding the threshold.

D. Use Amazon CloudWatch and AWS Lambda to periodically calculate the limits across all linked accounts using AWS Trusted Advisor, and use Amazon SNS for notifications if a limit is close to exceeding the threshold. Ensure that the accounts are using the AWS Business Support plan at a minimum.

Answer: **D** ([LEAVE A REPLY](#))

Explanation

<https://github.com/awslabs/aws-limit-monitor>

<https://aws.amazon.com/solutions/limit-monitor/>

NEW QUESTION: 57

A multimedia company with a single AWS account is launching an application for a global user base. The application storage and bandwidth requirements are unpredictable. The application will use Amazon EC2 instances behind an Application Load Balancer as the web tier and will use Amazon DynamoDB as the database tier. The environment for the application must meet the following requirements:

- * Low latency when accessed from any part of the world
- * WebSocket support
- * End-to-end encryption
- * Protection against the latest security threats
- * Managed layer 7 DDoS protection

Which actions should the solutions architect take to meet these requirements? (Select TWO)

- A. Use AWS WAF with AWS Shield Advanced to protect the application
- B. Use AWS Shield Standard to protect the application
- C. Use AWS WAF and Amazon Detective to protect the application
- D. Use Amazon Route 53 and AWS Transit Gateway for content distribution. Use an Amazon Elastic Block Store (Amazon EBS) volume to store static content
- E. Use Amazon Route 53 and Amazon CloudFront for content distribution. Use Amazon S3 to store static content

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 58

A company would like to implement a serverless application by using Amazon API Gateway, AWS Lambda, and Amazon DynamoDB. They deployed a proof of concept and stated that the average response time is greater than what their upstream services can accept. Amazon CloudWatch metrics did not indicate any issues with DynamoDB but showed that some Lambda functions were hitting their timeout.

Which of the following actions should the Solutions Architect consider to improve performance? (Choose two.)

- A. Configure the AWS Lambda function to reuse containers to avoid unnecessary startup time.
- B. Increase the amount of memory and adjust the timeout on the Lambda function. Complete performance testing to identify the ideal memory and timeout configuration for the Lambda function.
- C. Create an Amazon ElastiCache cluster running Memcached, and configure the Lambda function for VPC integration with access to the Amazon ElastiCache cluster.
- D. Enable API cache on the appropriate stage in Amazon API Gateway, and override the TTL for individual methods that require a lower TTL than the entire stage.
- E. Increase the amount of CPU, and adjust the timeout on the Lambda function. Complete performance testing to identify the ideal CPU and timeout configuration for the Lambda function.

Answer: B,D ([LEAVE A REPLY](#))

Explanation

<https://lumigo.io/blog/aws-lambda-timeout-best-practices/>

NEW QUESTION: 59

A company has a VPC with two domain controllers running Active Directory in the default configuration. The VPC DHCP options set is configured to use the IP addresses of the two domain controllers. There is a VPC interface endpoint defined; but instances within the VPC are not able to resolve the private endpoint addresses.

Which strategies would resolve this issue? (Select TWO)

- A.** Define an outbound Amazon Route 53 Resolver. Set a conditional forward rule for the Active Directory domain to the Active Directory servers. Update the VPC DHCP options set to AmazonProvidedDNS.
- B.** Update the DNS service on the Active Directory servers to forward all non-authoritative queries to the VPC Resolver
- C.** Update the DNS service on the Active Directory servers to forward all queries to the VPC Resolver.
- D.** Define an inbound Amazon Route 53 Resolver Set a conditional forward rule for the Active Directory domain to the Active Directory servers Update the VPC DHCP options set to AmazonProvidedDNS.
- E.** Update the DNS service on the client instances to split DNS queries between the Active Directory servers and the VPC Resolver

Answer: C,D ([LEAVE A REPLY](#))

NEW QUESTION: 60

A company has a complex web application that leverages Amazon CloudFront for global scalability and performance. Over time, Users report that the web application is slowing down. The company's operations team reports that the CloudFront cache hit ratio has been dropping steadily. The cache metrics report indicates that query strings on some URLs are inconsistently ordered and are specified sometimes in mixed-case letters and sometimes in lowercase letters. Which set of actions should the solutions architect take to increase the cache hit ratio as quickly as possible?

- A.** Update the CloudFront distribution to specify casing-insensitive query string processing
- B.** Deploy a Lambda@Edge function to sort parameters by name and force them to be lowercase. Select the CloudFront viewer request trigger to invoke the function
- C.** Update the CloudFront distribution to disable caching based on query string parameters
- D.** Deploy a reverse proxy after the load balancer to post-process the emitted URLs in the application to force the URL strings to be lowercase

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 61

A company is running a web application with On-Demand Amazon EC2 instances in Auto Scaling groups that scale dynamically based on custom metrics. After extensive testing, the company determines that the m5

2xlarge instance size is optimal for the workload. Application data is stored in db.r4.xlarge Amazon RDS instances that are confirmed to be optimal. The traffic to the web application spikes randomly during the day. What other cost-optimization methods should the company implement to further reduce costs without impacting the reliability of the application?

- A. Reserve capacity for the RDS database and the minimum number of EC2 instances that are constantly running
- B. Reduce the RDS instance size to db.r4.xlarge and add five equivalent-sized read replicas to provide reliability
- C. Double the instance count in the Auto Scaling groups and reduce the instance size to m5.large
- D. Reserve capacity for all EC2 instances and leverage Spot Instance pricing for the RDS database

Answer: ([SHOW ANSWER](#))

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (216 Q&As Dumps, **30%OFF** Special Discount: **freecram**)

NEW QUESTION: 62

A company is planning to host a three-tier application in the AWS Cloud. The application layer will use Amazon EC2 in an Auto Scaling group. A custom EC2 role named AppServer will be created and associated with the application instances. The entire application stack will be deployed using AWS CloudFormation. The company's security team requires encryption of all AMI snapshots and Amazon Elastic Block Store (Amazon EBS) volumes with an AWS Key Management Service (AWS KMS) CMK. Which action will deploy the stack correctly after the AMI snapshot is encrypted with the KMS key?

- A. Update the AppServer role to have the required permissions to access the KMS key
- B. Update the KMS key policy to provide the required permissions to the AWSServiceRoleForAutoScaling service-linked role
- C. Update the KMS key policy to provide the required permissions to the AppServer role
- D. Update the CloudFormation stack role to have the required permissions to access the KMS key

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 63

A company uses AWS Organizations with a single OU named Production to manage multiple accounts. All accounts are members of the Production OU. Administrators use deny list SCPs in the root of the organization to manage access to restricted services.

The company recently acquired a new business unit and invited the new unit's existing AWS account to the organization. Once onboarded, the administrators of the new business unit discovered that they are not able to update existing AWS Config rules to meet the company's policies.

Which option will allow administrators to make changes and continue to enforce the current policies without introducing additional long-term maintenance?

- A.** Convert the organization's root SCPs from deny list SCPs to allow list SCPs to allow the required services only. Temporally apply an SCP to the organization's root that allows AWS Config actions for principals only in the new account
- B.** Remove the organization's root SCPs that limit access to AWS Config. Create AWS Service Catalog products for the company's standard AWS Config rules and deploy them throughout the organization, including the new account
- C.** Create a temporary OU named Onboarding for the new account. Apply an SCP to the Onboarding OU to allow AWS Config actions. Move the new account to the Production OU when adjustments to AWS Config are complete
- D.** Create a temporary OU named Onboarding for the new account. Apply an SCP to the Onboarding OU to allow AWS Config actions. Move the organization's root SCP to the Production OU. Move the new account to the Production OU when adjustments to AWS Config are complete

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 64

A company has asked a Solutions Architect to design a secure content management solution that can be accessed by API calls by external customer applications. The company requires that a customer administrator must be able to submit an API call and roll back changes to existing files sent to the content management solution, as needed.

What is the MOST secure deployment design that meets all solution requirements?

- A.** Use Amazon EFS for object storage, using encryption at rest for the Amazon EFS volume and a customer managed key stored in AWS KMS. Use IAM roles and Amazon EFS access policies to specify separate encryption keys for each customer application. Deploy the content management application to store all new versions as new files in Amazon EFS and use a control API to revert a specific file to a previous version.
- B.** Use Amazon S3 for object storage with versioning and bucket access logging enabled, and an IAM role and access policy for each customer application. Encrypt objects using SSE-KMS. Develop the content management application to use a separate AWS KMS key for each customer.

- C.** Use Amazon S3 for object storage with versioning and enable S3 bucket access logging. Use an IAM role and access policy for each customer application. Encrypt objects using client-side encryption, and distribute an encryption key to all customers when accessing the content management application.
- D.** Use Amazon WorkDocs for object storage. Leverage WorkDocs encryption, user access management, and version control. Use AWS CloudTrail to log all SDK actions and create reports of hourly access by using the Amazon CloudWatch dashboard. Enable a revert function in the SDK based on a static Amazon S3 webpage that shows the output of the CloudWatch dashboard.

Answer: B (LEAVE A REPLY)

NEW QUESTION: 65

A company has developed a custom tool used in its workflow that runs within a Docker container. The company must perform manual steps each time the container code is updated to make the container image available to new workflow executions. The company wants to automate this process to eliminate manual effort and ensure a new container image is generated every time the tool code is updated. Which combination of actions should a solutions architect take to meet these requirements? (Select THREE.)

- A.** Configure an Amazon ECR repository for the tool. Configure an AWS CodeCommit repository containing code for the tool being deployed to the container image in Amazon ECR.
- B.** Configure an AWS CodePipeline pipeline that sources the tool code from the AWS CodeCommit repository and initiates an AWS CodeBuild build.
- C.** Configure an AWS CodeBuild project that pulls the latest tool container image from Amazon ECR, updates the container with code from the source AWS CodeCommit repository, and pushes the updated container image to Amazon ECR.
- D.** Configure an Amazon EventBridge rule that triggers on commits to the AWS CodeCommit repository for the tool. Configure the event to trigger an update to the tool container image in Amazon ECR. Push the updated container image to Amazon ECR.
- E.** Configure an AWS CodePipeline pipeline that sources the tool code from the AWS CodeCommit repository and initiates an AWS CodeDeploy application update.
- F.** Configure an AWS CodeDeploy application that triggers an application version update that pulls the latest tool container image from Amazon ECR, updates the container with code from the AWS CodeCommit repository, and pushes the updated container image to Amazon ECR.

Answer: B,D,E (LEAVE A REPLY)

NEW QUESTION: 66

A company operates pipelines across North America and South America. The company assesses pipeline inspection gauges with imagery and ultrasonic sensor data to monitor the condition of its pipelines. The pipelines are in areas with intermittent or unavailable internet connectivity. The imager data at each site requires terabytes of storage each month. The company wants a solution to collect the data at each site in monthly intervals and to store the data with high durability. The

imagery captured must be preprocessed and uploaded to a central location for persistent Storage.

Which actions should a solutions architect take to meet these requirements?

- A.** Deploy AWS IoT Greengrass on eligible hardware across the sites. Configure AWS Lambda on the devices for preprocessing. Ship the devices back to the closest AWS Region and store the data in Amazon S3 buckets
- B.** Deploy AWS Snowball devices at local sites in a cluster configuration. Configure AWS Lambda for preprocessing. Ship the devices back to the closest AWS Region and store the data in Amazon S3 buckets
- C.** Deploy AWS IoT Greengrass on eligible hardware across the sites. Configure AWS Lambda on the devices for preprocessing Upload the processed date to Amazon S3 buckets in AWS Regions closest to the sites
- D.** Deploy AWS Snowball Edge devices at local sites in a cluster configuration. Configure AWS Lambda for preprocessing Ship the devices back to the closest AWS Region and store the date in Amazon S3 buckets.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 67

A large company will be migrating to AWS. The company has 20 business units and anticipates another 10 coming online in the future. Each business unit will need its own IP range and will operate in its own AWS account. There will be a lot of communication between business units with very large data transfers. The company wants to make sure that the proposed solution will minimize data transfer costs and reduce complexity How should a solutions architect design the network to meet these requirements?

- A.** Create two subnets for each business unit in a networking account. Share the subnets with each business unit's AWS account using AWS Resource Access Manager.
- B.** Create a VPC for each business unit's AWS account Use VPC peering to route traffic between the VPCs in each account.
- C.** Create a transit gateway in a networking account. Share the transit gateway with each business unit's AWS account. Attach the VPC in each account to the transit gateway.
- D.** Create a transit VPC in a networking account. Within each business unit's AWS account create redundant VPN connections to the transit VPC.

Answer: ([SHOW ANSWER](#)**)**

NEW QUESTION: 68

A company that provides wireless services needs a solution to store and analyze log files about user activities.

Currently, log files are delivered daily to Amazon Linux on Amazon EC2 instance. A batch script is run once a day to aggregate data used for analysis by a third-party tool. The data pushed to the third-party tool is used to generate a visualization for end users. The batch script is cumbersome to maintain, and it takes several hours to deliver the ever-increasing data volumes to the third-

party tool. The company wants to lower costs, and is open to considering a new tool that minimizes development effort and lowers administrative overhead. The company wants to build a more agile solution that can store and perform the analysis in near-real time, with minimal overhead. The solution needs to be cost effective and scalable to meet the company's end-user base growth.

Which solution meets the company's requirements?

- A.** Develop a Python script to failover the data from Amazon EC2 in real time and store the data in Amazon S3. Use a copy command to copy data from Amazon S3 to Amazon Redshift. Connect a business intelligence tool running on Amazon EC2 to Amazon Redshift and create the visualizations.
- B.** Use an Amazon Kinesis agent running on an EC2 instance in an Auto Scaling group to collect and send the data to an Amazon Kinesis Data Firehose delivery stream. The Kinesis Data Firehose delivery stream will deliver the data directly to Amazon ES. Use Kibana to visualize the data.
- C.** Use an in-memory caching application running on an Amazon EBS-optimized EC2 instance to capture the log data in near real-time. Install an Amazon ES cluster on the same EC2 instance to store the log files as they are delivered to Amazon EC2 in near real-time. Install a Kibana plugin to create the visualizations.
- D.** Use an Amazon Kinesis agent running on an EC2 instance to collect and send the data to an Amazon Kinesis Data Firehose delivery stream. The Kinesis Data Firehose delivery stream will deliver the data to Amazon S3. Use an AWS Lambda function to deliver the data from Amazon S3 to Amazon ES. Use Kibana to visualize the data.

Answer: B ([LEAVE A REPLY](#))

Explanation

<https://docs.aws.amazon.com/firehose/latest/dev/writing-with-agents.html>

NEW QUESTION: 69

A solutions architect is designing a disaster recovery strategy for a three-tier application. The application has an RTO of 30 minutes and an RPO of 5 minutes for the data tier. The application and web tiers are stateless and leverage a fleet of Amazon EC2 instances. The data tier consists of an 50 TB Amazon Aurora database.

Which combination of steps satisfies the RTO and RPO requirements while optimizing costs? (Select Two.)

- A.** Create an Aws Backup job to replicate data to another Region.
- B.** Create daily snapshots of the EC2 instances and replicate the snapshots to another Region.
- C.** Deploy a hot standby of the application to another Region.
- D.** Create a cross-Region Aurora Replica of the database.
- E.** Create snapshots of the Aurora database every 5 minutes.

Answer: A,C ([LEAVE A REPLY](#))

NEW QUESTION: 70

A company has an application that runs a web service on Amazon EC2 instances and stores .jpg images in Amazon S3. The web traffic has a predictable baseline, but often demand spikes unpredictably for short periods of time. The application is loosely coupled and stateless. The .jpg images stored in Amazon S3 are accessed frequently for the first 15 to 20 days, they are seldom accessed thereafter but always need to be immediately available. The CIO has asked to find ways to reduce costs.

Which of the following options will reduce costs? (Choose two.)

- A.** Use On-Demand instances for baseline capacity requirements and use Spot Fleet instances for the demand spikes.
- B.** Configure a lifecycle policy to move the .jpg images on Amazon S3 to S3 IA after 30 days.
- C.** Configure a lifecycle policy to move the .jpg images on Amazon S3 to Amazon Glacier after 30 days.
- D.** Create a script that checks the load on all web servers and terminates unnecessary On-Demand instances.
- E.** Purchase Reserved instances for baseline capacity requirements and use On-Demand instances for the demand spikes.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 71

A company recently transformed its legacy infrastructure provisioning scripts to AWS CloudFormation templates. The newly developed templates are hosted in the company's private GitHub repository. Since adopting CloudFormation, the company has encountered several issues with updates to the CloudFormation templates, causing execution or creating environment Management is concerned by the increase in errors and has asked a Solutions architect to design the automated testing of CloudFormation template updates.

What should the Solution Architecture do to meet these requirements?

- A.** Use AWS CodePipeline to create a change set from the CloudFormation templates stored in the private GitHub repository. Execute the change set using AWS CodeDeploy. Include a CodePipeline action to test the deployment with testing scripts run by AWS CodeBuild.
- B.** Mirror the GitHub repository to AWS CodeCommit using AWS Lambda. Use AWS CodeDeploy to create a change set from the CloudFormation templates and execute it. Have CodeDeploy test the deployment with testing scripts run by AWS CodeBuild.
- C.** Use AWS CodePipeline to create and execute a change set from the CloudFormation templates stored in the GitHub repository. Configure a CodePipeline action to be deployment with testing scripts run by AWS CodeBuild.
- D.** Mirror the GitHub repository to AWS CodeCommit using AWS Lambda. Use AWS CodeBuild to create a change set from the CloudFormation templates and execute it. Have CodeBuild test the deployment with testing scripts.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 72

A company standardized its method of deploying applications to AWS using AWS CodePipeline and AWS Cloud Formation. The applications are in TypeScript and Python. The company has recently acquired another business that deploys applications to AWS using Python scripts. Developers from the newly acquired company are hesitant to move their applications under Cloud Formation because it would require that they learn a new domain-specific language and eliminate their access to language features, such as looping.

How can the acquired applications quickly be brought up to deployment standards while addressing the developers' concerns?

- A.** Create Cloud Formation templates and re-use parts of the Python scripts as Instance user data. Use the AWS Cloud Development Kit (AWS CDK) to deploy the application using these templates. Incorporate the AWS CDK into CodePipeline and deploy the application to AWS using these templates.
- B.** Define the AWS resources using TypeScript or Python Use the AWS Cloud Development Kit (AWS CDK) to create CloudFormation templates from the developers' code, and use the AWS CDK to create CloudFormation stacks Incorporate the AWS CDK as a CodeBuild job in CodePipeline
- C.** Standardize on AWS OpsWorks. Integrate OpsWorks with CodePipeline Have the developers create Chef recipes to deploy their applications on AWS.
- D.** Use a third-party resource provisioning engine inside AWS CodeBuild to standardize the deployment processes of the existing and acquired company Orchestrate the CodeBuild job using CodePipeline.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 73

During a security audit of a Service team's application a Solutions Architect discovers that a username and password for an Amazon RDS database and a set of AWSIAM user credentials can be viewed in the AWS Lambda function code. The Lambda function uses the username and password to run queries on the database and it uses the I AM credentials to call AWS services in a separate management account.

The Solutions Architect is concerned that the credentials could grant inappropriate access to anyone who can view the Lambda code The management account and the Service team's account are in separate AWS Organizations organizational units (OUs) Which combination of changes should the Solutions Architect make to improve the solution's security? (Select TWO)

- A.** Use an SCP on the management accounts OU to prevent IAM users from accessing resources in the Service team's account
- B.** Create a Lambda function to rotate the credentials every hour by deploying a new Lambda version with the updated credentials
- C.** Configure Lambda to use the stored database credentials in AWS Secrets Manager and enable automatic rotation
- D.** Enable AWS Shield Advanced on the management account to shield sensitive resources from unauthorized IAM access

E. Configure Lambda to assume a role in the management account with appropriate access to AWS

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 74

Company is developing a gene reporting device that will collect genomic information to assist researchers with collecting large samples of data from a diverse population. The device will push 8 KB of genomic data every second to a data platform that will need to process and analyze the data and provide Information back to researchers. The data platform must meet the following requirements:

- *Provide near-real-time analytics of the inbound genomic data
- *Ensure the data is flexible, parallel, and durable
- *Deliver results of processing to a data warehouse

Which strategy should a solutions architect use to meet these requirements?

- A.** Use Amazon Kinesis Data Streams to collect the inbound sensor data, analyze the data with Kinesis clients, and save the results to an Amazon Redshift duster using Amazon EMR
- B.** Use an Amazon API Gateway to put requests into an Amazon SQS queue, analyze the data with an AWS Lambda function, and save the results to an Amazon Redshift cluster using Amazon EMR.
- C.** Use Amazon Kinesis Data Firehose to collect the inbound sensor data, analyze the data with Kinesis clients, and save the results to an Amazon RDS instance.
- D.** Use Amazon S3 to collect the inbound device data, analyze the data from Amazon SOS with Kinesis, and save the results to an Amazon Redshift cluster.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 75

A company has developed a new release of a popular video game and wants to make it available for public download The new release package is approximately 5 GB in size The company provides downloads for existing releases from a Linux-based publicly facing FTP site hosted in an on-premises data center The company expects the new release will be downloaded by users worldwide The company wants a solution that provides improved download performance and low transfer costs regardless of a user's location Which solutions will meet these requirements?

- A.** Configure Amazon Route 53 and an Amazon S3 bucket for website hosting Upload the game files to the S3 bucket Set Requester Pays for the S3 bucket Publish the game download URL for users to download the package
- B.** Configure Amazon Route 53 and an Amazon S3 bucket for website hosting Upload the game files to the S3 bucket Use Amazon CloudFront for the website Publish the game download URL for users to download the package
- C.** Store the game files on Amazon EBS volumes mounted on Amazon EC2 instances within an Auto Scaling group Configure an FTP service on the EC2 instances Use an Application Load

Balancer in front of the Auto Scaling group. Publish the game download URL for users to download the package.

D. Store the game files on Amazon EFS volumes that are attached to Amazon EC2 instances within an Auto Scaling group. Configure an FTP service on each of the EC2 instances. Use an Application Load Balancer in front of the Auto Scaling group. Publish the game download URL for users to download the package.

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 76

A company runs a legacy system on a single m4.2xlarge Amazon EC2 instance with Amazon EBS2 storage.

The EC2 instance runs both the web server and a self-managed Oracle database. A snapshot is made of the EBS volume every 12 hours, and an AMI was created from the fully configured EC2 instance.

A recent event that terminated the EC2 instance led to several hours of downtime. The application was successfully launched from the AMI, but the age of the EBS snapshot and the repair of the database resulted in the loss of 8 hours of data. The system was also down for 4 hours while the Systems Operators manually performed these processes.

What architectural changes will minimize downtime and reduce the chance of lost data?

A. Create an Amazon CloudWatch alarm to automatically recover the instance. Create a script that will check and repair the database upon reboot. Subscribe the Operations team to the Amazon SNS message generated by the CloudWatch alarm.

B. Run the application on m4.xlarge EC2 instances behind an Elastic Load Balancer/Application Load Balancer. Run the EC2 instances in an Auto Scaling group across multiple Availability Zones with a minimum instance count of two. Migrate the database to an Amazon RDS Oracle Multi-AZ DB instance.

C. Run the application on m4.2xlarge EC2 instances behind an Elastic Load Balancer/Application Load Balancer. Run the EC2 instances in an Auto Scaling group across multiple Availability Zones with a minimum instance count of one. Migrate the database to an Amazon RDS Oracle Multi-AZ DB instance.

D. Increase the web server instance count to two m4.xlarge instances and use Amazon Route 53 round-robin load balancing to spread the load. Enable Route 53 health checks on the web servers.

Migrate the database to an Amazon RDS Oracle Multi-AZ DB instance.

Answer: B ([LEAVE A REPLY](#))

Explanation

Ensures that there are at least two EC instances, each of which is in a different AZ. It also ensures that the database spans multiple AZs. Hence this meets all the criteria.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html>

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (216 Q&As Dumps, **30%OFF** Special Discount: **freecram**)

NEW QUESTION: 77

A solutions architect is designing a network for a new cloud deployment. Each account will need autonomy to modify route tables and make changes. Centralized and controlled egress internet connectivity is also needed.

The cloud footprint is expected to grow to thousands of AWS accounts.

Which architecture will meet these requirements?

A. A shared services VPC to host central assets to include a fleet of firewalls with a route to the internet.

Each spoke VPC will peer to the central VPC.

B. A centralized shared VPC with a subnet for each account. Outbound internet traffic will be controlled through a fleet of proxy servers.

C. A shared transit gateway to which each VPC will be attached. Outbound internet access will route through a fleet of VPN-attached firewalls.

D. A centralized transit VPC with a VPN connection to a standalone VPC in each account. Outbound internet traffic will be controlled by firewall appliances.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 78

A company currently has data hosted in an IBM Db2 database. A web application calls an API that runs stored procedures on the database to retrieve user information data that is read-only. This data is historical in nature and changes on a daily basis. When a user logs in to the application, this data needs to be retrieved within 3 seconds. Each time a user logs in, the stored procedures run. Users log in several times a day to check stock prices.

Running this database has become cost-prohibitive due to Db2 CPU licensing. Performance goals are not being met. Timeouts from Db2 are common due to long-running queries. Which approach should a solutions architect take to migrate this solution to AWS?

A. Rehost the Db2 database in Amazon Fargate. Migrate all the data. Enable caching in Fargate. Refactor the API to use the Fargate Db2 database. Implement Amazon API Gateway and enable API caching.

B. Create a local cache on the mainframe to store query outputs. Use SFTP to sync to Amazon S3 on a daily basis. Refactor the API to use Amazon EFS. Implement Amazon API Gateway and enable API caching.

C. Extract data daily and copy the data to AWS Snowball for storage on Amazon S3. Sync daily. Refactor the API to use the S3 data. Implement Amazon API Gateway and enable API caching.

D. Use AWS DMS to migrate data to Amazon DynamoDB using a continuous replication task. Refactor the API to use the DynamoDB data. Implement the refactored API in Amazon API Gateway and enable API caching

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 79

A company's main intranet page has experienced degraded response times as its user base has increased although there are no reports of users seeing error pages. The application uses Amazon DynamoDB in read-only mode.

Amazon DynamoDB latency metrics for successful requests have been in a steady state even during times when users have reported degradation. The Development team has correlated the issue to ProvisionedThroughput Exceeded exceptions in the application logs when doing Scan and read operations. The team also identified an access pattern of steady spikes of read activity on a distributed set of individual data items. The Chief Technology Officer wants to improve the user experience. Which solutions will meet these requirements with the LEAST amount of changes to the application? (Select TWO)

- A.** Remove error retries and exponential backoffs in the application code to handle throttling errors
- B.** Change the data model of the DynamoDB tables to ensure that all Scan and read operations meet DynamoDB best practices of uniform data access, reaching the full request throughput provisioned for the DynamoDB tables
- C.** Provision Amazon ElastiCache for Redis with cluster mode enabled. The cluster should be provisioned with enough shards to spread the application load and provision at least one read replica node for each shard
- D.** Implement the DynamoDB Accelerator (DAX) client and provision a DAX cluster with the appropriate node types to sustain the application load. Tune the item and query cache configuration for an optimal user experience
- E.** Enable DynamoDB auto scaling to manage the throughput capacity as table traffic increases. Set the upper and lower limits to control costs and set a target utilization given the peak usage and how quickly the traffic changes.

Answer: D,E ([LEAVE A REPLY](#))

NEW QUESTION: 80

A company is configuring connectivity to a multi-account AWS environment to support application workloads that serve users in a single geographic region. The workloads depend on a highly available on-premises legacy system deployed across two locations. It is critical for the AWS workloads to maintain connectivity to the legacy system and a minimum of 5 Gbps of bandwidth is required. All application workloads within AWS must have connectivity with one another. Which solution will meet these requirements?

- A.** Configure multiple AWS Direct Connect (DX) 10 Gbps dedicated connections from two DX partners for each on-premises location Create a transit gateway and a DX gateway in a central network account Create a transit virtual interface for each DX interface and associate them with the DX gateway Create a gateway association between the DX gateway and the transit gateway
- B.** Configure multiple AWS Direct Connect (DX) 10 Gbps dedicated connections from two DX partners for each on-premises location Create and attach a virtual private gateway for each AWS account VPC Create a DX gateway in a central network account and associate it with the virtual private gateways Create a public virtual interface on each DX connection and associate the interface with the DX gateway
- C.** Configure multiple AWS Direct Connect (DX) 10 Gbps dedicated connections from a DX partner for each on-premises location Create private virtual interfaces on each connection for each AWS account VPC Associate the private virtual interface with a virtual private gateway attached to each VPC
- D.** Configure multiple AWS Direct Connect (DX) 10 Gbps dedicated connections from a DX partner for each on-premises location Create and attach a virtual private gateway for each AWS account VPC Create a transit gateway in a central network account and associate it with the virtual private gateways Create a transit virtual interface on each DX connection and attach the interface to the transit gateway

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 81

A company is migrating a subset of its application APIs from Amazon EC2 instances to run on a serverless infrastructure. The company has set up Amazon API Gateway, AWS Lambda, and Amazon DynamoDB for the new application. The primary responsibility of the Lambda function is to obtain data from a third-party Software as a Service (SaaS) provider. For consistency, the Lambda function is attached to the same virtual private cloud (VPC) as the original EC2 instances.

Test users report an inability to use this newly moved functionality, and the company is receiving 5xx errors from API Gateway. Monitoring reports from the SaaS provider shows that the requests never made it to its systems. The company notices that Amazon CloudWatch Logs are being generated by the Lambda functions.

When the same functionality is tested against the EC2 systems, it works as expected.

What is causing the issue?

- A.** API Gateway does not have the necessary permissions to invoke Lambda.
- B.** Lambda is in a subnet that does not have a NAT gateway attached to it to connect to the SaaS provider.
- C.** The end-user application is misconfigured to continue using the endpoint backed by EC2 instances.
- D.** The throttle limit set on API Gateway is too low and the requests are not making their way through.

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 82

A company is manually deploying its application to production and wants to move to a more mature deployment pattern. The company has asked a solutions architect to design a solution that leverages its current Chef tools and knowledge. The application must be deployed to a staging environment for testing and verification before being deployed to production. Any new deployment must be rolled back in 5 minutes if errors are discovered after a deployment. Which AWS service and deployment pattern should the solutions architect use to meet these requirements?

- A. Use AWS Elastic Beanstalk and deploy the application using a rolling update deployment strategy
- B. Use AWS CodeBuild and deploy the application using a canary deployment strategy
- C. Use AWS OpsWorks and deploy the application using a blue/green deployment strategy
- D. Use AWS CodePipeline and deploy the application using a rolling update deployment strategy

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 83

A company has a media metadata extraction pipeline running on AWS. Notifications containing a reference to a file in Amazon S3 are sent to an Amazon Simple Notification Service (Amazon SNS) topic. The pipeline consists of a number of AWS Lambda functions that are subscribed to the SNS topic. The Lambda functions extract the S3 file and write metadata to an Amazon RDS PostgreSQL DB instance. Users report that updates to the metadata are sometimes slow to appear. 01 are lost. During these times, the CPU utilization on the database is high and the number of failed Lambda invocations increases. Which combination of actions should a solutions architect take to help resolve this issue? (Select TWO)

- A. Enable the RDS Data API for the RDS instance. Update the Lambda functions to connect to the RDS instance using the Data API
- B. Create an RDS proxy for the RDS instance. Update the Lambda functions to connect to the RDS instance using the proxy
- C. Create an Amazon Simple Queue Service (Amazon SQS) FIFO queue and subscribe the queue to the SNS topic. Configure the Lambda functions to consume messages from the SQS queue
- D. Create an Amazon Simple Queue Service (Amazon SQS) standard queue for each Lambda function and subscribe the queues to the SNS topic. Configure the Lambda functions to consume messages from their respective SQS queue
- E. Enable message delivery status on the SNS topic. Configure the SNS topic delivery policy to enable retries with exponential backoff

Answer: B,D ([LEAVE A REPLY](#))

NEW QUESTION: 84

A three-tier web application runs on Amazon EC2 instances. Cron daemons are used to trigger scripts that collect the web server, application, and database logs and send them to a centralized location every hour.

Occasionally, scaling events or unplanned outages have caused the instances to stop before the latest logs were collected, and the log files were lost.

Which of the following options is the MOST reliable way of collecting and preserving the log files?

A. Update the cron jobs to run every 5 minutes instead of every hour to reduce the possibility of log messages being lost in an outage.

B. Use Amazon CloudWatch Events to trigger Amazon Systems Manager Run Command to invoke the log collection scripts more frequently to reduce the possibility of log messages being lost in an outage.

C. Use the Amazon CloudWatch Logs agent to stream log messages directly to CloudWatch Logs.

Configure the agent with a batch count of 1 to reduce the possibility of log messages being lost in an outage.

D. Use Amazon CloudWatch Events to trigger AWS Lambda to SSH into each running instance and invoke the log collection scripts more frequently to reduce the possibility of log messages being lost in an outage.

Answer: ([SHOW ANSWER](#))

Explanation

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/AgentReference.html>

NEW QUESTION: 85

A company has several Amazon EC2 instances in both public and private subnets within a VPC that is not connected to the corporate network. A security group associated with the EC2 instances allows the company to use the Windows remote desktop protocol (RDP) over the internet to access the instances. The security team has noticed connection attempts from unknown sources. The company wants to implement a more secure solution to access the EC2 instances.

Which strategy should a solutions architect implement?

A. Deploy AWS Systems Manager Agent on the EC2 instances. Access the EC2 instances using Session Manager, restricting access to users with permission.

B. Deploy a Linux bastion host with an Elastic IP address in the public subnet. Allow access to the bastion host from 0.0.0.0/0.

C. Establish a Site-to-Site VPN connecting the corporate network to the VPC. Update the security groups to allow access from the corporate network only.

D. Deploy a Linux bastion host on the corporate network that has access to all instances in the VPC.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 86

A Solutions Architect is working with a company that operates a standard three-tier web application in AWS.

The web and application tiers run on Amazon EC2 and the database tier runs on Amazon RDS. The company is redesigning the web and application tiers to use Amazon API Gateway and AWS Lambda, and the company intends to deploy the new application within 6 months. The IT Manager has asked the Solutions Architect to reduce costs in the interim.

Which solution will be MOST cost effective while maintaining reliability?

- A.** Use Spot Instances for the web and application tiers, and Reserved Instances for the database tier.
- B.** Use Reserved Instances for the web, application, and database tiers.
- C.** Use Spot Instances for the web tier, On-Demand Instances for the application tier, and Reserved Instances for the database tier.
- D.** Use On-Demand Instances for the web and application tiers, and Reserved Instances for the database tier.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 87

A group of research institutions and hospitals are in a partnership to study 2 PBs of genomic data. The institute that owns the data stores it in an Amazon S3 bucket and updates it regularly. The institute would like to give all of the organizations in the partnership read access to the data. All members of the partnership are extremely cost-conscious, and the institute that owns the account with the S3 bucket is concerned about covering the costs for requests and data transfers from Amazon S3.

Which solution allows for secure datasharing without causing the institute that owns the bucket to assume all the costs for S3 requests and data transfers?

- A.** Ensure that all organizations in the partnership have AWS accounts. In the account with the S3 bucket, create a cross-account role for each account in the partnership that allows read access to the data. Have the organizations assume and use that read role when accessing the data.
- B.** Ensure that all organizations in the partnership have AWS accounts. Create a bucket policy on the bucket that owns the data. The policy should allow the accounts in the partnership read access to the bucket. Enable Requester Pays on the bucket. Have the organizations use their AWS credentials when accessing the data.
- C.** Ensure that all organizations in the partnership have AWS accounts. Configure buckets in each of the accounts with a bucket policy that allows the institute that owns the data the ability to write to the bucket. Periodically sync the data from the institute's account to the other organizations. Have the organizations use their AWS credentials when accessing the data using their accounts.
- D.** Ensure that all organizations in the partnership have AWS accounts. In the account with the S3 bucket, create a cross-account role for each account in the partnership that allows read access to the data. Enable Requester Pays on the bucket. Have the organizations assume and use that read role when accessing the data.

Answer: (SHOW ANSWER)

Explanation

<https://docs.aws.amazon.com/AmazonS3/latest/dev/RequesterPaysBuckets.html>

NEW QUESTION: 88

A developer reports receiving an Error 403: Access Denied message when they try to download an object from an Amazon S3 bucket. The S3 bucket is accessed using an S3 endpoint inside a VPC, and is encrypted with an AWS KMS key. A solution architect has verified that the developer is assuming the correct IAM role in the account that allows the object to be downloaded. The S3 bucket policy and the NACL are also valid.

Which additional step should the solutions architect take to troubleshoot this issue?

- A. Ensure that blocking all public access has not been enabled in the S3 bucket.
- B. Verify that the IAM role has the correct trust relationship configured.
- C. Check that local firewall rules are not preventing access to the S3 endpoint.
- D. Verify that the IAM role has permission to decrypt the referenced KMS key.

Answer: A (LEAVE A REPLY)

NEW QUESTION: 89

A company provides auction services for artwork and has users across North America and Europe. The company hosts its application in Amazon EC2 instances in the us-east-1 Region. Artists upload photos of their work as large-size, high-resolution image files from their mobile phones to a centralized Amazon S3 bucket created in the us-east-1 Region. The users in Europe are reporting slow performance for their image uploads.

How can a solutions architect improve the performance of the image upload process?

- A. Redeploy the application to use S3 multipart uploads.
- B. Create an Amazon CloudFront distribution and point to the application as a custom origin.
- C. Configure the buckets to use S3 Transfer Acceleration.
- D. Create an Auto Scaling group for the EC2 instances and create a scaling policy.

Answer: C (LEAVE A REPLY)

Explanation

<https://docs.aws.amazon.com/AmazonS3/latest/dev/transfer-acceleration.html>

NEW QUESTION: 90

An organization has a write-intensive mobile application that uses Amazon API Gateway, AWS Lambda, and Amazon DynamoDB. The application has scaled well, however, costs have increased exponentially because of higher than anticipated Lambda costs. The application's use is unpredictable, but there has been a steady 20% increase in utilization every month.

While monitoring the current Lambda functions, the Solutions Architect notices that the execution-time averages 4.5 minutes. Most of the wait time is the result of a high-latency network call to a 3-TB MySQL database server that is on-premises. A VPN is used to connect to the VPC, so the Lambda functions have been configured with a five-minute timeout.

How can the Solutions Architect reduce the cost of the current architecture?

- A.** Migrate the MySQL database server into a Multi-AZ Amazon RDS for MySQL. Enable caching of the Amazon API Gateway results in Amazon CloudFront to reduce the number of Lambda function invocations. Monitor the Lambda function performance; gradually adjust the timeout and memory properties to lower values while maintaining an acceptable execution time. Enable DynamoDB Accelerator for frequently accessed records, and enable the DynamoDB Auto Scaling feature.
- B.** Migrate the MySQL database server into a Multi-AZ Amazon RDS for MySQL. Enable API caching on API Gateway to reduce the number of Lambda function invocations. Continue to monitor the AWS Lambda function performance; gradually adjust the timeout and memory properties to lower values while maintaining an acceptable execution time. Enable Auto Scaling in DynamoDB.
- C.** Replace the VPN with AWS Direct Connect to reduce the network latency to the on-premises MySQL database. Enable local caching in the mobile application to reduce the Lambda function invocation calls. Monitor the Lambda function performance; gradually adjust the timeout and memory properties to lower values while maintaining an acceptable execution time. Offload the frequently accessed records from DynamoDB to Amazon ElastiCache.
- D.** Replace the VPN with AWS Direct Connect to reduce the network latency to the on-premises MySQL database. Cache the API Gateway results to Amazon CloudFront. Use Amazon EC2 Reserved Instances instead of Lambda. Enable Auto Scaling on EC2, and use Spot Instances during peak times. Enable DynamoDB Auto Scaling to manage target utilization.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 91

A company's application is increasingly popular and experiencing latency because of high volume reads on the database server.

The service has the following properties:

- * A highly available REST API hosted in one region using Application Load Balancer (ALB) with auto scaling.
- * A MySQL database hosted on an Amazon EC2 instance in a single Availability Zone.

The company wants to reduce latency, increase in-region database read performance, and have multi-region disaster recovery capabilities that can perform a live recovery automatically without any data or performance loss (HA/DR).

Which deployment strategy will meet these requirements?

- A.** Use AWS CloudFormation StackSets to deploy the API layer in two regions. Add the database to an Auto Scaling group. Add a read replica to the database in the second region. Use Amazon Route 53 health checks on the database to trigger a DNS failover to the standby region if the health checks in the primary region fail. Promote the cross-region database replica to be the master and build out new read replicas in the standby region.
- B.** Use AWS CloudFormation StackSets to deploy the API layer in two regions. Migrate the database to an Amazon Aurora with MySQL database cluster with multiple read replicas in one

region and a read replica in a different region than the source database cluster. Use Amazon Route 53 health checks to trigger a DNS failover to the standby region if the health checks to the primary load balancer fail. In the event of Route 53 failover, promote the cross-region database replica to be the master and build out new read replicas in the standby region.

C. Use Amazon ElastiCache for Redis Multi-AZ with an automatic failover to cache the database read queries. Use AWS OpsWorks to deploy the API layer, cache layer, and existing database layer in two regions. Use Amazon Route 53 health checks on the ALB to trigger a DNS failover to the standby region if the health checks in the primary region fail. Back up the MySQL database frequently, and in the event of a failure in an active region, copy the backup to the standby region and restore the standby database.

D. Use Amazon ElastiCache for Redis Multi-AZ with an automatic failover to cache the database read queries. Use AWS OpsWorks to deploy the API layer, cache layer, and existing database layer in two regions. In the event of failure, use Amazon Route 53 health checks on the database to trigger a DNS failover to the standby region if the health checks in the primary region fail. Back up the MySQL database frequently, and in the event of a failure in an active region, copy the backup to the standby region and restore the standby database.

Answer: B ([LEAVE A REPLY](#))

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (216 Q&As Dumps, **30%OFF** Special Discount: **freecram**)

NEW QUESTION: 92

A company is creating a REST API to share information with six of its partners based in the United States.

The company has created an Amazon API Gateway Regional endpoint. Each of the six partners will access the API once per day to post daily sales figures.

After initial deployment, the company observes 1,000 requests per second originating from 500 different IP addresses around the world. The company believes this traffic is originating from a botnet and wants to secure its API while minimizing cost. Which approach should the company take to secure its API?

A. Create an AWS WAF web ACL with a rule to allow access to the IP addresses used by the six partners. Associate the web ACL with the API. Create a usage plan with a request limit and associate it with the API. Create an API key and add it to the usage plan.

- B.** Create an Amazon CloudFront distribution with the API as the origin Create an AWS WAF web ACL with a rule to block clients that submit more than five requests per day. Associate the web ACL with the CloudFront distribution Add a custom header to the CloudFront distribution populated with an API key Configure the API to require an API key on the POST method
- C.** Create an Amazon CloudFront distribution with the API as the origin Create an AWS WAF web ACL with a rule to block clients that submit more than five requests per day. Associate the web ACL with the CloudFront distribution Configure CloudFront with an origin access identity (OAI) and associate it with the distribution Configure API Gateway to ensure only the OAI can execute the POST method
- D.** Create an AWS WAF web ACL with a rule to allow access to the IP addresses used by the six partners Associate the web ACL with the API Create a resource policy with a request limit and associate it with the API Configure the API to require an API key on the POST method

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 93

A Solutions Architect has been asked to look at a company's Amazon Redshift cluster, which has quickly become an integral part of its technology and supports key business process. The Solutions Architect is to increase the reliability and availability of the cluster and provide options to ensure that if an issue arises, the cluster can either operate or be restored within four hours. Which of the following solution options BEST addresses the business need in the most cost-effective manner?

- A.** Ensure that the Amazon Redshift cluster has been set up to make use of Auto Scaling groups with the nodes in the cluster spread across multiple Availability Zones.
- B.** Ensure that the Amazon Redshift cluster creation has been template using AWS CloudFormation so it can easily be launched in another Availability Zone and data populated from the automated Redshift back-ups stored in Amazon S3.
- C.** Use Amazon Kinesis Data Firehose to collect the data ahead of ingestion into Amazon Redshift and create clusters using AWS CloudFormation in another region and stream the data to both clusters.
- D.** Create two identical Amazon Redshift clusters in different regions (one as the primary, one as the secondary). Use Amazon S3 cross-region replication from the primary to secondary). Use Amazon S3 cross-region replication from the primary to secondary region, which triggers an AWS Lambda function to populate the cluster in the secondary region.

Answer: B ([LEAVE A REPLY](#))

Explanation

https://aws.amazon.com/redshift/faqs/?nc1=h_ Is Q: What happens to my data warehouse cluster availability and data durability if my data warehouse cluster's Availability Zone (AZ) has an outage? If your Amazon Redshift data warehouse cluster's Availability Zone becomes unavailable, you will not be able to use your cluster until power and network access to the AZ are restored. Your data warehouse cluster's data is preserved so you can start using your Amazon Redshift data warehouse as soon as the AZ becomes available again. In addition, you can also

choose to restore any existing snapshots to a new AZ in the same Region. Amazon Redshift will restore your most frequently accessed data first so you can resume queries as quickly as possible.

FROM 37

NEW QUESTION: 94

A Solutions Architect must migrate an existing on-premises web application with 70 TB of static files supporting a public open-data initiative. The architect wants to upgrade to the latest version of the host operating system as part of the migration effort.

Which is the FASTEST and MOST cost-effective way to perform the migration?

- A.** Run a physical-to-virtual conversion on the application server. Transfer the server image over the internet, and transfer the static data to Amazon S3.
- B.** Re-platform the server by using the AWS Server Migration Service to move the code and data to a new Amazon EC2 instance.
- C.** Run a physical-to-virtual conversion on the application server. Transfer the server image over AWS Direct Connect, and transfer the static data to Amazon S3.
- D.** Re-platform the server to Amazon EC2, and use AWS Snowball to transfer the static data to Amazon S3.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 95

A healthcare company runs a production workload on AWS that stores highly sensitive personal information.

The security team mandates that, for auditing purposes, any AWS API action using AWS account root user credentials must automatically create a high-priority ticket in the company's ticketing system. The ticketing system has a monthly 3-hour maintenance window when no tickets can be created.

To meet security requirements, the company enabled AWS CloudTrail logs and wrote a scheduled AWS Lambda function that uses Amazon Athena to query API actions performed by the root user. The Lambda function submits any actions found to the ticketing system API. During a recent security audit, the security team discovered that several tickets were not created because the ticketing system was unavailable due to planned maintenance.

Which combination of steps should a solutions architect take to ensure that the incidents are reported to the ticketing system even during planned maintenance? (Select TWO.)

- A.** Create an Amazon EventBridge rule that triggers on all API events where the invoking user identity is root. Configure the EventBridge rule to write the event to an Amazon SQS queue.
- B.** Create an Amazon SNS topic to which Amazon CloudWatch alarms will be published. Configure a CloudWatch alarm to invoke the Lambda function.
- C.** Modify the Lambda function to be triggered when there are messages in the Amazon SQS queue and to return successfully when the ticketing system API has processed the request.

D. Create an Amazon SQS queue to which Amazon CloudWatch alarms will be published. Configure a CloudWatch alarm to publish to the SQS queue.

E. Modify the Lambda function to be triggered by messages published to an Amazon SNS topic. Update the existing application code to retry every 5 minutes if the ticketing system's API endpoint is unavailable.

Answer: C,D ([LEAVE A REPLY](#))

NEW QUESTION: 96

A company is migrating its on-premises build artifact server to an AWS solution. The current system consists of an Apache HTTP server that serves artifacts to clients on the local network, restricted by the perimeter firewall. The artifact consumers are largely build automation scripts that download artifacts via anonymous HTTP, which the company will be unable to modify within its migration timetable.

The company decides to move the solution to Amazon S3 static website hosting. The artifact consumers will be migrated to Amazon EC2 instances located within both public and private subnets in a virtual private cloud (VPC).

Which solution will permit the artifact consumers to download artifacts without modifying the existing automation scripts?

A. Create a NAT gateway within a public subnet of the VPC. Add a default route pointing to the NAT gateway into the route table associated with the subnets containing consumers. Configure the bucket policy to allow the s3:ListBucket and s3:GetObject actions using the condition IpAddress and the condition key aws:SourceIp matching the elastic IP address of the NAT gateway.

B. Create a VPC endpoint and add it to the route table associated with subnets containing consumers.

Configure the bucket policy to allow s3:ListBucket and s3:GetObject actions using the condition IpAddress and the condition key aws:SourceIp matching the VPC CIDR block.

C. Create an IAM role and instance profile for Amazon EC2 and attach it to the instances that consume build artifacts. Configure the bucket policy to allow the s3:ListBucket and s3:GetObjects actions for the principal matching the IAM role created.

D. Create a VPC endpoint and add it to the route table associated with subnets containing consumers.

Configure the bucket policy to allow s3:ListBucket and s3:GetObject actions using the condition StringEquals and the condition key aws:sourceVpce matching the identification of the VPC endpoint.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 97

A large company recently experienced an unexpected increase in Amazon RDS and Amazon DynamoDB costs. The company needs to increase visibility into delays of AWS Billing and Cost

Management. There are various accounts associated with AWS Organizations, including many development and production accounts.

There is no consistent tagging strategy across the organization, but there are guidelines in place that require all infrastructure to be deployed using AWS CloudFormation with consistent tagging. Management requires cost center numbers and project ID numbers for all existing and future DynamoDB tables and RDS instances.

Which strategy should the solutions architect provide to meet these requirements?

- A.** Use an AWS Config rule to alert the finance team of untagged resources. Create a centralized AWS Lambda based solution to tag untagged RDS databases and DynamoDB resources every hour using a cross-account role
- B.** Use Tag Editor to tag existing resources. Create cost allocation tags to define the cost center and project ID. Use SCPs to restrict resource creation that do not have the cost center and project ID on the resource.
- C.** Use Tag Editor to tag existing resources. Create cost allocation tags to define the cost center and project ID and allow 24 hours for tags to propagate to existing resources.
- D.** Create cost allocation tags to define the cost center and project ID and allow 24 hours for tags to propagate to existing resources. Update existing federated roles to restrict privileges to provision resources that do not include the cost center and project ID on the resource

Answer: A ([LEAVE A REPLY](#))

NEW QUESTION: 98

A company has developed a new billing application that will be released in two weeks.

Developers are testing the application running on 10 EC2 instances managed by an Auto Scaling group in subnet 172.31.0.0/24 within VPC A with CIDR block 172.31.0.0/16. The Developers noticed connection timeout errors in the application logs while connecting to an Oracle database running on an Amazon EC2 instance in the same region within VPC B with CIDR block 172.50.0.0/16. The IP of the database instance is hard-coded in the application instances.

Which recommendations should a Solutions Architect present to the Developers to solve the problem in a secure way with minimal maintenance and overhead?

- A.** Create an additional Amazon EC2 instance for each VPC as a customer gateway; create one virtual private gateway (VGW) for each VPC, configure an end-to-end VPC, and advertise the routes for 172.50.0.0/16
- B.** Create a VPC peering connection between the two VPCs and add a route to the routing table of VPC A that points to the IP address range of 172.50.0.0/16
- C.** Disable the SrcDestCheck attribute for all instances running the application and Oracle Database.

Change the default route of VPC A to point ENI of the Oracle Database that has an IP address assigned within the range of 172.50.0.0/26

- D.** Create and attach internet gateways for both VPCs. Configure default routes to the Internet gateways for both VPCs. Assign an Elastic IP for each Amazon EC2 instance in VPC A

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 99

A company runs an IoT platform on AWS. IoT sensors in various locations send data to the company's Node.js API servers on Amazon EC2 instances running behind an Application Load Balancer. The data is stored in an Amazon RDS MySQL DB instance that uses a 4 TB General Purpose SSD volume.

The number of sensors the company has deployed in the field has increased over time, and is expected to grow significantly. The API servers are consistently overloaded and RDS metrics show high write latency.

Which of the following steps together will resolve the issues permanently and enable growth as new sensors are provisioned, while keeping this platform cost-efficient? (Choose two.)

- A.** Re-architect the database tier to use Amazon Aurora instead of an RDS MySQL DB instance and add read replicas
- B.** Re-architect the database tier to use Amazon DynamoDB instead of an RDS MySQL DB instance
- C.** Leverage Amazon Kinesis Data Streams and AWS Lambda to ingest and process the raw data
- D.** Use AWS-X-Ray to analyze and debug application issues and add more API servers to match the load
- E.** Resize the MySQL General Purpose SSD storage to 6 TB to improve the volume's IOPS

Answer: B,C ([LEAVE A REPLY](#))

NEW QUESTION: 100

A company has an application written using an in-house software framework. The framework installation takes 30 minutes and is performed with a user data script. Company Developers deploy changes to the application frequently. The framework installation is becoming a bottleneck in this process.

Which of the following would speed up this process?

- A.** Create a pipeline to build a custom AMI with the framework installed and use this AMI as a baseline for application deployments.
- B.** Employ a user data script to install the framework but compress the installation files to make them smaller.
- C.** Create a pipeline to parallelize the installation tasks and call this pipeline from a user data script.
- D.** Configure an AWS OpsWorks cookbook that installs the framework instead of employing user data. Use this cookbook as a base for all deployments.

Answer: ([SHOW ANSWER](#))

Explanation

<https://aws.amazon.com/codepipeline/features/?nc=sn&loc=2>

NEW QUESTION: 101

A company wants to follow its website on AWS using serverless architecture design patterns for global customers. The company has outlined its requirements as follow:

- * The website should be responsive.
- * The website should offer minimal latency.
- * The website should be highly available.
- * Users should be able to authenticate through social identity providers such as Google, Facebook, and Amazon.
- * There should be baseline DDoS protections for spikes in traffic.

How can the design requirements be met?

A. Use AWS Direct Connect with Amazon CloudFront and Amazon S3 for hosting static web resources.

Use Amazon Cognito to provide user management and authentication functions. Use AWS Lambda to build an API.

B. Use Amazon Route 53 latency routing with an Application Load Balancer and AWS Fargate in different regions for hosting the website. Use Amazon Cognito to provide user management and authentication functions. Use Amazon EKS containers to build an API.

C. Use Amazon CloudFront with Amazon ECS for hosting the website. Use AWS Secrets Manager to provide user management and authentication functions. Use ECS Docker containers to build an API.

D. Use Amazon CloudFront with Amazon S3 for hosting static web resources. Use Amazon Cognito to provide user management and authentication functions. Use Amazon API Gateway with AWS Lambda to build an API.

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 102

An organization has two Amazon EC2 instances:

- * The first is running an ordering application and an inventory application.
- * The second is running a queuing system.

During certain times of the year, several thousand orders are placed per second. Some orders were lost when the queuing system was down. Also, the organization's inventory application has the incorrect quantity of products because some orders were processed twice.

What should be done to ensure that the applications can handle the increasing number of orders?

A. Put the ordering and inventory applications into their own AWS Lambda functions. Have the ordering application write the messages into an Amazon SQS FIFO queue.

B. Put the ordering and inventory applications into their own Amazon ECS containers and create an Auto Scaling group for each application. Then, deploy the message queuing server in multiple Availability Zones.

C. Put the ordering and inventory applications into their own Amazon EC2 instances, and create an Auto Scaling group for each application. Use Amazon SQS standard queues for the incoming orders, and implement idempotency in the inventory application.

D. Put the ordering and inventory applications into their own Amazon EC2 instances. Write the incoming orders to an Amazon Kinesis data stream. Configure AWS Lambda to poll the stream and update the inventory application.

Answer: C ([LEAVE A REPLY](#))

Explanation

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/standard-queues.html>

NEW QUESTION: 103

A company wants to refactor its retail ordering web application that currently has a load-balanced Amazon EC2 instance fleet for web hosting, database API services, and business logic. The company needs to create a decoupled, scalable architecture with a mechanism for retaining failed orders while also minimizing operational costs. Which solution will meet these requirements?

A. Use Amazon Lightsail for web hosting with AWS AppSync for database API services. Use Amazon Simple Email Service (Amazon SES) for order queuing. Use Amazon Elastic Kubernetes Service (Amazon EKS) for business logic with Amazon Elasticsearch Service (Amazon ES) for retaining failed orders.

B. Use AWS Elastic Beanstalk for web hosting with Amazon API Gateway for database API services. Use Amazon MQ for order queuing. Use AWS Step Functions for business logic with Amazon S3 Glacier Deep Archive for retaining failed orders.

C. Use Amazon S3 for web hosting with Amazon API Gateway for database API services. Use Amazon Simple Queue Service (Amazon SQS) for order queuing. Use Amazon Elastic Container Service (Amazon ECS) for business logic with Amazon SQS long polling for retaining failed orders.

D. Use Amazon S3 for web hosting with AWS AppSync for database API services. Use Amazon Simple Queue Service (Amazon SQS) for order queuing. Use AWS Lambda for business logic with an Amazon SQS dead-letter queue for retaining failed orders.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 104

A company runs a video processing platform. Files are uploaded by users who connect to a web server, which stores them on an Amazon EFS share. This web server is running on a single Amazon EC2 instance. A different group of instances, running in an Auto Scaling group, scans the EFS share directory structure for new files to process and generates new videos (thumbnails, different resolution, compression, etc.) according to the instructions file, which is uploaded along with the video files. A different application running on a group of instances managed by an Auto Scaling group processes the video files and then deletes them from the EFS share. The results are stored in an S3 bucket. Links to the processed video files are emailed to the customer. The company has recently discovered that as they add more instances to the Auto Scaling Group, many files are processed twice, so image processing speed is not improved. The maximum size of these video files is

2GB.

What should the Solutions Architect do to improve reliability and reduce the redundant processing of video files?

A. Rewrite the web application to run from Amazon S3 and upload the video files to an S3 bucket. Each time a new file is uploaded, trigger an AWS Lambda function to put a message in an SQS queue containing the link and the instructions. Modify the video processing application to read from the SQS queue and the S3 bucket. Use the queue depth metric to adjust the size of the Auto Scaling group for video processing instances.

B. Rewrite the web application to run directly from Amazon S3 and use Amazon API Gateway to upload the video files to an S3 bucket. Use an S3 trigger to run an AWS Lambda function each time a file is uploaded to process and store new video files in a different bucket. Using CloudWatch Events, trigger an SES job to send an email to the customer containing the link to the processed file.

C. Modify the web application to upload the video files directly to Amazon S3. Use Amazon CloudWatch Events to trigger an AWS Lambda function every time a file is uploaded, and have this Lambda function put a message into an Amazon SQS queue. Modify the video processing application to read from SQS queue for new files and use the queue depth metric to scale instances in the video processing Auto Scaling group.

D. Set up a cron job on the web server instance to synchronize the contents of the EFS share into Amazon S3. Trigger an AWS Lambda function every time a file is uploaded to process the video file and store the results in Amazon S3. Using Amazon CloudWatch Events trigger an Amazon SES job to send an email to the customer containing the link to the processed file.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 105

A company is using Amazon Aurora MySQL for a customer relationship management (CRM) application. The application requires frequent maintenance on the database and the Amazon EC2 instances on which the application runs. For AWS Management Console access, the system administrators authenticate against AWS Identity and Access Management (IAM) using an internal identity provider. For database access, each system administrator has a user name and password that have previously been configured within the database.

A recent security audit revealed that the database passwords are not frequently rotated. The company wants to replace the passwords with temporary credentials using the company's existing AWS access controls. Which set of options will meet the company's requirements?

A. Enable IAM database authentication on the database. Attach an IAM policy to each system administrator's role to map the role to the database user name. Install the Amazon Aurora SSL certificate bundle to the system administrators' certificate trust store. Use the AWS CLI to generate an authentication token used when connecting to the database.

B. Create a new AWS Systems Manager Parameter Store entry for each database password. Enable parameter expiration to invoke an AWS Lambda function to perform password rotation by updating the parameter value. Create an IAM policy allowing each system administrator to

retrieve their current password from the Parameter Store. Use the AWS CLI to retrieve credentials when connecting to the database

C. Enable 1AM database authentication on the database Configure the database to use the 1AM identity provider to map the administrator roles to the database user Install the Amazon Aurora SSL certificate bundle to the system administrators' certificate trust store Use the AWS CLI to generate an authentication token used when connecting to the database.

D. Create a new AWS Secrets Manager entry for each database password Configure password rotation for each secret using an AWS Lambda function in the same VPC as the database cluster Create an 1AM policy allowing each system administrator to retrieve their current password Use the AWS CLI to retrieve credentials when connecting to the database.

Answer: A ([LEAVE A REPLY](#))

NEW QUESTION: 106

A company is using AWS for production and development workloads. Each business unit has its own AWS account for production, and a separate AWS account to develop and deploy its applications. The Information Security department has introduced new security policies that limit access for terminating certain Amazon EC2 instances in all accounts to a small group of individuals from the Security team.

How can the Solutions Architect meet these requirements?

A. Create an organizational unit under AWS Organizations. Move all the accounts into this organizational unit and use SCP to apply a whitelist policy to allow access to these EC2 instances for the Security team only.

B. Create a new IAM policy that allows access to those EC2 instances only for the Security team. Apply this policy to the AWS Organizations master account.

C. Set up SAML federation for all accounts in AWS. Configure SAML so that it checks for the service API call before authenticating the user. Block SAML from authenticating API calls if anyone other than the Security team accesses these instances.

D. Create a new tag-based IAM policy that allows access to these EC2 instances only for the Security team.

Tag the instances appropriately, and apply this policy in each account.

Answer: D ([LEAVE A REPLY](#))

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with

Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (216 Q&As Dumps, **30%OFF** Special Discount: **freecram**)

NEW QUESTION: 107

A utility company wants to collect usage data every 5 minutes from its smart meters to facilitate time-of-use metering. When a meter sends data to AWS, the data is sent to Amazon API Gateway, processed by an AWS Lambda function, and stored in an Amazon DynamoDB table. During the pilot phase, the Lambda functions took from 3 to 5 seconds to complete. As more smart meters are deployed, the Engineers notice the Lambda functions are taking from 1 to 2 minutes to complete. The functions are also increasing in duration as new types of metrics are collected from the devices. There are many `ProvisionedThroughputExceededException` errors while performing `PUT` operations on DynamoDB, and there are also many `TooManyRequestsException` errors from Lambda.

Which combination of changes will resolve these issues? (Select TWO)

- A. Increase the memory available to the Lambda functions
- B. Stream the data into an Amazon Kinesis data stream from API Gateway and process the data in batches
- C. Increase the write capacity units to the DynamoDB table
- D. Collect data in an Amazon SQS FIFO queue, which triggers a Lambda function to process each message
- E. Increase the payload size from the smart meters to send more data

Answer: B,C ([LEAVE A REPLY](#))

NEW QUESTION: 108

A company plans to move regulated and security-sensitive businesses to AWS. The Security team is developing a framework to validate the adoption of AWS best practice and industry-recognized compliance standards. The AWS Management Console is the preferred method for teams to provision resources.

Which strategies should a Solutions Architect use to meet the business requirements and continuously assess, audit, and monitor the configurations of AWS resources? (Choose two.)

- A. Use AWS Config rules to periodically audit changes to AWS resources and monitor the compliance of the configuration. Develop AWS Config custom rules using AWS Lambda to establish a test-driven development approach, and further automate the evaluation of configuration changes against the required controls.
- B. Use Amazon CloudWatch Logs agent to collect all the AWS SDK logs. Search the log data using a pre-defined set of filter patterns that machines mutating API calls. Send notifications using Amazon CloudWatch alarms when unintended changes are performed. Archive log data by using a batch export to Amazon S3 and then Amazon Glacier for a long-term retention and auditability.
- C. Use AWS CloudTrail events to assess management activities of all AWS accounts. Ensure that CloudTrail is enabled in all accounts and available AWS services. Enable trails, encrypt

CloudTrail event log files with an AWS KMS key, and monitor recorded activities with CloudWatch Logs.

D. Use the Amazon CloudWatch Events near-real-time capabilities to monitor system events patterns, and trigger AWS Lambda functions to automatically revert non-authorized changes in AWS resources. Also, target Amazon SNS topics to enable notifications and improve the response time of incident responses.

E. Use CloudTrail integration with Amazon SNS to automatically notify unauthorized API activities.

Ensure that CloudTrail is enabled in all accounts and available AWS services. Evaluate the usage of Lambda functions to automatically revert non-authorized changes in AWS resources.

Answer: ([SHOW ANSWER](#))

Explanation

<https://docs.aws.amazon.com/awsccloudtrail/latest/userguide/cloudwatch-alarms-for-cloudtrail.html>

https://docs.aws.amazon.com/en_pv/awsccloudtrail/latest/userguide/best-practices-security.html

The AWS Config console shows the compliance status of your rules and resources. You can see how your AWS resources comply overall with your desired configurations, and learn which specific resources are noncompliant. You can also use the AWS CLI, the AWS Config API, and AWS SDKs to make requests to the AWS Config service for compliance information.

<https://docs.aws.amazon.com/config/latest/developerguide/evaluate-config.html>

NEW QUESTION: 109

A company maintains a restaurant review website. The website is a single-page application where files are stored in Amazon S3 and delivered using Amazon CloudFront. The company receives several fake postings every day that are manually removed. The security team has identified that most of the fake posts are from bots with IP addresses that have a bad reputation within the same global region. The team needs to create a solution to help restrict the bots from accessing the website. Which strategy should a solutions architect use?

A. Use AWS Firewall Manager to control the CloudFront distribution security settings. Select the managed Amazon IP reputation rule group and associate it with Firewall Manager with a deny action.

B. Use AWS Firewall Manager to control the CloudFront distribution security settings. Create a geographical block rule and associate it with Firewall Manager.

C. Associate an AWS WAF web ACL with the CloudFront distribution. Select the managed Amazon IP reputation rule group for the web ACL with a deny action.

D. Associate an AWS WAF web ACL with the CloudFront distribution. Create a rule group for the web ACL with a geographical match statement with a deny action.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 110

A company runs a memory-intensive analytics application using on-demand Amazon EC2 compute optimized instances. The application is used continuously and application demand

doubles during working hours. The application currently scales based on CPU usage. When scaling in occurs, a lifecycle hook is used because the instance requires 4 minutes to clean the application state before terminating.

Because users reported poor performance during working hours, scheduled scaling actions were implemented so additional instances would be added during working hours. The Solutions Architect has been asked to reduce the cost of the application.

Which solution is MOST cost-effective?

A. Use the existing launch configuration that uses C5 instances, and update the application AMI to include the Amazon CloudWatch agent. Change the Auto Scaling policies to scale based on memory utilization.

Use Reserved Instances for the number of instances required after working hours, and use Spot Instances to cover the increased demand during working hours.

B. Update the existing launch configuration to use R5 instances, and update the application AMI to include SSM Agent. Change the Auto Scaling policies to scale based on memory utilization.

Use Reserved instances for the number of instances required after working hours, and use Spot Instances with on-Demand instances to cover the increased demand during working hours.

C. Use the existing launch configuration that uses C5 instances, and update the application AMI to include SSM Agent. Leave the Auto Scaling policies to scale based on CPU utilization. Use scheduled Reserved Instances for the number of instances required after working hours, and use Spot Instances to cover the increased demand during work hours.

D. Create a new launch configuration using R5 instances, and update the application AMI to include the Amazon CloudWatch agent. Change the Auto Scaling policies to scale based on memory utilization. use Reserved Instances for the number of instances required after working hours, and use Standard Reserved Instances with On-Demand Instances to cover the increased demand during working hours.

Answer: D (LEAVE A REPLY)

Explanation

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/monitoring_ec2.html

NEW QUESTION: 111

During an audit a Security team discovered that a Development team was putting IAM user secret access keys in their code and then committing it to an AWS CodeCommit repository The Security team wants to automatically find and remediate instances of this security vulnerability Which solution will ensure that the credentials are appropriately secured automatically?

A. Configure Amazon Macie to scan for credentials in CodeCommit repositories If credentials are found, trigger an AWS Lambda function to disable the credentials and notify the user

B. Use a scheduled AWS Lambda function to download and scan the application code from CodeCommit If credentials are found generate new credentials and store them in AWS KMS

C. Configure a CodeCommit trigger to invoke an AWS Lambda function to scan new code submissions for credentials If credentials are found, disable them in AWS IAM and notify the user

D. Run a script rightly using AWS Systems Manager Run Command to search (or credentials on the development instances It found, use AWS Secrets Manager to rotate the credentials

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 112

A company uses AWS Organizations to manage one parent account and nine member accounts. The number of member accounts is expected to grow as the business grows. A security engineer has requested consolidation of AWS CloudTrail logs into the parent account for compliance purposes. Existing logs currently stored in Amazon S3 buckets in each individual member account should not be lost. Future member accounts should comply with the logging strategy. Which operationally efficient solution meets these requirements?

A. Configure CloudTrail in each member account to deliver log events to a central S3 bucket. Ensure the central bucket policy allows Put Object access from the member accounts. Migrate existing logs to the central S3 bucket. Set up an Amazon CloudWatch alarm to alert if CloudTrail is not configured properly.

B. Configure an organization-level CloudTrail in the parent account to deliver log events to a central S3 bucket. Configure CloudTrail in each member account to deliver log events to the central S3 bucket.

C. Create an AWS Lambda function in each member account with a cross-account role. Trigger the Lambda functions when new CloudTrail logs are created and copy the CloudTrail logs to a centralized S3 bucket. Set up an Amazon CloudWatch alarm to alert if CloudTrail is not configured properly.

D. Configure an organization-level CloudTrail in the parent account to deliver log events to a central S3 bucket. Migrate the existing CloudTrail logs from each member account to the central S3 bucket. Delete the existing CloudTrail and logs in the member accounts.

Answer: **A** ([LEAVE A REPLY](#))

NEW QUESTION: 113

A company's CISO has asked a Solutions Architect to re-engineer the company's current CI/CD practices to make sure patch deployments to its applications can happen as quickly as possible with minimal downtime if vulnerabilities are discovered. The company must also be able to quickly roll back a change in case of errors.

The web application is deployed in a fleet of Amazon EC2 instances behind an Application Load Balancer.

The company is currently using GitHub to host the application source code and has configured an AWS CodeBuild project to build the application. The company also intends to use AWS CodePipeline to trigger builds from GitHub commits using the existing CodeBuild project.

What CI/CD configuration meets all of the requirements?

A. Configure CodePipeline with a deploy stage using AWS CodeDeploy configured for in-place deployment. Monitor the newly deployed code, and if there are any issues, push another code update.

B. Configure the CodePipeline with a deploy stage using AWS OpsWorks and in-place deployments.

Monitor the newly deployed code and if there are any issues, push another code update.

C. Configure CodePipeline with a deploy stage using AWS CloudFormation to create a pipeline for test and production stacks. Monitor the newly deployed code and if there are any issues push another code update.

D. Configure CodePipeline with a deploy stage using AWS CodeDeploy configure for blue/green deployments. Monitor the new deployed code and if there are any issues, trigger a manual rollback using CodeDeploy.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 114

A Solutions Architect is working with a company that is extremely sensitive to its IT costs and wishes to implement controls that will result in a predictable AWS spend each month.

Which combination of steps can help the company control and monitor its monthly AWS usage to achieve a cost that is as close as possible to the target amount? (Choose three.)

A. Place conditions in the users' IAM policies that limit the number of instances they are able to launch.

B. Contact AWS Support and ask that they apply limits to the account so that users are not able to launch more than a certain number of instance types.

C. Define 'workload' as a cost allocation tag in the AWS Billing and Cost Management console.

D. Implement an IAM policy that requires users to specify a 'workload' tag for cost allocation when launching Amazon EC2 instances.

E. Purchase all upfront Reserved Instances that cover 100% of the account's expected Amazon EC2 usage.

F. Set up AWS Budgets to alert and notify when a given workload is expected to exceed a defined cost.

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 115

A weather service provides high-resolution weather maps from a web application hosted on AWS in the eu-west-1 Region. The weather maps are updated frequently and stored in Amazon S3 along with static HTML content. The web application is fronted by Amazon CloudFront.

The company recently expanded to serve users in the us-east Region and these new users report that viewing their respective weather maps is slow from time to time Which combination of steps will resolve the us-east performance issues? (Select TWO)

A. Use Lambda@Edge to modify requests from North America to use the S3 Transfer Acceleration endpoint in us-east-1

B. Use Lambda@Edge to modify requests from North America to use the S3 bucket in us-east-1

C. Configure the AWS Global Accelerator endpoint for the S3 bucket in eu-west-1 Configure endpoint groups for TCP ports 80 and 443 in us-east-1

D. Create a new S3 bucket in us-east-1. Configure S3 cross-Region replication to synchronize from the S3 bucket in eu-west-1.

E. Configure the AWS Global Accelerator endpoint for us-east-1 as an origin on the CloudFront distribution. Use `lambda@Edge` to modify requests from North America to use the new origin.

Answer: D,E ([LEAVE A REPLY](#))

NEW QUESTION: 116

A company is running a two-tier web-based application in an on-premises data center. The application user consists of a single server running a stateful application. The application connects to a PostgreSQL database running on a separate server. The application's user base is expected to grow significantly, so the company is migrating the application and database to AWS. The solution will use Amazon Aurora PostgreSQL, Amazon EC2 Auto Scaling, and Elastic Load Balancing.

Which solution will provide a consistent user experience that will allow the application and database tiers to scale?

A. Enable Aurora Auto Scaling for Aurora writers. Use an Application Load Balancer with the round robin routing algorithm and sticky sessions enabled.

B. Enable Aurora Scaling for Aurora writers. Use a Network Load Balancer with the least outstanding requests routing algorithm and sticky sessions enabled.

C. Enable Aurora Auto Scaling for Aurora Replicas. Use a Network Load Balancer with the outstanding requests routing algorithm and sticky sessions enabled.

D. Enable Aurora Auto Scaling for Aurora Replicas. Use an Application Load Balancer with the round robin routing and sticky sessions enabled.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 117

A company uses AWS Organizations to manage more than 1,000 AWS accounts. The company has created a new developer organization. There are 540 developer members in that must be moved to the new developer organization. All accounts are set up with all the required information so that each account can be operated as a standalone account.

Which combination of steps should a solutions architect take to move all of the developer accounts to the new developer organization? (Select THREE)

A. Call the `InviteAccountToOrganization` operation in the Organizations API from the new developer organization's master account to send invitations to the developer accounts.

B. Have each developer sign in to their account and confirm to join the new developer organization.

C. Sign in to the new developer organization's master account and create a placeholder member account that acts as a target for the developer account migration.

D. From the master account, remove each developer account from the old organization using the `RemoveAccountFromOrganization` operation in the Organizations API.

E. From each developer account, remove the account from the old organization using the RemoveAccountFromOrganization operation In the Organization API.

F. Call the MoveAccount operation in the Organizations API from the old organization's master account to migrate the developer accounts to the new developer organization

Answer: A,C,D ([LEAVE A REPLY](#))

NEW QUESTION: 118

The following AWS Identity and Access Management (IAM) customer managed policy has been attached to an IAM user:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "s3:*",
      "Resource": [
        "arn:aws:s3:::prod-data",
        "arn:aws:s3:::prod-data/*"
      ]
    },
    {
      "Effect": "Deny",
      "NotAction": "s3:*",
      "NotResource": [
        "arn:aws:s3:::prod-data",
        "arn:aws:s3:::prod-data/*"
      ]
    }
  ]
}
```

which statement describes the access that this policy provides to the user?

A. This policy denies access to all Amazon S3 actions, excluding all actions in the prod-data S3 bucket.

B. This policy grants access to all Amazon S3 actions including all actions in the prod-data S3 bucket.

C. This policy denies access to the Amazon S3 bucket and objects not having prod-data in the bucket name.

Answer: B ([LEAVE A REPLY](#))

NEW QUESTION: 119

A company wants to replace its call system with a solution built using AWS managed services. The company call center would like the solution to receive calls, create contact flows, and scale to handle growth projections. The call center would also like the solution to use deep learning capabilities to recognize the intent of the callers and handle basic tasks, reducing the need to speak an agent. The solution should also be able to query business applications and provide relevant information back to calls as requested.

Which services should the Solution Architect use to build this solution? (Choose three.)

- A. Amazon Connect to create a cloud-based contact center.
- B. Amazon Rekognition to identity who is calling.
- C. Amazon Alexa for Business to build conversational interface.
- D. Amazon SQS to add incoming callers to a queue.
- E. Amazon Lex to recognize the intent of the caller.
- F. AWS Lambda to integrate with internal systems.

Answer: A,E,F ([LEAVE A REPLY](#))

NEW QUESTION: 120

A solutions architect is designing a solution that consists of a fleet of Amazon EC2 Reserved Instances (RIs) in an Auto Scaling group that will grow over time as usage increases. The solution needs to maintain 80% RI coverage to maintain cost control with an alert to the DevOps team using an email distribution list when coverage drops below 30% The solution must also include the ability to generate a report to easily track and manage coverage. The company has a policy that allows only one workload for each AWS account Which set of steps should the solutions architect take to create the report and alert the DevOps team?

- A. Create an Amazon Simple Notification Service (Amazon SNS) topic and subscribe the DevOps email distribution list Use the Cost Explorer console to configure the report for RI utilization set the utilization target to 30% and link to the SNS topic created in the alert configuration
- B. Enable cost allocation tags and ensure instances populate a customer-managed cost allocation tag at startup Use the Cost Explorer console to configure the report for RI coverage, filter using the customer-managed cost allocation tag and set the threshold to 80% and enter the email distribution list in the alert configuration.
- C. Create an Amazon Simple Notification Service (Amazon SNS) topic and subscribe the DevOps email distribution list Enable cost allocation tags and ensure instances populate a customer-managed cost allocation tag at startup Use the AWS Billing and Cost Management console to create a budget for RI coverage, filter using the customer-managed cost allocation tag and set the threshold to 80% and link to the SNS topic created in the alert configuration
- D. Use the AWS Billing and Cost Management console to create a reservation budget for RI utilization set the utilization to 80% and enter the email distribution list in the alert configuration

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 121

A company built an ecommerce website on AWS using a three-tier web architecture. The application is Java-based and composed of an Amazon CloudFront distribution, an Apache web server layer of Amazon EC2 instances in an Auto Scaling group, and a backend Amazon Aurora MySQL database. Last month, during a promotional sales event, users reported errors and timeouts while adding items to their shopping carts. The operations team recovered the logs created by the web servers and reviewed Aurora DB cluster performance metrics. Some of the web servers were terminated before logs could be collected and the Aurora metrics were not

sufficient for query performance analysis. Which combination of steps must the solutions architect take to improve application performance visibility during peak traffic events? (Select THREE.)

- A.** Configure the Aurora MySQL DB cluster to publish slow query and error logs to Amazon CloudWatch Logs.
- B.** Install and configure an Amazon CloudWatch Logs agent on the EC2 instances to send the Apache logs to CloudWatch Logs
- C.** Enable and configure AWS CloudTrail to collect and analyze application activity from Amazon EC2 and Aurora.
- D.** Configure the Aurora MySQL DB cluster to stream slow query and error logs to Amazon Kinesis.
- E.** Enable Aurora MySQL DB cluster performance benchmarking and publish the stream to AWS X-Ray.
- F.** Implement the AWS X-Ray SDK to trace incoming HTTP requests on the EC2 instances and implement tracing of SQL queries with the X-Ray SDK for Java.

Answer: A,B,F (LEAVE A REPLY)

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (216 Q&As Dumps, **30%OFF** Special Discount: **freecram**)

NEW QUESTION: 122

A Solutions Architect needs to design a highly available application that will allow authenticated users to stay connected to the application even when there are underlying failures.

Which solution will meet these requirements?

- A.** Deploy the application on Amazon EC2 instances in an Auto Scaling group Use an internet-facing Application Load Balancer on the front end Use EC2 instances to save the authenticated connection details
- B.** Deploy the application on Amazon EC2 instances Use Amazon Route 53 to forward requests to the EC2 Instances. Use Amazon DynamoDB to save the authenticated connection details.
- C.** Deploy the application on Amazon EC2 instances in an Auto Scaling group Use an internet-facing Application Load Balancer to handle requests Use Amazon DynamoDB to save the authenticated connection details
- D.** Deploy the application on Amazon EC2 instances in an Auto Scaling group Use an internet-facing Application Load Balancer on the front end Use EC2 instances hosting a MySQL database to save the authenticated connection details

Answer: ([SHOW ANSWER](#))

NEW QUESTION: 123

A company has an application that sells tickets online and experiences bursts of demand even/ 7 days. The application has a stateless presentation layer running on Amazon EC2 an Oracle database to store unstructured data catalog information and a backend API layer. The front-end layer uses an Elastic Load Balancer to distribute the load across nine On-Demand instances over three Availability Zones (AZs). The Oracle database is running on a single EC2 instance The company is experiencing performance issues when running more than two concurrent campaigns. A solutions architect must design a solution that meets the following requirements

- * Address scalability issues
- * Increase the level of concurrency
- * Eliminate licensing costs
- * improve reliability

Which set of steps should the solutions architect take?

- A.** Create an Auto Scaling group for the front end with a combination of On-Demand and Spot Instances to reduce costs Create two additional copies of the database instance, then distribute the databases in separate AZs
- B.** Create an Auto Scaling group for the front end with a combination of On-Demand and Spot instances to reduce costs Convert the Oracle database into a single Amazon RDS reserved DB instance
- C.** Create an Auto Scaling group foe the front end with a combination of On-Demand and Spot instances to reduce costs Convert the tables in the Oracle database into Amazon DynamoDB tables
- D.** Convert the On-Demand instances into Spot instances to reduce costs for the front end Convert the tables in the Oracle database into Amazon DynamoDB tables.

Answer: C ([LEAVE A REPLY](#))

NEW QUESTION: 124

A finance company is running its business-critical application on current-generation Linux EC2 instances. The application includes a self-managed MySQL database performing heavy I/O operations. The application is working fine to handle a moderate amount of traffic during the month. However, it slows down during the final three days of each month due to month-end reporting, even though the company is using Elastic Load Balancers and Auto Scaling within its infrastructure to meet the increased demand.

Which of the following actions would allow the database to handle the month-end load with the LEAST impact on performance?

- A.** Using Amazon CloudWatch with AWS Lambda to change the type, size, or IOPS of Amazon EBS volumes in the cluster based on a specific CloudWatch metric.

- B.** Replacing all existing Amazon EBS volumes with new PIOPS volumes that have the maximum available storage size and I/O per second by taking snapshots before the end of the month and reverting back afterwards.
- C.** Pre-warming Elastic Load Balancers, using a bigger instance type, changing all Amazon EBS volumes to GP2 volumes.
- D.** Performing a one-time migration of the database cluster to Amazon RDS, and creating several additional read replicas to handle the load during end of month.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 125

A company wants to replace its call system with a solution built using AWS managed services. The company call center would like the solution to receive calls, create contact flows, and scale to handle growth projections. The call center would also like the solution to use deep learning capabilities to recognize the intent of the callers and handle basic tasks, reducing the need to speak an agent. The solution should also be able to query business applications and provide relevant information back to calls as requested.

Which services should the Solution Architect use to build this solution? (Choose three.)

- A.** Amazon Lex to recognize the intent of the caller.
- B.** Amazon Rekognition to identify who is calling.
- C.** Amazon Alexa for Business to build conversational interface.
- D.** Amazon SQS to add incoming callers to a queue.
- E.** AWS Lambda to integrate with internal systems.
- F.** Amazon Connect to create a cloud-based contact center.

Answer: ([SHOW ANSWER](#)**)**

NEW QUESTION: 126

An ecommerce company has an order processing application it wants to migrate to AWS. The application has inconsistent data volume patterns, but needs to be available at all times. Orders must be processed as they occur and in the order that they are received.

Which set of steps should a solutions architect take to meet these requirements?

- A.** Use Amazon SNS with FIFO and send orders as they occur. Use a single large Reserved Instance for processing.
- B.** Use Amazon SQS with FIFO and send orders as they occur. Use Spot Instances in multiple Availability Zones for processing.
- C.** Use AWS Transfer for SFTP and upload orders as they occur. Use On-Demand Instances in multiple Availability Zones for processing.
- D.** Use Amazon SQS with FIFO and send orders as they occur. Use Reserved Instances in multiple Availability Zones for processing.

Answer: D ([LEAVE A REPLY](#))

NEW QUESTION: 127

A photo-sharing and publishing company receives 10,000 to 150,000 images daily. The company receives the images from multiple suppliers and users registered with the service. The company is moving to AWS and wants to enrich the existing metadata by adding data using Amazon Rekognition.

The following is an example of the additional data:

`list celebrities [name of the personality, wearing (color), looking (happy, sad), near (location example Eiffel Tower in Paris)]`

As part of the cloud migration program, the company uploaded existing image data to Amazon S3 and told users to upload images directly to Amazon S3.

What should the Solutions Architect do to support these requirements?

- A.** Trigger AWS Lambda based on an S3 event notification to create additional metadata using Amazon Rekognition. Use Amazon DynamoDB to store the metadata and Amazon ES to create an index. Use a web front-end to provide search capabilities backed by Amazon ES.
- B.** Use Amazon Kinesis to stream data based on an S3 event. Use an application running in Amazon EC2 to extract metadata from the images. Then store the data on Amazon DynamoDB and Amazon CloudSearch and create an index. Use a web front-end with search capabilities backed by CloudSearch.
- C.** Start an Amazon SQS queue based on S3 event notifications. Then have Amazon SQS send the metadata information to Amazon DynamoDB. An application running on Amazon EC2 extracts data from Amazon Rekognition using the API and adds data to DynamoDB and Amazon ES. Use a web front-end to provide search capabilities backed by Amazon ES.
- D.** Trigger AWS Lambda based on an S3 event notification to create additional metadata using Amazon Rekognition. Use Amazon RDS MySQL Multi-AZ to store the metadata information and use Lambda to create an index. Use a web front-end with search capabilities backed by Lambda.

Answer: ([SHOW ANSWER](#))

Explanation

<https://github.com/aws-samples/lambda-refarch-imagerecognition>

Valid AWS-Solutions-Architect-Professional Dumps shared by Fast2test.com for Helping Passing AWS-Solutions-Architect-Professional Exam! Fast2test.com now offer the **newest AWS-Solutions-Architect-Professional exam dumps**, the Fast2test.com AWS-Solutions-Architect-Professional exam **questions have been updated** and **answers have been corrected** get the **newest** Fast2test.com AWS-Solutions-Architect-Professional dumps with Test Engine here: <https://www.fast2test.com/AWS-Solutions-Architect-Professional-premium-file.html> (216 Q&As Dumps, **30%OFF** Special Discount: **freecram**)