

KONRAD CŁAPA | BRIAN GERRARD

PROFESSIONAL CLOUD ARCHITECT GOOGLE CLOUD CERTIFICATION GUIDE

Second Edition

Build a solid foundation in Google Cloud Platform
to achieve the most lucrative IT certification

Foreword by Yujun Liang,
Fully Certified on GCP, Cloud Certification King on LinkedIn



Professional Cloud Architect – Google Cloud Certification Guide

Second Edition

Build a solid foundation in Google Cloud Platform
to achieve the most lucrative IT certification

Konrad Cłapa

Brian Gerrard



BIRMINGHAM—MUMBAI

Professional Cloud Architect – Google Cloud Certification Guide

Second Edition

Copyright © 2022 Packt Publishing

All rights reserved. No part of this book may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, without the prior written permission of the publisher, except in the case of brief quotations embedded in critical articles or reviews.

Every effort has been made in the preparation of this book to ensure the accuracy of the information presented. However, the information contained in this book is sold without warranty, either express or implied. Neither the authors, nor Packt Publishing or its dealers and distributors, will be held liable for any damages caused or alleged to have been caused directly or indirectly by this book.

Packt Publishing has endeavored to provide trademark information about all of the companies and products mentioned in this book by the appropriate use of capitals. However, Packt Publishing cannot guarantee the accuracy of this information.

Group Product Manager: Rahul Nair

Publishing Product Manager: Rahul Nair

Senior Editor: Sangeeta Purkayastha

Content Development Editor: Nihar Kapadia

Technical Editor: Nithik Cheruvakodan

Copy Editor: Safis Editing

Project Coordinator: Shagun Saini

Proofreader: Safis Editing

Indexer: Hemangini Bari

Production Designer: Aparna Bhagat

Marketing Coordinator: Nimisha Dua and Sanjana Gupta

First published: October 2019

Second edition: January 2022

Production reference: 1101221

Published by Packt Publishing Ltd.

Livery Place

35 Livery Street

Birmingham

B3 2PB, UK.

ISBN 978-1-80181-229-0

www.packt.com

To my wife and parents, for their constant support.

– Konrad Clapa

To Linda, for her constant support and understanding.

– Brian Gerrard

Foreword

I came to know Konrad Cłapa on LinkedIn. We both have an interest in Google Cloud and have both become fully certified on Google Cloud. Then we met at Google Cloud Next '19, where Konrad invited me to co-author a book to teach people how to use Google Cloud – this book.

I had had a failed writing experience before and I knew it takes a lot of effort to write a book, so I didn't agree to co-author it, but I did write the foreword for the first edition of the book, published in 2019.

Those 2 years passed quickly and so many new services have been released on Google Cloud since then. Google Cloud Training and Certification released an upgraded version of the exam for Google Cloud Certified Professional Cloud Architect in March 2021. Even though my certification won't expire until September 2021, I took the beta exam and passed.

As expected, a lot of new content has been added to the exam, and new services such as Google Kubernetes Engine and Anthos are tested in the new exam. These changes made the first edition of the book irrelevant.

I am glad Konrad and Brian have spent time refreshing their book to reflect the changes in the exam syllabus. This will be the first book ever published on the topic since the new exam was released and it will provide much-needed guidance for people to learn about the most up-to-date services from Google in order to pass the exam and advance their careers.

A better job, a better life. I achieved so much more after I became certified on a number of cloud vendors. I encourage people to learn the cloud and plan for a better career. This book gives you what you need to pass one of the most sought-after certifications, and also the highest-paid certification, according to some sources online.

We didn't get a chance to meet again after Google Cloud Next '19. Google Cloud Next '20 and '21 were online only due to COVID-19. But we kept in touch on LinkedIn and I am aware that Konrad has become a Google Cloud Certified Fellow, the most prestigious certification from Google, and one of my own career goals. I am happy about his achievement.

Even though I don't know Brian Gerrard in person, since Konrad and Brian are colleagues and they co-authored the first edition of the book, we connected on LinkedIn. Brian has decade-long experience on VMware and he claimed to have many VMware-related certifications. Brian just moved to the public cloud in recent years. His transition from data center-centric technology to the public cloud is an inspiration for people facing the same dilemma – stop doubting the public cloud; rather, embrace it to further your career advancement. There are many similar success stories on LinkedIn for people to use as a reference and as inspiration for their own career moves. LinkedIn is a professional social media platform where you can connect with people with a common interest, learn from each other, and grow together to reach your own professional goals.

I feel honored to be given this opportunity again to write the foreword to this second edition. And I understand why Konrad wanted me to write it. For the past 5 years, I have inspired thousands of people to learn the cloud and find their purpose. I am the right person to introduce this book to you, which will help you learn Google Cloud and become certified. And great things will happen after becoming certified in Google Cloud.

Yujun Liang

Fully Certified on GCP

Cloud Certification King on LinkedIn

Contributors

About the authors

Konrad Clapa is a Certified Google Trainer and Lead Cloud Architect working for Atos R&D. He has over 12 years' experience in the IT industry. He holds over 40 IT certification. This includes all 12/12 active Google Cloud Platform certifications making him the first in the world to achieve it. He is also listed among 40 individuals who hold Google Cloud Certified Fellow certificate. Sharing knowledge has always been important to him, so he contributes to the community by acting as a leader for a local Google Cloud Developer group. Funny fact about Konrad is that till now he never failed a GCP exam attempt including multiple beta exams that he has taken.

First of all, I would like to thank my lovely wife Kinga for the patience and understanding of my passion. It is difficult to be a wife of a geek. I would also like to thank my parents that always supported me in any decisions I have taken in my life. Last but not the least I thank Brian for being a fellow for multiple journeys and above all being a great friend!

Brian Gerrard is currently a DevOps Engineer from Scotland with over 15 years' experience in the IT industry. As well as holding the Google Certified Professional Architect and Google Certified Associate Engineer certifications, he is certified in Azure, AWS, Terraform and VMware. Brian is a firm believer in lifelong learning, and you will regularly find him contributing to his local user groups.

I would like to thank my wife and children for their support. Also, thanks to Konrad for constantly pushing the Google community to achieve more!

About the reviewers

Dr. Artem Nikulchenko is a chief software architect at Teamwork Commerce, a Google Cloud Platform Partner that specializes in the development of retail management systems. Artem spearheaded the usage of Google Cloud Platform in teamwork products from the very first publicly available version of the platform.

Artem has a Ph.D. in computer science and holds an associate professor position at National Technical University Kharkiv Polytechnic Institute. He teaches courses on databases, software architecture, and public clouds.

As a big fan of Google Cloud, Artem leads his local Google Developer Group and delivers talks about Google Cloud Platform at specialized conferences around the world.

Having worked with GCP for over 10 years, I still found this book very useful while preparing for my Google Cloud Architect certification exam.

Then, I was invited to become a reviewer of that same book's newest edition. It's a great honor for me and I'd like to thank the authors, as well as Packt Publishing, for this exceptional opportunity.

Artur Zejfer has been working in the IT industry for more than 20 years, ranging across application support, database administration, automation, and architect roles. This has given him a broad base of skills and the ability to work with a diverse range of clients. He holds a master's degree in computer science and several industry-recognized certifications from Oracle, Google, The Open Group, and so on. He currently works for Atos R&D as a database architect focusing on automating daily operations. Artur has a passion for learning and information that he continues to pursue daily.

Above all the previously stated, he is a father, husband, son, brother, and friend.

I'd like to thank my wife, Kamila, and our son, Franek, for their daily support and patience. To my parents, siblings, relatives, friends, and mentors (you know who you are), thank you for guiding and supporting me.

Shouvik Basak is a solution architect at a leading Fortune 500 global IT services organization. He specializes in architecting cloud-based solutions and data center infrastructure technologies. In a career spanning more than 20 years in technology, he has worked extensively on public, private, and hybrid cloud delivery, traditional infrastructure, and operations, primarily for global enterprise customers. He has led multiple programs in the technology, service management, and integration areas and has worked hands-on in multiple technical roles.

To my mentors and friends, who have kept me motivated and curious and helped me navigate my journey and find joy in learning and doing what I do. I would also like to thank Packt Publishing for the opportunity to review this wonderful book.

Vijaykumar Jangamashetti (VJ) is a cloud solutions architect at Rackspace Technology, with a specialization in Google Cloud Platform, and is 7X Google Cloud Platform-certified. VJ is passionate about data and analytics and a strong customer advocate of educating on and promoting the power of data. As a Google Cloud Platform expert, VJ has led multiple efforts in strengthening alliance partnerships with Google Cloud through the achievement of MSP credentials, specialization awards, and expertise badges for the respective organizations he has worked in. With 15+ years of IT experience working with customers across the globe with multiple positions as an ETL architect, enterprise data architect, and cloud architect, he has built a strong reputation as a customer champion and a trusted advisor. He goes by the name @vijaykumarpj on LinkedIn and Twitter.

I'd like to thank my wife, Jyoti, and our daughter, Shriya, for their continued support and encouragement with everything that I do with great patience. I genuinely appreciate what you both have done for me and I love you both. I'd also like to thank my grandparents, parents, siblings, relatives, friends, colleagues, and mentors for helping me to be who I am today, and I greatly appreciate you all. Thanks to Packt for providing me with the opportunity to review this wonderful book.

Jarosław Gajewski holds the lead architect and distinguished expert positions in Atos. He is responsible for designing multi- and hybrid cloud solutions for cloud-agnostic and cloud-native services.

His technical knowledge and experience are backed by multiple industry-standard certificates. He is already part of an elite group of Google Cloud Certified Fellows. He holds the Google Cloud Certified Professional and other Google, VMware, Dell, Microsoft, and AWS certifications.

As someone who is passionate about the cloud, outside of his work, he is also an active community speaker and one of the Google Developer group leads for GDG Cloud Bydgoszcz.

In his free time, he loves to enjoy time with his wife, two daughters, and son, play board games, and constantly increase his knowledge.

I'd like to thank my wife Weronika for her patience and constant support.

Maciej Stopa is a multi-certified cloud solution architect experienced in Azure, Google Cloud Platform, AWS, Red Hat, IBM, Oracle, and VMware. With more than 17 years of experience in IT, he has gained deep on-premises, HPC, and distributed systems experience while working for a leading systems integrator in Poland. He transformed through virtualization and hybrid cloud, then architected a custom-tailored public cloud offering for the Polish market. After moving to Microsoft, he enjoyed working in a corporate environment, spending a couple of years as a cloud solution architect for enterprise customers. He now supports a multi-national organization called Digital and Cloud Services in Deutsche Telekom as a lead cloud solution architect, also leading an internal community of practice called Cloud Architect Ninjas.

I'd like to thank my beloved wife, Beata, and our children, Oliwia and Oskar, for their daily support, patience, and encouragement with everything that I do. I'd like to thank my parents, for giving me everything I needed to become who I am, my Mom, for her endless patience, and my Dad, for starting his IT journey living in a small Polish city by finishing offline snail-mailed courses almost 35 years ago! A spark that ignited my interests.

Dr. Nabil Hadj-Ahmed is a Google Cloud authorized trainer, a cloud architect, and a Google Cloud developer expert with over 20 years of experience working in real-world, challenging, enterprise, and demanding environments. Nabil facilitates Google Cloud training and brings a wealth of experience to enhance the learning experience of Google Cloud.

Nabil holds a Ph.D. and an MSc from Leeds Beckett University, where he also taught during his doctoral research. He also manages and speaks at Google Developer groups, organizing monthly meetup events and workshops focusing on Google technologies.

Pawel Piwowarek is a highly technical IT professional with a master's degree in computer science and over 18 years of industry experience designing, implementing, and supporting a wide range of IT solutions. He is an experienced technical leader and communicator who works with executives, managers, suppliers, and technical peers on large-scale projects with critical business impact. He maintains a keen interest in emerging and disruptive technologies and continuously adopts skillsets backed by relevant vendors' accreditations to match evolving business needs.

Table of Contents

Preface

Section 1: Introduction to GCP

1

GCP Cloud Architect Professional

Understanding the benefits of being a certified architect	4	Suggested hints and tips	11
Registering for the exam	5	Summary	12
What to expect from the exam	10	Further reading	12
Online proctored exams	11		

2

Getting Started with Google Cloud Platform

Introducing the cloud	14	Resource manager	20
Understanding GCP	15	Summary	41
GCP differentiators	17	Further reading	41
GCP locations	18		

3

Google Cloud Platform Core Services

Understanding computing and hosting services	44	Getting to know about networking services	49
Exploring storage services	47	Going through big data services	51

Understanding machine learning services	52	Summary	54
Learning about identity services	53	Further reading	54

Section 2: Manage, Design, and Plan a Cloud Solution Architecture

4

Working with Google Compute Engine

Deploying our first GCE instance	58	Instance templates and instance groups	81
Deployment options	64	Setting the location	84
Region	64	Port name mapping	85
Zone	65	Autoscaling	85
Boot disk	66	Autohealing	92
Snapshots	66	Containers on Compute Engine	93
Existing disks	69	Running a simple Windows container on Compute Engine	94
Labels	70	A quick look at Google Cloud VMware Engine	96
Confidential VM service	71	Use cases	96
Container	71	Quotas and limits	98
Deletion protection	71	IAM roles	98
Reservations	71	Pricing	100
Metadata	72	Summary	103
Startup scripts	72	Further reading	103
Preemptibility	73		
Availability policy	74		
Automatic restart	74		
Shielded VM	75		
Deletion rule	75		
Node affinity labels	76		
GPUs and TPUs	79		

5

Exploring Google App Engine as a Compute Option

App Engine components	106	Working with App Engine	108
Choosing the right location	107	Environment types	108

App Engine standard environment	109	Scaling	121
Flexible environment	109	Cron jobs	122
Deploying an App Engine application	110	Memcache	124
Versions	116	IAM	125
Splitting traffic	117	Quotas and limits	125
Migrating traffic	118	Pricing	126
Firewall rules	119	Summary	126
Settings	120	Further reading	126
Custom domain	120		
SSL certificates	121		

6

Managing Kubernetes Clusters with Google Kubernetes Engine

An introduction to microservices	130	GKE cluster management	151
Containers	131	Creating a GKE cluster	151
Kubernetes	132	IAM	179
Kubernetes architecture	133	Kubernetes role-based access control	180
Kubernetes objects	135	Artifact Registry	180
Google Kubernetes Engine	147	Cloud Build	181
GKE Autopilot	148	Best practices for building containers	182
Node pools	148	Quotas and limits	182
Container-Optimized OS (COS)	149	Pricing	183
Storage	149	Summary	183
		Further reading	183

7

Deploying Cloud-Native Workloads with Cloud Run

Using Cloud Run	186	Deploying a new revision	201
Cloud Run containers	188	Mapping custom domains	204
Cloud Run triggers	188	IAM roles	205
Deploying to Cloud Run	190	Quotas and limits	205

Pricing	206	Further reading	207
Summary	207		

8

Managing Cloud-Native Workloads with Anthos

Anthos components	210	Service management with Anthos Service Mesh	227
Anthos clusters	211	Istio	227
Anthos Connect Agent	212	Using Anthos Service Mesh	230
Fleets (formerly environs)	212	Summary of Anthos Service Mesh	232
Anthos cluster deployment options	213	Anthos Binary Authorization	234
Anthos clusters on VMware (GKE on-prem)	214	Migrate for Anthos and GKE	235
Anthos on AWS	217	Cloud Run for Anthos	236
Anthos on Azure	219	Quotas and limits	238
Anthos clusters on bare metal	221	Pricing	239
Anthos Config Management	224	Summary	239
		Further reading	240

9

Running Serverless Functions with Google Cloud Functions

Main Cloud Functions characteristics	242	Triggers	247
Use cases	243	Other considerations	247
Application backends	243	Cloud SQL connectivity	248
Real-time data processing systems	244	Connecting to internal resources in a VPC network	248
Smart applications	244	Environmental variables	248
Runtime environments	245	Cold start	248
Types of Cloud Functions	245	Local emulator	249
HTTP functions	246	Deploying Cloud Functions	249
Background functions	246	Deploying Cloud Functions with the Google Cloud Console	249
Events	247	Deploying functions with the gcloud command	258

IAM	260	Summary	261
Quotas and limits	260	Further reading	261
Pricing	261		

10

Networking Options in GCP

Exploring GCP networking	264	Interconnects	281
Understanding Virtual Private Clouds	265	Peering	282
Connectivity	267	Choosing the right connectivity method	282
Cost	268	Network Connectivity Center	283
VPC Flow Logs	269	DNS	284
Cross-VPC connectivity	269	DNS resolution	284
		Cloud DNS	285
Load balancing	272	DNSSEC	285
Global versus regional load balancing	273	VPC firewall rules	285
External versus internal	274	Default rules	286
Proxy versus load balancer	274	Implied rules	287
Load balancer types	274	Always allowed traffic rules	287
Comparison	276	Always denied rules	288
Choosing the right load balancer	277	User-defined rules	288
NAT	278	Firewall logging	289
NAT gateway	278	Hierarchical firewall policies	289
Cloud NAT	279	Private access options	290
Hybrid connectivity	279	Summary	291
VPN	279	Further reading	292

11

Exploring Storage and Database Options in GCP – Part 1

Choosing the right storage option	294	Storage classes	299
Data consistency	297	Data consistency	301
		Cloud Storage FUSE	302
Understanding Cloud Storage	297	Creating and using a bucket	302
Bucket locations	299	Versioning and life cycle management	306
		Transferring data	312

Cloud Storage Transfer Service	313	Data consistency	329
Understanding IAM	315	Creating and managing Cloud SQL	330
Quotas and limits	315	Read replicas	335
Pricing	316	High availability	338
Understanding Cloud Firestore	316	Backup and recovery	338
Creating and using Cloud Datastore	318	Database migration service	340
Summary of Datastore mode		Instance cloning	340
versus Native mode	323	IAM	340
IAM	323	Quotas and limits	341
Quotas and limits	324	Pricing	341
Pricing	325	Summary	342
Understanding Cloud SQL	325	Further reading	342

12

Exploring Storage and Database Options in GCP – Part 2

Cloud Spanner	346	Pricing	356
Instance configurations	346	Bigtable	356
Node count	347	Bigtable configuration	358
Processing units	347	Data consistency	362
Replication	347	Planning capacity	362
TrueTime	348	Creating a Bigtable instance and table	362
Data consistency	348	IAM	365
Encryption	348	Quotas and limits	365
Point-in-time recovery	349	Pricing	366
Creating a Cloud Spanner instance	349	Summary	366
IAM	354	Further reading	366
Quotas and limits	355		

13

Analyzing Big Data Options

End-to-end big data solution	368	Pub/Sub Lite	374
Cloud Pub/Sub	368	IAM	374
Creating a topic and a subscription	370	Quotas and limits	374
		Pricing	375

Cloud Dataflow	375	Cloud IoT Core	397
IAM	380	IAM	398
Quotas and limits	381	Quotas and limits	398
Pricing	381	Pricing	399
BigQuery	381	Data Fusion	399
BigQuery features	382	Core concepts	399
Using BigQuery	386	IAM	405
IAM	392	Quotas and limits	405
Quotas and limits	393	Pricing	405
Pricing	393	Datastream API	405
Dataproc	393	Additional considerations	406
Architecture	394	Summary	407
IAM	396	Further reading	408
Quotas and limits	397		

14

Putting Machine Learning to Work

An introduction to AI and ML	412	Cloud TPUs	421
The seven steps of ML	413	Pretrained ML models	422
Gathering and preparing the data	414	The Cloud Speech-to-Text API	422
Choosing a model	414	The Cloud Text-to-Speech API	423
Training	415	The Cloud Translation API	423
Evaluation	416	The Cloud Natural Language API	423
Hyperparameter tuning	416	The Cloud Vision API	428
Prediction	416	The Cloud Video Intelligence API	431
Learning models	417	Dialogflow	432
GCP ML options	418	AutoML	433
TensorFlow	418	Summary	435
Cloud Vertex AI	420	Further reading	435
Using Vertex AI	420		

Section 3: Secure, Manage and Monitor a Google Cloud Solution

15

Security and Compliance

Introduction to security	440	Industry regulations	472
Shared responsibility model	442	PCI compliance	472
Cloud Identity	443	Data loss prevention (DLP)	472
Resource Manager	446	Penetration testing in GCP	473
Identity and Access Management (IAM)	448	CI/CD security overview	473
Service accounts	454	Identity-Aware Proxy (IAP)	474
Cloud Storage access management	458	TCP forwarding	475
Organization Policy Service	460	Access Context Manager	475
Firewall rules and load balancers	462	Chronicle	475
Cloud Web Security Scanner	464	Summary	476
Monitoring and logging	466	Additional security services	476
Encryption	467	Security Command Center (SCC)	476
Data encryption keys versus key encryption keys	467	Forseti	477
CMEs versus CSEs	467	Cloud Armor	477
		Secret Manager	479
		Summary	479
		Further reading	479

16

Google Cloud Management Options

Using APIs to manage resources	482	Exploring Cloud Deployment Manager	498
Using Google Cloud Shell	487	Terraform	514
The GCP SDK	489	Cloud Foundation Toolkit	516
gcloud	489	Pricing Calculator	517
gsutil	492	Additional things to consider	519
bq	494	Summary	520
cbt	496	Further reading	521

17

Monitoring Your Infrastructure

Technical requirements	524	Logging agent	547
Introduction to Google Cloud's operations suite	524	Ops Agent	547
Cost	526	Log-based metrics	548
Configuring Google Cloud's operations suite	526	Cloud Audit Logs	549
Cloud Monitoring	529	Activity	550
Groups	530	Retention	551
Dashboards	531	Google Cloud's operations suite for GKE	552
Alerting policies	532	GKE dashboard	553
Uptime checks	534	APM	554
Monitoring agents	537	Trace	554
Cloud Logging	538	Debugger	554
Legacy Logs Viewer	538	Profiler	555
Logs Explorer	539	Error Reporting	555
Exporting logs	544	Summary	556
		Further reading	556

Section 4: Exam Focus

18

Case Studies

Understanding how to approach the exam case studies	560	Forming a solution	566
What are we looking for in the case studies?	561	Summary of Mountkirk	569
Company overview	561	Additional case studies	570
Solution concept	563	EHR Healthcare	570
Existing technical environment	564	TerramEarth	574
Business requirements	564	Helicopter Racing League	577
Technical requirements	565	Summary	580
Executive summary	565	Further reading	580

19

Test Your Knowledge

Mock test one	581	Answers to mock test one	602
Mock test two	587	Answers to mock test two	606
Mock test three	592	Answers to mock test three	608
Mock test four	597	Answers to mock test four	612

Index

Other Books You May Enjoy
