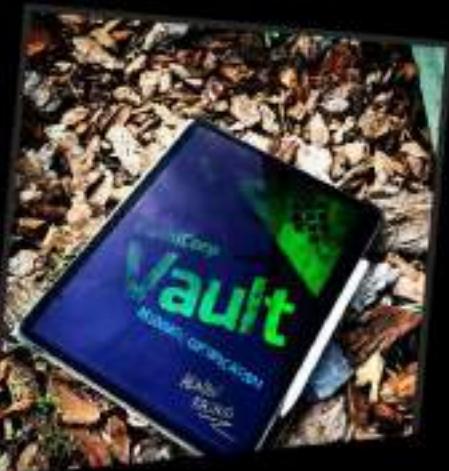


# CLOUD NATIVE CERTIFIED

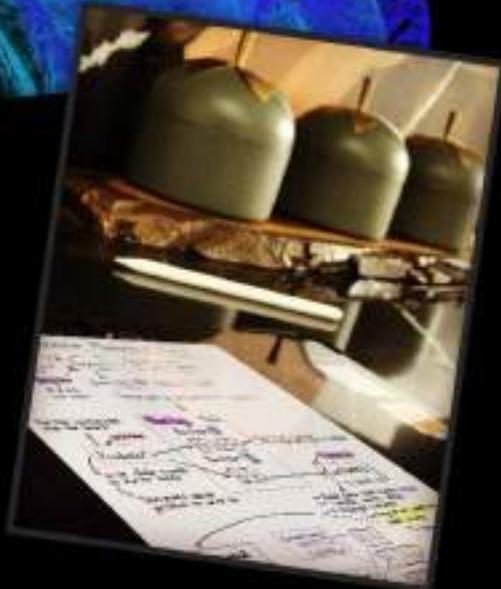


# KUBERNETES ADMINISTRATOR

ADNAN RASHID

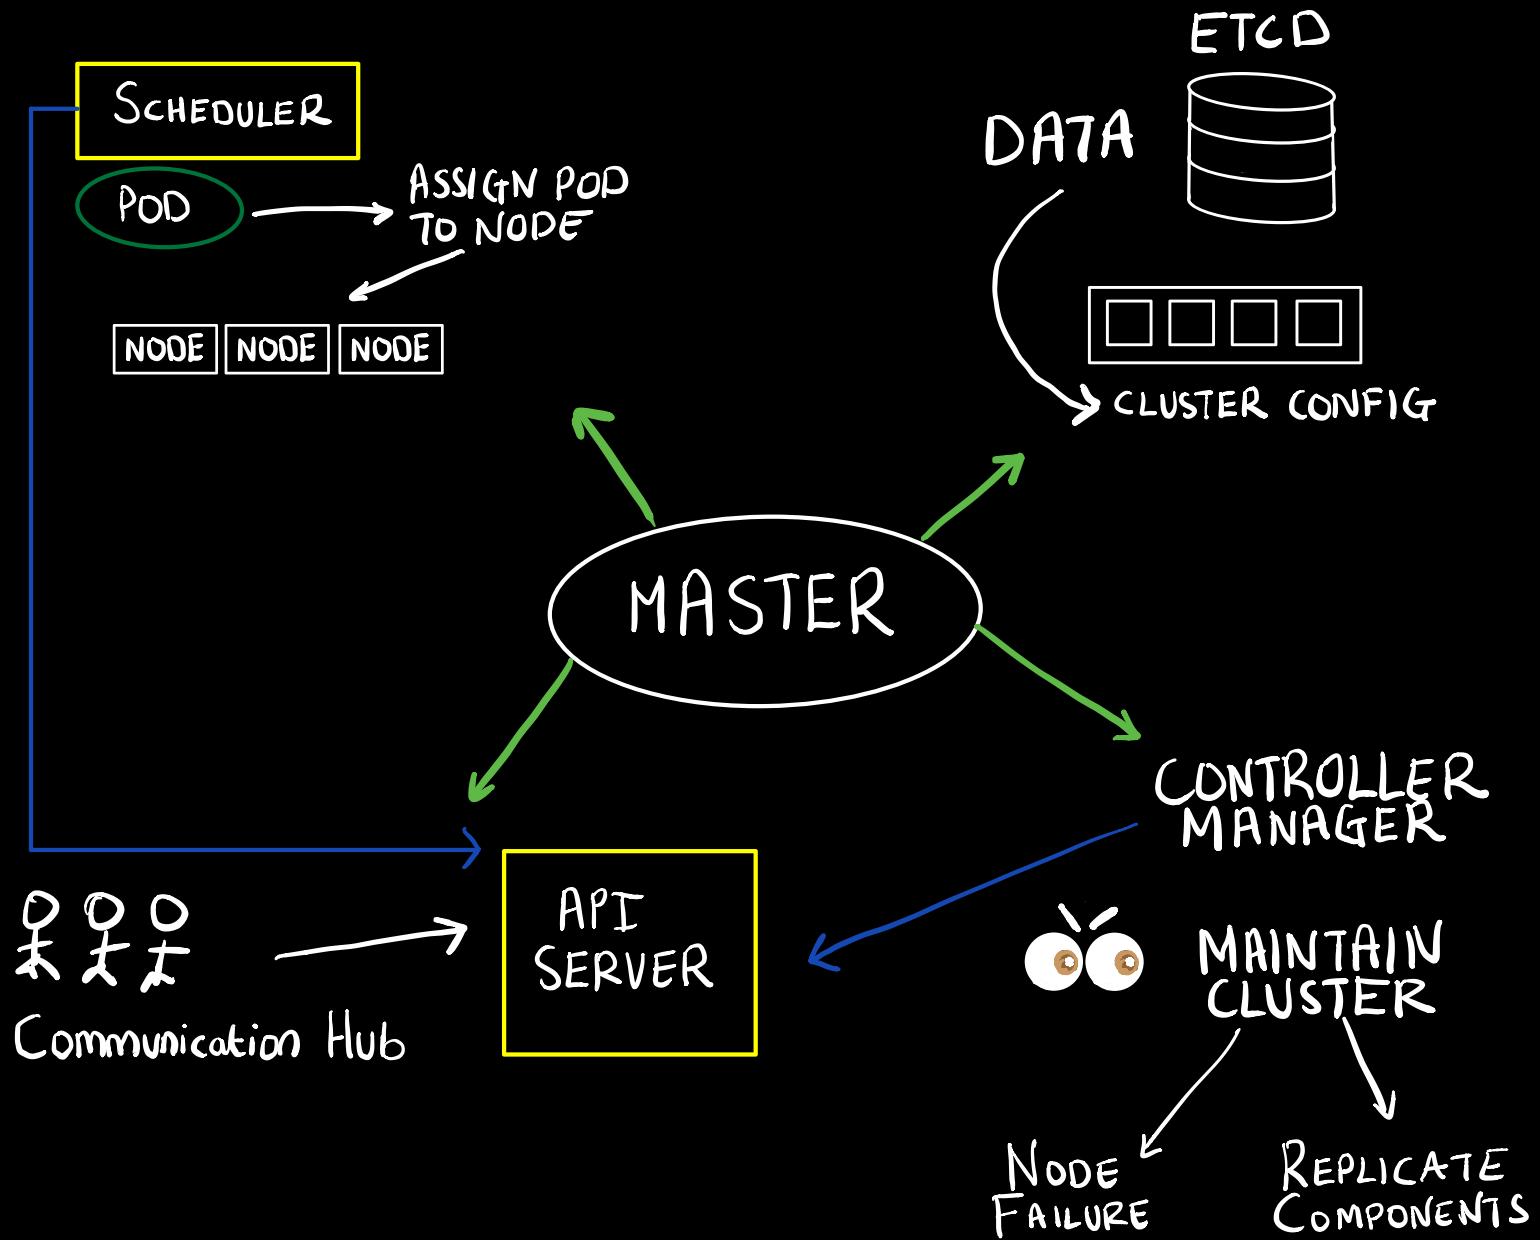


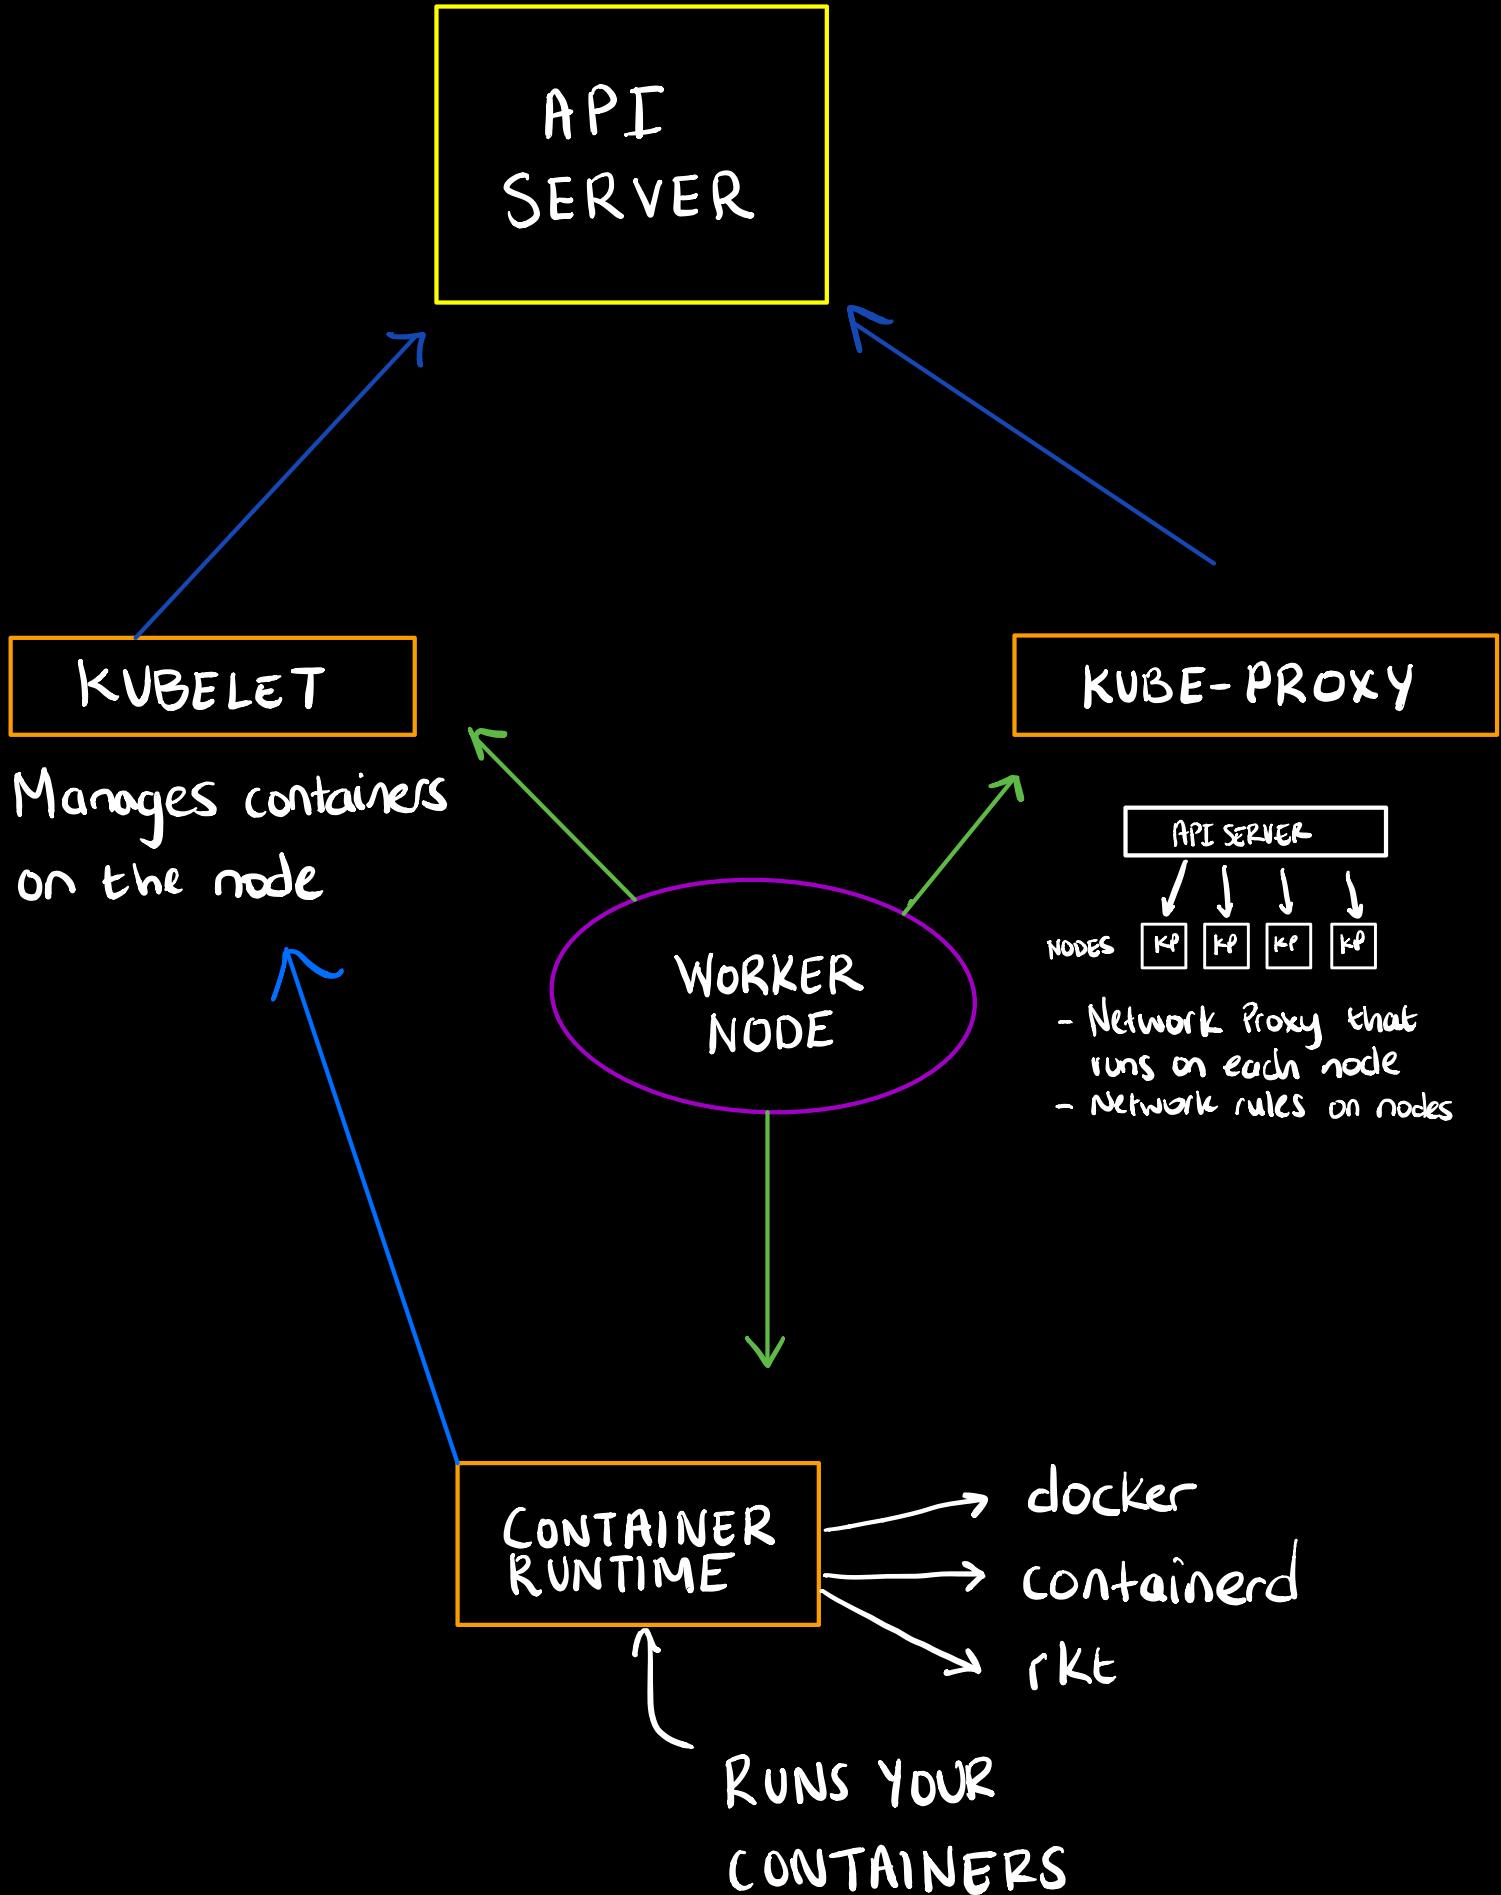
[instagram.com/  
adnans\\_techie\\_studies](https://instagram.com/adnans_techie_studies)



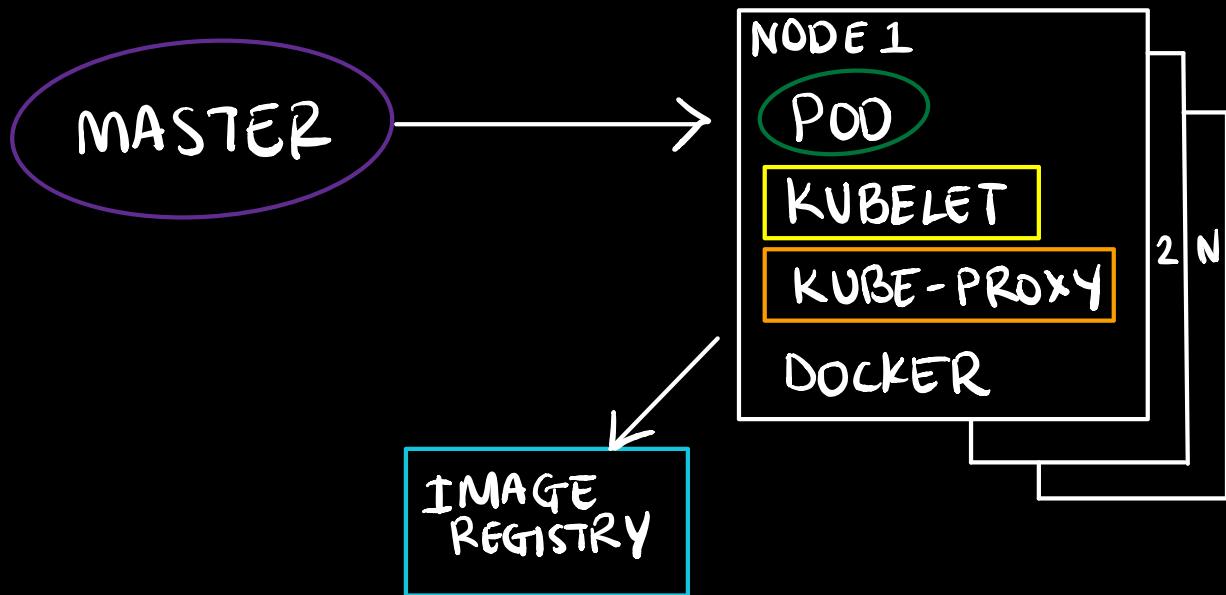
# UNDERSTANDING THE KUBERNETES ARCHITECTURE

# CLUSTER ARCHITECTURE



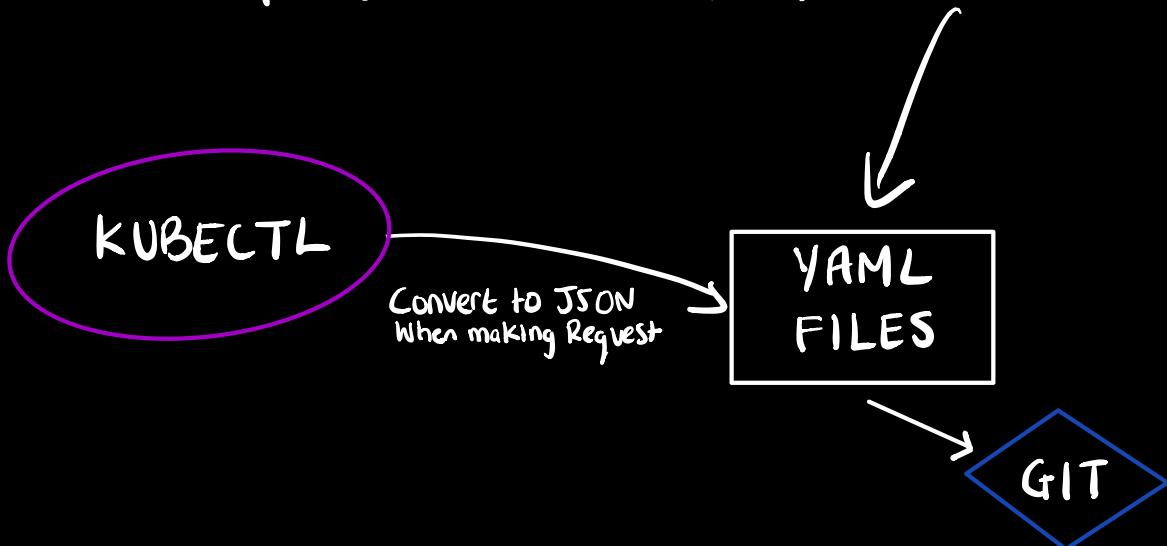


# APPLICATION RUNNING ON K8S



# API PRIMITIVES

API Server is the only one that communicates with etcd  
Every single component speak with API server only not to each other  
Objects like pods and services are declarative intents



# YAML FILE COMPOSITION

**API VERSION**

→ clear consistent view  
of resources

**KIND**

→ The kind of object you want to create

- Pod
- Deployment
- Job

**METADATA**

→ uniquely identify the object

Name  
String

↓  
UID

↓  
namespace

**SPEC**

Container image  
volume exposed ports

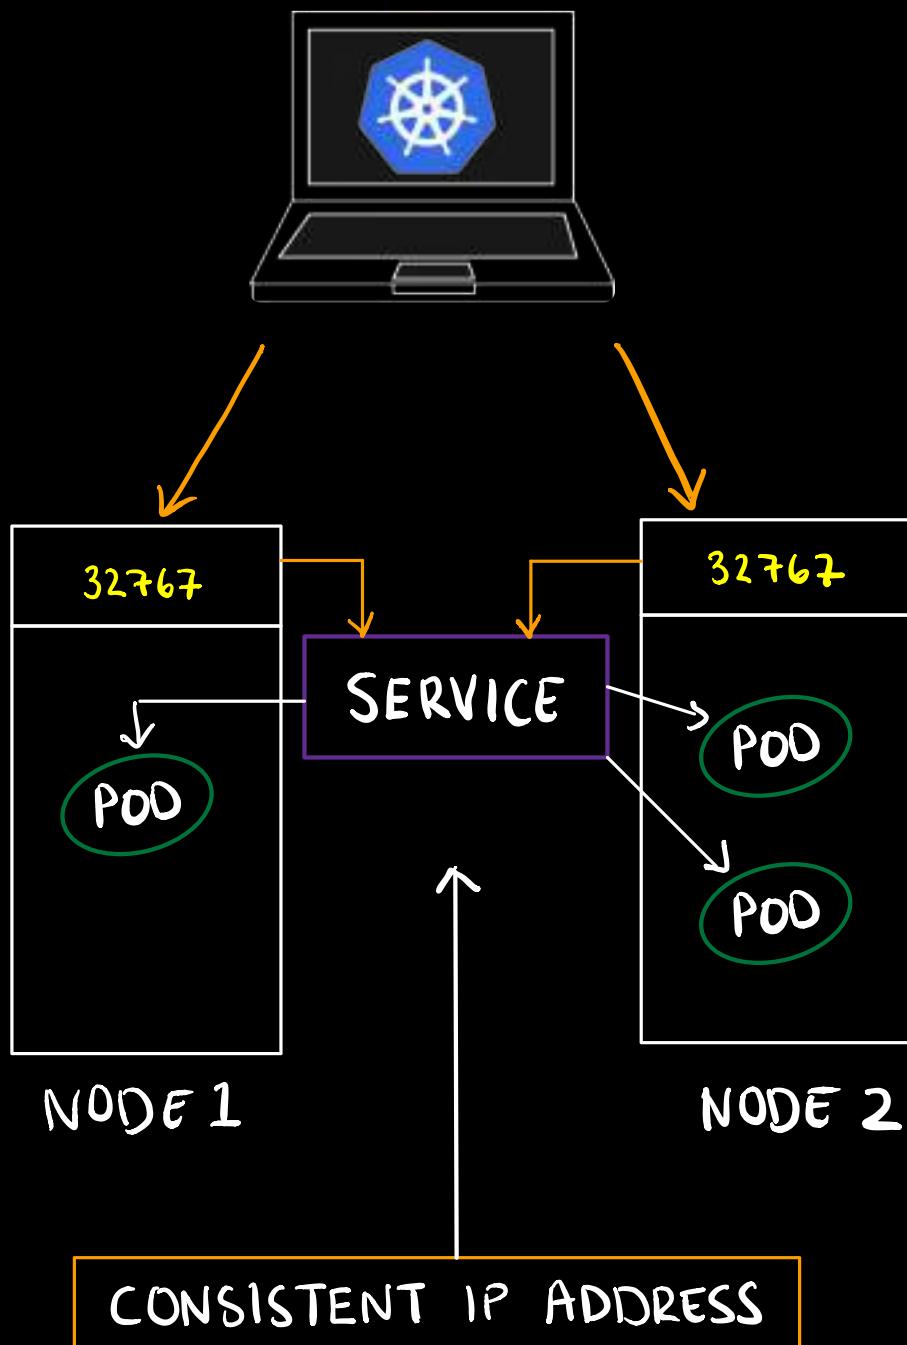
**STATUS**

→ State of the object

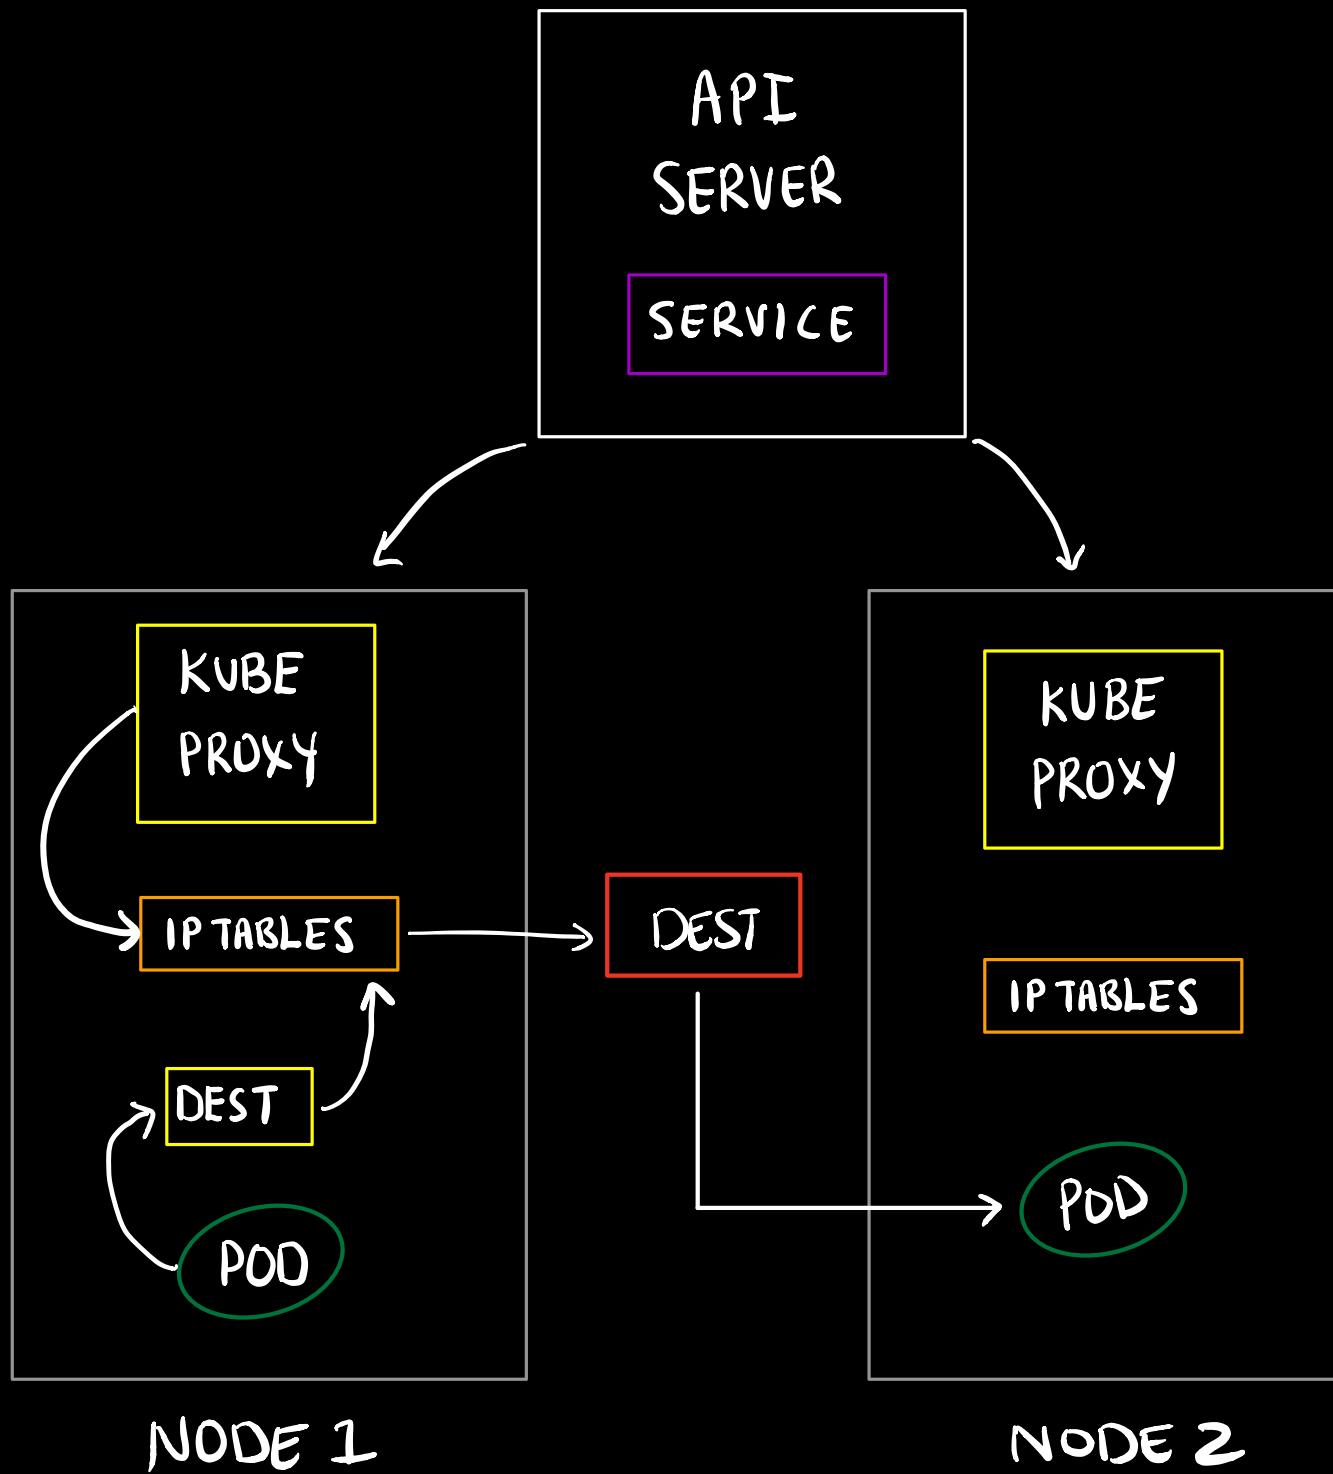
↓  
match desired states

# SERVICES AND NETWORK PRIMITIVES

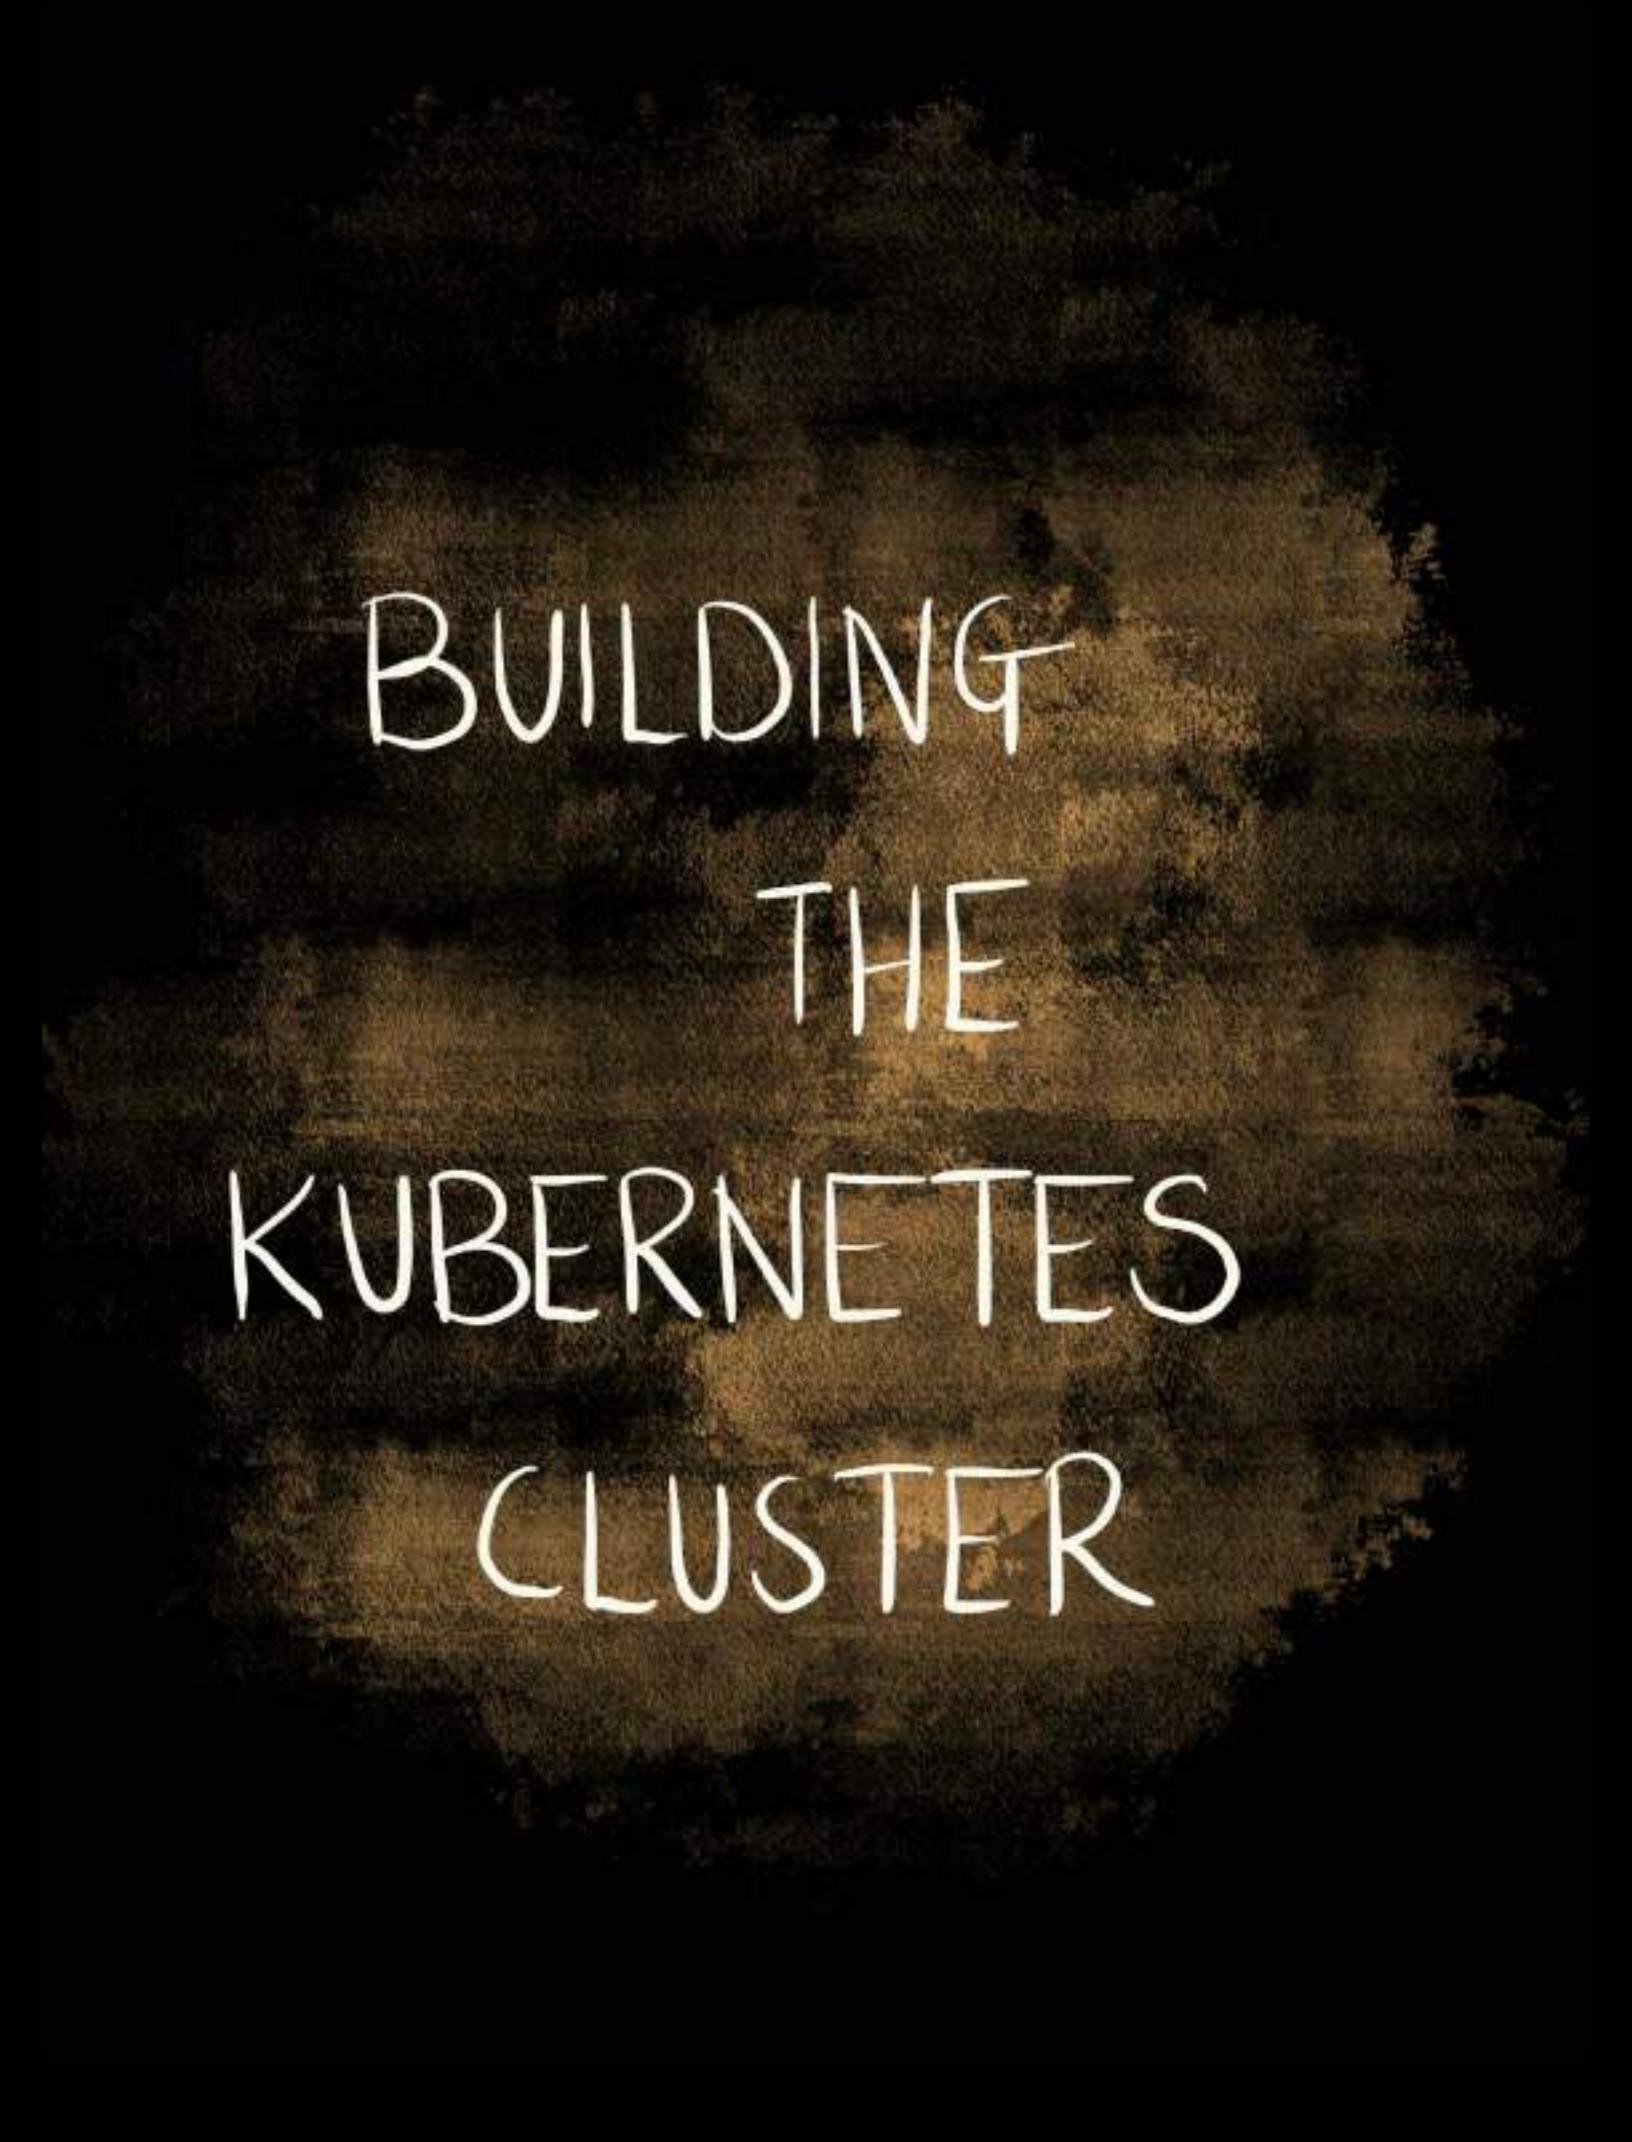
SERVICES ALLOW YOU TO DYNAMICALLY  
ACCESS A GROUP OF REPLICA PODS



# KUBE-PROXY



KUBE-PROXY HANDLES THE TRAFFIC ASSOCIATED WITH A SERVICE BY CREATING IP TABLE RULES



# BUILDING THE KUBERNETES CLUSTER

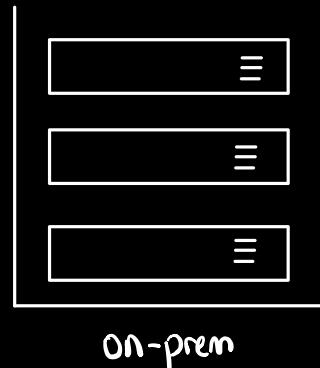
# RELEASE BINARIES, PROVISIONING & TYPES OF CLUSTERS

## Picking the right Solution



CLOUD

VS



on-prem

ALSO

### CUSTOM

- Install Manually
- Configure your own network fabric
- Locate the release binaries
- Build your own Images
- Secure cluster config

VS

### Pre-Built

- Minikube
- Minishift
- Microk8s
- Ubuntu on LXD
- AWS, Azure, GCP

✓ using minikube locally on MacOS

# INSTALLING KUBE MASTER AND NODES

## MASTER + WORKERS

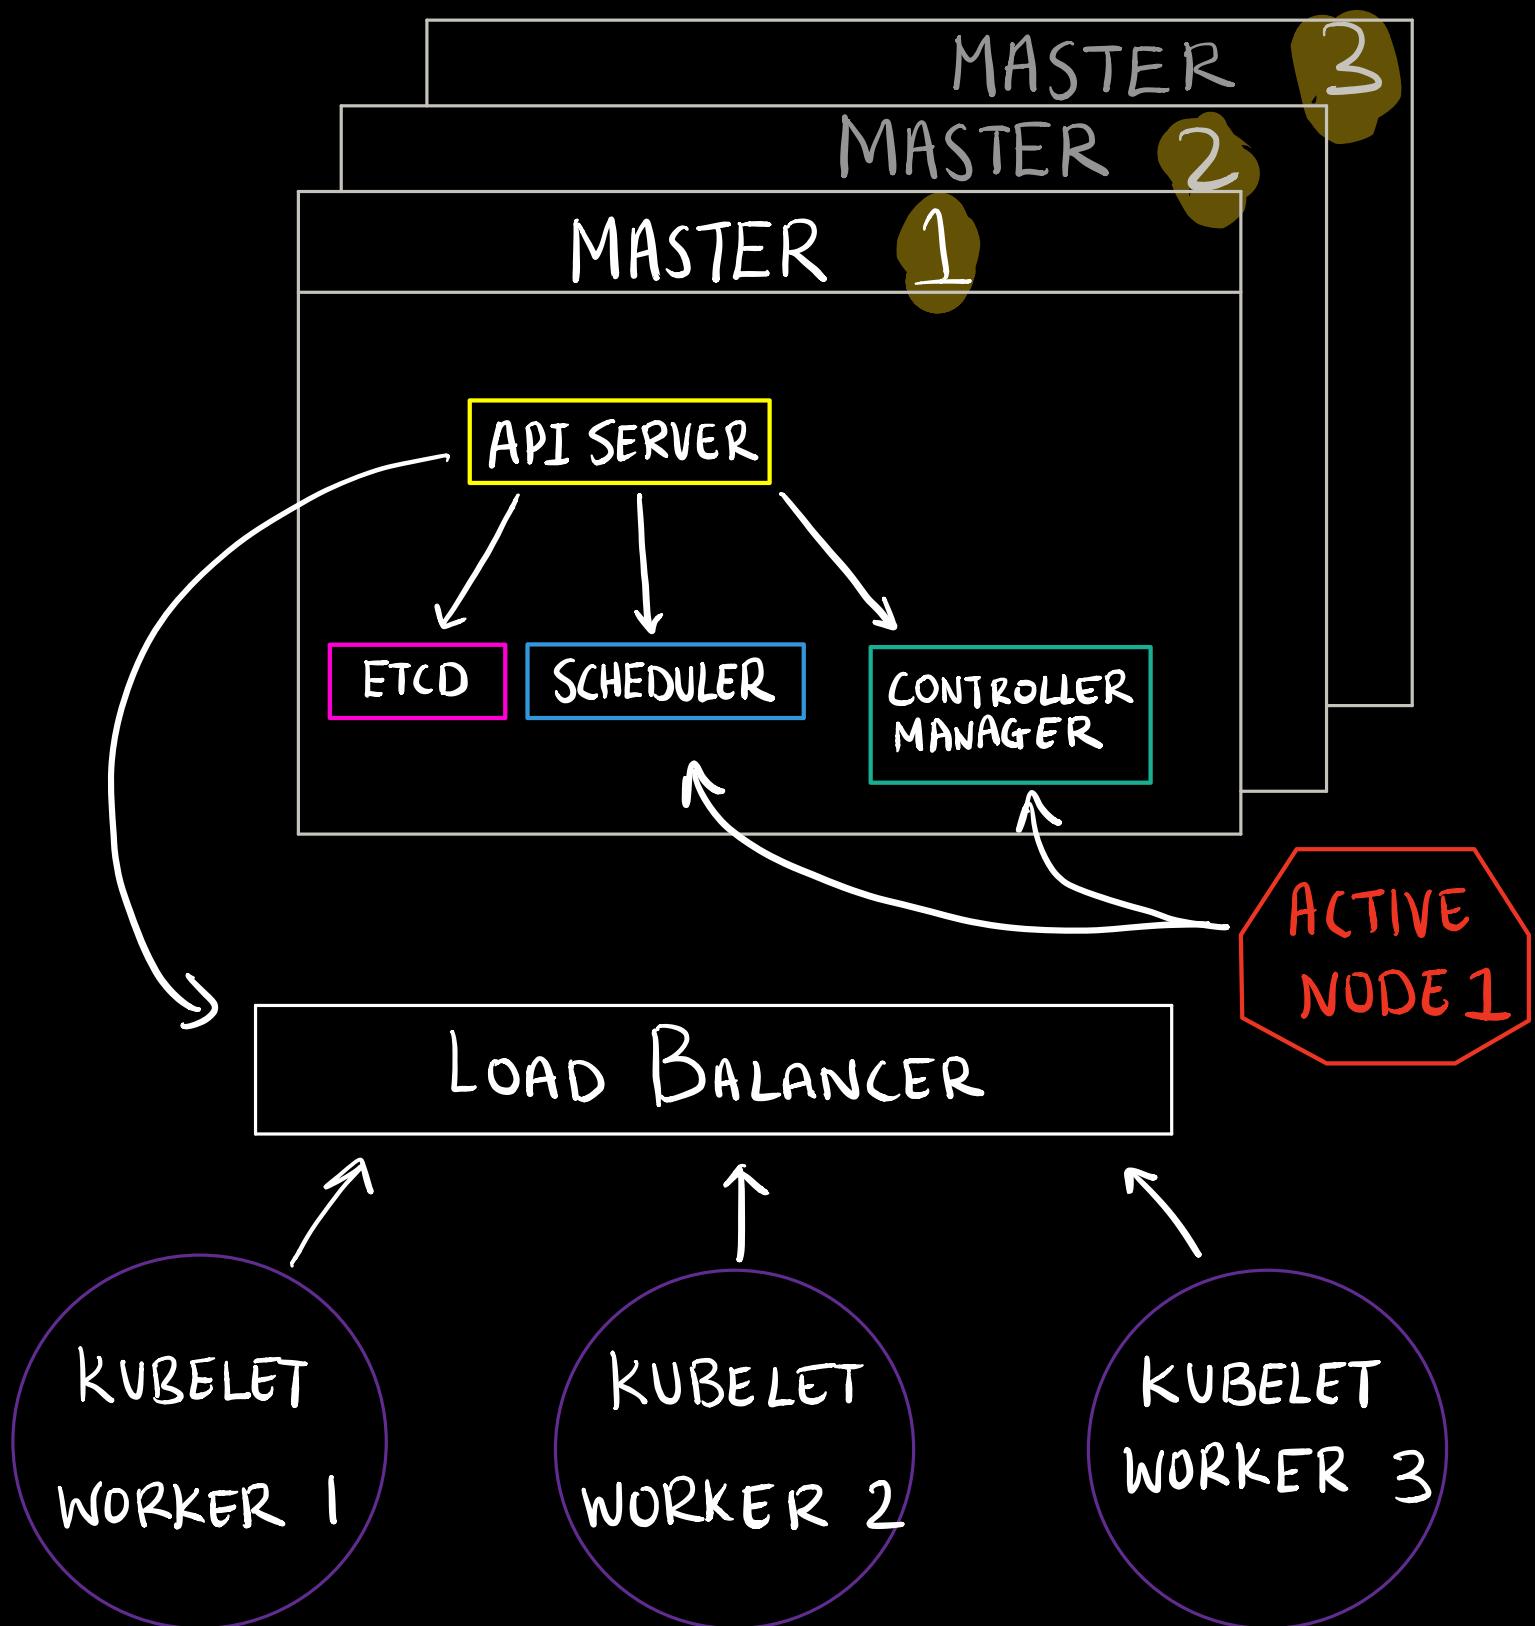
- ① DOCKER + KUBE → GPG KEY  
→ ADD REPOS
- ② UPDATE PACKAGES
- ③ INSTALL Docker, kubelet, kubeadm, kubectl
- ④ Modify bridge adapter settings

## MASTER ONLY

- ① Initialise Cluster
- ② Make directory for kss
- ③ Copy kube config
- ④ Change ownership of config
- ⑤ Apply Flannel CNI

# BUILDING HIGHLY AVAILABLE CLUSTER

ALL components can be replicated, but only certain can operate simultaneously



The controller manager and the scheduler actively watch the cluster state and take action when it changes

SCHEDULER



CM



CLUSTER

ARE WE IN

CHARGE?

SCHEDULER



CM



NO WE ARE!

LEADER  
ELECT  
OPTION

HOW DO WE  
DECIDE?

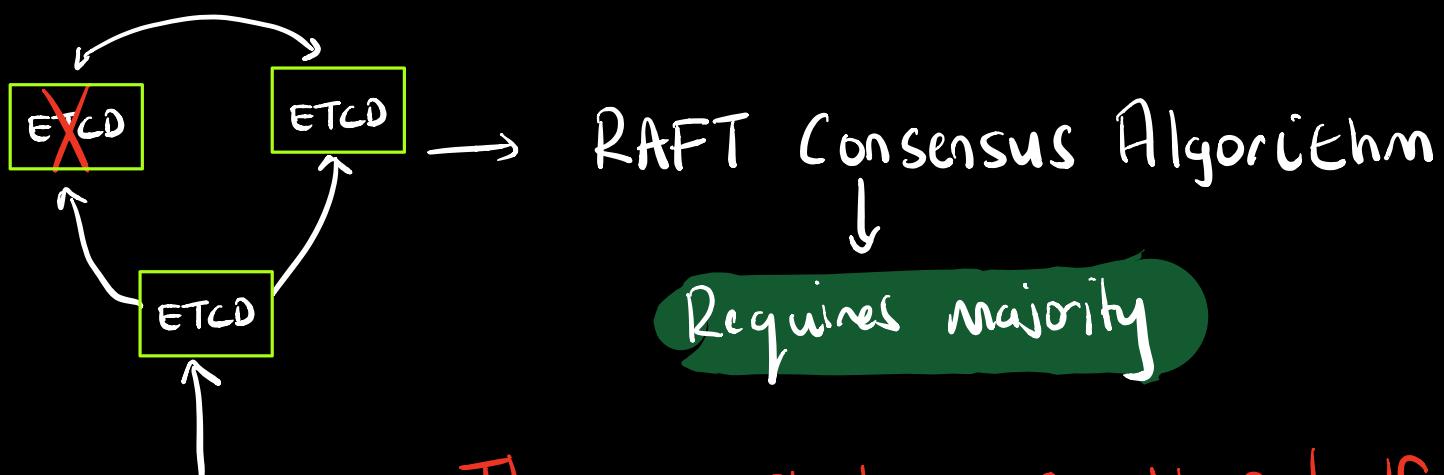
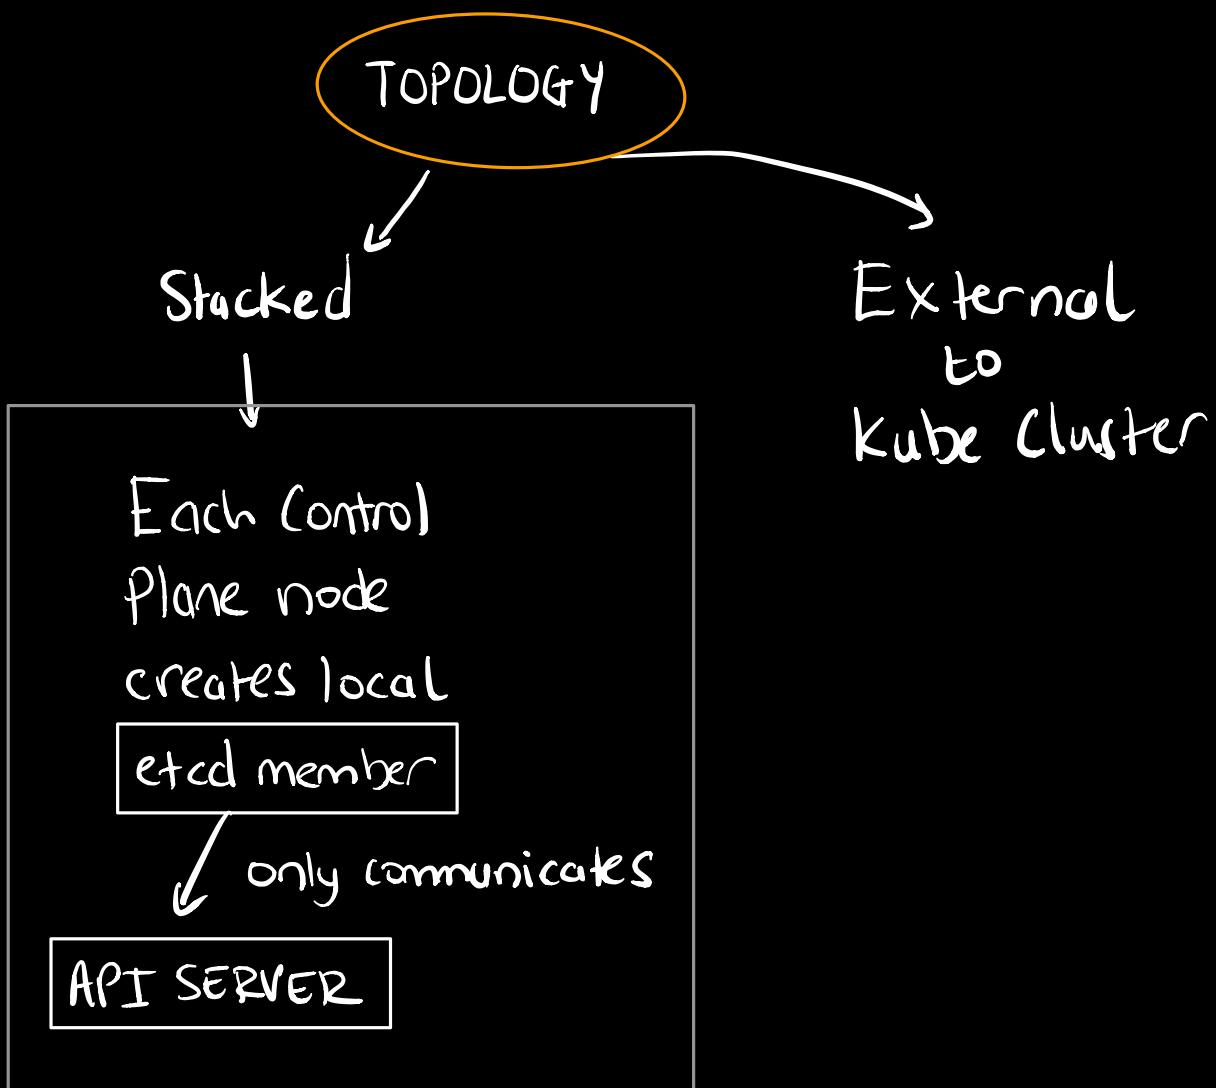
creates endpoint resource

↑ see in scheduler YAML

↑ holderIdentity

→ Leading to duplicate resources or corruption.

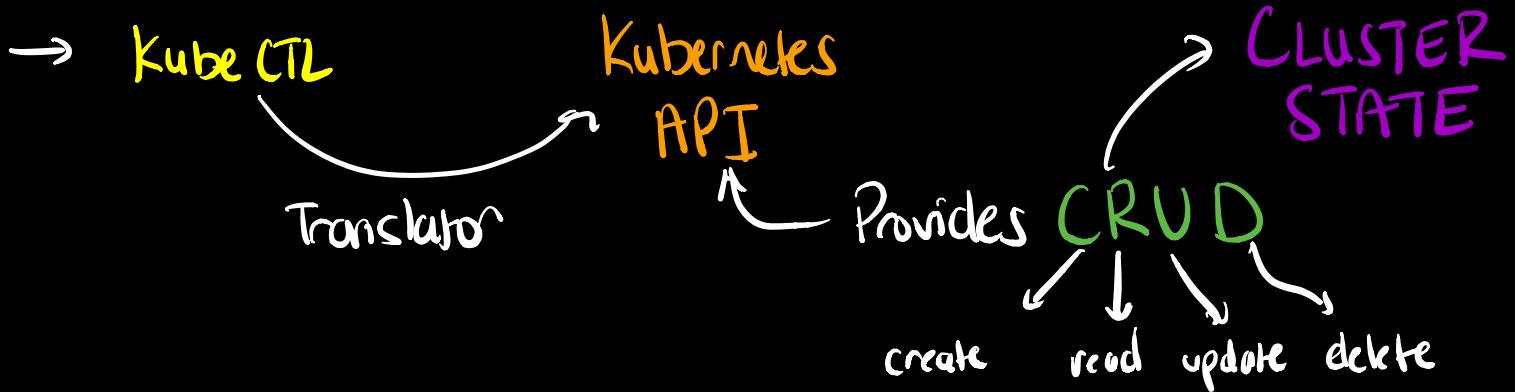
# REPLICATING ETCD



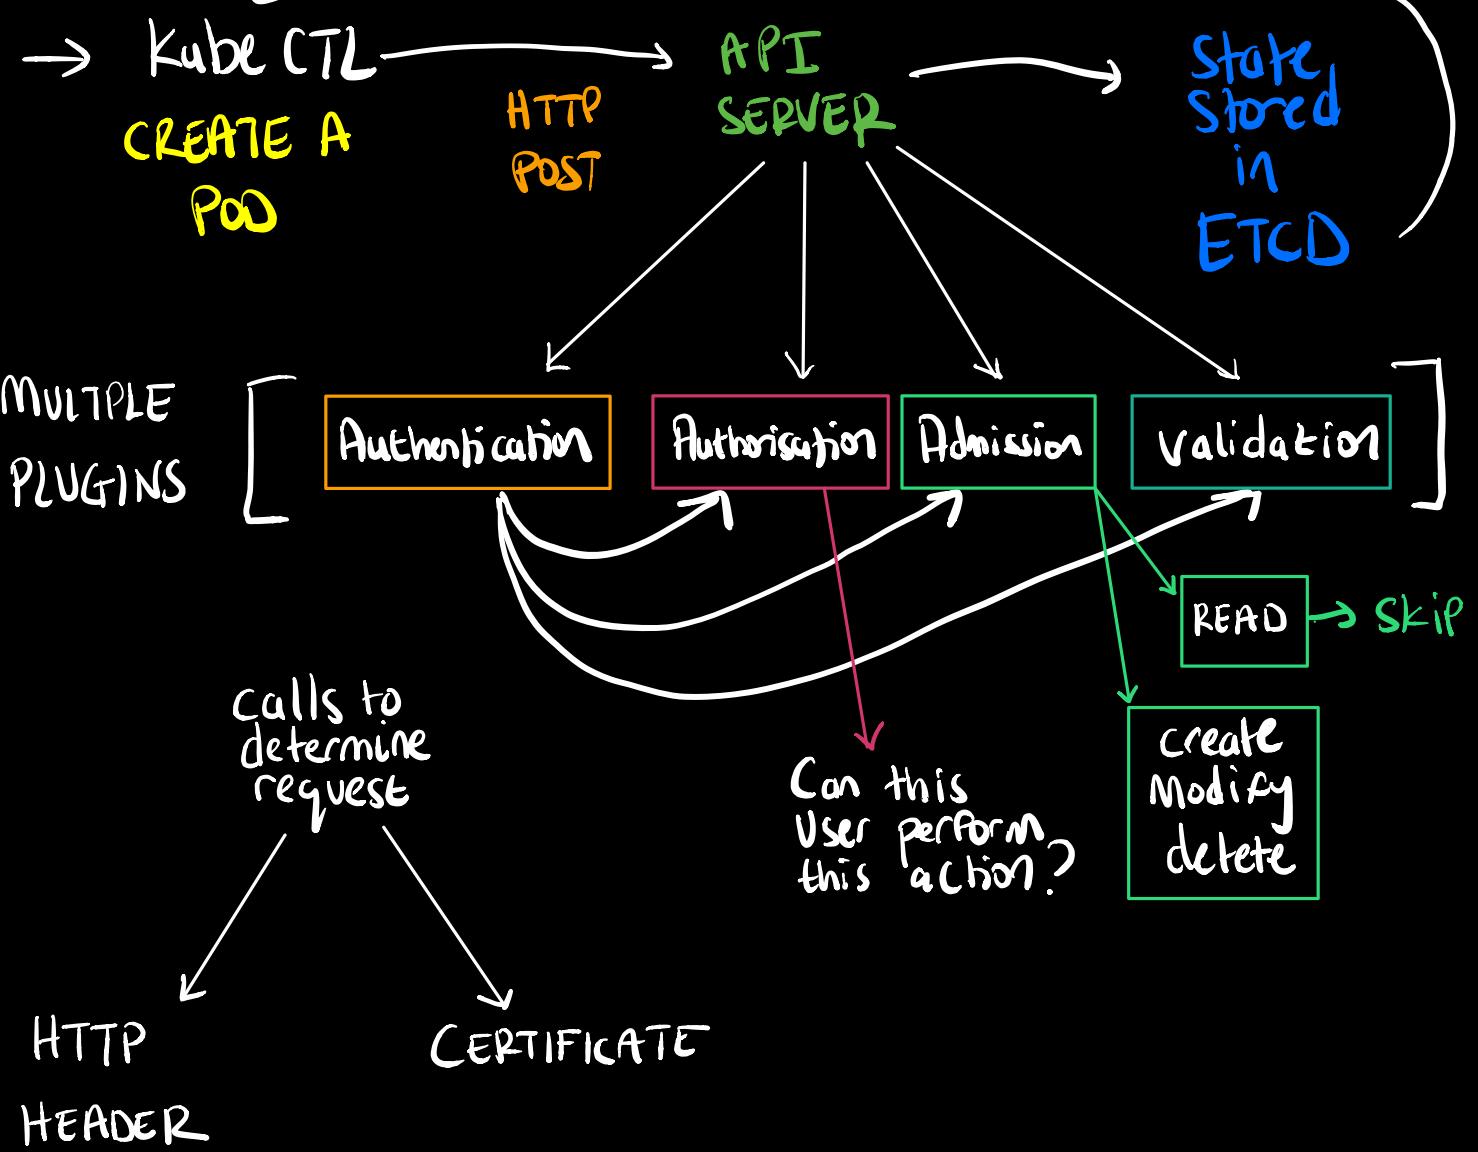
There must be more than half taking place in the state change  
∴ there must be odd number of nodes

# CONFIGURING SECURE CLUSTER COMMUNICATIONS

→ All communication via HTTPS

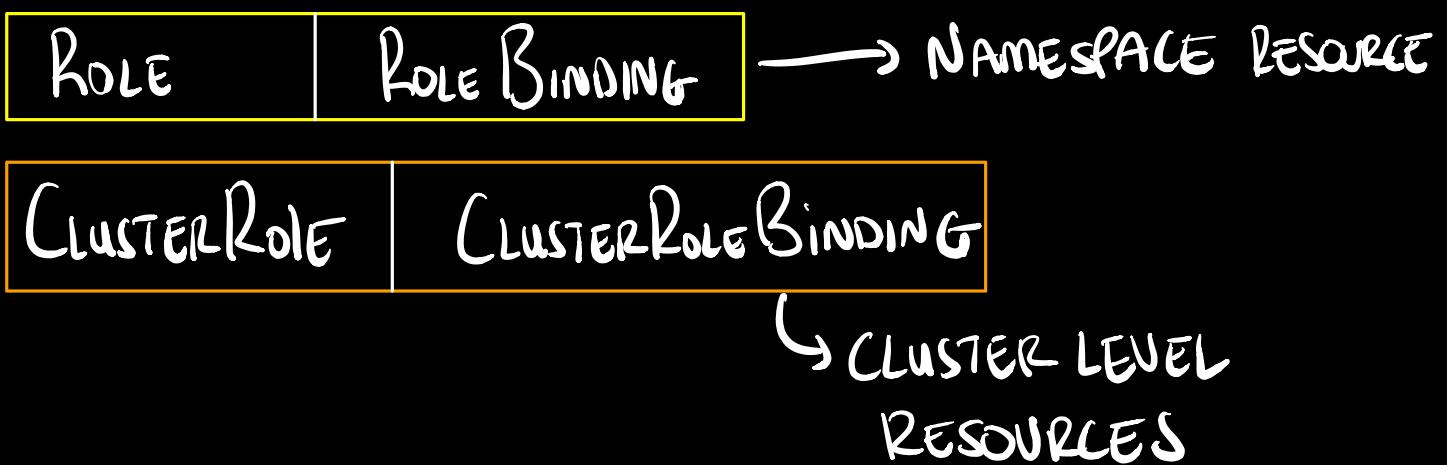
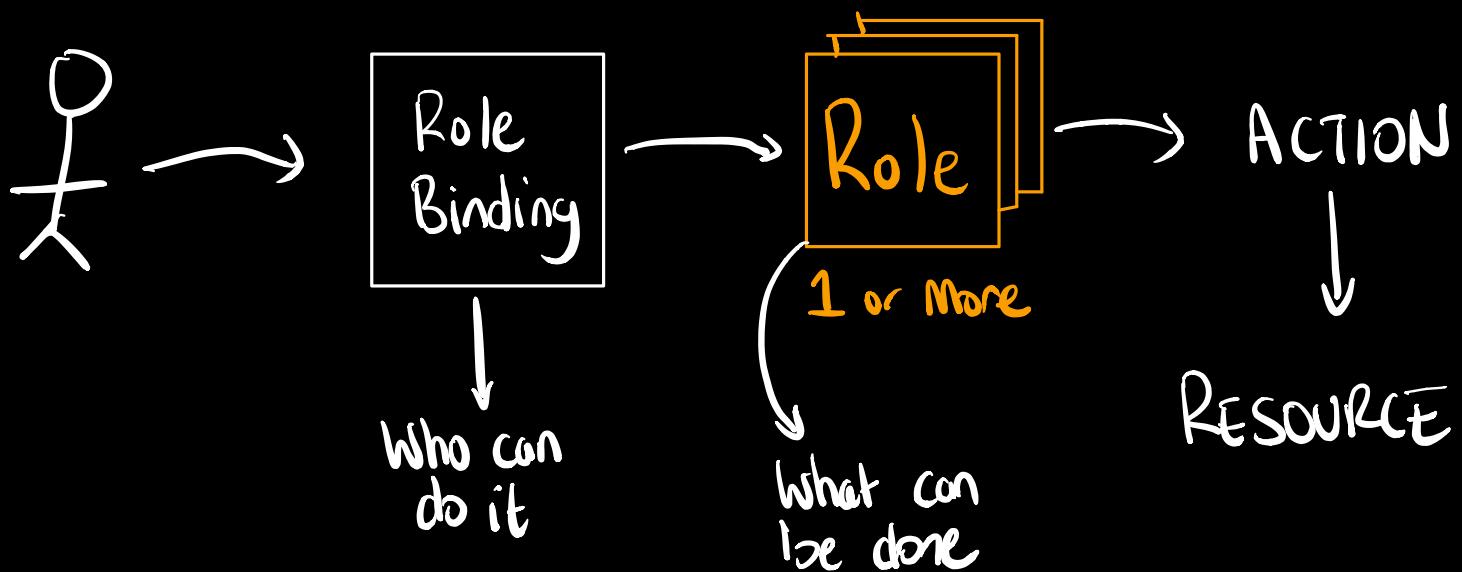


Return Response

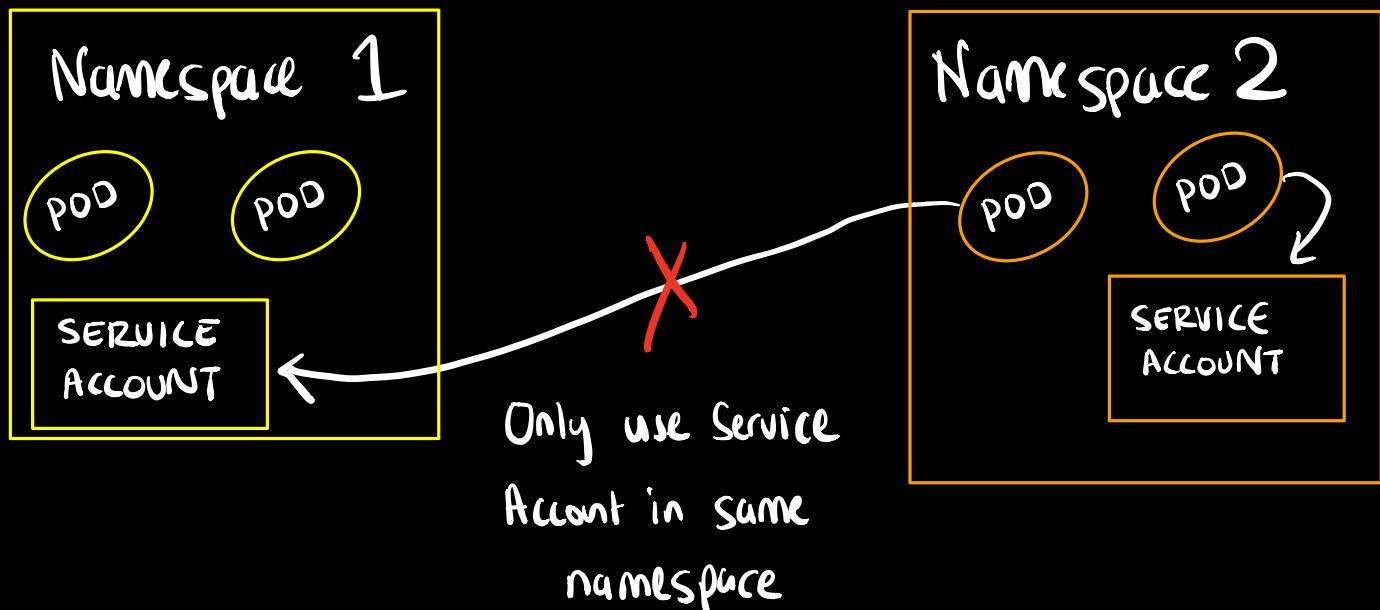
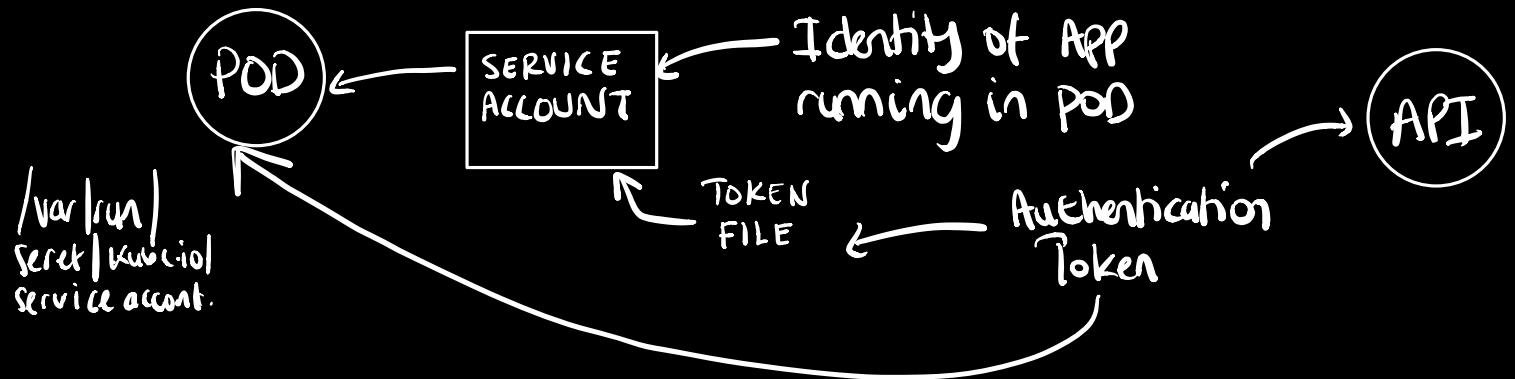


# ROLES AND ACCESS

RBAC is used to prevent unauthorised users from modifying the cluster state



# SERVICE ACCOUNT



# RUNNING END TO END TESTS ON CLUSTER

Performance  
and  
Response of  
Application

Poor Cluster  
Performance

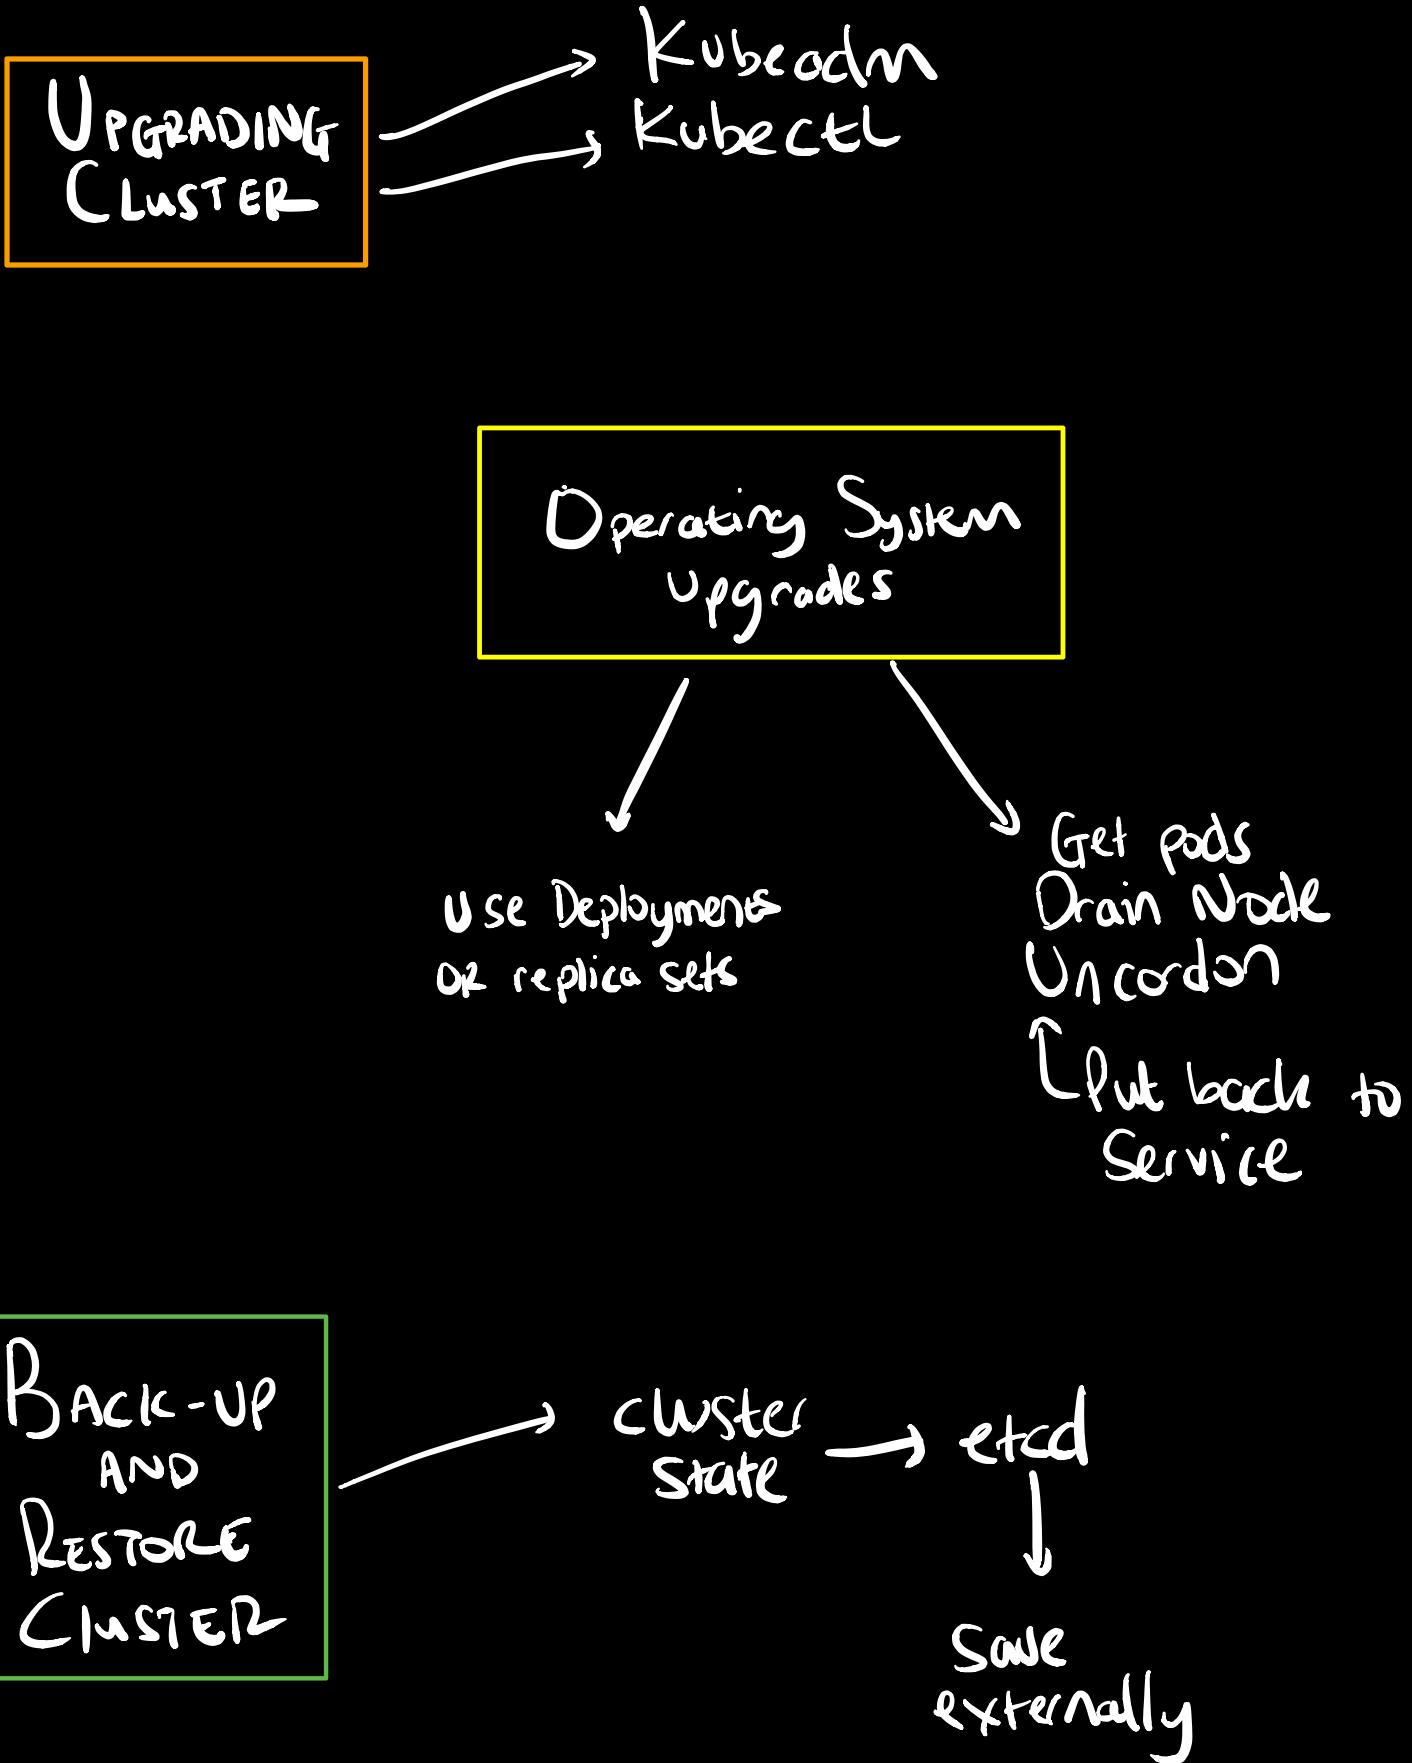


kubeTest

Example  
Tests

- ✓ Deployments can run
- ✓ Pods can run
- ✓ Pods can be directly accessed
- ✓ Logs can be collected
- ✓ Commands run from pod
- ✓ Services can provide access
- ✓ Nodes are healthy
- ✓ Pods are healthy

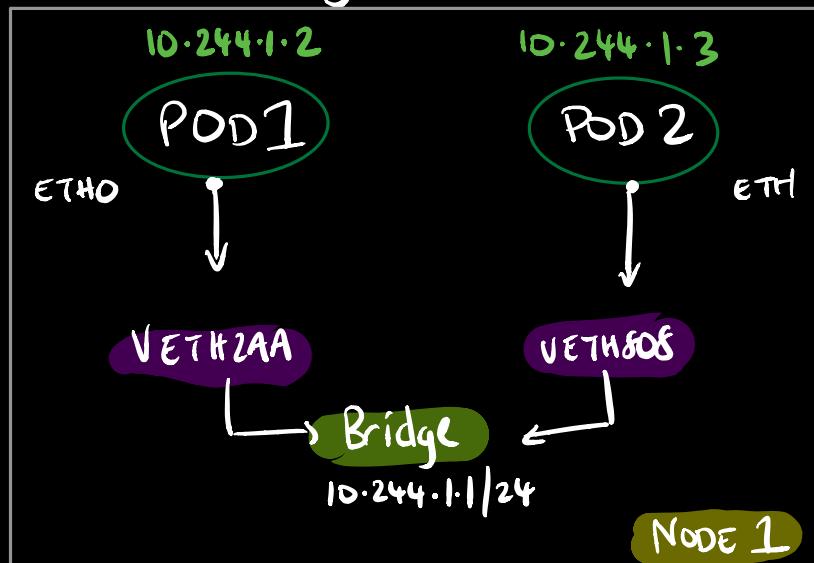
# MANAGING CLUSTER



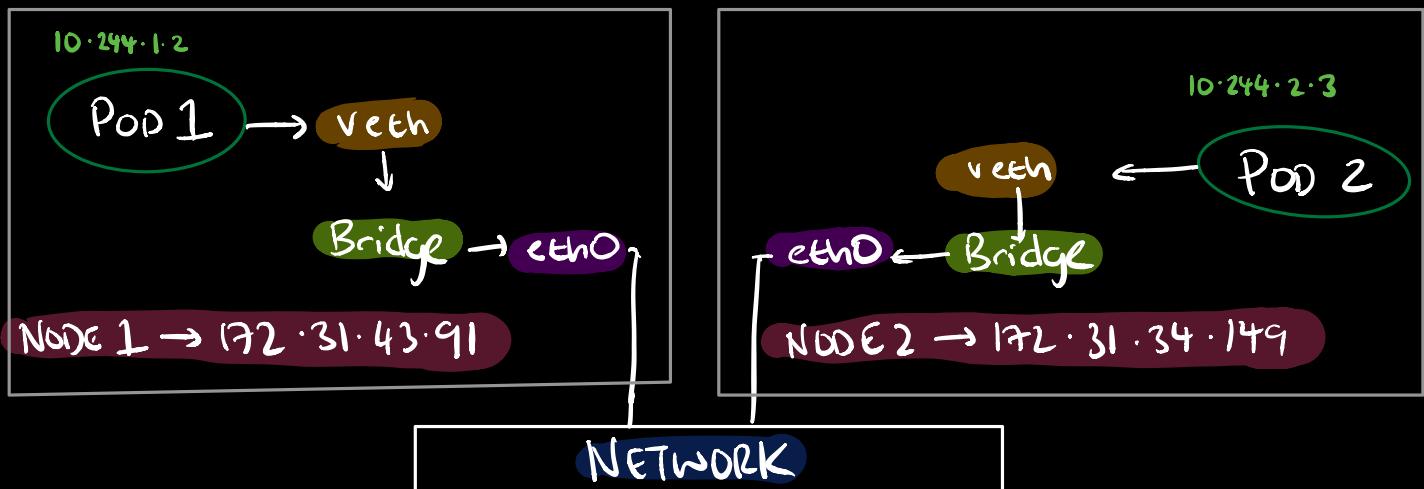
NETWORK  
CLUSTER  
COMMUNICATIONS

# POD AND NODE NETWORKING

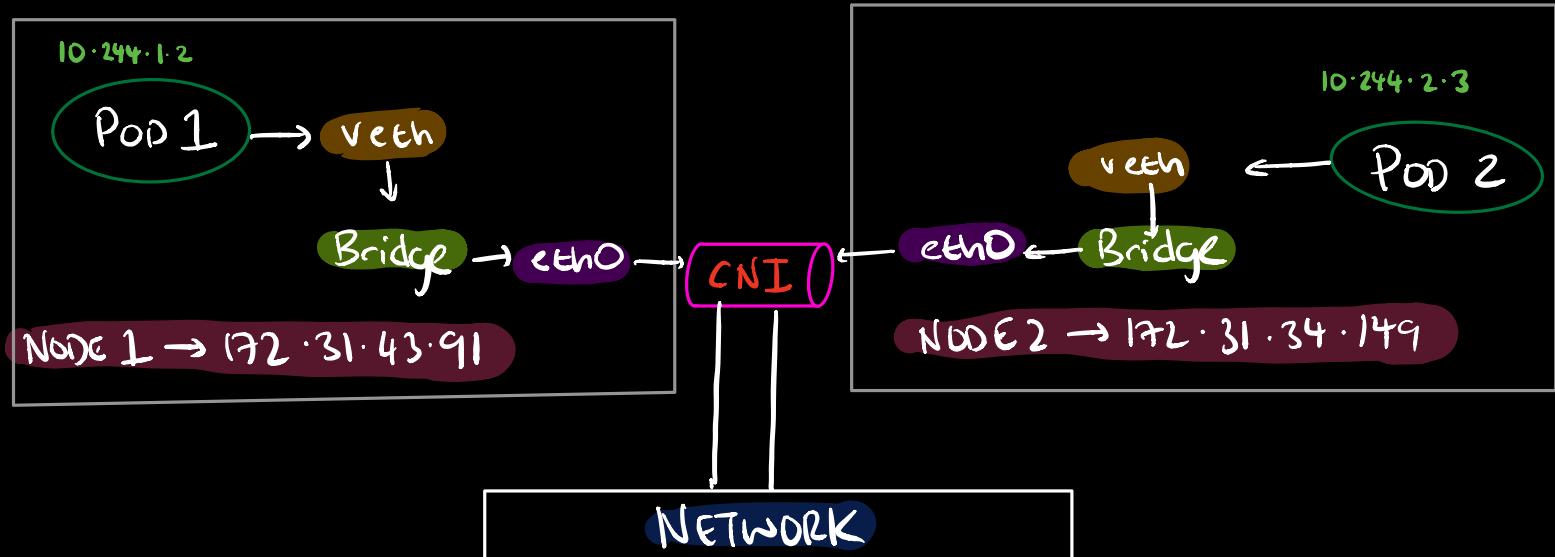
## Networking within a node



## Networking Outside of the Node



# CONTAINER NETWORK INTERFACE



## CNI is a network Overlay

- ↳ Allows building tunnel between nodes
  - ↳ Sits on top of existing networks
    - ↳ Encapsulates Packet
      - ↳ changes SRC & DST



How  
CNI  
DO  
THIS?

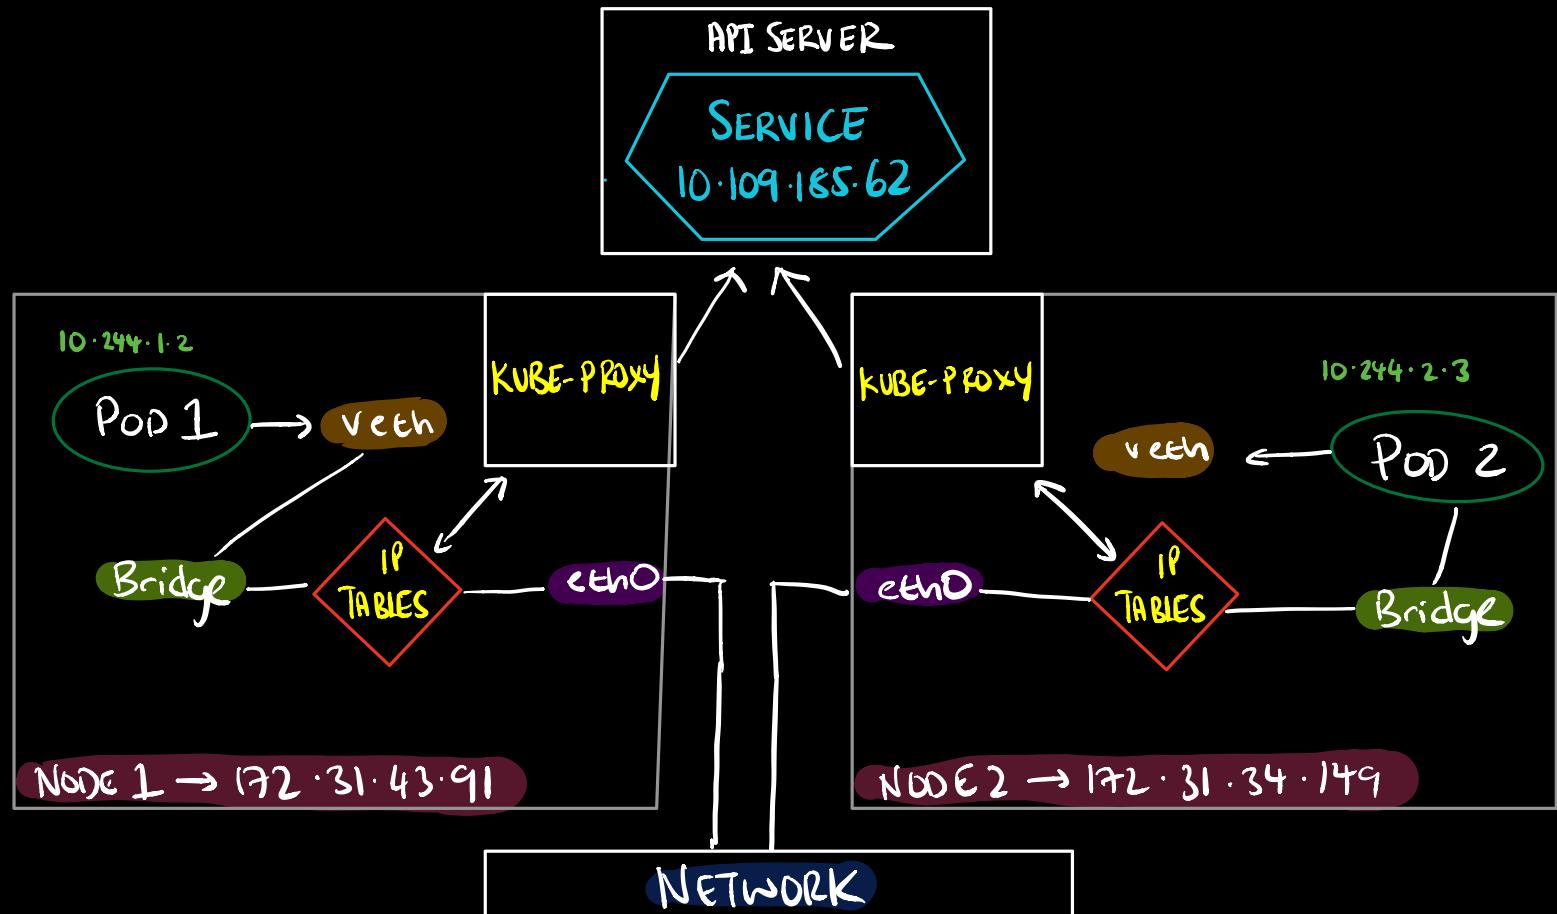
- There is a mapping associated in user space
  - ↳ Program all pod ip addresses to node IP, when reach other node, de-encapsulate packet and give to bridge

Example CNI

Calico      Flannel

Appears Local To Node!

# SERVICE NETWORKING



→ Pods come and go.

↳ How does cluster keep track?

↳ Services!

↳ Provides virtual interface → Auto assigned to pods behind interface.

Example

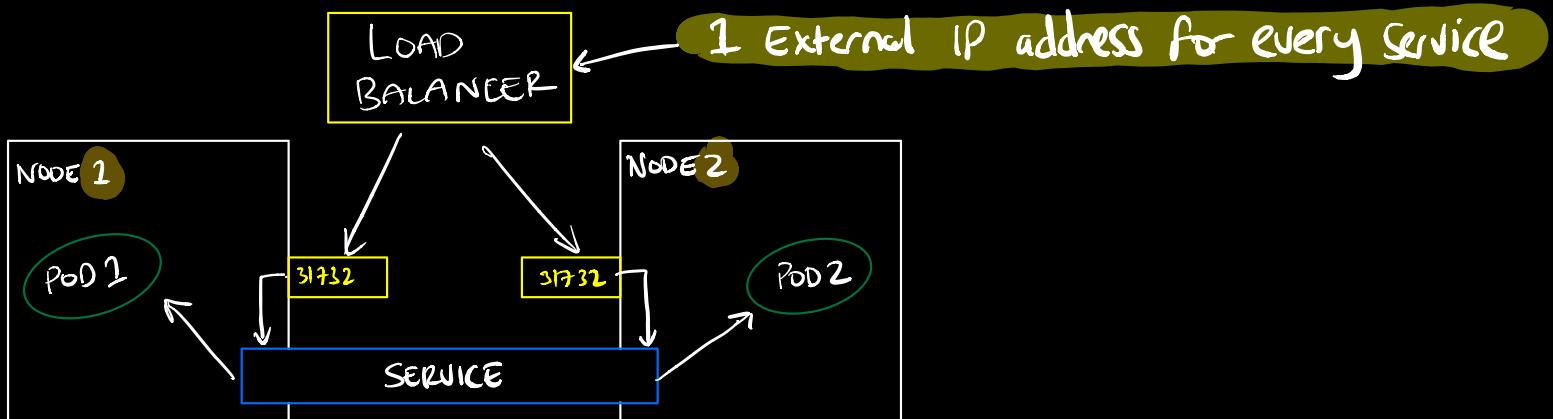
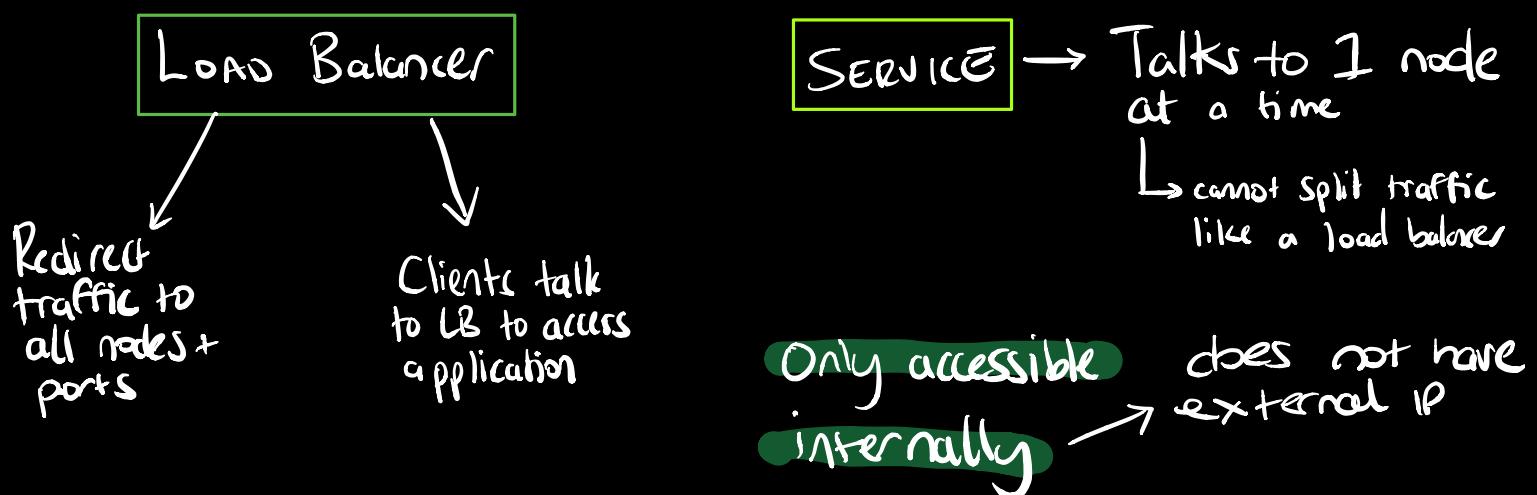
Cluster IP Services

→ Auto created on cluster creation

→ Takes care of internal routing

→ No matter where moves other pods know how to communicate to it

# INGRESS RULES AND LOAD BALANCERS



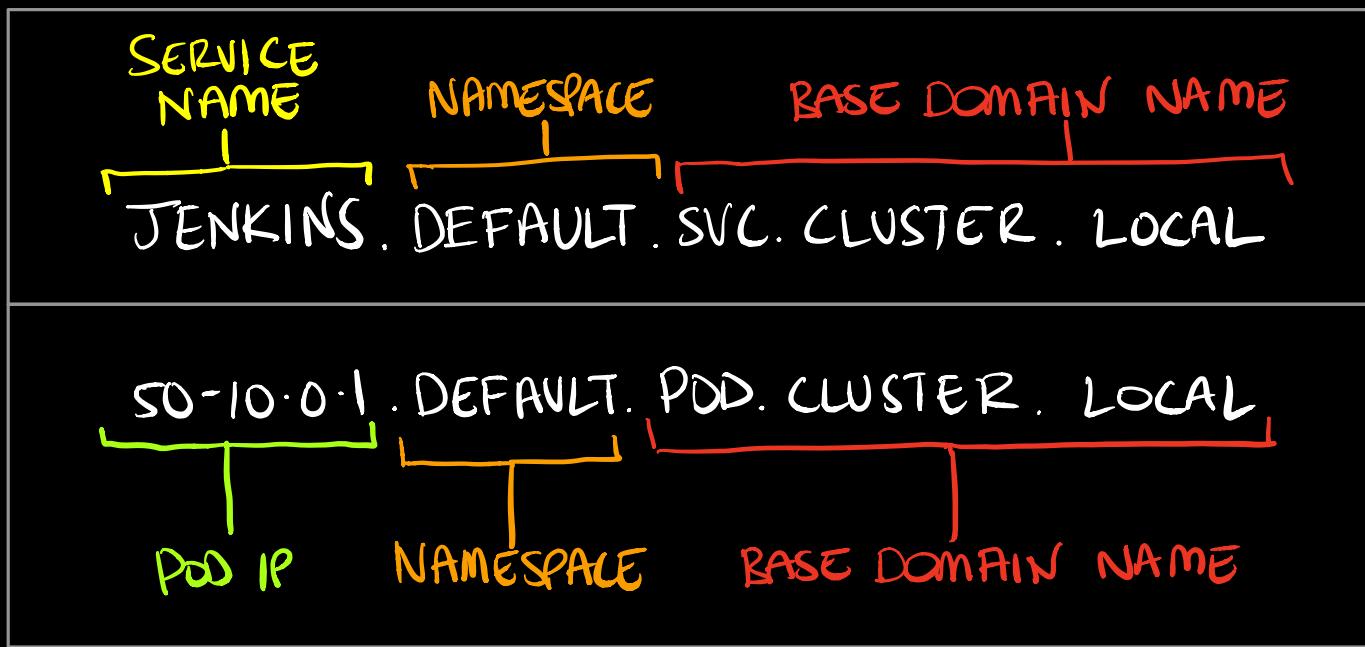
## INGRESS



Access multiple Services with Single IP Address

# CLUSTER DNS

EVERY SERVICE DEFINED IN THE CLUSTER IS ASSIGNED A DNS NAME



A PODS DNS SEARCH WILL INCLUDE THE PODS OWN NAMESPACE AND THE CLUSTERS DEFAULT DOMAIN

# POD SCHEDULING WITHIN THE KUBERNETES CLUSTER

# CONFIGURING THE KUBERNETES SCHEDULER

SCHEDULER RESPONSIBLE FOR ASSIGNING POD TO NODE  
BASED ON RESOURCE REQUIREMENTS OF THE POD

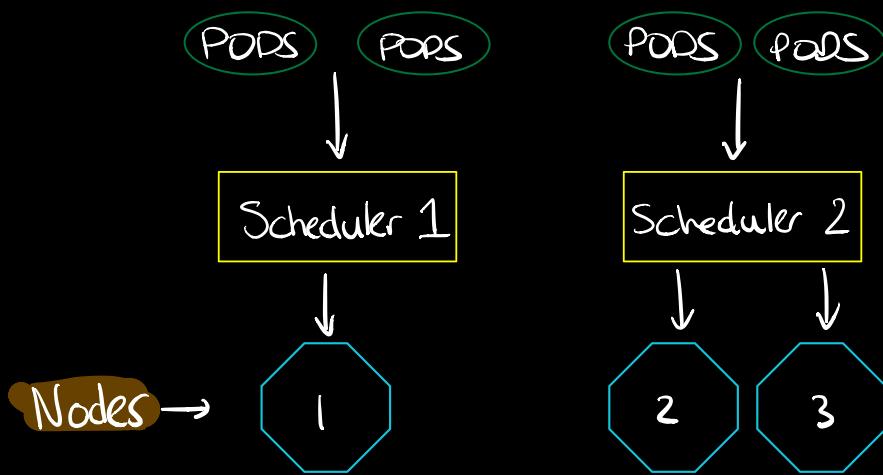


SCHEDULER

- 1 Does the node have adequate hardware resources?
- 2 Is the node running out of resources?
- 3 Does the pod request a specific node?
- 4 Does the node have a matching label?
- 5 If pod requests a port, is it available?
- 6 If pod requests a volume, can it be mounted?
- 7 Does the pod tolerate the taints of the node?
- 8 Does the pod specify node or pod affinity?

# RUNNING MULTIPLE SCHEDULERS FOR MULTIPLE PODS

It is possible to have 2 schedulers working alongside each other.

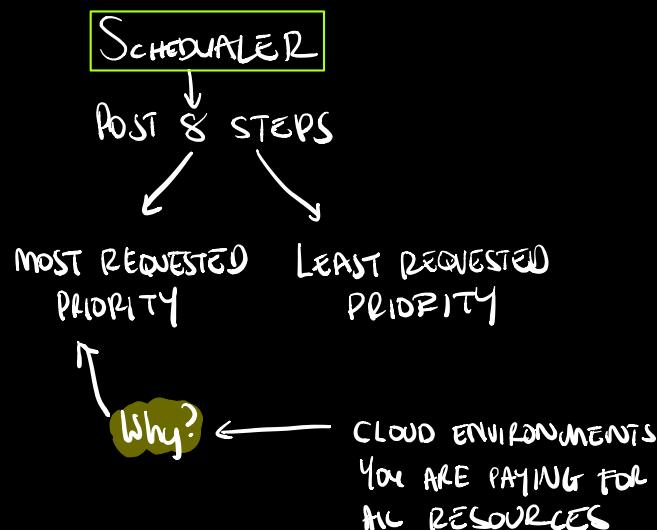


## SCHEDULING PODS WITH LIMITS AND LABEL SELECTORS

TRAINTS → REPEL WORK → EXAMPLE  
MASTER NODE NO SCHEDUAL

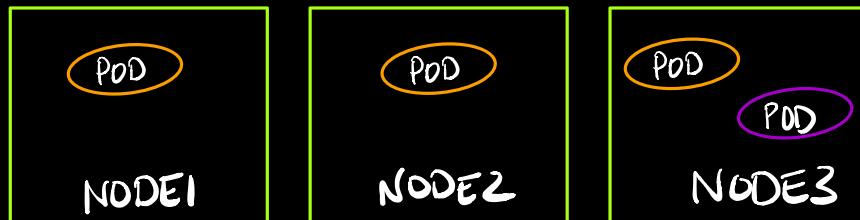
TOLERATIONS → ALLOW YOU TO TOLERATE A TRAINT → EXAMPLE  
KUBE-PROXY ← DAEMON SET POD MUST RUN ON ALL NODES

CPU/MEMORY  
↓  
POD MAY NOT BE USING ALL REQUESTED RESOURCE AT A GIVEN TIME.  
SCHEDULER LOOKS AT THE SUM OF RESOURCES REQUESTED BY EXISTING PODS



# DAEMONSETS

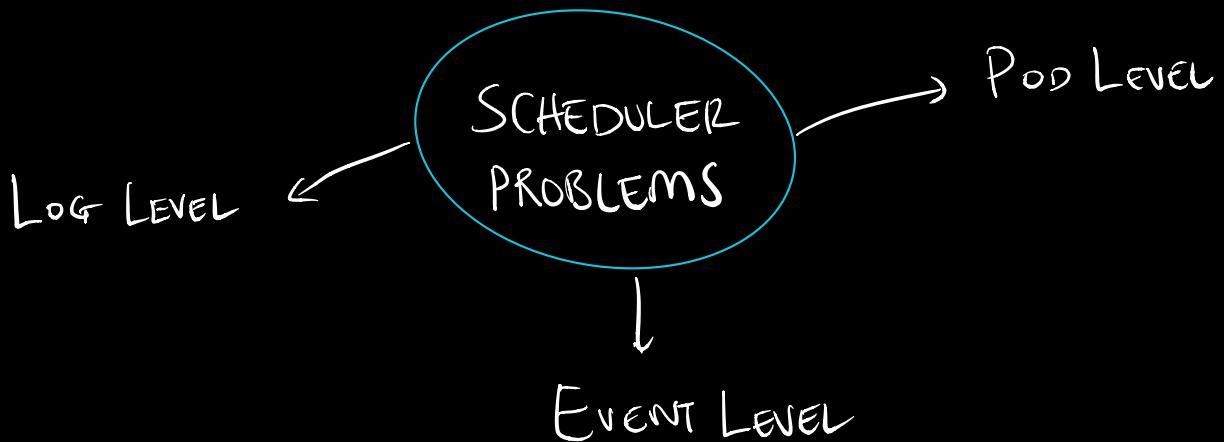
DaemonSets ensure that a single replica of a pod is running on each node at all times



POD DAEMONSET POD  
POD REPLICASET POD

If you try delete a daemonset pod, it will simply recreate it.

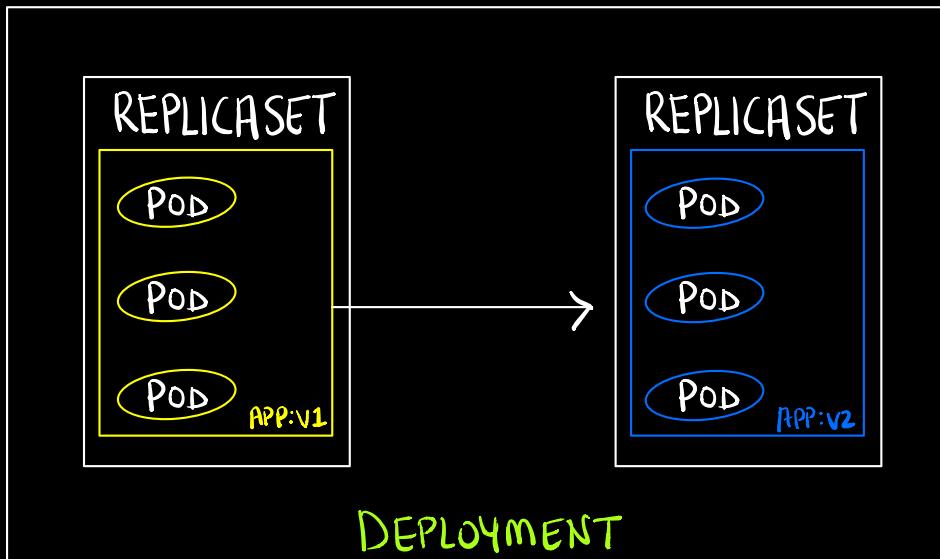
# DISPLAY SCHEDULER EVENTS



# DEPLOYING APPLICATIONS IN THE KUBERNETES CLUSTER

# DEPLOYING AN APPLICATION, ROLLING UPDATES, AND ROLLBACKS

DEPLOYMENTS → HIGH LEVEL RESOURCE FOR  
DEPLOYING AND UPDATING APPS



KUBECTL APPLY → MODIFY OBJECTS TO EXISTING YAML  
IF DEPLOYMENT NOT CREATED → ALSO CREATE

KUBECTL REPLACE → REPLACES OLD WITH NEW AND OBJECT MUST  
EXIST.

ROLLING UPDATE → PREFERRED WAY → SERVICE NOT INTERRUPTED  
→ FASTEST WAY

KUBECTL ROLLOUT → ROLL BACK PREVIOUS VERSION

# CONFIGURING AN APP FOR HA AND SCALE

AVOIDING  
BAD  
VERSIONS

→ BLOCK BAD VERSION  
RELEASE

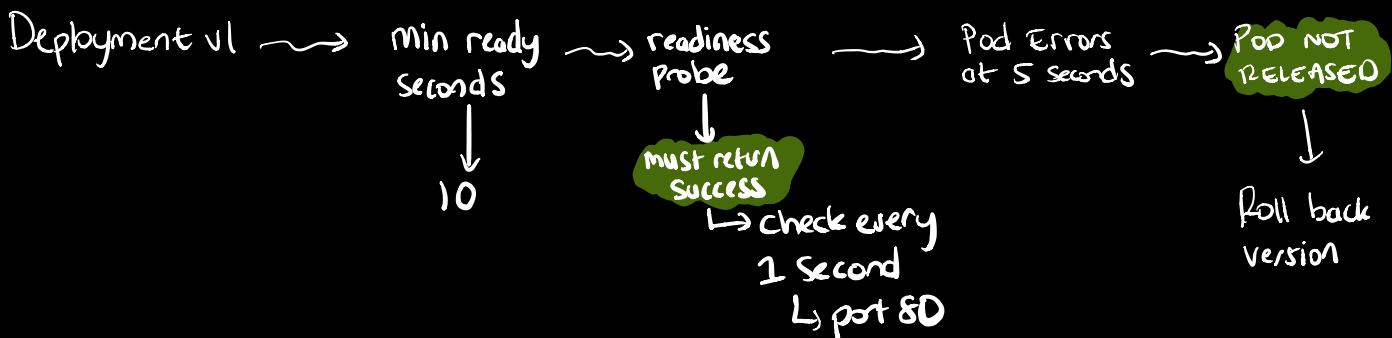
MIN READY  
SECONDS

How long a  
newly created  
pod should be  
ready before  
considered  
available



READINESS  
PROBE

Determines  
if a specific  
pod should  
receive  
client request  
or not.



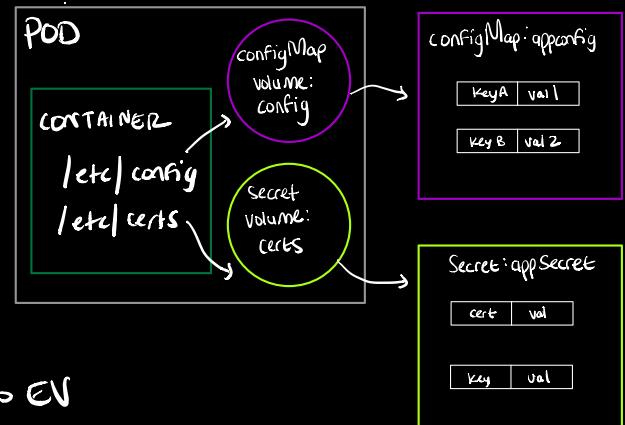
PASSING CONFIGURATION OPTIONS TO APP

ENVIRONMENT VARIABLES

STORE IN CONFIG MAP  
CREATE SECRET & PASS TO EV

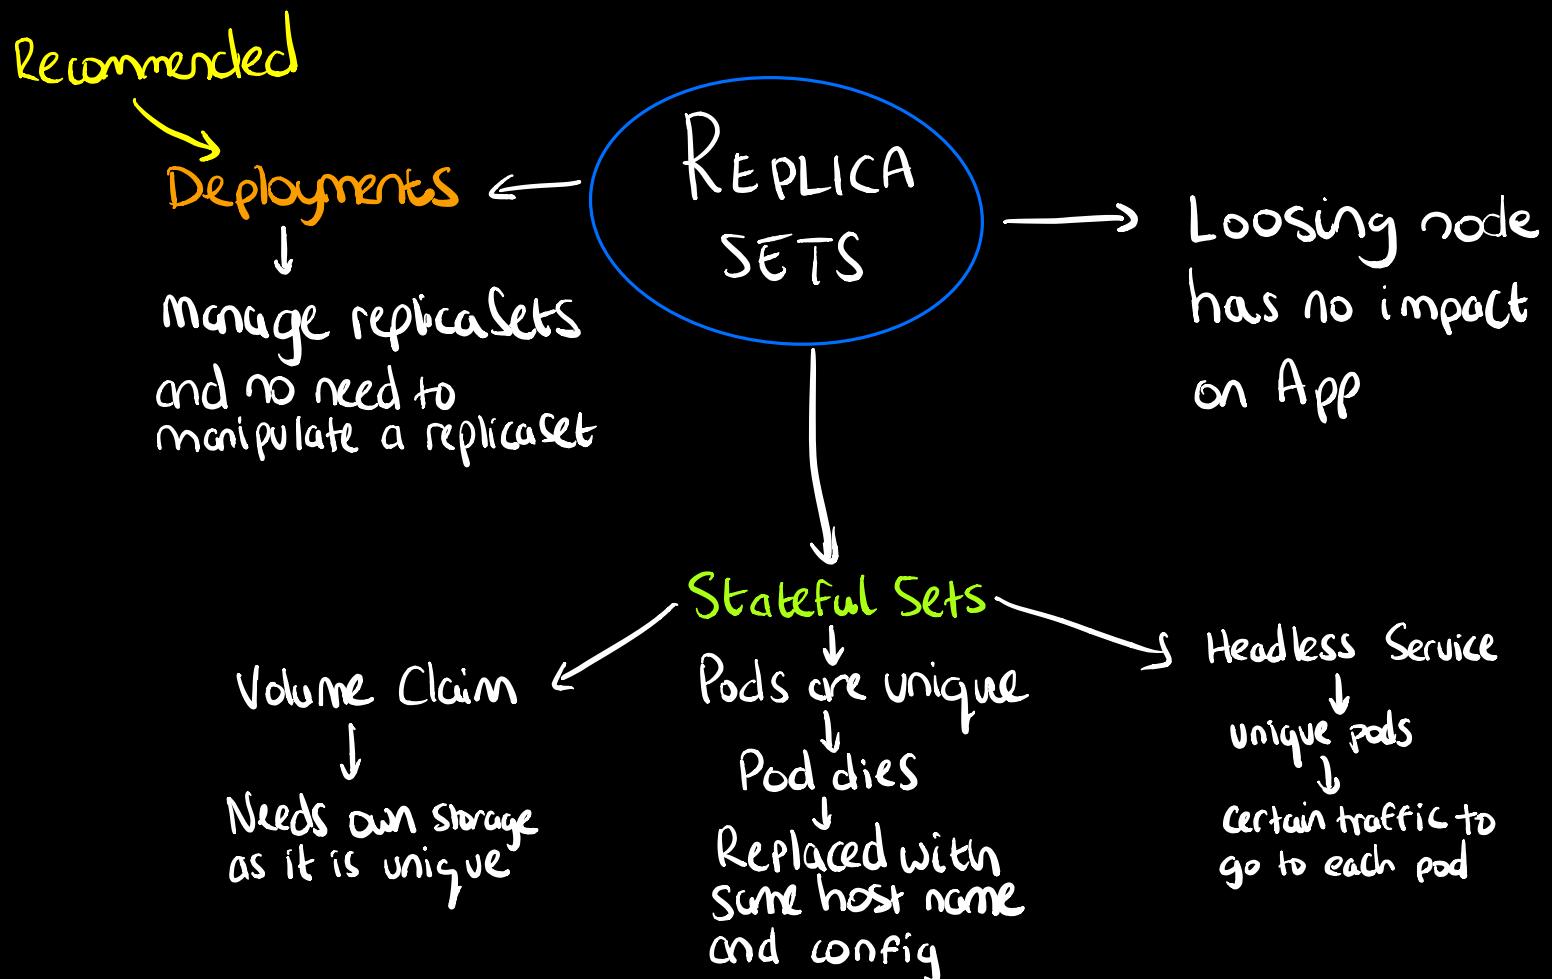
JUST UPDATE → NO NEED TO REBUILD  
IMAGE

MULTIPLE  
CONTAINERS  
CAN USE  
SAME



# CREATING A SELF-HEALING APP

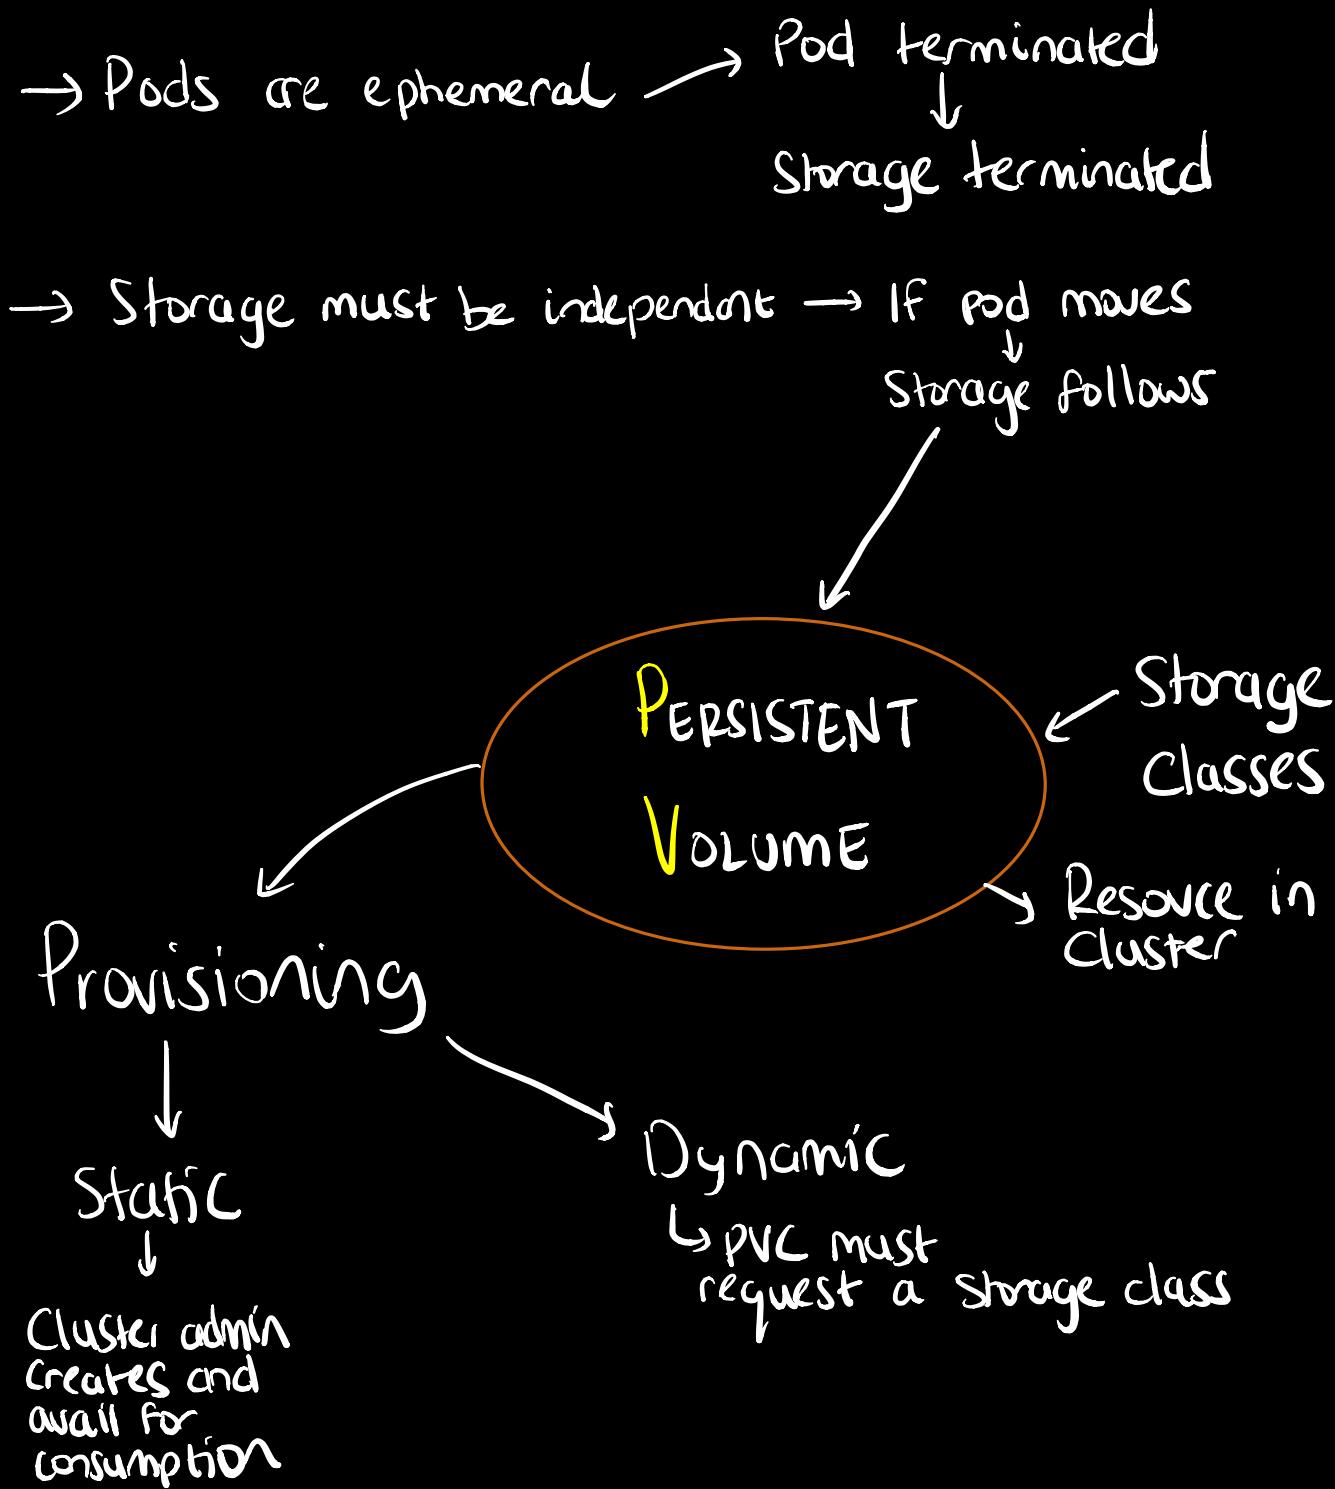
ReplicaSets ensure that a identically configured pods are running at the desired replica count





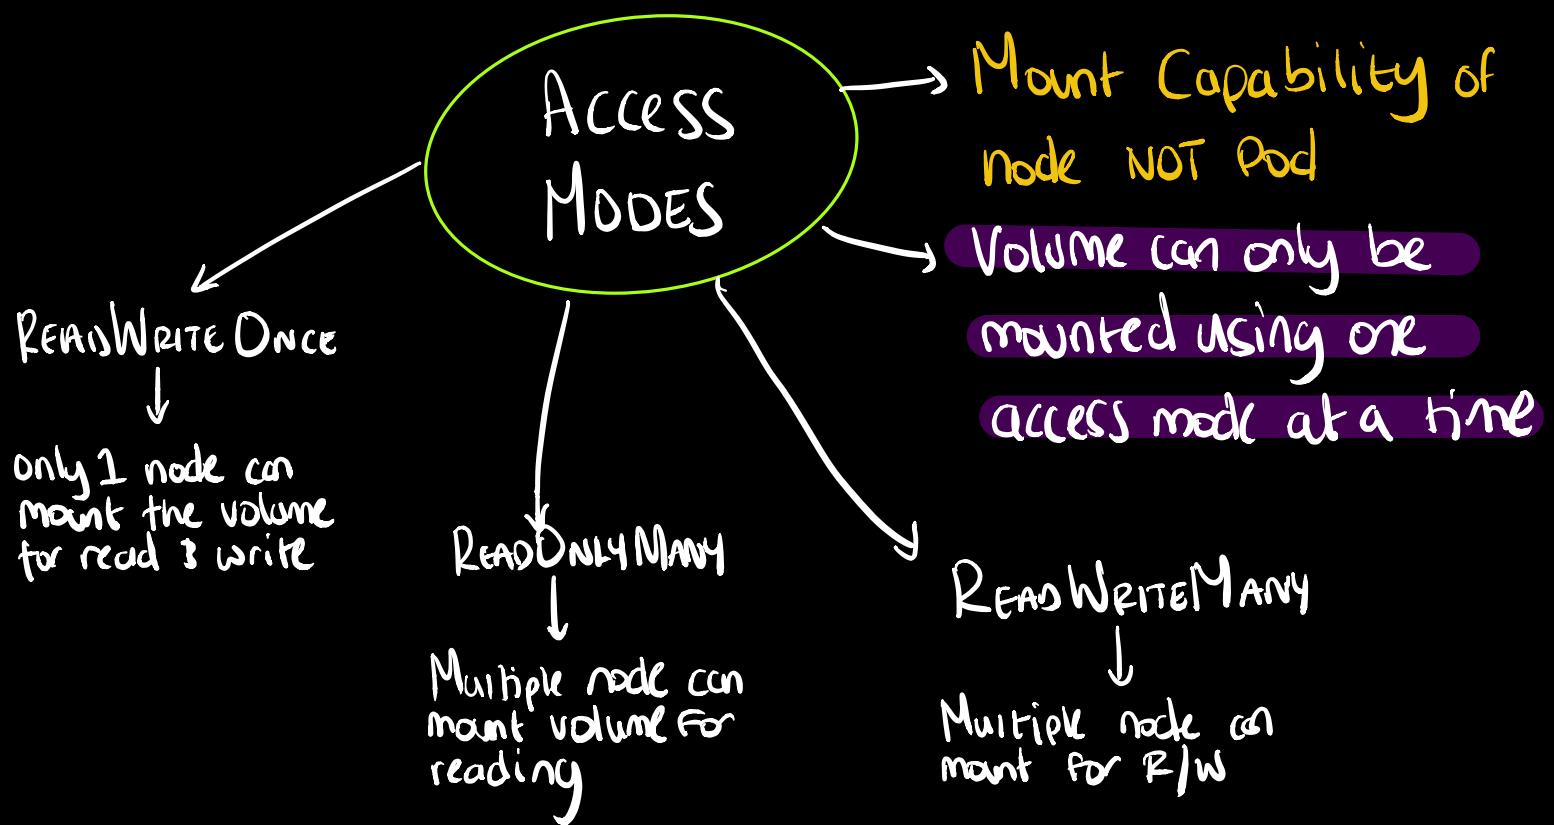
# MANAGING DATA IN KUBERNETES CLUSTER

# PERSISTENT VOLUMES



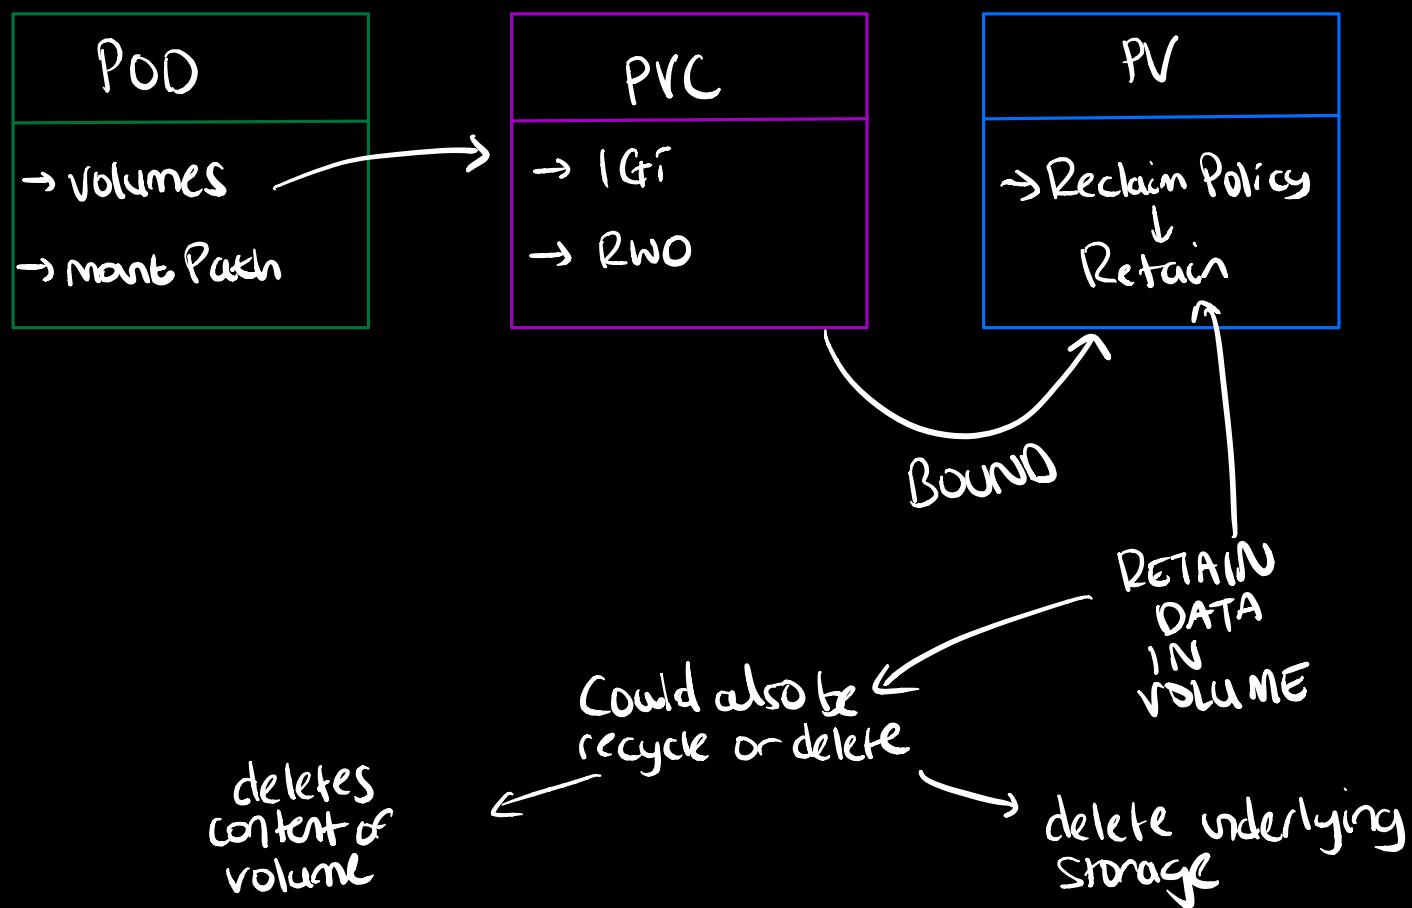
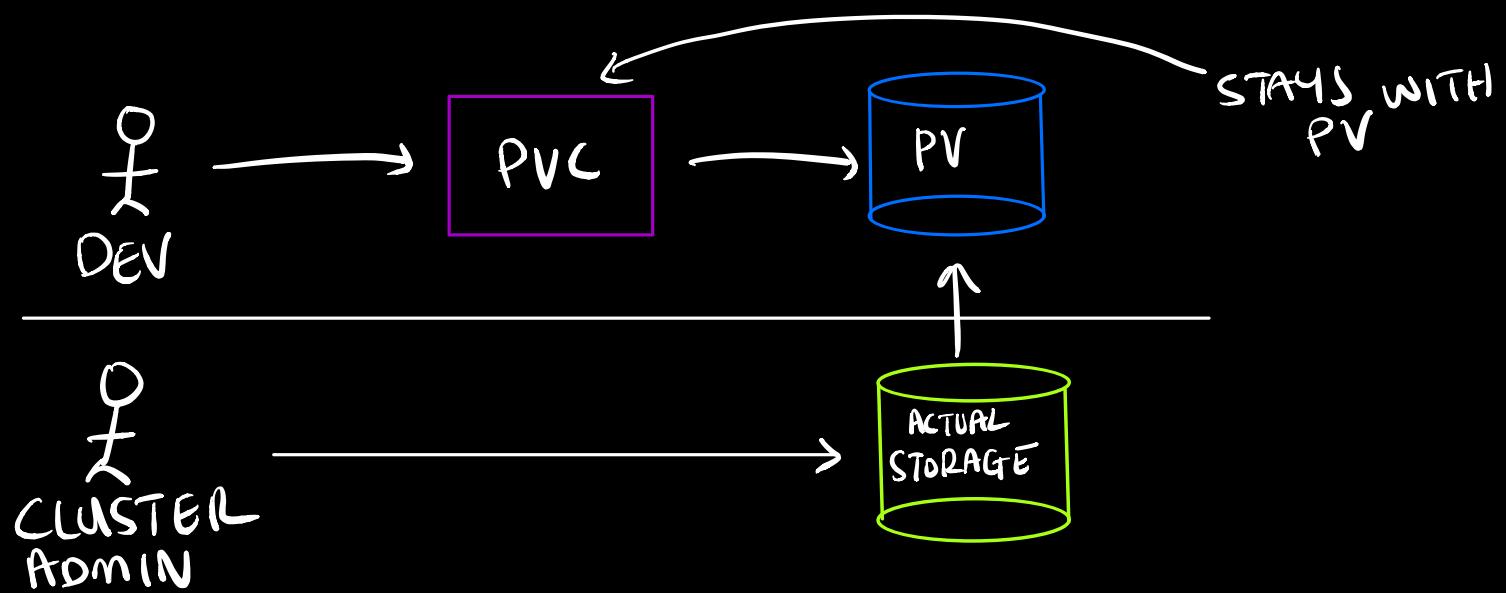
# VOLUME ACCESS MODES

By specifying an access mode with your PV,  
you allow the volume to be mounted to one  
or many nodes, as well as read by one or many



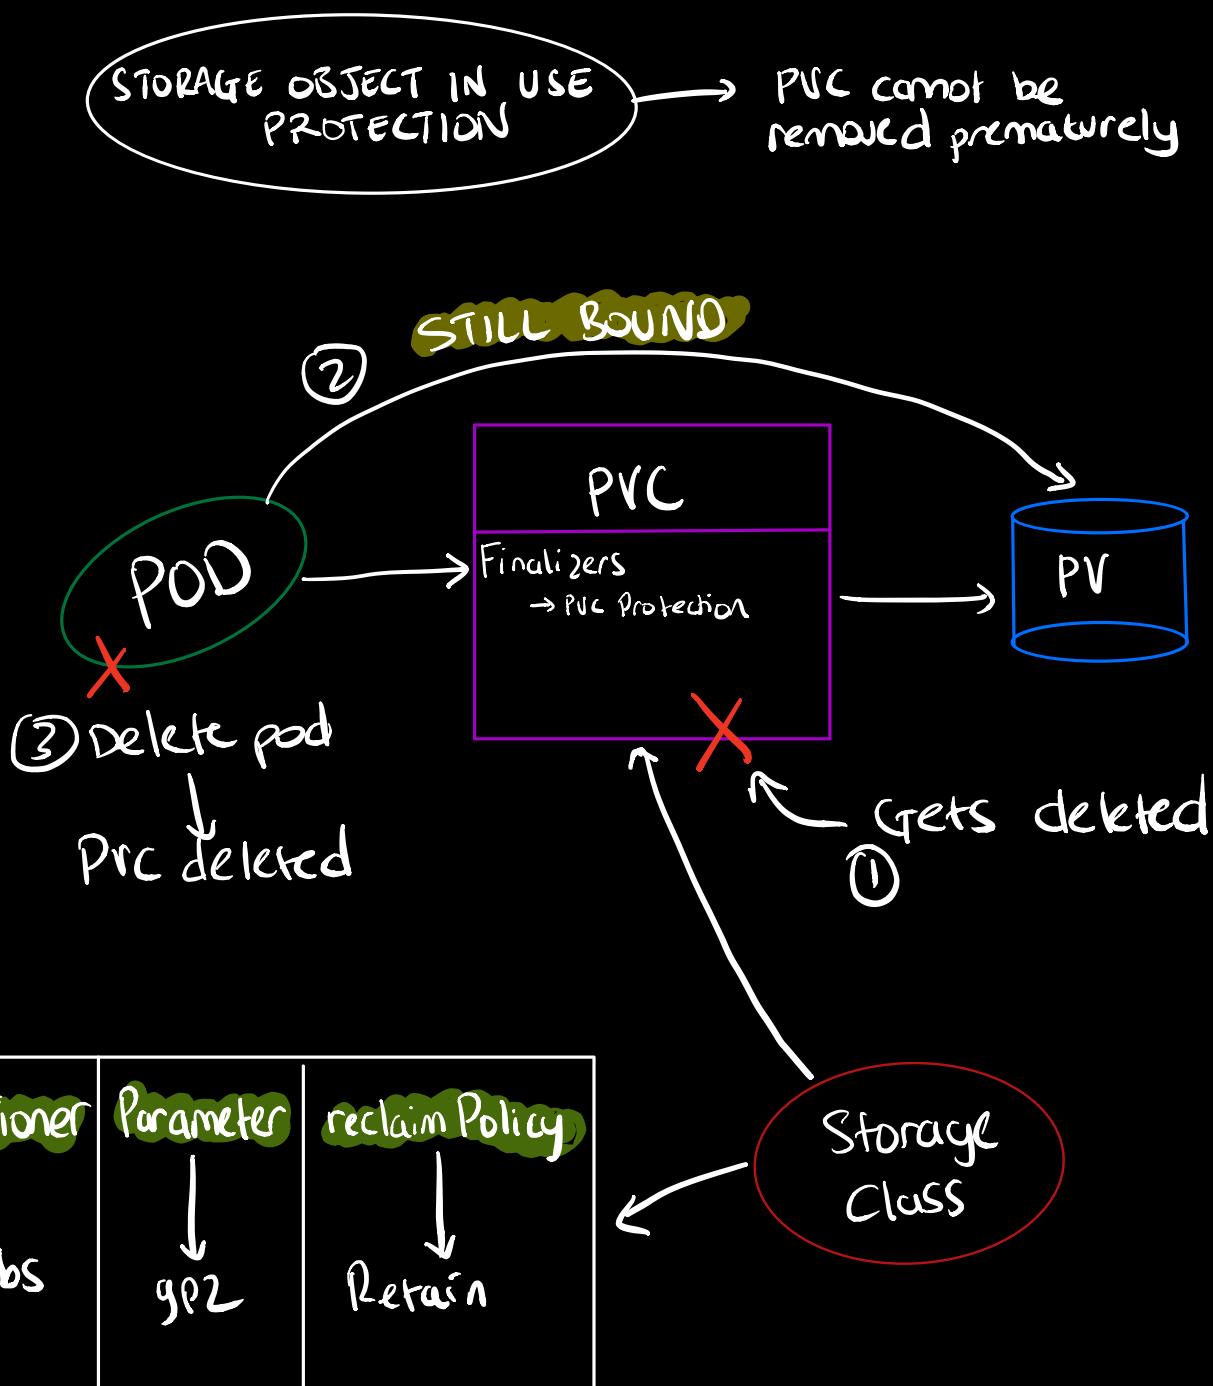
# PERSISTENT VOLUME CLAIMS (PVC)

PVC allows the application developer to request storage for the application, without having to know underlying infra.



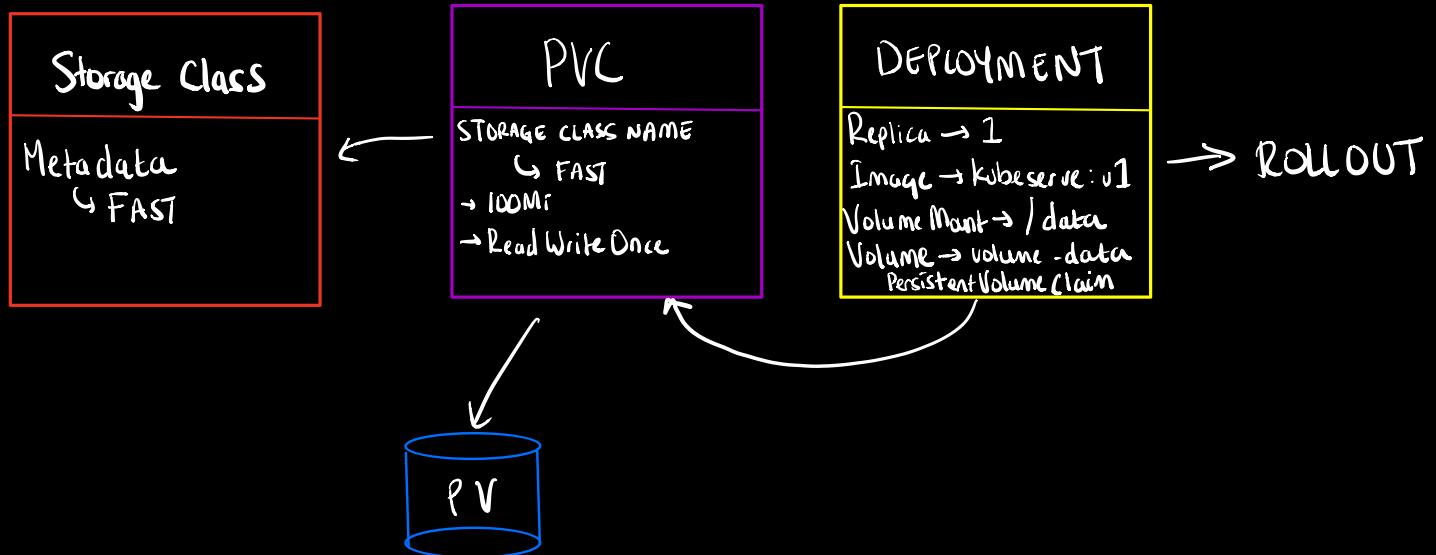
# STORAGE OBJECTS

VOLUMES that are already in use by a pod are protected against data loss. This means even if you delete a PVC, you can still access volume from pod.



# APPLICATIONS WITH PERSISTENT STORAGE

Example →

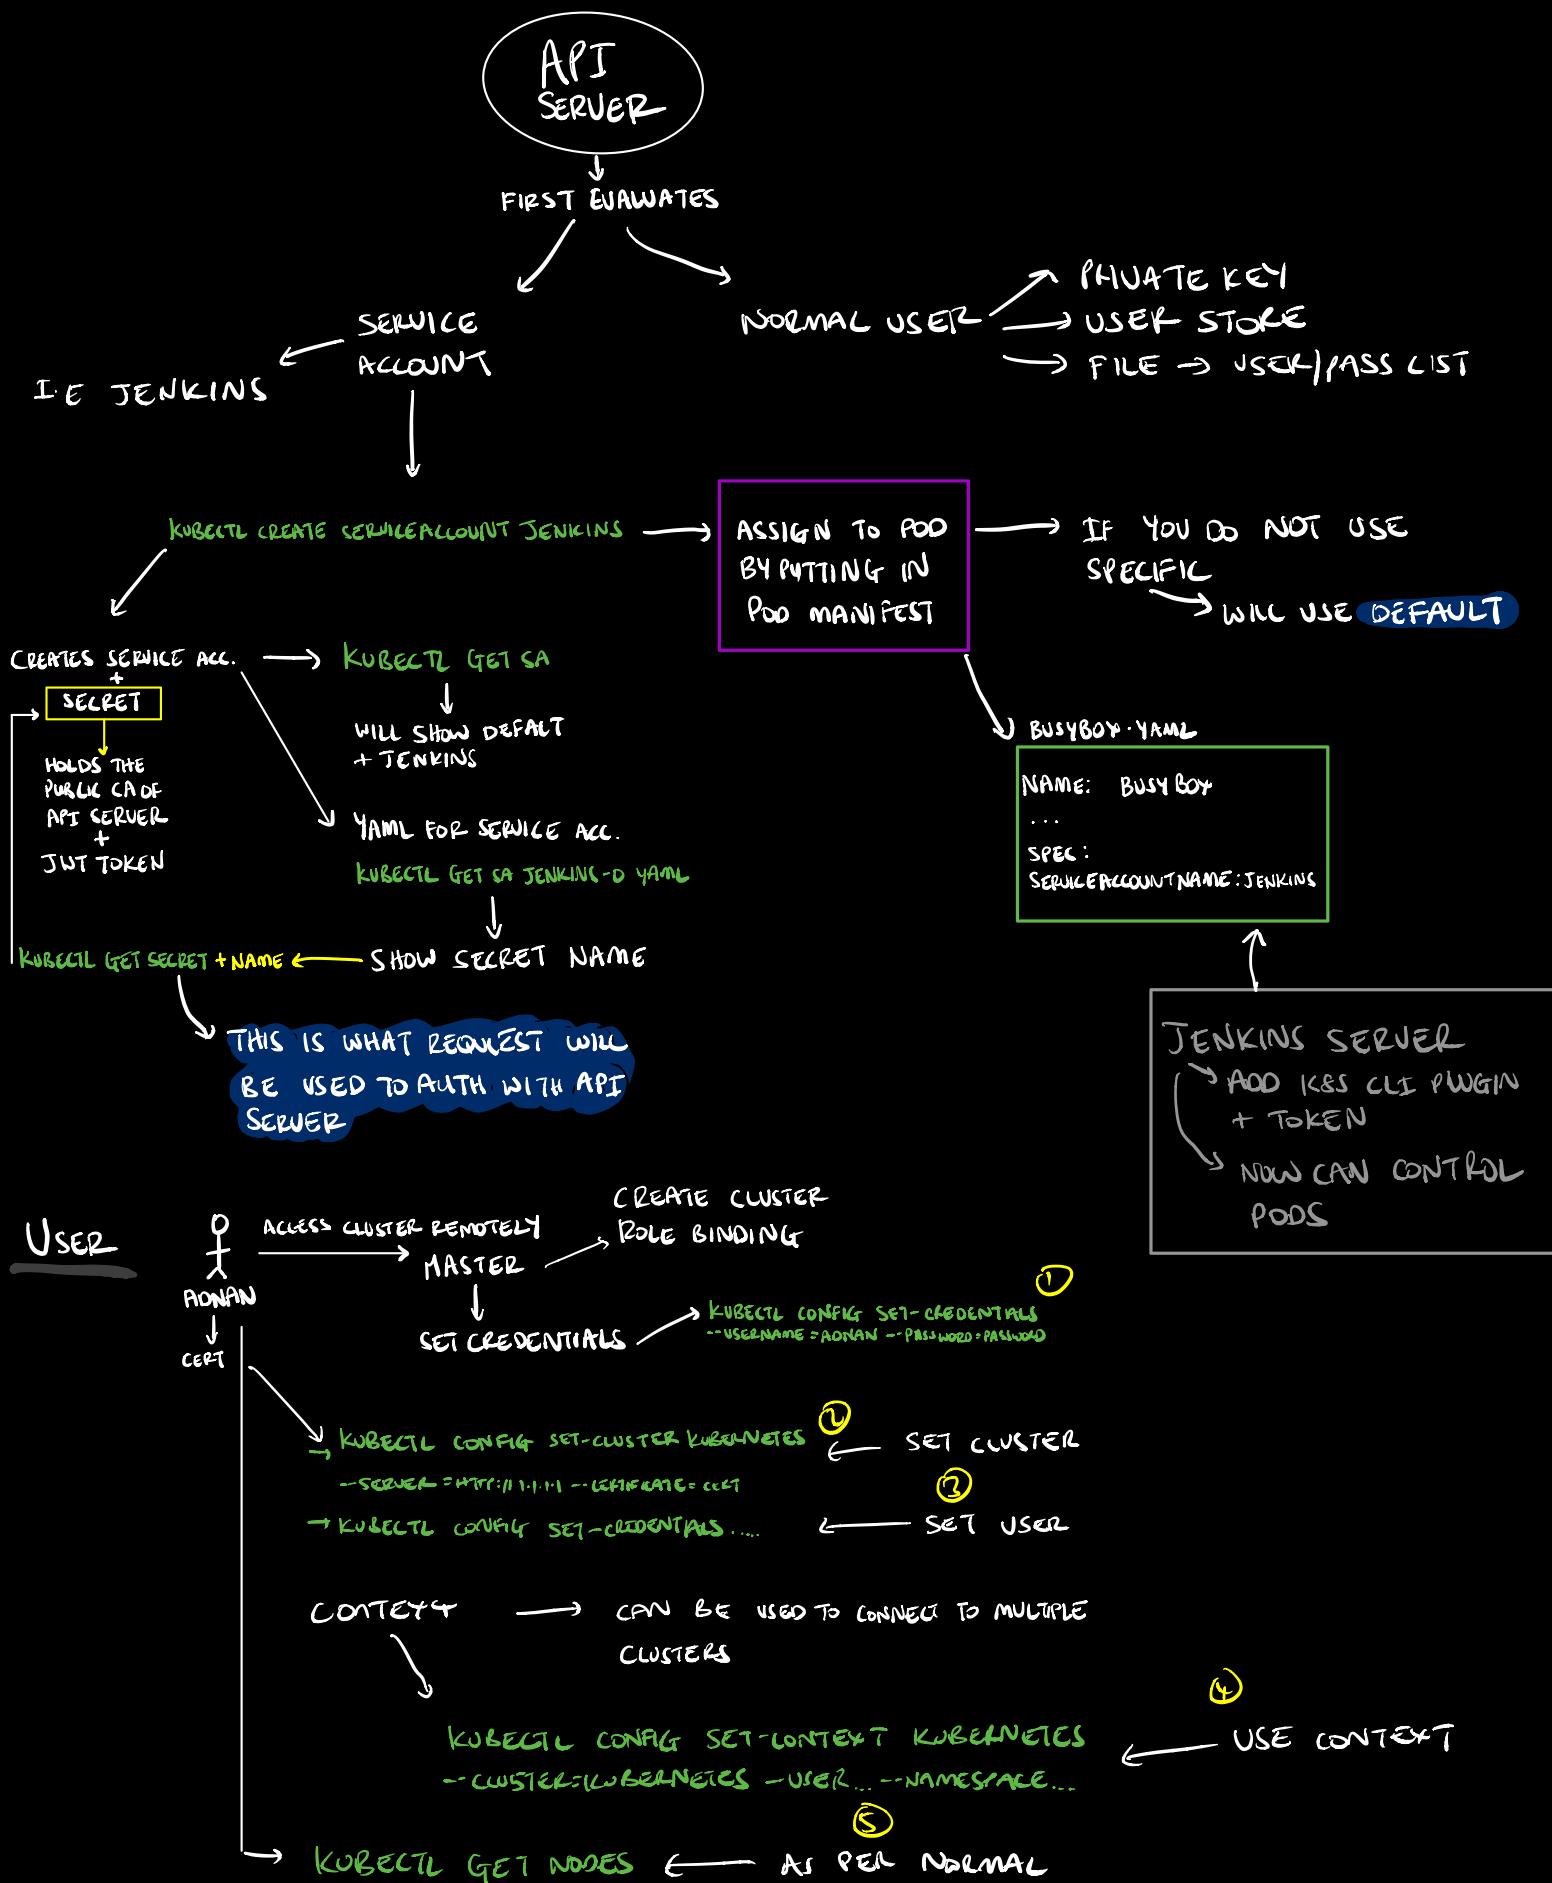


- ① create storage class object
- ② create PVC object
- ③ create deployment
- ④ rollout deployment
- ⑤ check pods
- ⑥ create file on mount
- ⑦ List contents.

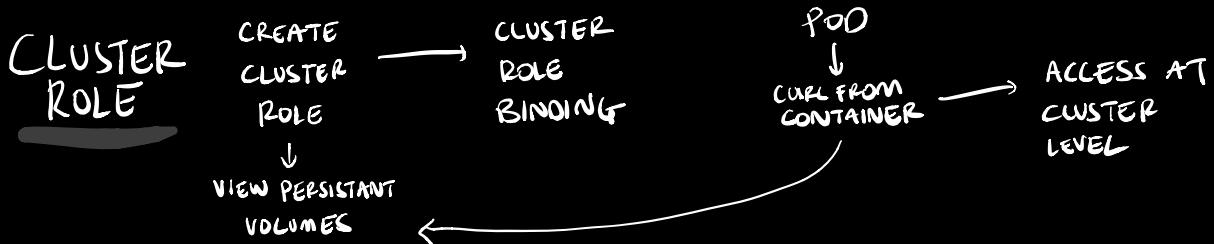
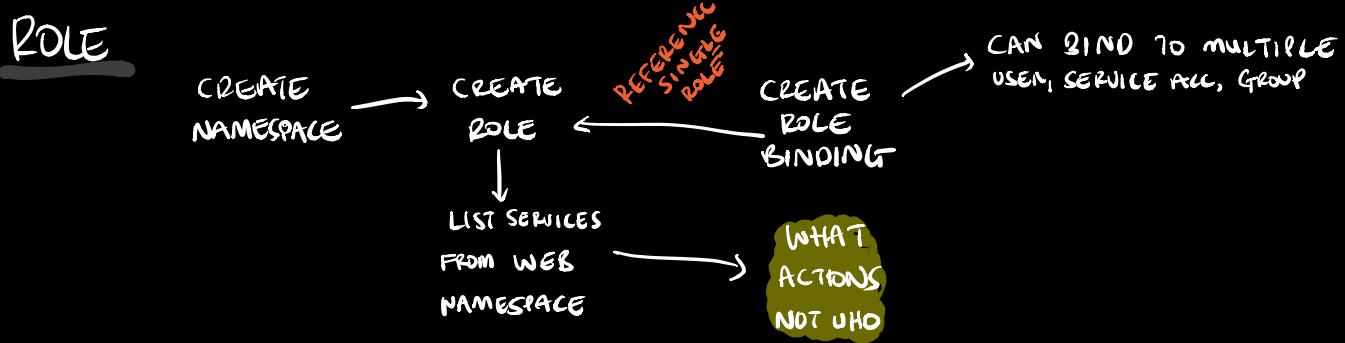
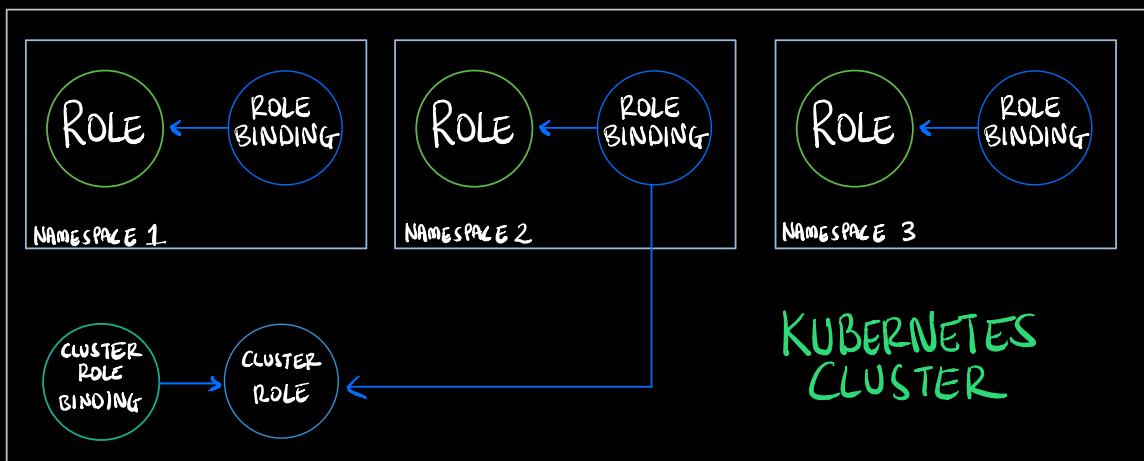
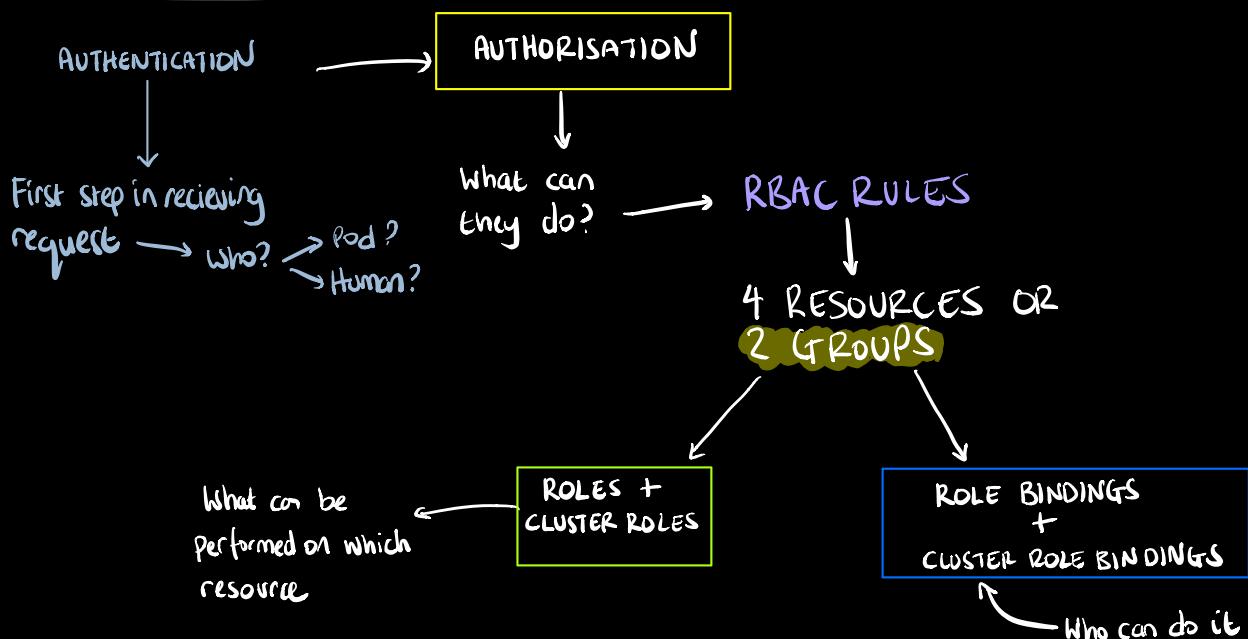


# SECURING THE KUBERNETES CLUSTER

# SERVICE ACCOUNTS AND USERS

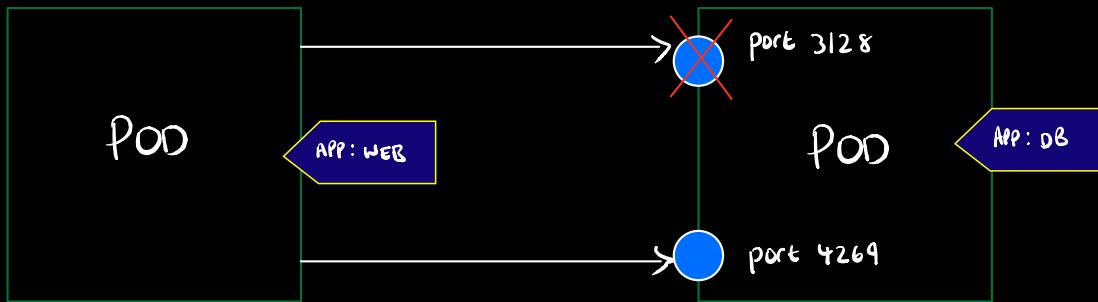


# CLUSTER AUTHENTICATION AND AUTHORISATION

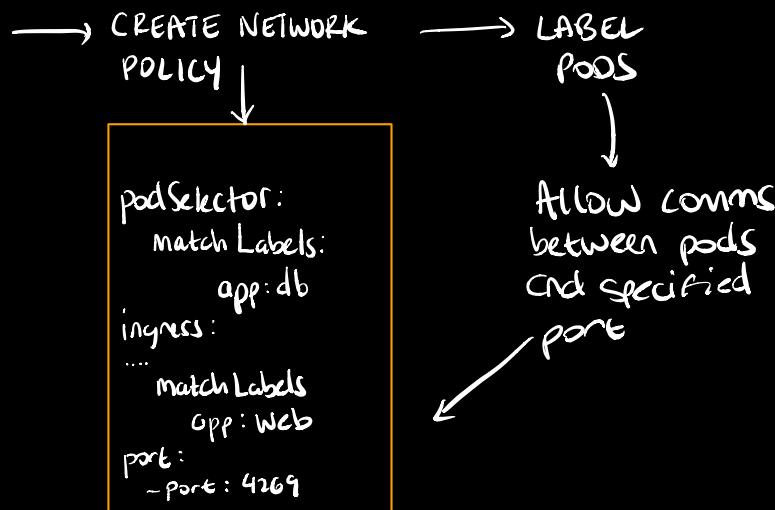
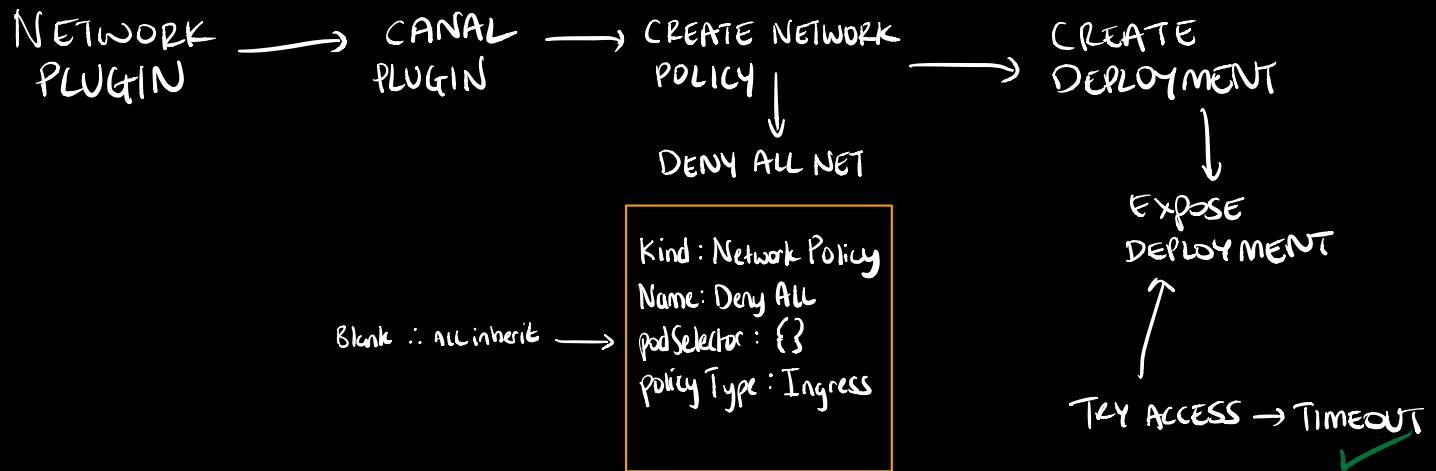


# CONFIGURING NETWORK POLICIES

Network policies use selectors to apply rules to pods for communication throughout the cluster

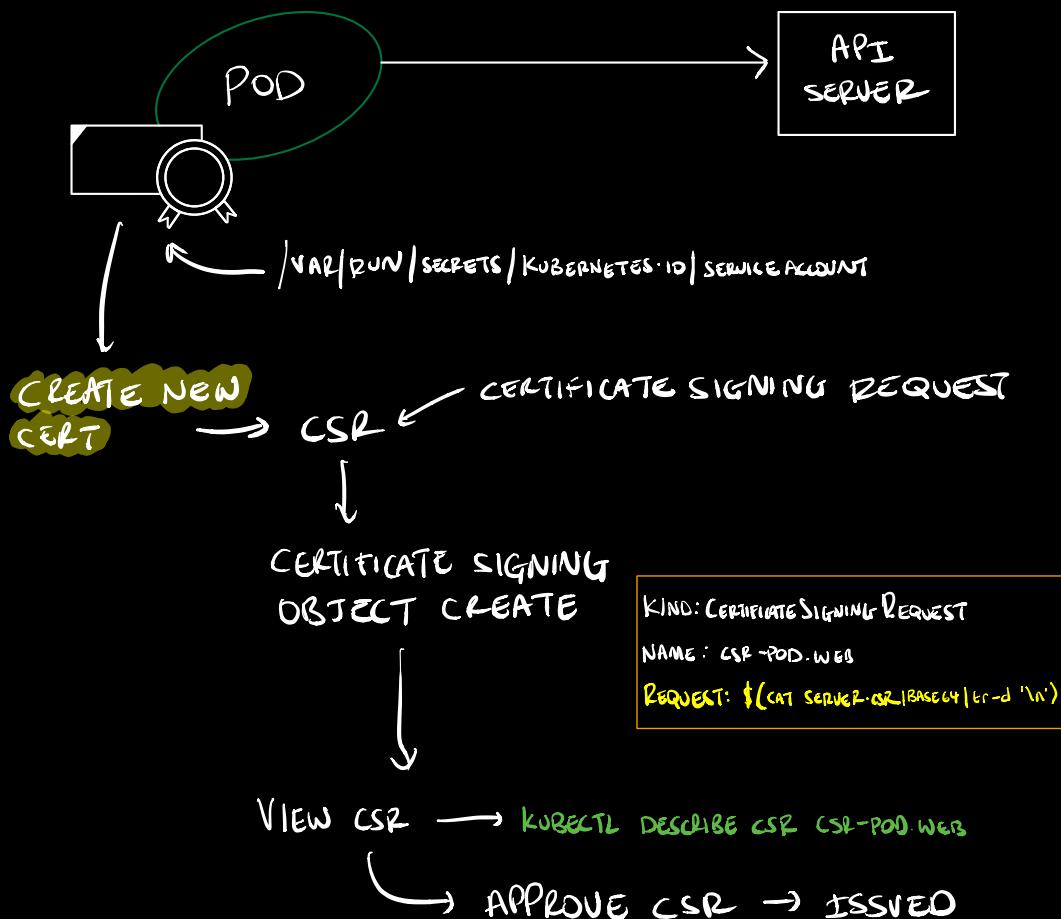


How?



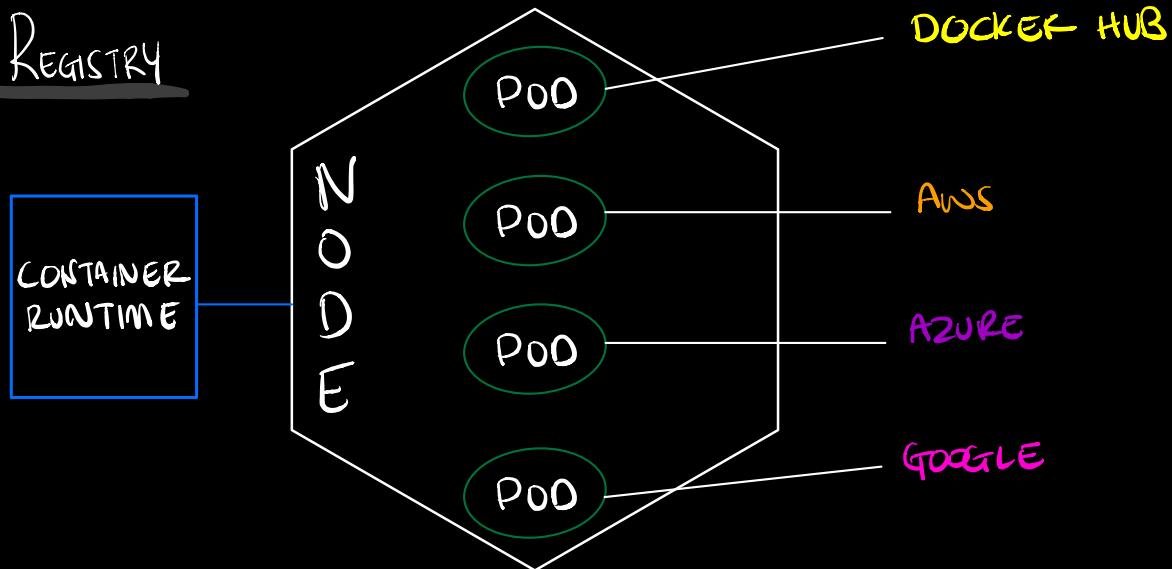
# CREATING TLS CERTIFICATES

The CA is used to generate a TLS certificate and authenticate with the API Server.



# SECURE IMAGES

## PRIVATE REGISTRY



CONTROLLING  
IMAGES THAT  
GO INTO  
PRODUCTION

VULNERABILITIES → CLAIR (SCANNING)

SOMETHING ON IT  
CAUSE NODE CRASH

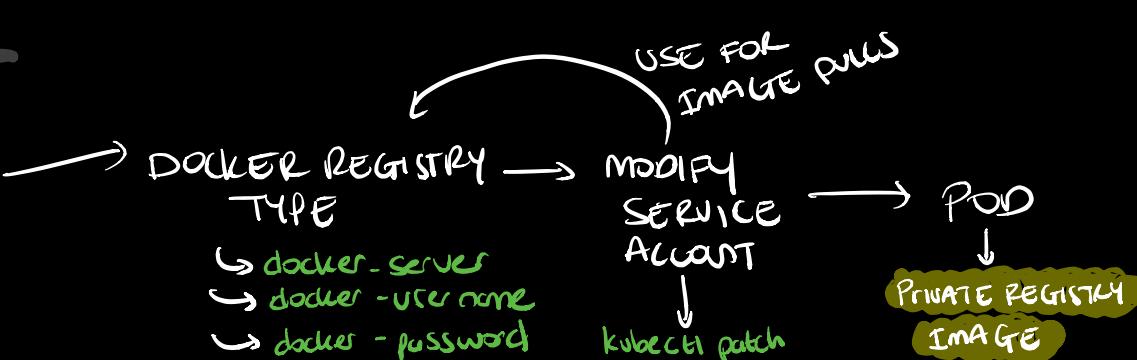
LOGIN TO PRIVATE REGISTRY

↳ TAG DOCKER IMAGE

↳ PUSH TO PRIVATE REGISTRY

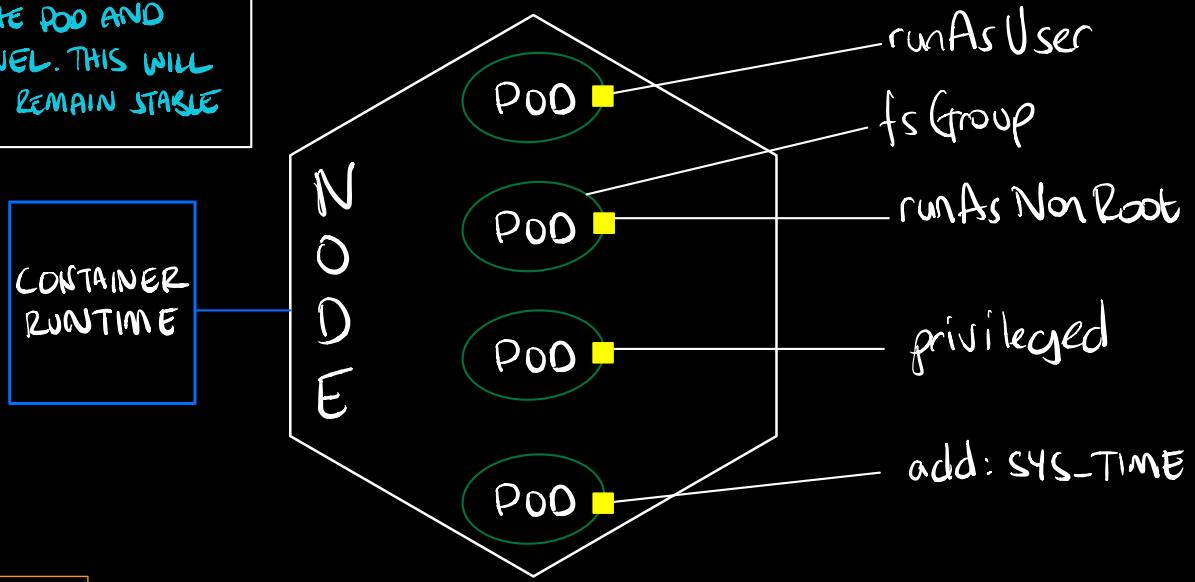
## How?

KUBERNETES  
CREATE  
SECRET



# DEFINING SECURITY CONTEXT

LIMIT ACCESS TO CERTAIN OBJECTS AT THE POD AND CONTAINER LEVEL. THIS WILL ALLOW IMAGES TO REMAIN STABLE



Kind: Pod  
image: alpine  
Security Context:  
runAsUser: 405

Run Pod as 405

Can also put 'runAsNonRoot'

Ability to run as privileged → 'privileged: true'

## CONTAINER LEVEL

ABILITY TO LOCK DOWN KERNEL LEVEL FEATURES ON CONTAINER

SETTING CAPABILITIES ON POD LEVEL

ADD

Security Context:  
add:  
- SYS\_TIME  
- NET\_ADMIN

REMOVE

Security Context:  
drop:  
CHOWN

# SECURING PERSISTENT KEY/VALUE STORE

SECRETS ALLOW YOU TO EXPOSE ENTRIES AS FILES  
IN A VOLUME. KEEPING THIS DATA  
SECURE IS CRITICAL TO CLUSTER SECURITY

DATA MUST  
LIVE BEYOND  
LIFE OF POD → SECRETS → KEY/VALUE PAIR

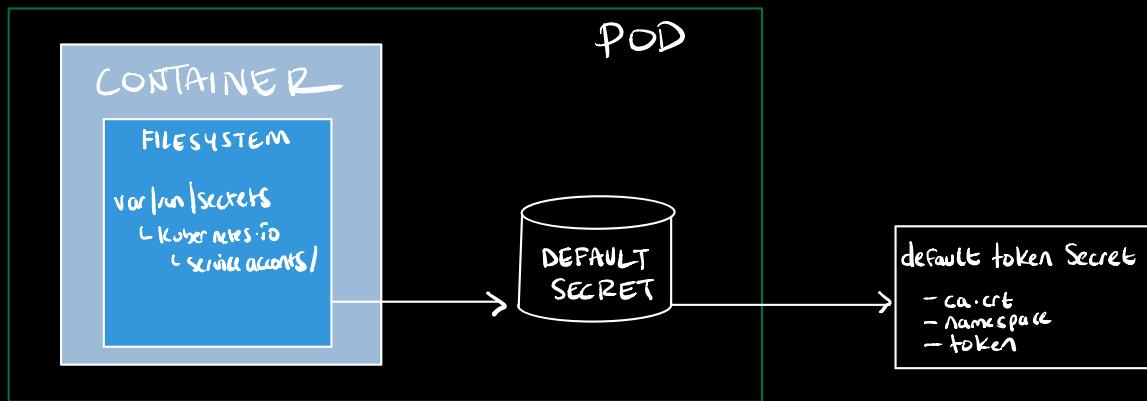
↓  
PASS AS ENV VAR

OR

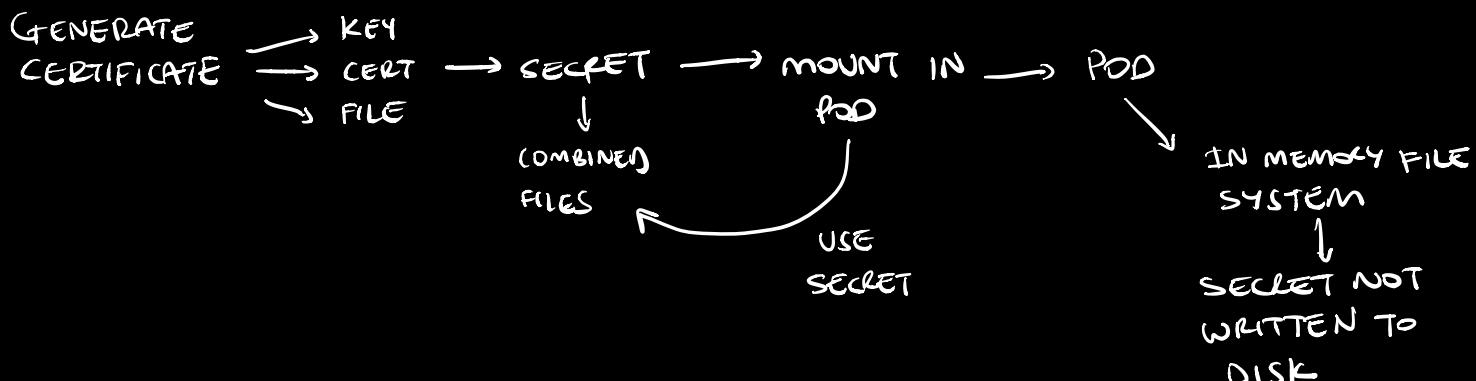
EXPOSE AS FILES  
IN VOLUME

NOT BEST  
PRACTICE

↓  
MAY BE  
OUTPUT TO  
LOG FILES



## HTTPS TO WEBSITE

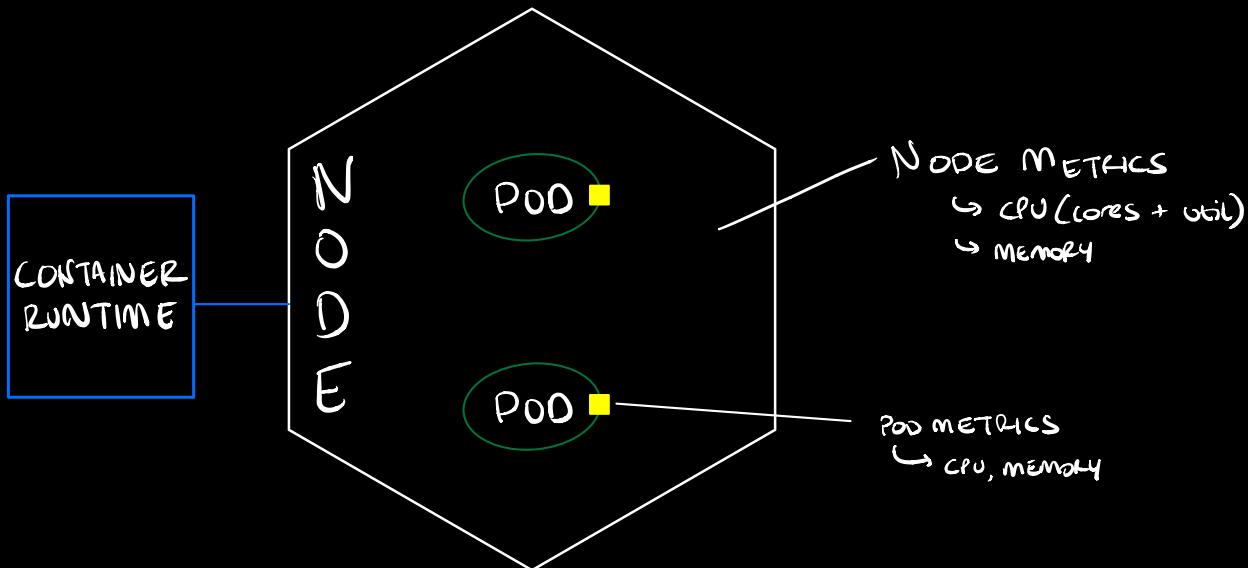




# MONITORING CLUSTER COMPONENTS

# MONITORING THE CLUSTER COMPONENTS

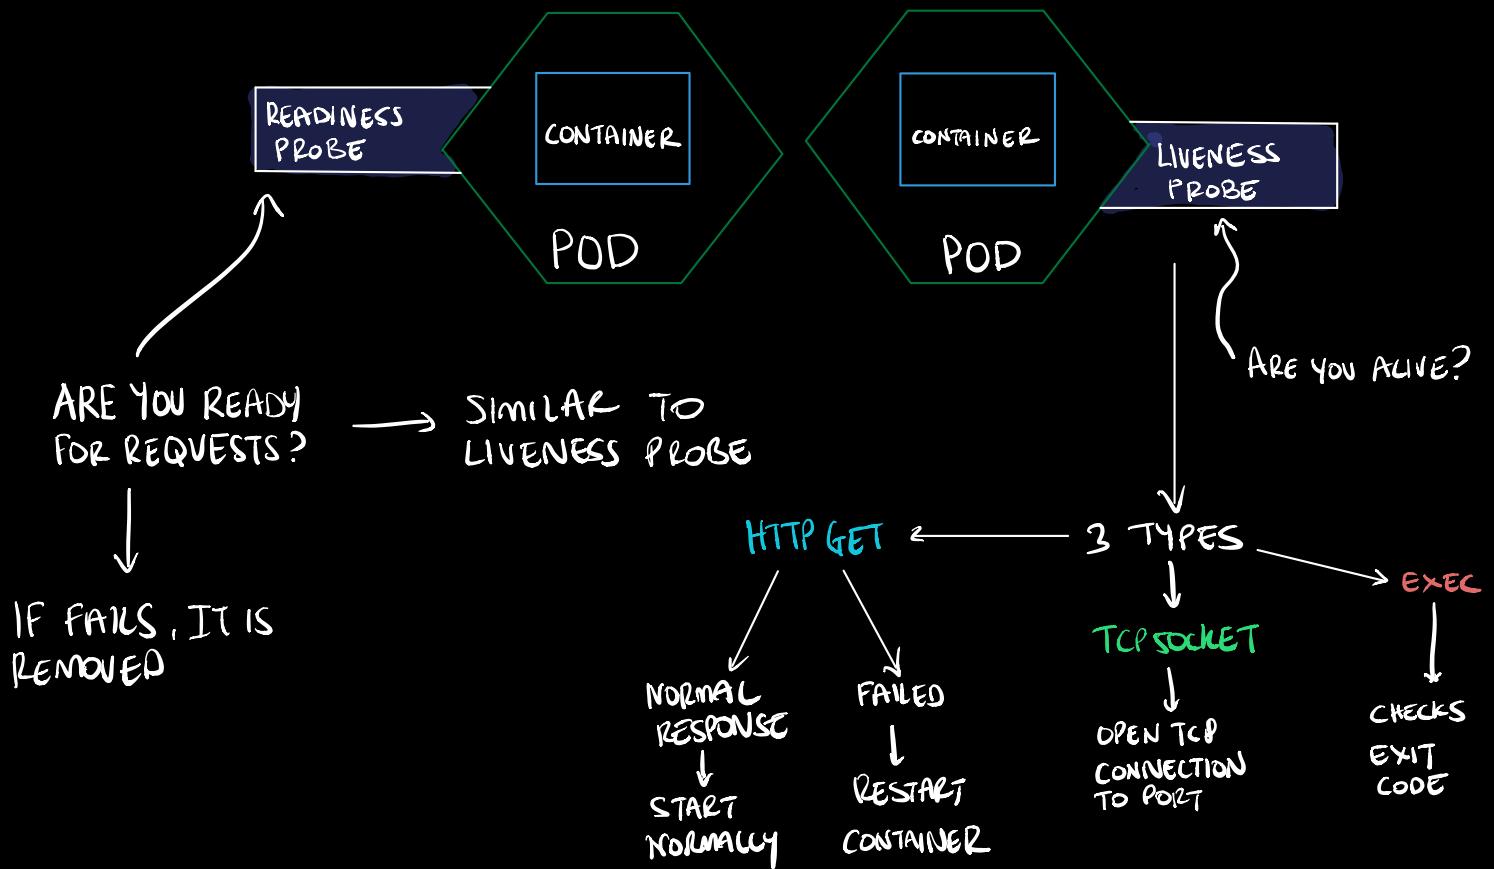
THE METRIC SERVER ALLOWS YOU TO COLLECT CPU AND MEMORY DATA FROM THE NODES AND PODS IN YOUR CLUSTER



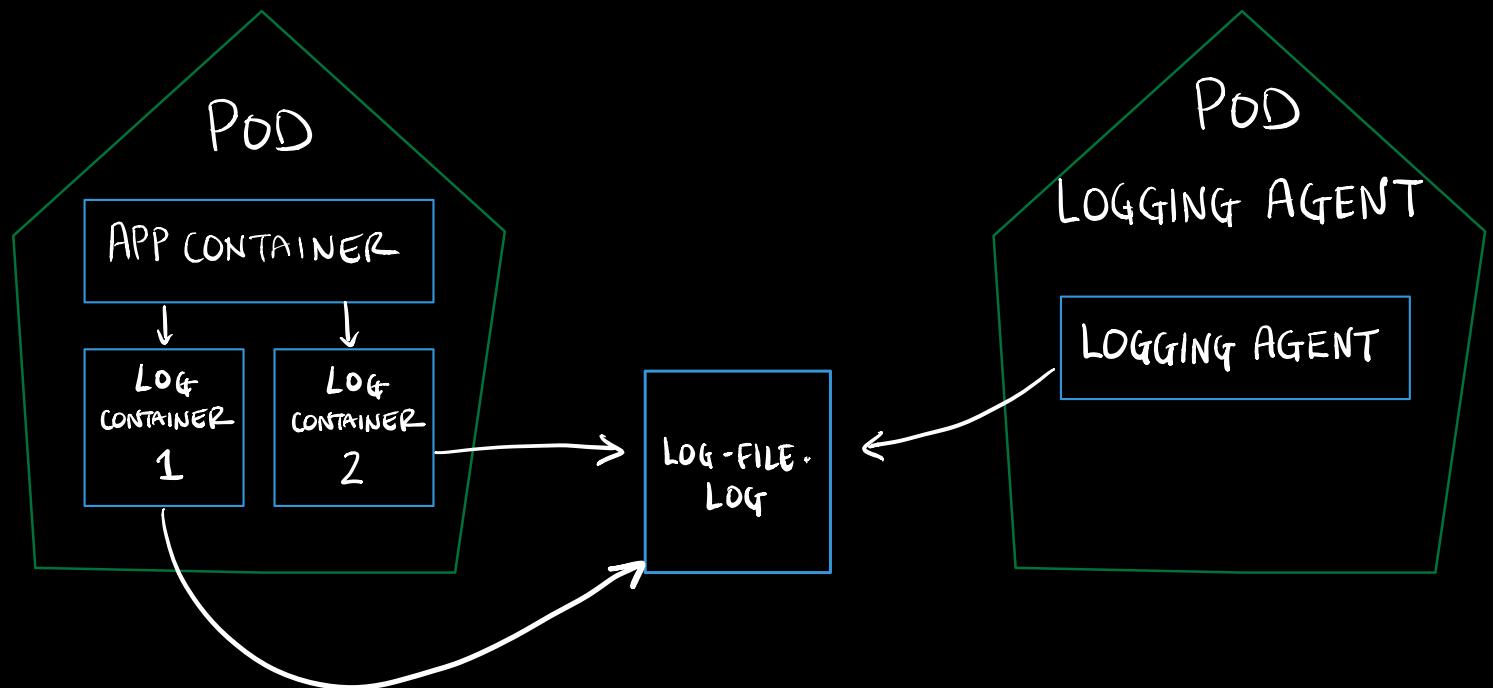
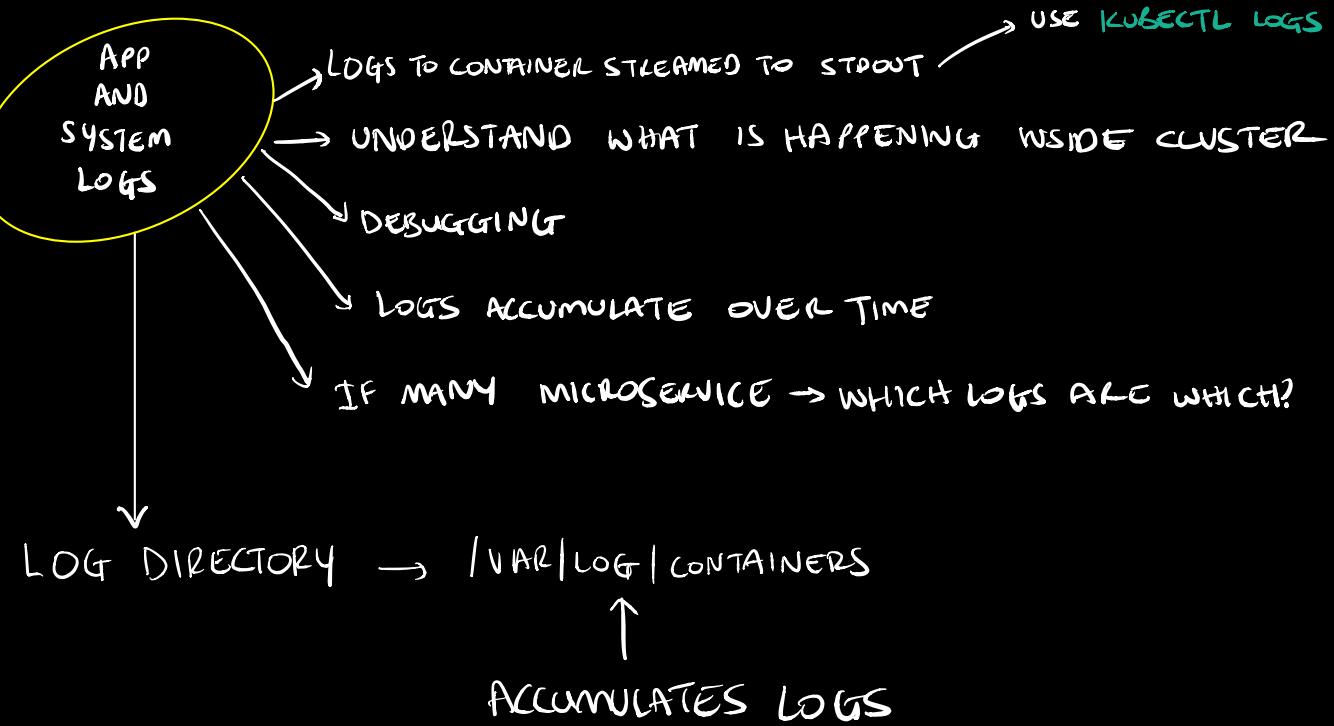
INSTALL METRIC SERVER → KUBECTL TOP NODE → CPU/MEMORY FOR ALL THE NODES  
KUBECTL TOP POD → CPU/MEMORY FOR ALL THE PODS  
KUBECTL TOP POD --ALL-NAMESPACES → ALL NAMESPACES  
KUBECTL TOP POD -N KUBE-SYSTEM → KUBE-SYSTEM NAMESPACE  
KUBECTL TOP GROUP-CONTEXT --CONTAINERS → PODS CONTAINERS

# MONITORING THE APPS RUNNING WITHIN A CLUSTER

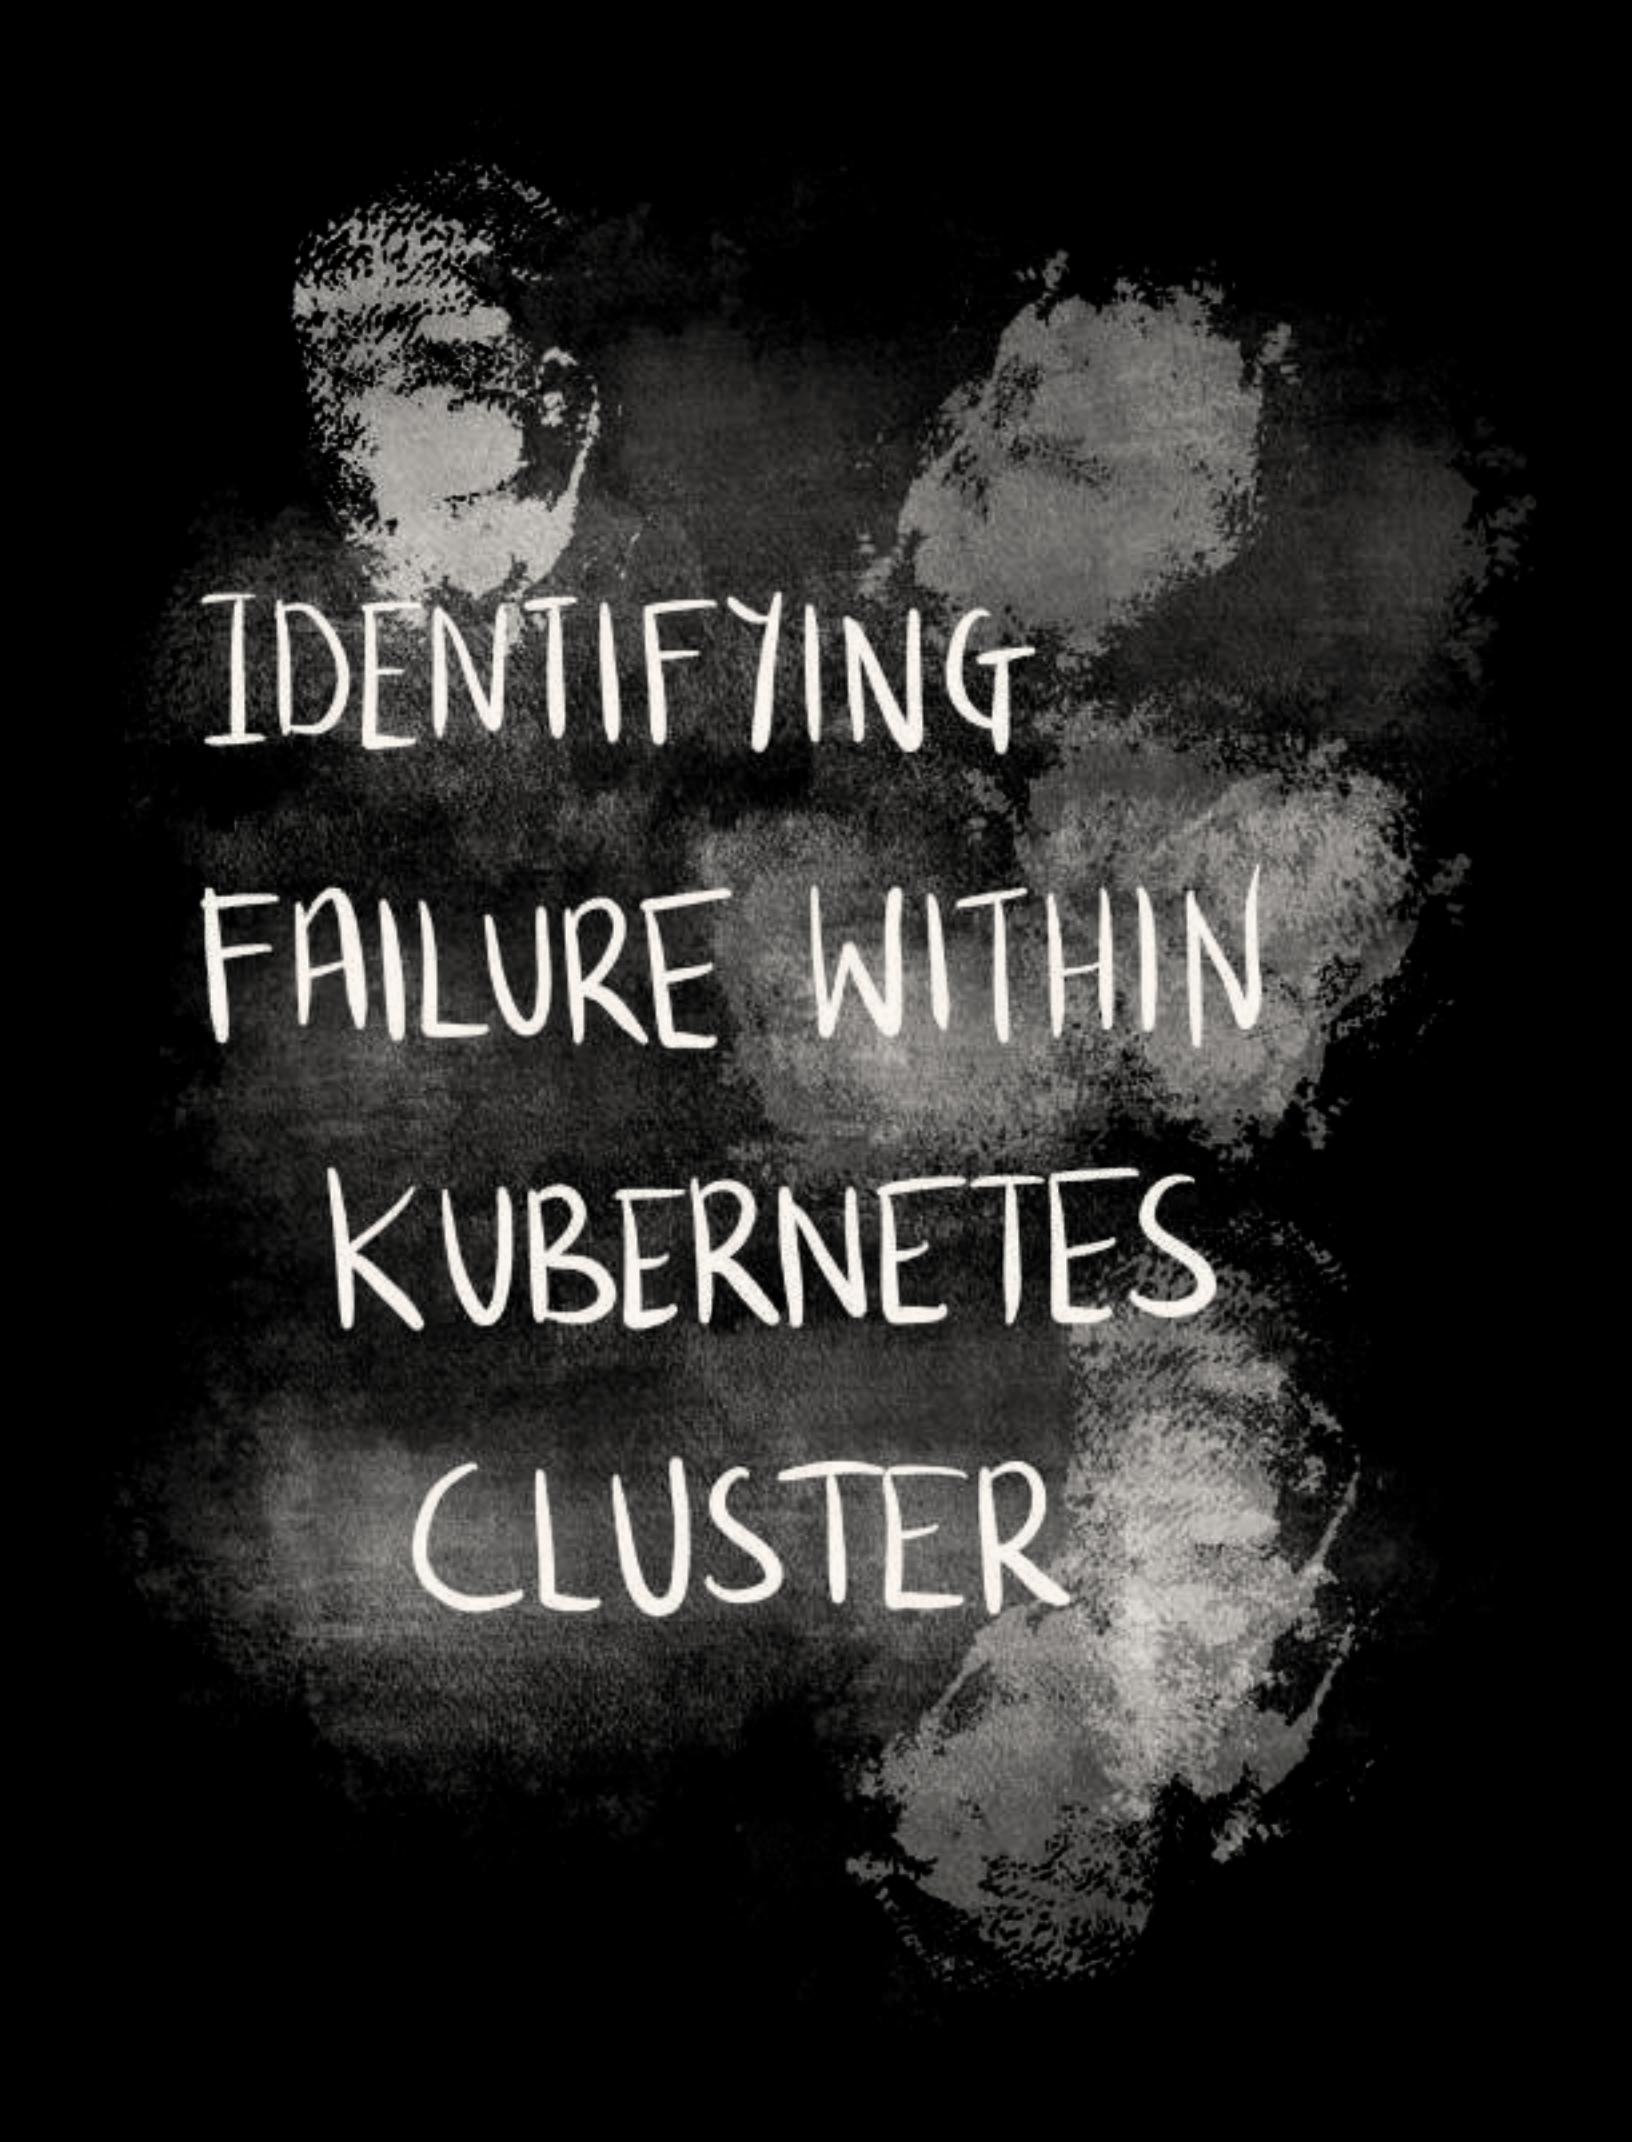
## LIVENESS AND READINESS PROBE



# MANAGING CLUSTER COMPONENT LOGS



- HAVE SIDEKAR CONTAINER TO DO LOGGING SO YOU CAN ACCESS SPECIFIC LOGS
- ABLE TO ROTATE LOGS USING OTHER TOOLS → NO NATIVE



# IDENTIFYING FAILURE WITHIN KUBERNETES CLUSTER

# TROUBLESHOOTING APPLICATION FAILURE

ABILITY TO WRITE  
TERMINATION MESSAGE  
TO SPECIFIC FILE ON  
CONTAINER



POD1.YAML

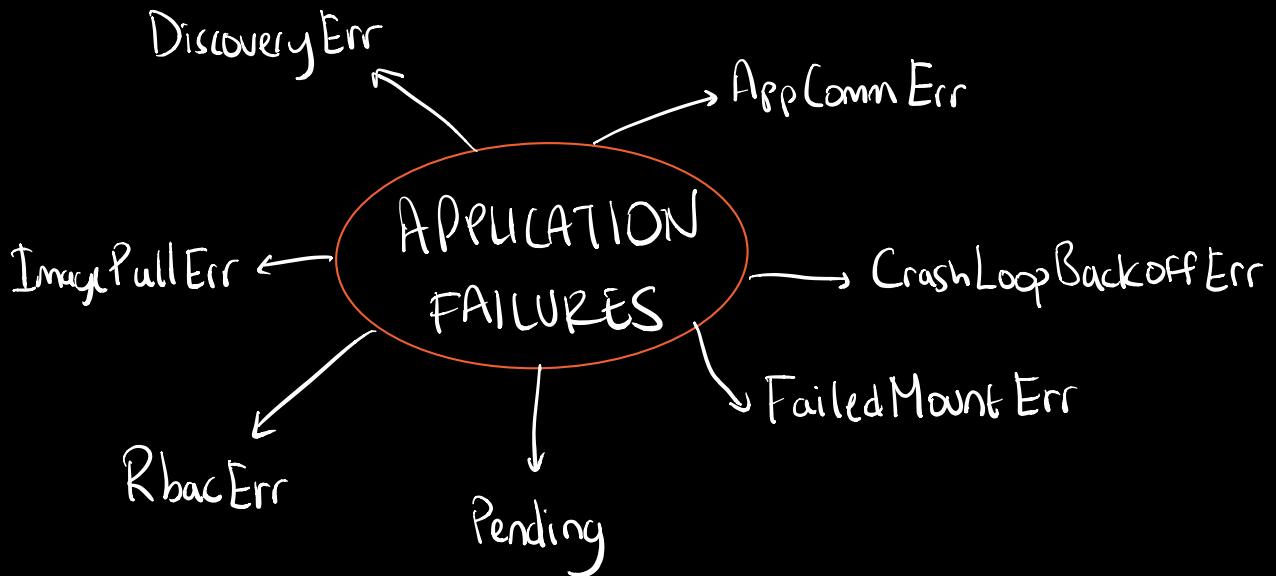
```
KIND: POD
NAME: POD1
IMAGE: BusyBox
COMMAND:
...
TERMINATION MESSAGE PATH
```

→ KUBECTL  
DESCRIBE  
↓  
WILL SHOW  
ERROR MESSAGE

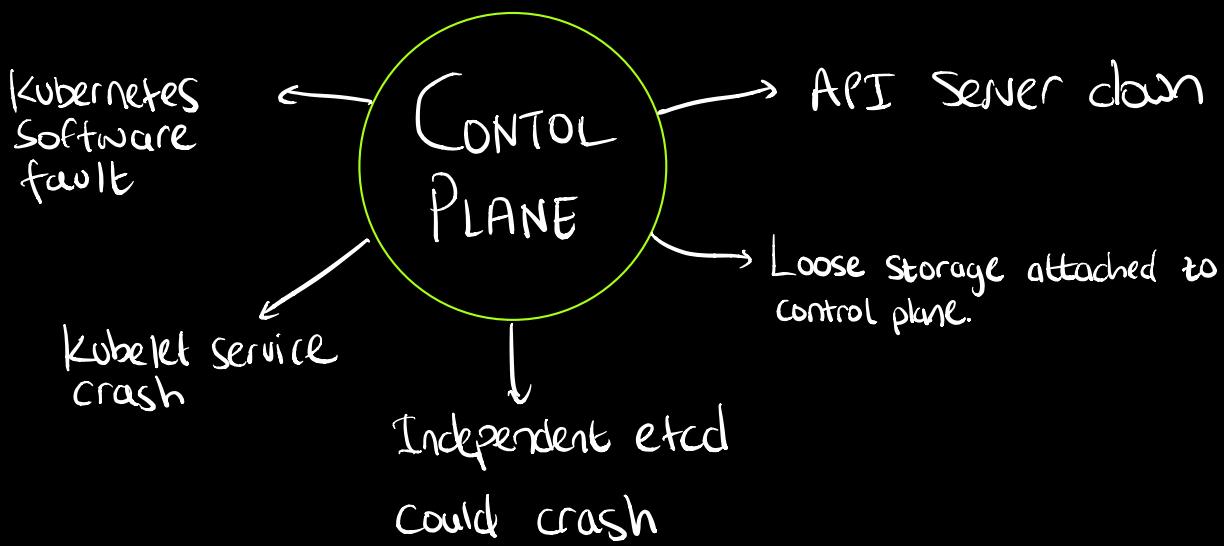
→ Only particular fields can be changed i.e Image

→ To change other fields of failed pod

↓  
export configuration → -o yaml --export  
modify yaml to change memory request

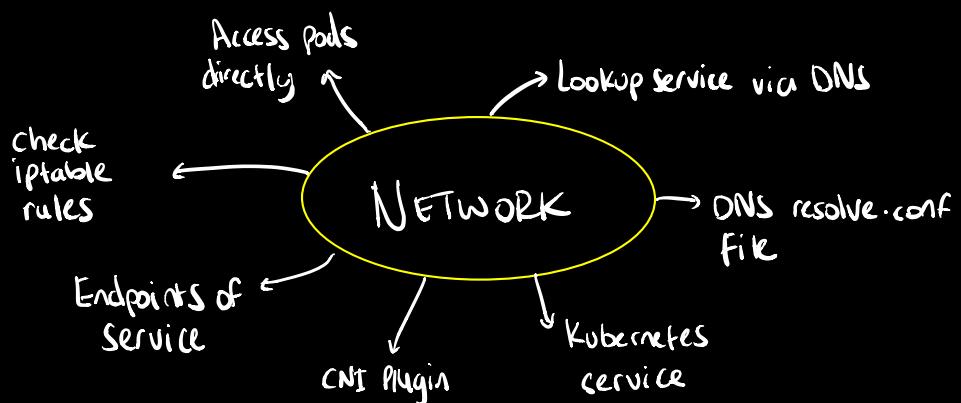
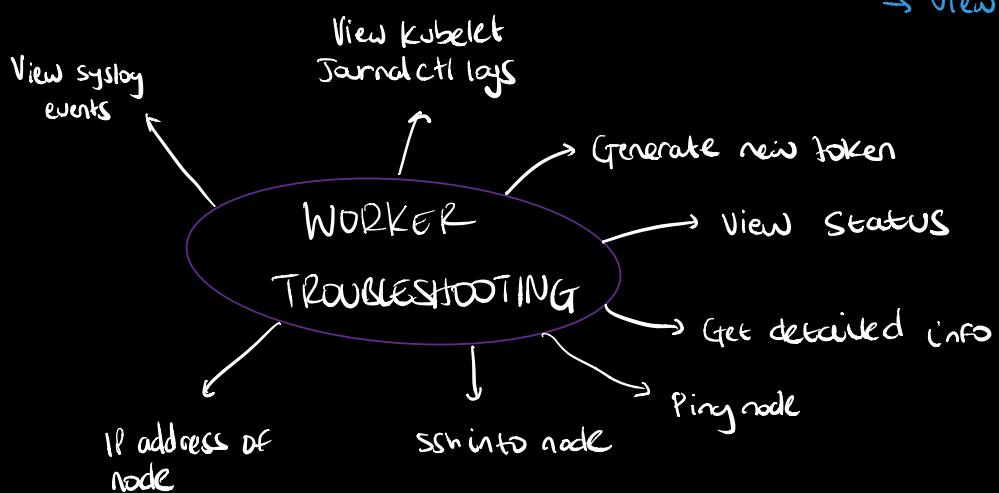


# TROUBLESHOOT FAILURES



- View the events from control plane components
- View logs for control plane pods
- check status of docker service

- check status of kubelet service
- Disable swap
- check firewalld service
- View kube config



THANK You FOR READING

PLEASE LIKE AND SHARE FOR MORE

LATEST NOTES AVAILABLE ON

INSTAGRAM → adnans\_techie\_studies

ADNAN  
RASHID  
Z