




WELCOME TO NOMAD HANDS-ON TRAINING!




oscon.com
#oscon



Anubhav Mishra

Developer Advocate, HashiCorp
@anubhavm





HashiCorp

Nomad

Agenda



1. Architecture Overview
2. Installing and Configuring Nomad
3. Creating and Running Jobs
4. Service Discovery with Consul
5. Operating Nomad
6. Interacting with the HTTP API
7. Advanced Job Strategies

Workstations




Claim your workstation!

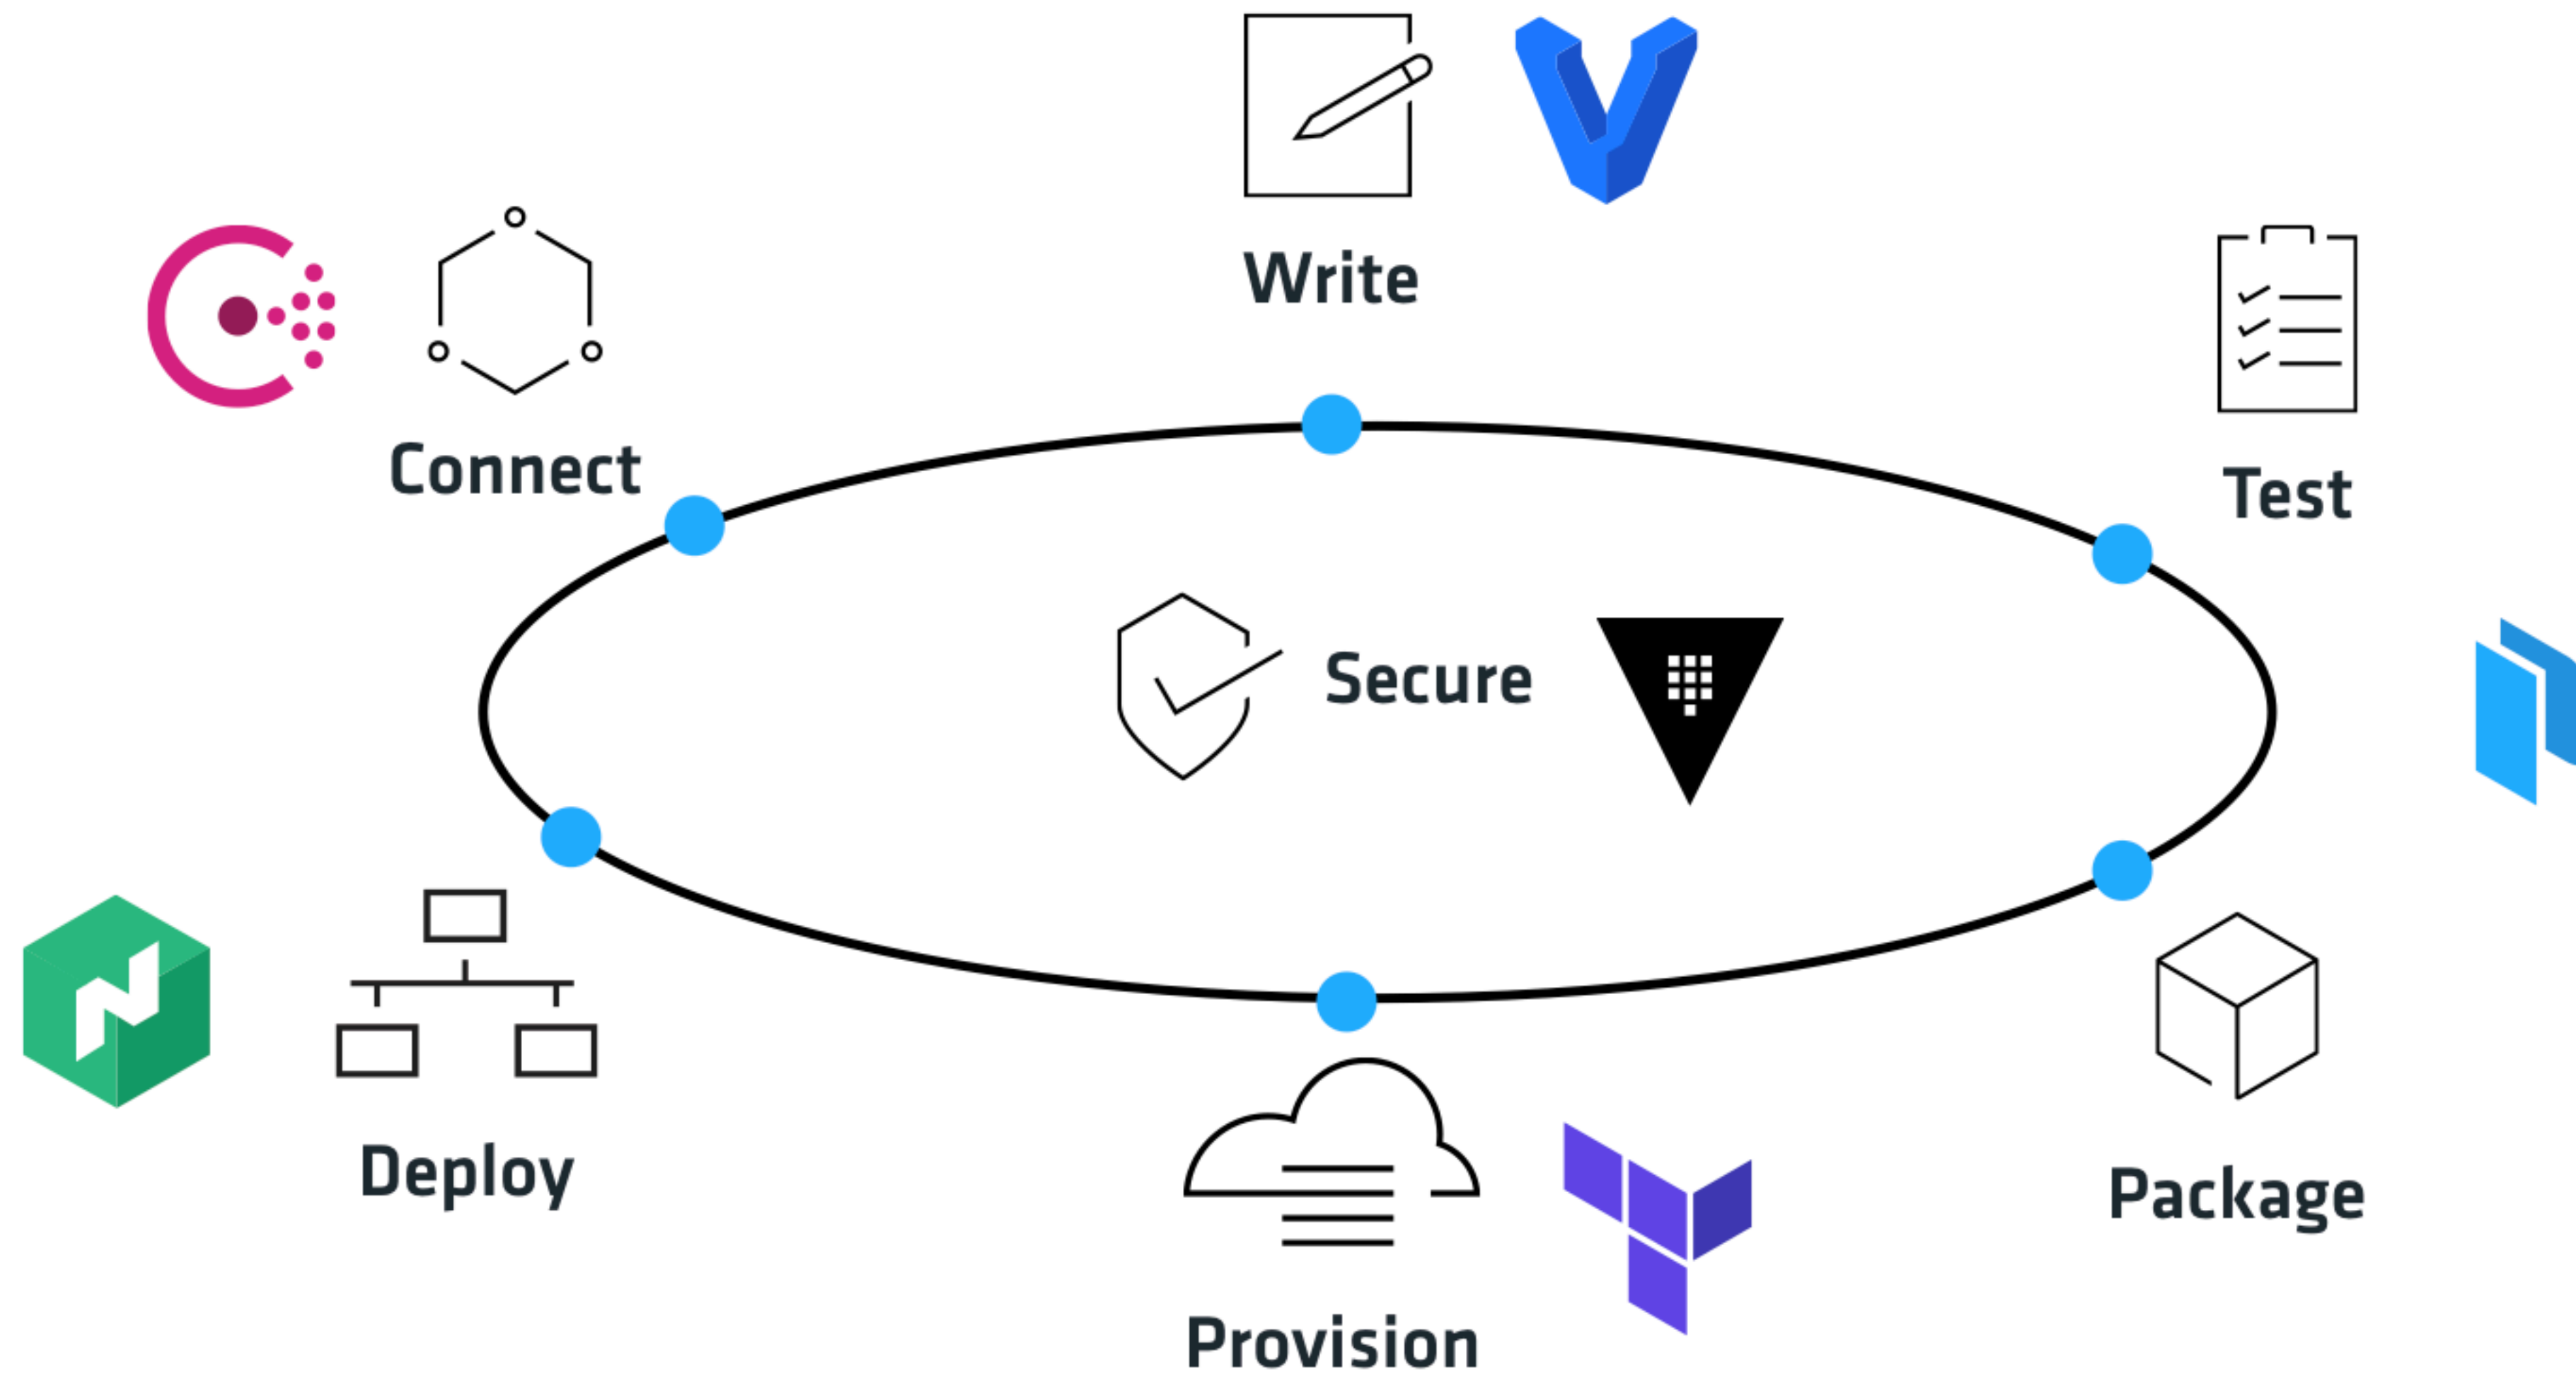
<https://hashi.co/nomad-hands-on-oscon>

Introduction

Nomad Powers Application Deployment



Application Delivery Lifecycle



Nomad's Goals



Treat entire data center as a collection of resources

Support non-containerized workloads

Achieve massive scale and speed

Cross-platform portability

Support heterogeneous workflows



Nomad

Million Container Challenge

1,000 Jobs

1,000 Tasks per Job

5,000 Hosts on GCE


1,000,000 Containers



==> Nomad agent started! Log data will stream in below:

```
2016/03/14 17:34:26 [DEBUG] client: available drivers [raw_exec java qemu docker exec]
2016/03/14 17:34:26 [INFO] client: server address list: [127.0.0.1:4647]
2016/03/14 17:34:26 [DEBUG] client: checking for node changes at duration 5s
2016/03/14 17:34:27 [DEBUG] client: complete
2016/03/14 17:34:28 [DEBUG] client:
2016/03/14 17:35:05 [DEBUG] client: 8 (pulled 1) (filtered 0)
2016/03/14 17:35:05 [DEBUG] client: updated 0) (ignore 0)
2016/03/14 17:35:05 [DEBUG] client: 76cb4259-a20f-1581-877f-1a
2016/03/14 17:35:05 [DEBUG] client: 76cb4259-a20f-158
2016/03/14 17:35:05 [DEBUG] client: (filtered 0)
2016/03/14 17:35:05 [DEBUG] client: 99-94
2016/03/14 17:35:05 [DEBUG] client: 8-975
```

Nomad C1M (1 Million Docker Containers)



“640 KB ought to be enough for anybody.”

- Bill Gates

2nd Largest Hedge Fund

18K Cores

5 Hours

2,200 Containers/second



Nomad vs. Other Tools



Schedule multiple workload types: VMs, containers, binaries, etc

Optimistically concurrent to schedule faster and at greater scale

Globally-aware with multi-region and multi-datacenter support

Operationally simple

Integration into the HashiCorp ecosystem

Glossary



Node

Physical or virtual machine in the cluster. In the context of Nomad, a node is a machine running the Nomad agent.



Agent

Long-running daemon on every member of the Nomad cluster. The agent is able to run in either client or server mode.



Agent (Client)

Agent that fingerprints the host to determine capabilities, resources, and available drivers.



Agent (Server)

Agent that holds the global state of the cluster and participates in scheduling decisions.



Job

Definition of how a workload should be scheduled. The job specification is composed of one or more task groups, and each task defines a series of resource constraints and configuration.



Job File

HCL or JSON configuration file on disk which describes how a workload should be scheduled.



Driver

Pluggable components that execute a task and provide resource isolation. Example drivers include *docker*, *java*, and *raw-exec*.



Task

A command, service, application or "set of work" to execute. Tasks are executed by their driver.



Task Group

A collection of individual tasks that should be co-located on the same node. This is especially useful for applications that require low latency or have high throughput to another application.



Evaluation

A calculation performed by the Nomad servers to determine what action(s) need to take place to execute a job.



Allocation

An instance of a task group placed on a node. Allocations can fail (not enough memory, node is down, etc).



Datacenter

Networking environment that is private, low latency, and high bandwidth. Example datacenters include the `us-east-1` or `us-west-1`.



Region

Collection of multiple datacenters, typically grouped geographically. For example, the `north-america` region might include the `us-east-1` and `us-west-1` datacenters.



Consensus

Agreement upon the elected leader.



Gossip

Random node-to-node communication primarily over UDP that provides membership, failure detection, and event broadcast information to the cluster. Built on Serf.

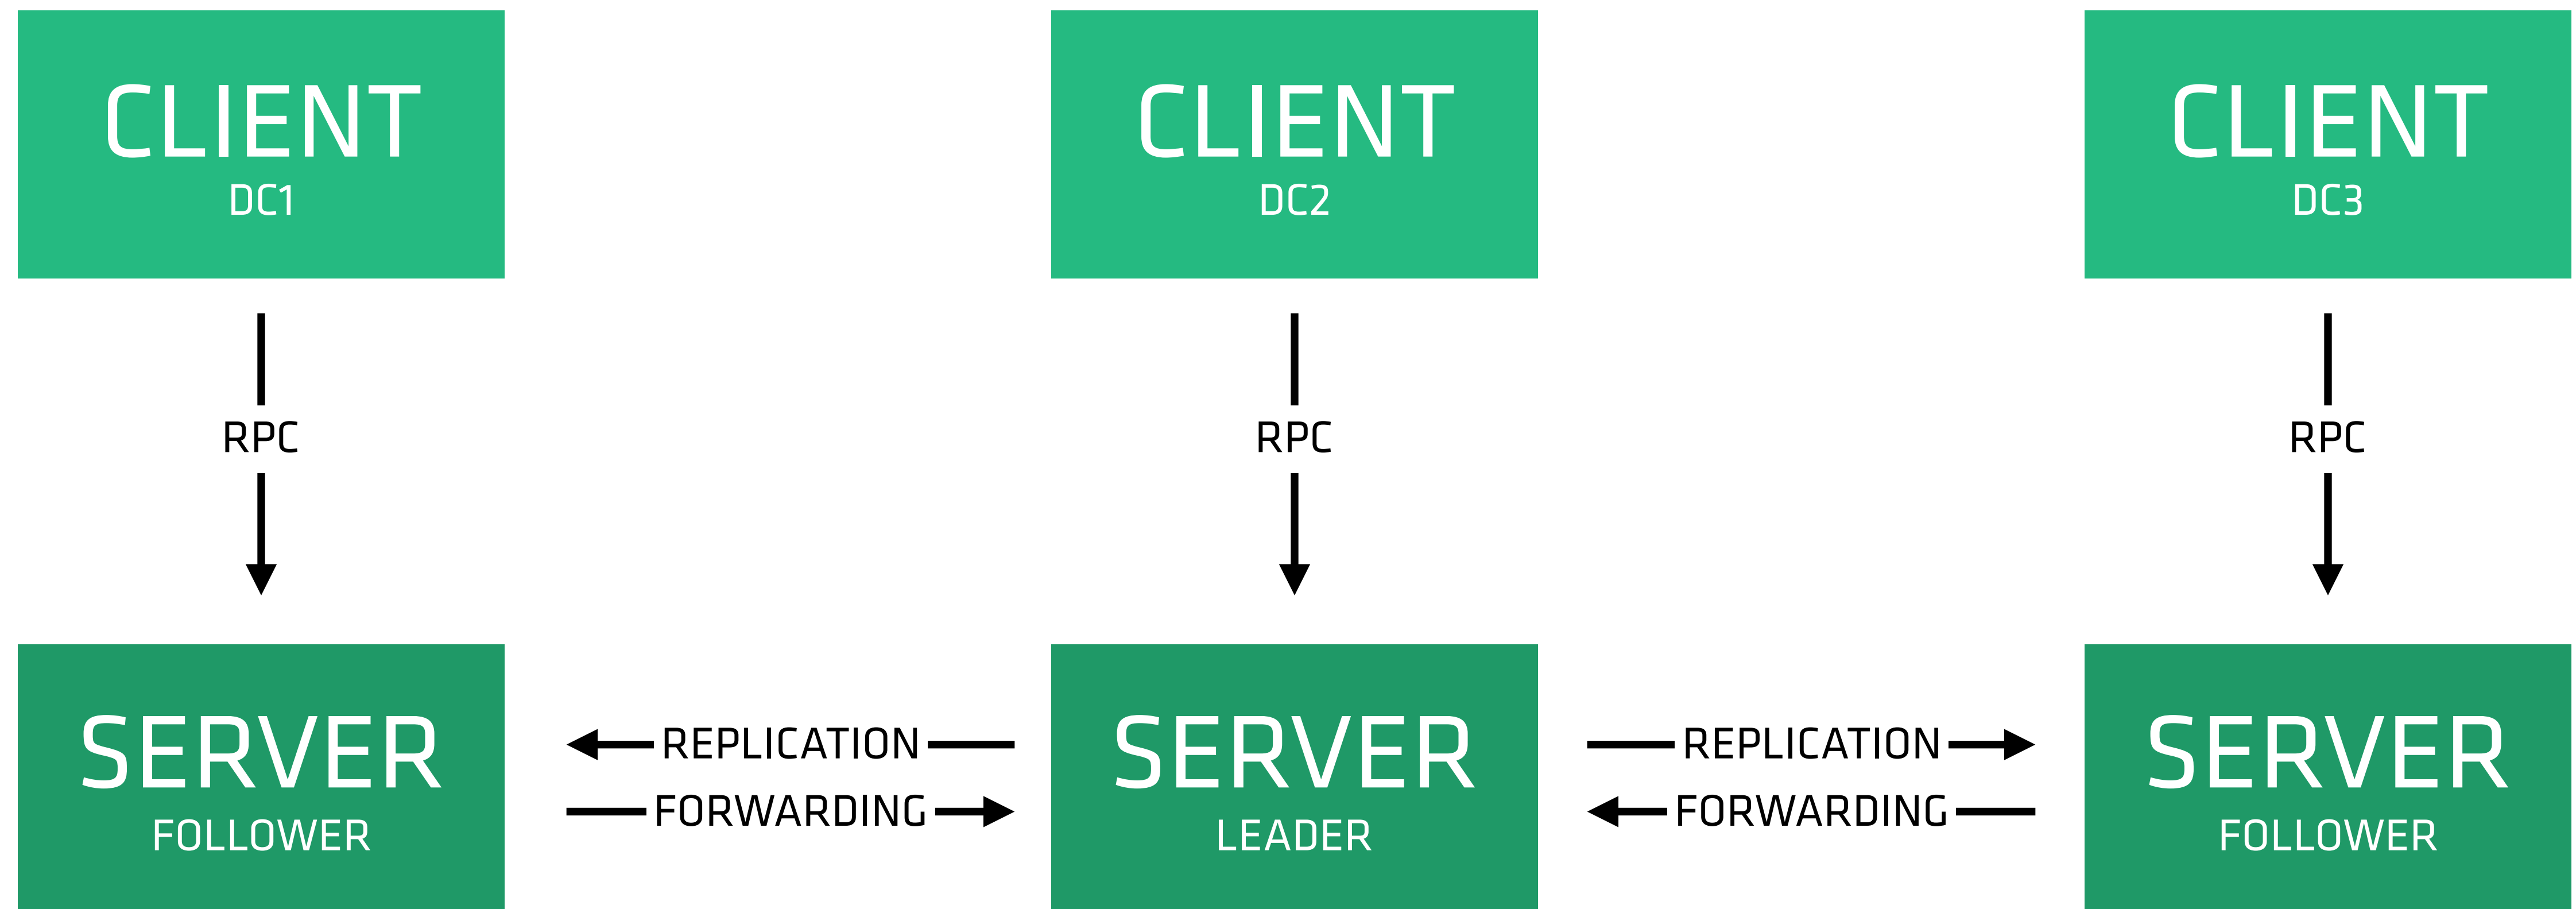


Bin-Packing

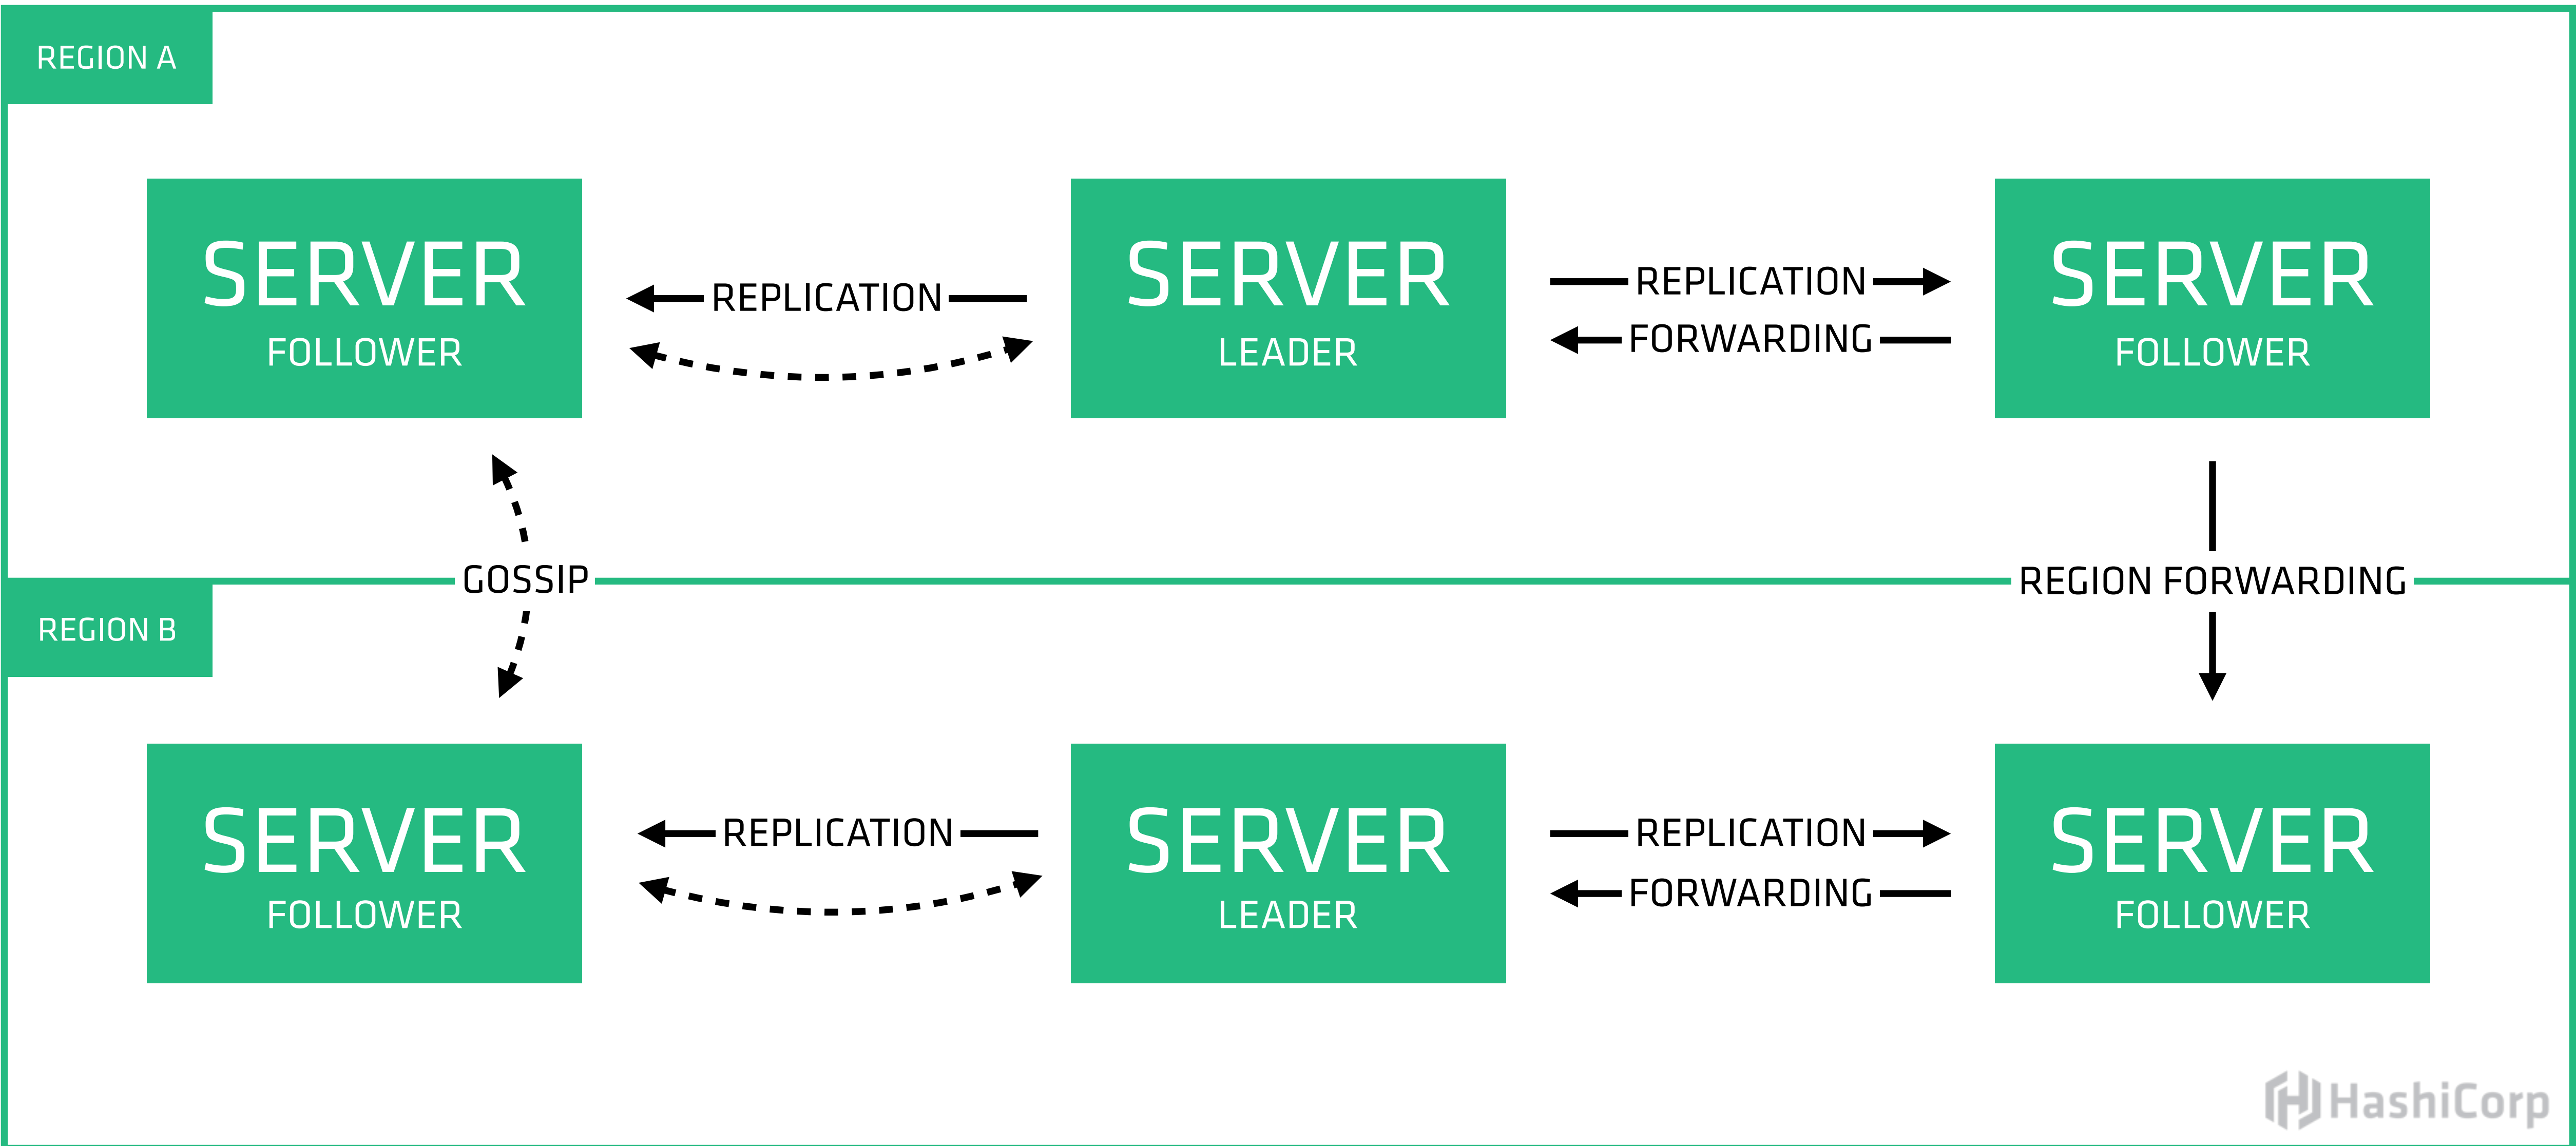
A algorithm which optimizes the resource utilization and density of applications, but is also augmented by affinity and anti-affinity rules.

Architecture

Single Region Architecture



Multi-Region Architecture




100's of Regions

10,000's of Clients per Region

1000's of Jobs per Region

Installing & Configuring

Single Binary



Agent Functionality (Client)



Fingerprint to determine resources and capabilities of each host

Send node information to the central server cluster

Heartbeat to provide liveness

Run any tasks assigned by the central server cluster

Agent Functionality (Server)



Store resource, capability, and availability of each host

Schedule workloads

Participate in leader election and state replication (consensus)

Discover other Nomad servers (gossip)



Terminal

```
$ nomad agent -config=/path/to/config.hcl
```





Terminal

```
$ nomad agent -config=/path/to/config.hcl -config=/path/to/more/config.hcl
```



Bootstrapping Nomad

Bootstrapping Nomad



There are two common strategies for bootstrapping a Nomad cluster: *automatic* and *manual*.

We will walk through the manual steps first, then automatic.

Bootstrapping Nomad: Manual



Chicken-and-egg problem: requires one server IP

The **bootstrap_expect** field tells how many servers to reach quorum

This server IP is specified in a configuration file or used with the **nomad server join** command

Clients specify server IPs via configuration file

Requires human involvement



/etc/nomad.d/config.hcl

```
server {  
  enabled          = true  
  bootstrap_expect = 3  
  
  # This is the IP address of the first server we provisioned  
  retry_join = ["<known-address>:4648"]  
}
```



Terminal


```
$ nomad server join <known-address>
```





client-config.hcl

```
client {  
  enabled = true  
  servers = ["<known-address>:4647"]  
}
```



Bootstrapping Nomad: Automatic



Leverages another open source tool - Consul

The `bootstrap_expect` field tells how many servers to reach quorum


Nomad assumes the Consul agent is accessible on the local IP and port, but this is configurable

Fully automated and scalable



server-config.hcl


```
server {  
  enabled          = true  
  bootstrap_expect = 3  
}
```





client-config.hcl

```
client {  
  enabled = true  
}
```





consul.hcl

```
consul {  
  # The address to the Consul agent.  
  address = "127.0.0.1:8500"  
  
  # The service name to register the server and client with Consul.  
  server_service_name = "nomad"  
  client_service_name = "nomad-client"  
  
  # Enables automatically registering the services.  
  auto_advertise = true  
  
  # Enabling the server and client to bootstrap using Consul.  
  server_auto_join = true  
  client_auto_join = true  
}
```

Exercise: Connect to Workstation



Go to: <https://hashi.co/nomad-hands-on-oscon> and claim a workstation IP.

SSH into your workstation using the provided credentials.

```
ssh nomaduser@<your.ip.address>  
password: oscon2018
```

Change directory into `/workstation/nomad`.

Creating & Running Jobs

Nomad Jobs



Jobs specifications are defined in HCL or JSON, but we will use HCL for this training.

An example can be generated by running `nomad init`.

Nomad Job Types



Nomad has three scheduler types that can be used when creating your job: *service*, *batch*, and *system*.

Nomad Job Types: Service



Service Scheduler Job Type

The service scheduler is designed for scheduling long-lived services that should never go down. The service scheduler ranks a large portion of the nodes that meet the jobs constraints and selects the optimal node to place a task group on.

Examples: *webapp*, *redis*

Nomad Job Types: Batch



Batch Scheduler Job Type

Batch jobs are less sensitive to short-term performance fluctuations and are short lived, finishing after some period.

Examples: *billing, data replication*

Nomad Job Types: System



System Scheduler Job Type

The system scheduler is used to register jobs that should be run on all clients that meet the job's constraints. The system scheduler is also invoked when clients join the cluster or transition into the ready state.

Examples: *logging agent, security auditing tool*

Exercise: Run `nomad init`



Inside the `/workstation/nomad` folder, run the `nomad init` command to generate a sample job file.

Open the resulting file on the workstation in an editor.



Terminal

```
$ nomad init
```

```
Example job file written to example.nomad
```





Terminal

```
$ vi example.nomad
```

Nomad Jobs: Hierarchy



Job

(Group)

Task


Resources

Constraint



example.nomad

```
job "example" {  
  # ...  
}
```





example.nomad

```
job "example" {  
  # Run the job in the global region, which is the default.  
  # region = "global"  
  
  # ...  
}
```



example.nomad

```
job "example" {  
  # ...  
  
  # Specify the datacenters within the region this job can run in.  
  datacenters = ["dc1"]  
  
  # ...  
}
```




example.nomad

```
job "example" {  
  # ...  
  
  # Service type jobs optimize for long-lived services. This is  
  # the default but we can change to batch for short-lived tasks.  
  # type = "service"  
  
  # ...  
}
```



example.nomad

```
job "example" {  
  # ...  
  
  # Restrict our job to only linux. We can specify multiple  
  # constraints as needed.  
  # constraint {  
  #   attribute = "${attr.kernel.name}"  
  #   value      = "linux"  
  # }  
  
  # ...  
}
```



example.nomad

```
job "example" {  
  # ...  
  
  # Configure the job to do rolling updates  
  update {  
    # Stagger updates every 10 seconds  
    stagger = "10s"  
  
    # Update a single task at a time  
    max_parallel = 1  
  }  
  
  # ...  
}
```



example.nomad

```
job "example" {  
  # ...  
  
  # Create a 'cache' group. Each task in the group will be  
  # scheduled onto the same machine.  
  group "cache" {  
    # Control the number of instances of this group.  
    # Defaults to 1  
    count = 1  
  
    # ...  
  }  
  
  # ...  
}
```



example.nomad

```
job "example" {  
  # ...  
  
  group "cache" {  
    # ...  
    restart {  
      # ...  
    }  
  
    # Define a task to run  
    task "redis" {  
      # ...  
    }  
  }  
}
```

example.nomad

```
job "example" {  
  # ...  
  group "cache" {  
    task "redis" {  
      # Use Docker to run the task.  
      driver = "docker"  
  
      # Configure Docker driver with the image  
      config {  
        image = "redis:latest"  
        port_map {  
          db = 6379  
        }  
      }  
    }  
  
    service {  
      name = "${TASKGROUP}-redis"  
      tags = ["global", "cache"]  
    }  
  }  
}
```

example.nomad

```
job "example" {  
  # ...  
  group "cache" {  
    # ...  
    task "redis" {  
      # We must specify the resources required for  
      # this task to ensure it runs on a machine with  
      # enough capacity.  
      resources {  
        cpu      = 500 # 500 MHz  
        memory = 256 # 256MB  
        network {  
          mbits = 10  
          port "db" {}  
        }  
      }  
    }  
  }  
}
```

Exercise: Delete Example Job File



This job file is just an example, so delete it.



Terminal

```
$ rm example.nomad
```



Exercise: Inspect Job File



There is already a job file on the system named `http-echo.nomad`.

Open and inspect this job file.

About: http-echo



`http-echo` is a small application that accepts text as a command-line flag and renders that text as an HTML webpage.

It accept `-listen` and `-text` flags.

Example invocation:

```
http-echo -text="hello world"
```



Terminal

```
$ http-echo -text="hello world"  
Server is listening on :5678
```

```
$ curl localhost:5678  
hello world
```

About: http-echo Docker Container



The `hashicorp/http-echo` Docker container packages the application in a distributable format.

It accepts the same arguments and flags as `http-echo`.

Example invocation:

```
docker run hashicorp/http-echo -text="hello world"
```

About: http-echo Nomad Job



Requests Nomad to download and run a Docker container from the public docker registry (`hashicorp/http-echo`).

Configures the container to bind to port 80 on the host.

Registers the service for discovery with Consul with an integrated health check (more on this later).

http-echo.nomad

```
config {
  image = "hashicorp/http-echo:x.y.z"
  args = [
    "-listen", ":80",
    "-text", "hello world",
  ]
}

resources {
  network {
    mbits = 10
    port "http" {
      static = 80
    }
  }
}
```



http-echo.nomad

```
service {  
  name = "http-echo"  
  port = "http"  
  
  tags = [  
    "anaconda",  
    "urlprefix-/http-echo",  
  ]  
  
  check {  
    type      = "http"  
    path      = "/health"  
    interval  = "10s"  
    timeout   = "5s"  
  }  
}
```


Exercise: Run Nomad Job



Submit the `http-echo.nomad` job to the server for evaluation and scheduling.



Terminal

```
$ nomad run http-echo.nomad
==> Monitoring evaluation "89164c43"
      Evaluation triggered by job "http-echo-testing-mongrel"
      Allocation "1c1057f3" created: node "ec255cce", group "echo"
      Evaluation status changed: "pending" -> "complete"
==> Evaluation "89164c43" finished with status "complete"
```

Exercise: Contact Service



Your instance of `http-echo` was deployed on the cluster. It may have been scheduled on any host, so we need to use service discovery (Consul) to address it.

Use `curl` to query your instance

```
$(identity).http-echo.service.consul
```



Terminal

```
$ curl $(identity).http-echo.service.consul  
hello world
```



Exercise: Change Job



Modify the job file to render your name (or anything you choose) instead of "hello world".

Resubmit the modified job to Nomad.

Use `curl` to query the updated instance.

● ● ● http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args  = [  
    "-listen", ":80",  
    "-text", "hello world",  
  ]  
}
```



http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args  = [  
    "-listen", ":80",  
    "-text", "smiling banana",  
  ]  
}
```



Terminal

```
$ nomad run http-echo.nomad
==> Monitoring evaluation "c12830a2"
      Evaluation triggered by job "http-echo-testing-mongrel"
      Allocation "30c3ee48" created: node "ec277448", group "echo"
      Evaluation status changed: "pending" -> "complete"
==> Evaluation "c12830a2" finished with status "complete"
```




Terminal

```
$ curl $(identity).http-echo.service.consul  
smiling banana
```



Exercise: Scale `http-echo` Job



Increase the `count` attribute on the group to "5" to run five instances of our application.

Resubmit this job for evaluation.

HINT: You may need to add a `count` attribute or look online at the documentation.



http-echo.nomad

```
group "echo" {  
  count = "5"  
  
  task "server" {  
    driver = "docker"  
  
    config {  
      image = "hashicorp/http-echo:x.y.z"  
      args = [  
        "-listen", ":80",  
        "-text", "smiling banana",  
      ]  
    }  
  
    # ...  
  }  
}
```



Terminal

```
$ nomad run http-echo.nomad
==> Monitoring evaluation "78bd7efb"
      Evaluation triggered by job "http-echo-testing-mongrel"
      Allocation "0c5724a0" created: node "ec255589", group "echo"
      Allocation "5d4dab42" created: node "ec261569", group "echo"
      Allocation "8290a411" created: node "ec255cce", group "echo"
      Allocation "a29935f1" created: node "ec2ee7db", group "echo"
      Allocation "f400c498" created: node "ec277448", group "echo"
      Allocation "781ca69d" modified: node "ec274728", group "echo"
      Evaluation status changed: "pending" -> "complete"
==> Evaluation "78bd7efb" finished with status "complete" but failed to place
all allocations:
      Task Group "echo" (failed to place 4 allocations):
        * Resources exhausted on 6 nodes
        * Dimension "network: reserved port collision" exhausted on 6 nodes
      Evaluation "546caa43" waiting for additional capacity to place remainder
```

About: Static Ports



Static ports are generally a bad idea in scheduled environments as they restrict exactly one instance of the job to running per host.

Letting Nomad choose dynamic ports allows for better scale.

http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args = [  
    "-listen", ":80",  
    "-text", "hello world",  
  ]  
}  
  
resources {  
  network {  
    mbits = 10  
    port "http" {  
      static = 80  
    }  
  }  
}
```

← Remove this line



http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args  = [  
    "-listen", ":80",  
    "-text", "hello world",  
  ]  
}  
  
resources {  
  network {  
    mbits = 10  
    port "http" {}  
  }  
}
```

● ● ● http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args = [  
    "-listen", ":??",  
    "-text", "hello world",  
  ]  
}  
  
resources {  
  network {  
    mbits = 10  
    port "http" {}  
  }  
}
```


Nomad Interpolation



Nomad populates certain variables with information about the job.

Values between `${ }` are analyzed by the parser.



http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args  = [  
    "-listen", ":",  
    "-text", "hello world",  
  ]  
}  
  
resources {  
  network {  
    mbits = 10  
    port "http" {}  
  }  
}
```

http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args = [  
    "-listen", ":%${NOMAD_PORT_http}",  
    "-text", "hello world",  
  ]  
}  
  
resources {  
  network {  
    mbits = 10  
    port "http" {}  
  }  
}
```

http-echo.nomad

```
config {
  image = "hashicorp/http-echo:x.y.z"
  args  = [
    "-listen", ":%${NOMAD_PORT_http}",
    "-text", "hello world",
  ]
}

resources {
  network {
    mbits = 1
    port "http" {}
  }
}
```



● ● ● http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args  = [  
    "-listen", ":%${NOMAD_PORT_banana}",  
    "-text", "hello world",  
  ]  
}  
  
resources {  
  network {  
    mbits = 10  
    port "banana" {}  
  }  
}
```

Exercise: Resubmit Job



Make these required changes to the job.

Resubmit the job to the Nomad server for evaluation.

Use `curl` to query the updated instance.

HINT: Google "nomad interpolation".



http-echo.nomad

```
config {  
  image = "hashicorp/http-echo:x.y.z"  
  args  = [  
    "-listen", ":${NOMAD_PORT_http}",  
    "-text", "hello world",  
  ]  
}  
  
resources {  
  network {  
    mbits = 10  
    port "http" {}  
  }  
}
```



Terminal

```
$ nomad run http-echo.nomad
==> Monitoring evaluation "4cf66353"
Evaluation triggered by job "http-echo-testing-mongrel"
Allocation "2d1b84a1" created: node "ec255589", group "echo"
Allocation "5a87dcf3" created: node "ec277448", group "echo"
Allocation "c8bd019c" created: node "ec261569", group "echo"
Allocation "64db5ade" created: node "ec255cce", group "echo"
Allocation "67e19895" created: node "ec2ee7db", group "echo"
Allocation "7e5803cb" created: node "ec277448", group "echo"
Allocation "8a82a6b5" created: node "ec274728", group "echo"
Allocation "2117cf22" created: node "ec2ee7db", group "echo"
Allocation "52e56b0d" created: node "ec255cce", group "echo"
Allocation "53528b48" created: node "ec261569", group "echo"
Allocation "52e56b0d" status changed: "pending" -> "running"
Allocation "5a87dcf3" status changed: "pending" -> "running"
Allocation "67e19895" status changed: "pending" -> "running"
Allocation "7e5803cb" status changed: "pending" -> "running"
Allocation "c8bd019c" status changed: "pending" -> "running"
```




Terminal

```
$ curl $(identity).http-echo.service.consul  
curl: (7) Failed to connect to anaconda.http-echo.service.consul port 80:  
Connection refused
```



Why is the Job Inaccessible?



Previously our job was hard-bound to port 80 (default HTTP port).

Now Nomad is dynamically allocating a high-numbered port, so our service could be listening on any port.

Must rely on service discovery to find the port (which we will discuss later).

Monitoring Jobs

Command: `nomad status`



The `nomad status` command lists the status of all jobs in the system.

If supplied an optional argument, the `nomad status` command lists detailed information about the job name.

Exercise: Run `nomad status`



Run the `nomad status` command to see the status of all jobs in the system.

Run the `nomad status` command with the name of *your* job to get detailed job information.



Terminal

```
$ nomad status
```

ID	Type	Priority	Status	Submit Date
fabio	system	75	running	01/01/17 01:30:10 UTC
hashi-ui	system	75	running	01/01/17 01:30:10 UTC
http-echo-llama	service	50	running	01/01/17 01:30:10 UTC





Terminal

```
$ nomad status http-echo-$(identity)
ID           = http-echo-llama
Name         = http-echo-llama
Submit Date  = 01/01/17 01:30:10 UTC
Type         = service
Priority      = 50
Datacenters  = dc1
Status       = running
Periodic     = false
Parameterized = false
```

Summary

Task Group	Queued	Starting	Running	Failed	Complete	Lost
echo	0	0	5	0	6	0

Allocations

ID	Node ID	Task Group	Version	Desired	Status	Created At
404add9c	fb3eff16	echo	4	run	running	01/01/17 01:30:10



Terminal

```
$ nomad status fabio
ID           = fabio
Name         = fabio
Submit Date  = 01/01/17 01:30:10 UTC
Type         = system
Priority      = 75
Datacenters  = dc1
Status       = running
Periodic     = false
Parameterized = false
```

Summary

Task Group	Queued	Starting	Running	Failed	Complete	Lost
fabio	0	0	6	0	0	0

Allocations

ID	Node ID	Task Group	Version	Desired	Status	Created At
10a9306c	e1ae29b8	fabio	0	run	running	01/01/17 01:30:10



Terminal

```
$ nomad status fabio
ID           = fabio
Name         = fabio
Submit Date  = 01/01/17 01:30:10 UTC
Type         = system
Priority      = 75
Datacenters  = dc1
Status       = running
Periodic     = false
Parameterized = false
```

Summary

Task Group	Queued	Starting	Running	Failed	Complete	Lost
fabio	0	0	6	0	0	0

Allocations

ID	Node ID	Task Group	Version	Desired	Status	Created At
10a9306c	e1ae29b8	fabio	0	run	running	01/01/17 01:30:10

Exercise: Verify Fabio is Running



We see fabio is running under Nomad.

Fabio is running as a local-exec job (outside of Docker).

Manually verify fabio is running by running **ps**.



Terminal

```
$ ps aux | grep fabio
```

```
nobody      1640  0.3  0.4  37372 33324 ?        S1   02:57   0:00 fabio
```

Command: `nomad alloc-status`



The `nomad alloc status` displays information about the given allocation ID, including run status, metadata, and failure messages.

Exercise: Query Nomad Allocation Status



Find the ID of a running allocation using the `nomad status` command.

Query that allocation status using the `nomad alloc status` command with the allocation ID.



Terminal

```
$ nomad status http-echo-$(identity)
ID           = http-echo-llama
Name         = http-echo-llama
Submit Date  = 01/01/17 01:30:10 UTC
Type         = service
Priority      = 50
Datacenters  = dc1
Status       = running
Periodic     = false
Parameterized = false
```

Summary

Task Group	Queued	Starting	Running	Failed	Complete	Lost
echo	0	0	5	0	0	0

Allocations

ID	Node ID	Task Group	Version	Desired	Status	Created At
<u>1e6892d3</u>	744e4d82	echo	0	run	running	01/01/17 01:30:10



Terminal

```
$ nomad alloc status 1e6892d3
ID                = 1e6892d3
Eval ID           = 8a5e22a0
Name              = http-echo-llama.echo[2]
Node ID           = 744e4d82
Job ID            = http-echo-llama
Job Version       = 0
Client Status     = running
Client Description = <none>
Desired Status    = run
Desired Description = <none>
Created At        = 01/01/17 01:30:10 UTC
```

Task "server" is "running"

Task Resources

CPU	Memory	Disk	IOPS	Addresses
0/100 MHz	2.2 MiB/10 MiB	300 MiB	0	http: 10.1.1.14:21136



Terminal

```
$ nomad alloc status 1e6892d3
ID                = 1e6892d3
Eval ID           = 8a5e22a0
Name              = http-echo-llama.echo[2]
Node ID           = 744e4d82
Job ID            = http-echo-llama
Job Version       = 0
Client Status     = running
Client Description = <none>
Desired Status    = run
Desired Description = <none>
Created At        = 01/01/17 01:30:10 UTC
```

Task "server" is "running"

Task Resources

CPU	Memory	Disk	IOPS	Addresses
0/100 MHz	2.2 MiB/10 MiB	300 MiB	0	<u>http: 10.1.1.14:21136</u>



Terminal

```
$ curl 10.1.1.14:21136  
smiling banana
```



Exercise: Query Nomad Allocation Stats



Find and display the detailed resource statistics on your allocation.

HINT: You may need to pass an extra flag to the `alloc status` command.



Terminal

```
$ nomad alloc status -stats 1e6892d3
```





Terminal

```
$ nomad alloc status -stats 1e6892d3
```

```
# ...
```

Memory Stats

Cache	Max Usage	RSS	Swap
16 KiB	1.3 MiB	1.0 MiB	0 B

CPU Stats

Percent	Throttled Periods	Throttled Time
0.00%	0	0

```
# ...
```

Command: `nomad logs`



The `nomad logs` command can query the stdout and stderr from your task

Requires an allocation ID

Exercise: Run `nomad logs`



Run `nomad status` to get an allocation ID of a running job (or use the previous one).

Run the `nomad logs` command to view the most recent output for that allocation.



Terminal

```
$ nomad logs 1e6892d3
```

```
10.1.1.14:21136 10.1.1.14:50336 "GET / HTTP/1.1" 200 15 "curl/7.47.0" 26.075μs
```



Resources, Constraints, and Planning

Command: `nomad plan`



The `nomad plan` command invokes the scheduler in a dry-run mode to show you what will happen if the job was submitted.

The resulting index can be specified when running the job to ensure no changes have happened.



Terminal

```
$ nomad plan http-echo.nomad
```



Terminal

```
$ nomad plan http-echo.nomad
```

```
+/- Job: "http-echo-testing-iguana"
```

```
Task Group: "echo" (5 in-place update)
```

```
Task: "server"
```

```
Scheduler dry-run:
```

```
- All tasks successfully allocated.
```

```
Job Modify Index: 287
```

```
To submit the job with version verification run:
```

```
nomad run -check-index 287 http-echo.nomad
```

When running the job with the check-index flag, the job will only be run if the

server side version matches the job modify index returned. If the index has changed, another user has modified the job and the plan's results are potentially invalid.



Terminal

```
$ nomad run -check-index 287 http-echo.nomad
```





Terminal

```
$ nomad run -check-index 287 http-echo.hcl
==> Monitoring evaluation "8a87cd9c"
      Evaluation triggered by job "http-echo-testing-iguana"
      Allocation "7b34edb0" modified: node "ec2ee7db", group "echo"
      Allocation "b9d5a1ff" modified: node "ec255cce", group "echo"
      Allocation "f00e188f" modified: node "ec2ee7db", group "echo"
      Allocation "1d843f7c" modified: node "ec255589", group "echo"
      Allocation "259496bf" modified: node "ec261569", group "echo"
      Evaluation status changed: "pending" -> "complete"
==> Evaluation "8a87cd9c" finished with status "complete"
```

About: Constraints



Constraints are requirements the scheduler must evaluate *about the client* such as operating system, architecture, kernel version, etc.

Constraint requirements can be specified at the job, group, or task level.

Exercise: Add Constraint Requirement



Add a constraint requirement to the `http-echo` job which requires the client kernel to be linux.

Plan this job to see the changes that will take place.

Submit this job to the Nomad server.



http-echo.nomad

```
job "http-echo-anaconda" {  
  #...  
  
  constraint {  
    attribute = "${attr.kernel.name}"  
    value      = "linux"  
  }  
  
  group "echo" {  
    #...  
  }  
}
```




Terminal

```
$ nomad plan http-echo.nomad
+/- Job: "http-echo-testing-iguana"
+ Constraint {
  + LTarget: "${attr.kernel.name}"
  + Operand: "="
  + RTarget: "linux"
}
Task Group: "echo" (5 in-place update)
  Task: "server"

Scheduler dry-run:
- All tasks successfully allocated.

Job Modify Index: 361
To submit the job with version verification run:

nomad run -check-index 361 http-echo.nomad
```



Terminal

```
$ nomad run -check-index 361 http-echo.nomad
==> Monitoring evaluation "58394788"
      Evaluation triggered by job "http-echo-testing-iguana"
      Allocation "6704f93f" modified: node "ec277448", group "echo"
      Allocation "b854d9f0" modified: node "ec261569", group "echo"
      Allocation "fc790941" modified: node "ec255589", group "echo"
      Allocation "0d9eb8a6" modified: node "ec274728", group "echo"
      Allocation "28c581c3" modified: node "ec274728", group "echo"
      Evaluation status changed: "pending" -> "complete"
==> Evaluation "58394788" finished with status "complete"
```

About: Resources



Resources are minimum requirements the *task must have* to run on the client such as memory or cpu.

Resource constraints can only be specified on the task.

Exercise: Add Resource Requirement



Add a resource requirement to the `http-echo` job on the `server` task which allocates 50GB of memory.

Plan this job to see the changes that will take place.

Submit this job to the Nomad server.

Inspect the allocation status.



http-echo.nomad

```
job "http-echo-${identity}" {  
  group "echo" {  
    task "server" {  
  
      # ...  
  
      resources {  
        memory = 50000 # 50GB of RAM  
  
        network {  
          mbits = 10  
          port "http" {}  
        }  
      }  
  
      # ...  
    }  
  }  
}
```



Terminal

```
$ nomad plan http-echo.nomad
+/- Job: "http-echo-testing-iguana"
+/- Task Group: "echo" (5 create/destroy update)
  +/- Task: "server" (forces create/destroy update)
    +/- Resources {
      CPU:      "100"
      DiskMB:   "0"
      IOPS:     "0"
      +/- MemoryMB: "10" => "5000"
    }
```

Scheduler dry-run:

- **WARNING: Failed to place all allocations.**
Task Group "echo" (failed to place 4 allocations):
 - * Resources exhausted on 6 nodes
 - * Dimension "memory exhausted" exhausted on 6 nodes



Terminal

```
$ nomad run -check-index 373 http-echo.nomad
==> Monitoring evaluation "372f378d"
      Evaluation triggered by job "http-echo-testing-iguana"
      Allocation "1072afb4" created: node "ec255589", group "echo"
      Allocation "1072afb4" status changed: "pending" -> "running"
      Allocation "14e2b099" status changed: "pending" -> "running"
      Allocation "6e6d2c9e" status changed: "pending" -> "running"
      Allocation "d2fd22c1" status changed: "pending" -> "running"
      Evaluation status changed: "pending" -> "complete"
==> Evaluation "372f378d" finished with status "complete" but failed to place
all allocations:
      Task Group "echo" (failed to place 4 allocations):
        * Resources exhausted on 6 nodes
```



Terminal

```
$ nomad alloc status 1e6892d3
ID                = 1e6892d3
Eval ID           = fff87c39
Name              = http-echo-llama.echo[2]
Node ID           = 744e4d82
Job ID            = http-echo-llama
Job Version       = 3
Client Status     = running
Client Description = <none>
Desired Status    = run
Desired Description = <none>
Created At        = 01/01/17 01:30:10 UTC

Task "server" is "running"
Task Resources
CPU      Memory      Disk    IOPS  Addresses
1/100 MHz 2.4 MiB/10 MiB 300 MiB 0      http: 10.1.1.14:21136
```


Revert Changes



Update the `http-echo.nomad` file to use the default resource constraints.

Plan and submit the job.

Service Discovery

Service Discovery



Provide a unified mechanism for addressing services in a microservices-oriented architecture.

Various techniques exist, but DNS is usually easiest as it requires zero-touch integration.

Recall: Previously



Previously our `http-echo` job was hard-bound to port 80, and we changed that to allow for scale.

We used `nomad alloc status` to "cheat" and see the port Nomad chose.

Service discovery is a better solution for identifying and addressing these microservices as they **move throughout the system**.

Move Throughout the System?



As jobs are scheduled, their host and port are unpredictable.

Moreover, as hosts join and leave the cluster, jobs may move throughout the cluster.

Service discovery will adapt to the moving jobs over time, without human intervention.

About: Consul



Consul is a free and open-source tool by HashiCorp that implements service discovery.

It uses the RAFT and gossip protocols to reach massive scale.

It has integrations with health checks, so unhealthy services are not added to the service discovery layer.

Similar client-server model to Nomad.

About: Consul



Even though this is not a Consul course, service discovery is a key component of a scheduled architecture.

Consul is already configured and running on your workstation on `127.0.0.1:8500`.

All Consul queries go through the local agent (do not query the service directly).

Exercise: Run `consul members`



Execute the `consul members` command to list all the cluster members.



Terminal

```
$ consul members
```

Node	Address	Status	Type	Build	Protocol	DC	Segment
server-0	10.1.1.104:8301	alive	server	0.9.3	2	dc1	<all>
server-1	10.1.2.135:8301	alive	server	0.9.3	2	dc1	<all>
server-2	10.1.1.106:8301	alive	server	0.9.3	2	dc1	<all>
goldfish	10.1.1.174:8301	alive	client	0.9.3	2	dc1	<default>
grasshopper	10.1.2.96:8301	alive	client	0.9.3	2	dc1	<default>
llama	10.1.1.14:8301	alive	client	0.9.3	2	dc1	<default>

About: Consul DNS Service Discovery



Randomized round-robin to all services that match the query.

Filters based on health checks (unhealthy hosts are not returned from the query).

"Health" is determined by the application.

About: Consul DNS Service Discovery



Previously we queried `$(identity).http-echo.service.consul`.

`$(identity)` is a tag (services can have zero or more tags).

`http-echo` is the logical service name (could have more than one).

`service` is the DNS namespace which queries Consul services.

`consul` is the DNS suffix which delegates system DNS to Consul.



Terminal

```
$ curl <tag>.<logical-service-name>.service.consul
```





Terminal

```
$ curl $(identity).http-echo.service.consul  
curl: (7) Failed to connect to anaconda.http-echo.service.consul port 80:  
Connection refused
```





Terminal

```
$ dig +short $(identity).http-echo.service.consul  
10.1.1.207  
10.1.1.116  
10.1.2.187
```





Terminal

```
$ dig +short SRV $(identity).http-echo.service.consul  
1 1 27196 testing-server-0.node.dc1.consul.  
1 1 33946 testing-server-0.node.dc1.consul.  
1 1 43831 testing-server-2.node.dc1.consul.
```





Terminal

```
$ dig +short SRV $(identity).http-echo.service.consul  
1 1 27196 testing-server-0.node.dc1.consul.  
1 1 33946 testing-server-0.node.dc1.consul.  
1 1 43831 testing-server-2.node.dc1.consul.
```





Terminal

```
$ curl $(identity).http-echo.service.consul:27196  
smiling banana
```





Terminal

```
$ curl $(identity).http-echo.service.consul:40633  
smiling banana
```

```
$ curl $(identity).http-echo.service.consul:40633  
smiling banana
```

```
$ curl $(identity).http-echo.service.consul:40633  
smiling banana
```

```
$ curl $(identity).http-echo.service.consul:40633  
smiling banana
```

Lessons: Service Discovery



Scheduled architectures make heavy use of service discovery.

Service discovery is integrated with the health of the application.

Nomad uses Consul for service discovery by default.

Load Balancing

Load Balancing



Load balancing is a close cousin of service discovery.

Allows providing a known URL or path to other services.

Round robin and integrates health checks.

Load Balancing in Nomad



Load balancing in Nomad is possible through the use of fabio, an open source tool

Fabio integrates with Consul and acts as a load balancer for all healthy services in a given name.

Consul Connect support coming soon!



http-echo.nomad

```
job "http-echo-anaconda" {  
  
    service {  
        name = "http-echo"  
        port = "http"  
  
        tags = [  
            "testing-iguana",  
            "urlprefix-/http-echo",  
        ]  
  
        check {  
            type      = "http"  
            path      = "/health"  
            interval  = "10s"  
            timeout   = "5s"  
        }  
    }  
}
```

Exercise: Load Balance



Using your local machine's public IP address, visit the `/http-echo` URL on port 9999 in your browser.

(9999 is the default fabio port)

HINT: Execute `public_ip`




Terminal

```
$ public_ip  
52.206.242.2
```



Visit in Browser



Exercise: Add URL



Add an additional URL suffix where your service should be addressable.


Resubmit the nomad job.

Visit the URL in your browser.



http-echo.nomad

```
tags = [  
  "testing-iguana",  
  "urlprefix-/http-echo",  
  "urlprefix-/my-path",  
]
```





Terminal

```
$ nomad run http-echo.hcl
==> Monitoring evaluation "430abac8"
      Evaluation triggered by job "http-echo-testing-iguana"
      Allocation "f222da9b" created: node "ec255589", group "echo"
      Allocation "7082864e" created: node "ec274728", group "echo"
      Allocation "7b8f25ef" created: node "ec274728", group "echo"
      Allocation "8e45b28d" created: node "ec255589", group "echo"
      Allocation "acddd1ba" created: node "ec255cce", group "echo"
      Allocation "043678c3" created: node "ec277448", group "echo"
      Allocation "092a782c" created: node "ec261569", group "echo"
      Allocation "3cba85a7" created: node "ec261569", group "echo"
      Allocation "6b5efd95" created: node "ec2ee7db", group "echo"
      Allocation "c353d506" created: node "ec2ee7db", group "echo"
      Evaluation status changed: "pending" -> "complete"
==> Evaluation "430abac8" finished with status "complete"
```

Downloading Artifacts

Nomad Artifacts



Many jobs need external files to run such as configuration or the even the binary itself.

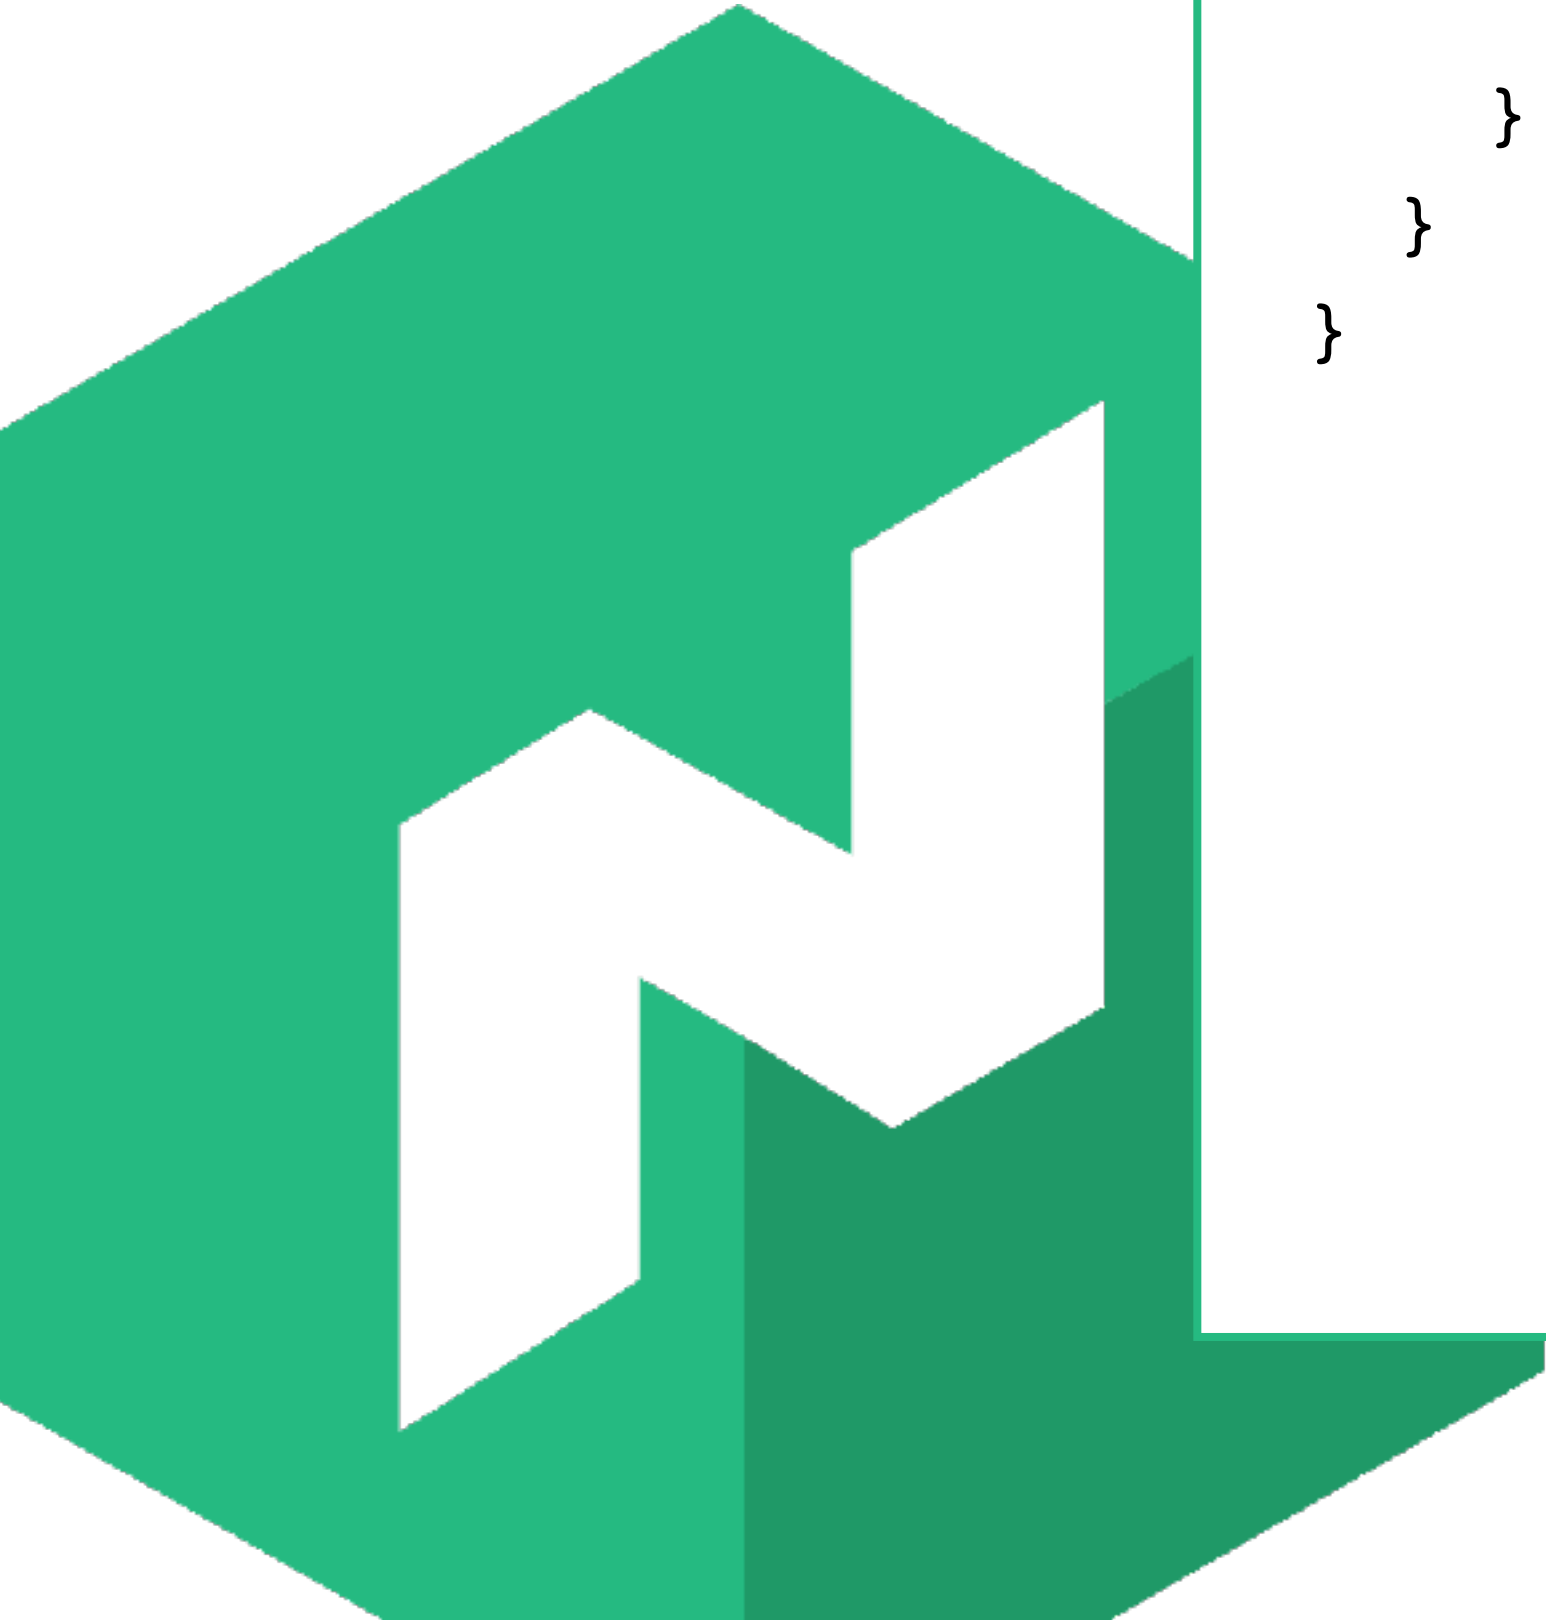
Nomad supports downloading artifacts from many different sources including S3 buckets, git, and more.

Optional checksum verification can ensure integrity.

Automatically extracts commonly-known archive types.



```
job "job" {  
  group "group" {  
    task "task" {  
      artifact {  
        source = "https://example.com/file.tar.gz"  
      }  
    }  
  }  
}
```

```
job "job" {
  group "group" {
    task "task" {
      artifact {
        source      = "https://example.com/file.tar.gz"
        destination = "/tmp"
      }
    }
  }
}
```

Templating

Nomad Templating



Nomad can take an input template, apply interpolations, and produce an output template for a task.

The template responds to signals and handles upstream updates.

Can query data from Nomad, Consul, and Vault.



job.nomad

```
job "job" {
  group "group" {
    task "task" {
      template {
        data = <<EOH
        bind_port: {{ env "NOMAD_PORT_db" }}
        service_id: {{ key "service/my-key" }}
        EOH
        destination = "local/foo"
      }
    }
  }
}
```



job.nomad

```
job "job" {
  group "group" {
    task "task" {
      artifact {
        source      = "https://example.com/file.tpl"
        destination = "local/"
      }

      template {
        source      = "local/file.tpl"
        destination = "local/file.yml"
      }
    }
  }
}
```

Advanced Job Strategies

Advanced Job Strategies



Nomad currently supports two advanced job strategies:

- Rolling Upgrades
- Blue/Green & Canary Deployments

Advanced Job Strategies: Rolling Upgrades



Nomad supports rolling updates as a first class feature.

To enable rolling updates a job or task group is annotated with a high-level description of the update strategy using the **update** stanza.



http-echo.nomad

```
job "http-echo-anaconda" {  
  
    # Add an update stanza to enable rolling updates of the service  
    update {  
        max_parallel = 1  
        min_healthy_time = "30s"  
        healthy_deadline = "10m"  
    }  
  
    ...  
}
```

Exercise: Add update stanza to http-echo job



```
job "http-echo-anaconda" {  
  
    # Add an update stanza to enable rolling updates of the service  
    update {  
        max_parallel = X  
        min_healthy_time = "Xs"  
    }  
  
    .....  
}
```

Open **<http://nomad.hashicorp.live/>** in your browser.

Advanced Job Strategies: Blue/Green & Canary



Sometimes rolling upgrades do not offer the required flexibility for updating an application in production.

Often organizations prefer to put a "canary" build into production or utilize a technique known as a "blue/green" deployment to ensure a safe application rollout to production while minimizing downtime.



http-echo.nomad

```
job "http-echo-anaconda" {  
  
  update {  
    max_parallel      = 1  
    canary            = 5  
    min_healthy_time  = "30s"  
    healthy_deadline  = "10m"  
    auto_revert       = true  
  }  
  
  ...  
}
```



Terminal

```
$ nomad plan http-echo.nomad
+/- Job: "http-echo"
+/- Task Group: "http-echo" (1 canary, 1 ignore)
.....
```





Terminal

```
$ nomad status http-echo
```

```
ID           = http-echo
```

```
Type         = service
```

```
Status       = running
```

```
Periodic     = false
```

```
Parameterized = false
```

```
Latest Deployment
```

```
ID           = 32a080c1
```

```
Status       = running
```

```
Description = Deployment is running but requires promotion
```

```
Deployed
```

Task Group	Auto Revert	Promoted	Desired	Canaries	Placed	Healthy
------------	-------------	----------	---------	----------	--------	---------

Unhealthy						
-----------	--	--	--	--	--	--

api	true	false	1	1	1	1	0
-----	------	-------	---	---	---	---	---



Terminal

```
$ nomad deployment promote 32a080c1
==> Monitoring evaluation "61ac2be5"
      Evaluation triggered by job "docs"
      Evaluation within deployment: "32a080c1"
      Evaluation status changed: "pending" -> "complete"
==> Evaluation "61ac2be5" finished with status "complete"
```



Terminal


```
$ nomad status http-echo
```



Nomad UI

<http://nomad.hashicorp.live/>

Demo



Q&A